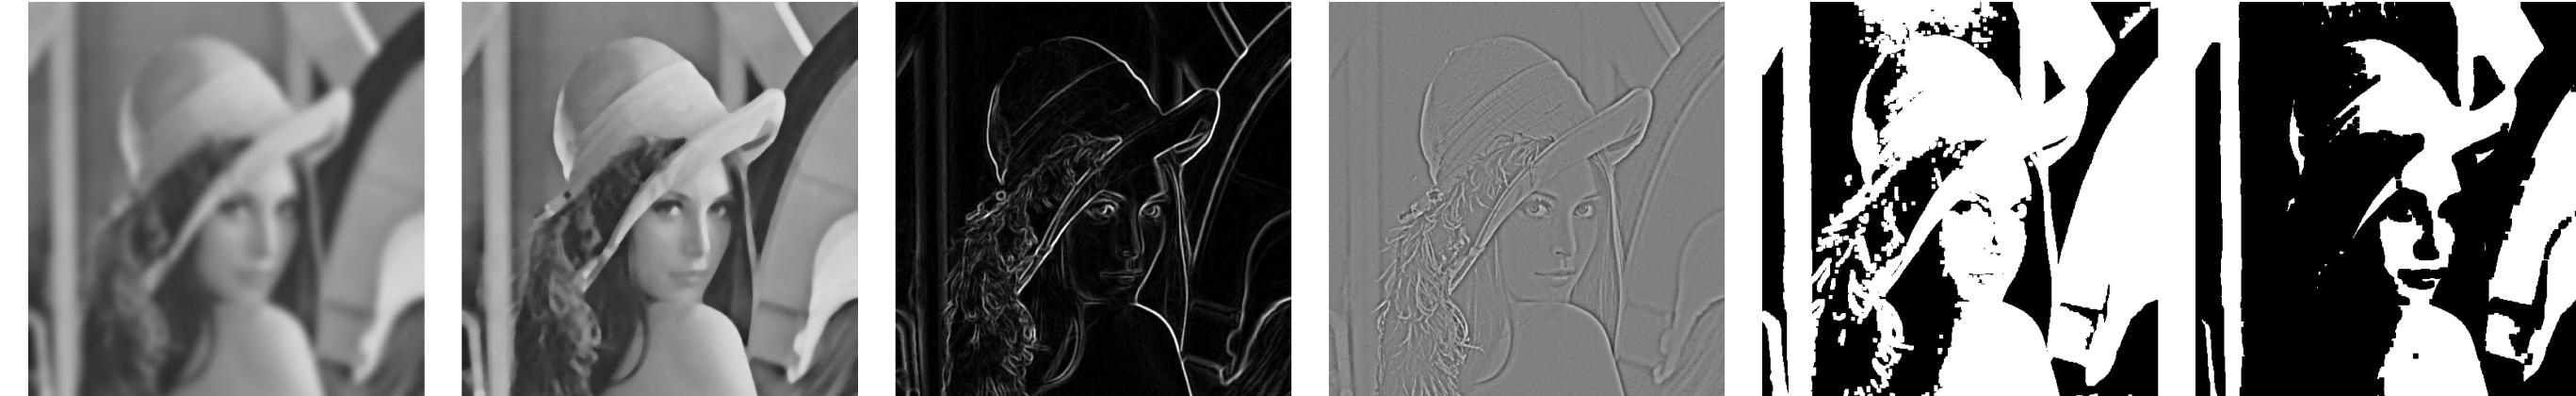




# CPSC 425: Computer Vision



## Lecture 5: Image Filtering (final)

( unless otherwise stated slides are taken or adopted from **Bob Woodham, Jim Little and Fred Tung** )

# Menu for Today

## Topics:

- **Linear Filtering** recap
- Efficient convolution, Fourier aside
- **Quiz 1**
- **Non-linear Filters:**  
Median, ReLU, Bilateral Filter

## Readings:

- **Today's Lecture:** Szeliski 3.3-3.4, Forsyth & Ponce (2nd ed.) 4.4

## Reminders:

- **Assignment 1:** Image Filtering and Hybrid Images due **September 28th**

# New TA Minglong + Updated Office Hours

Bereket Guta



Rayat Hossain



Minglong Li



Mon 3-4, In Person

Wed 3-4, In Person

Fri 11-12, Zoom

See **Canvas** for Links and Locations

# Minor Error in Last week's notes

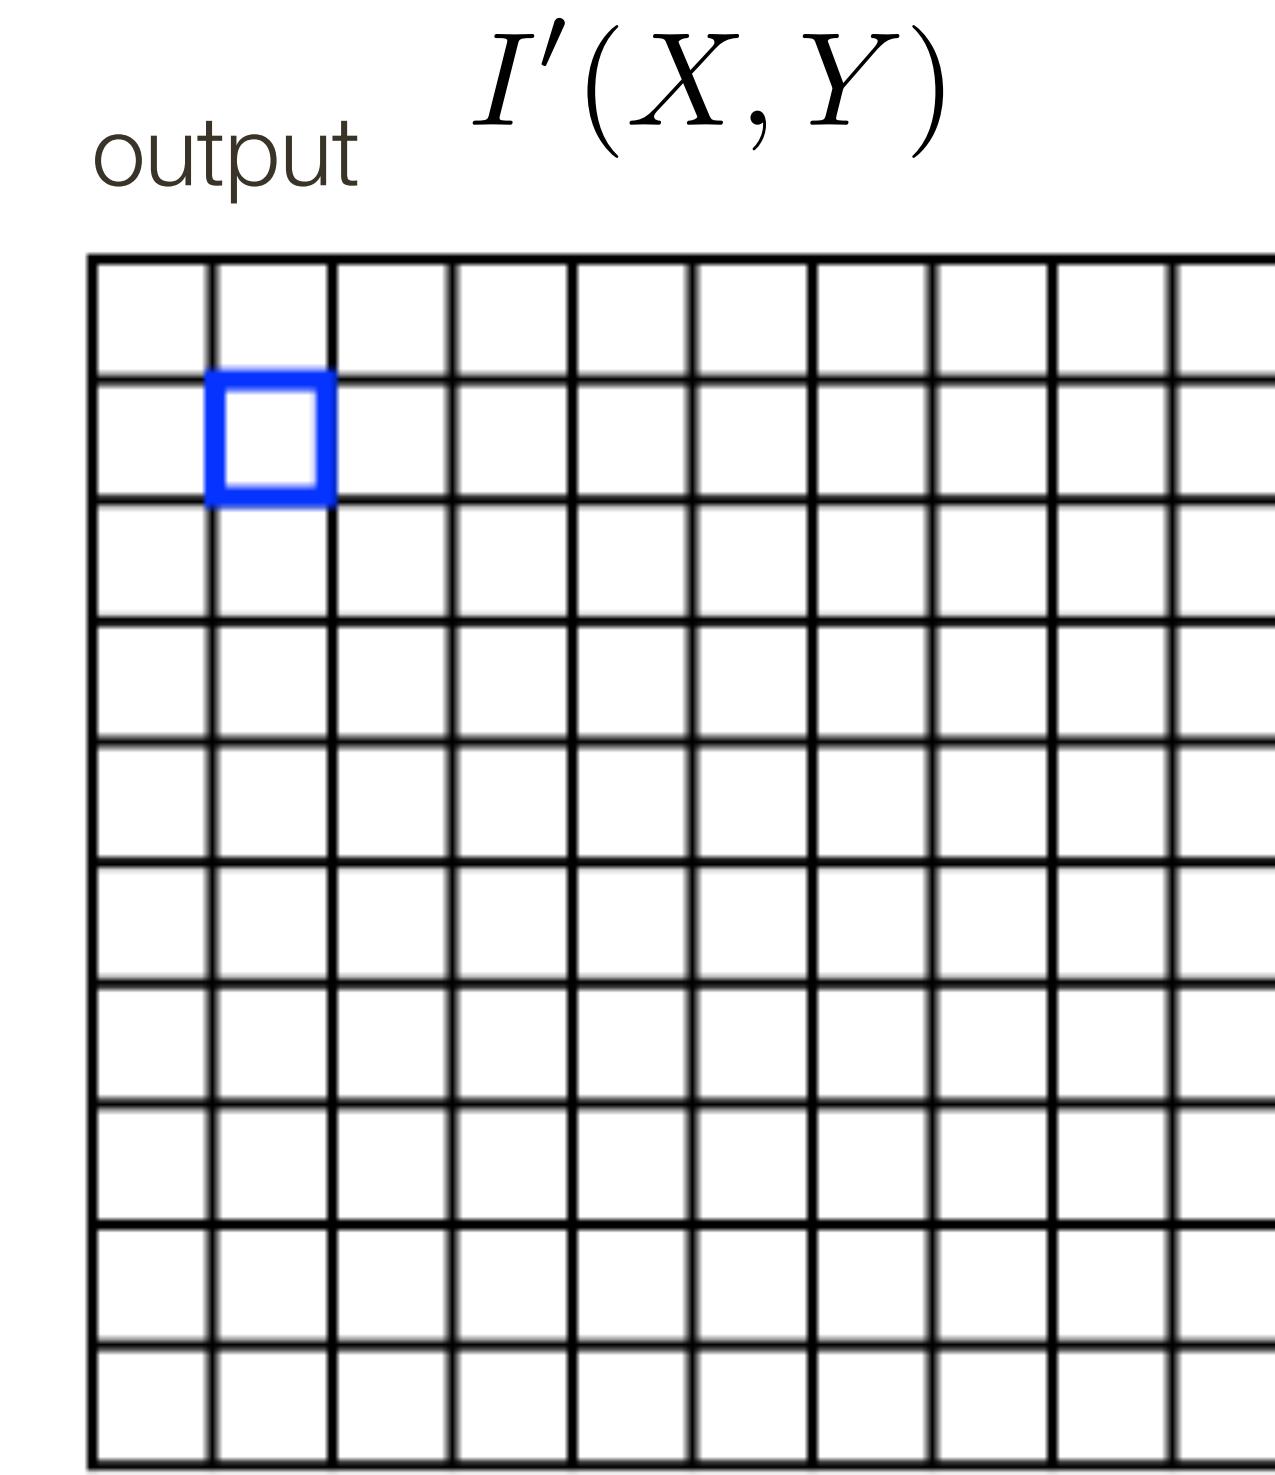
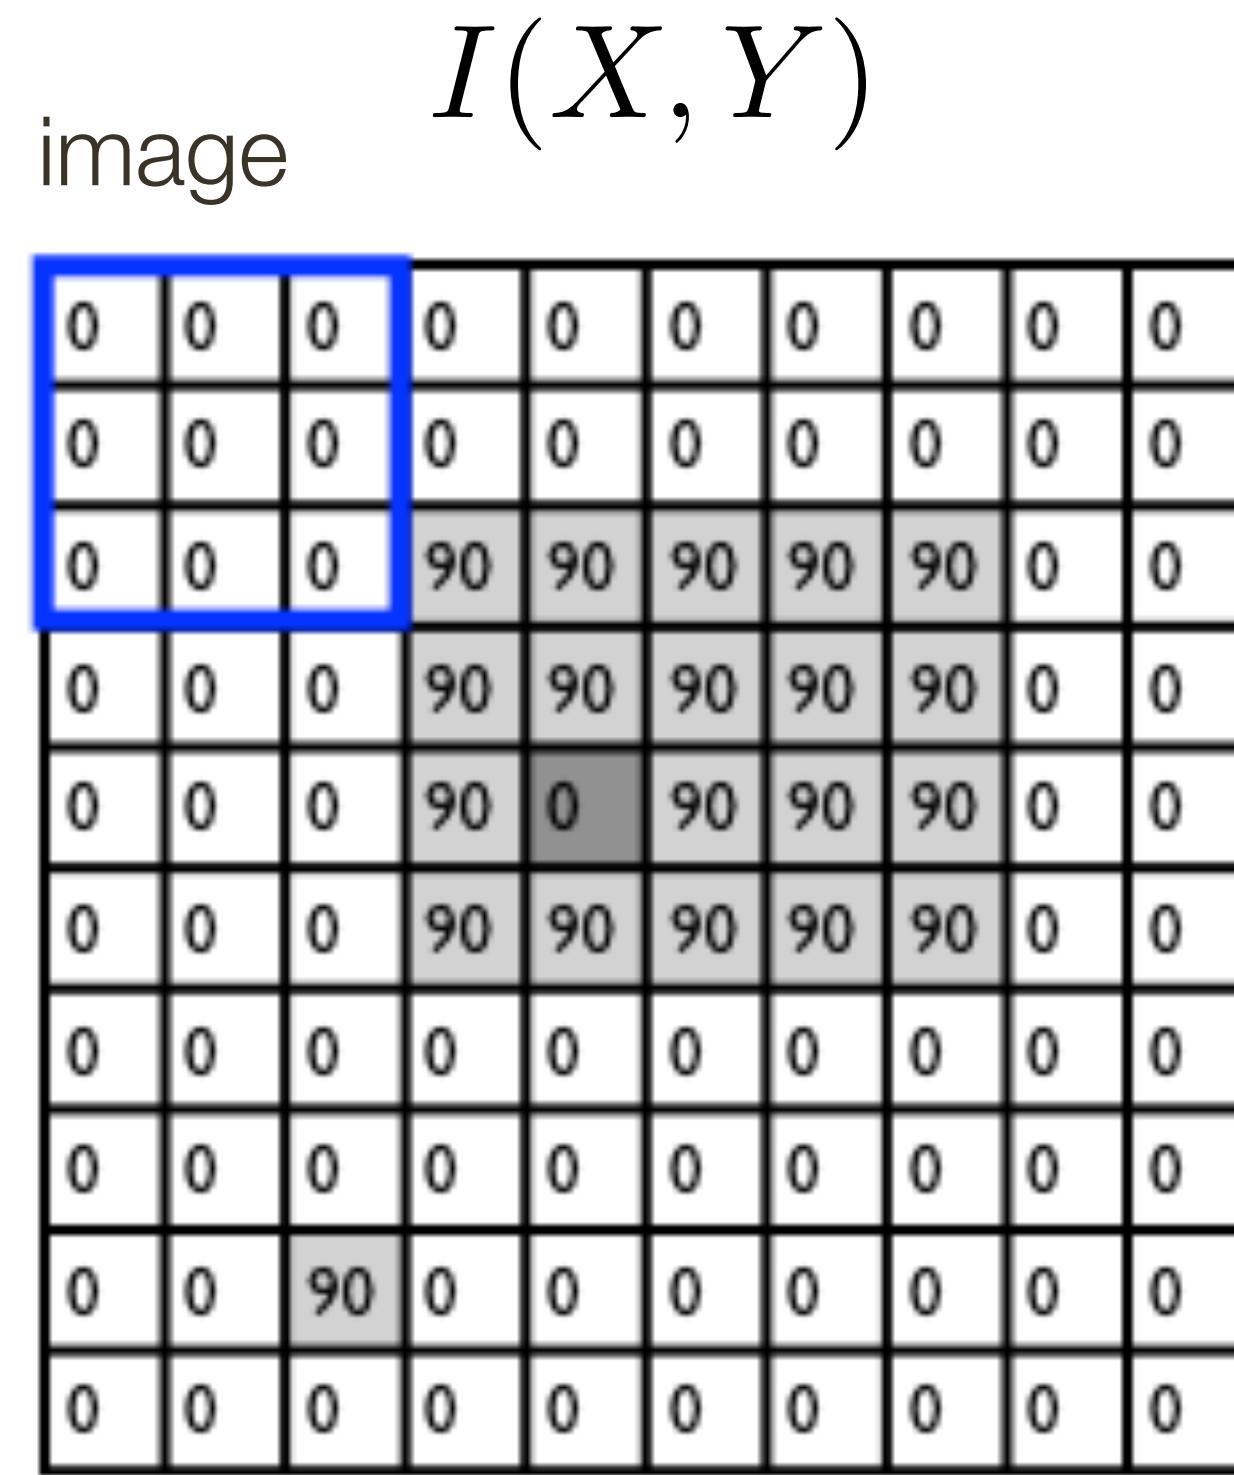
$$\begin{aligned} & \frac{4.2}{\text{corr}(I, F)} = \sum_i \sum_j \frac{I(x+i, y+j) F(i, j)}{F(i) F(j)} \quad F(i, j) = F(i) F(j) \\ &= \sum_i F(i) \sum_j I(x+i, y+j) F(j) \\ & \quad \underbrace{\qquad \qquad}_{\text{corr}(I(x+i, y), F(y)) = I^*(x+i, y)} = I^*(x+i, y) \\ &= (I(x, y) * F(y)) * F(x) \end{aligned}$$

# Linear Filter Example

$$F(X, Y)$$

filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



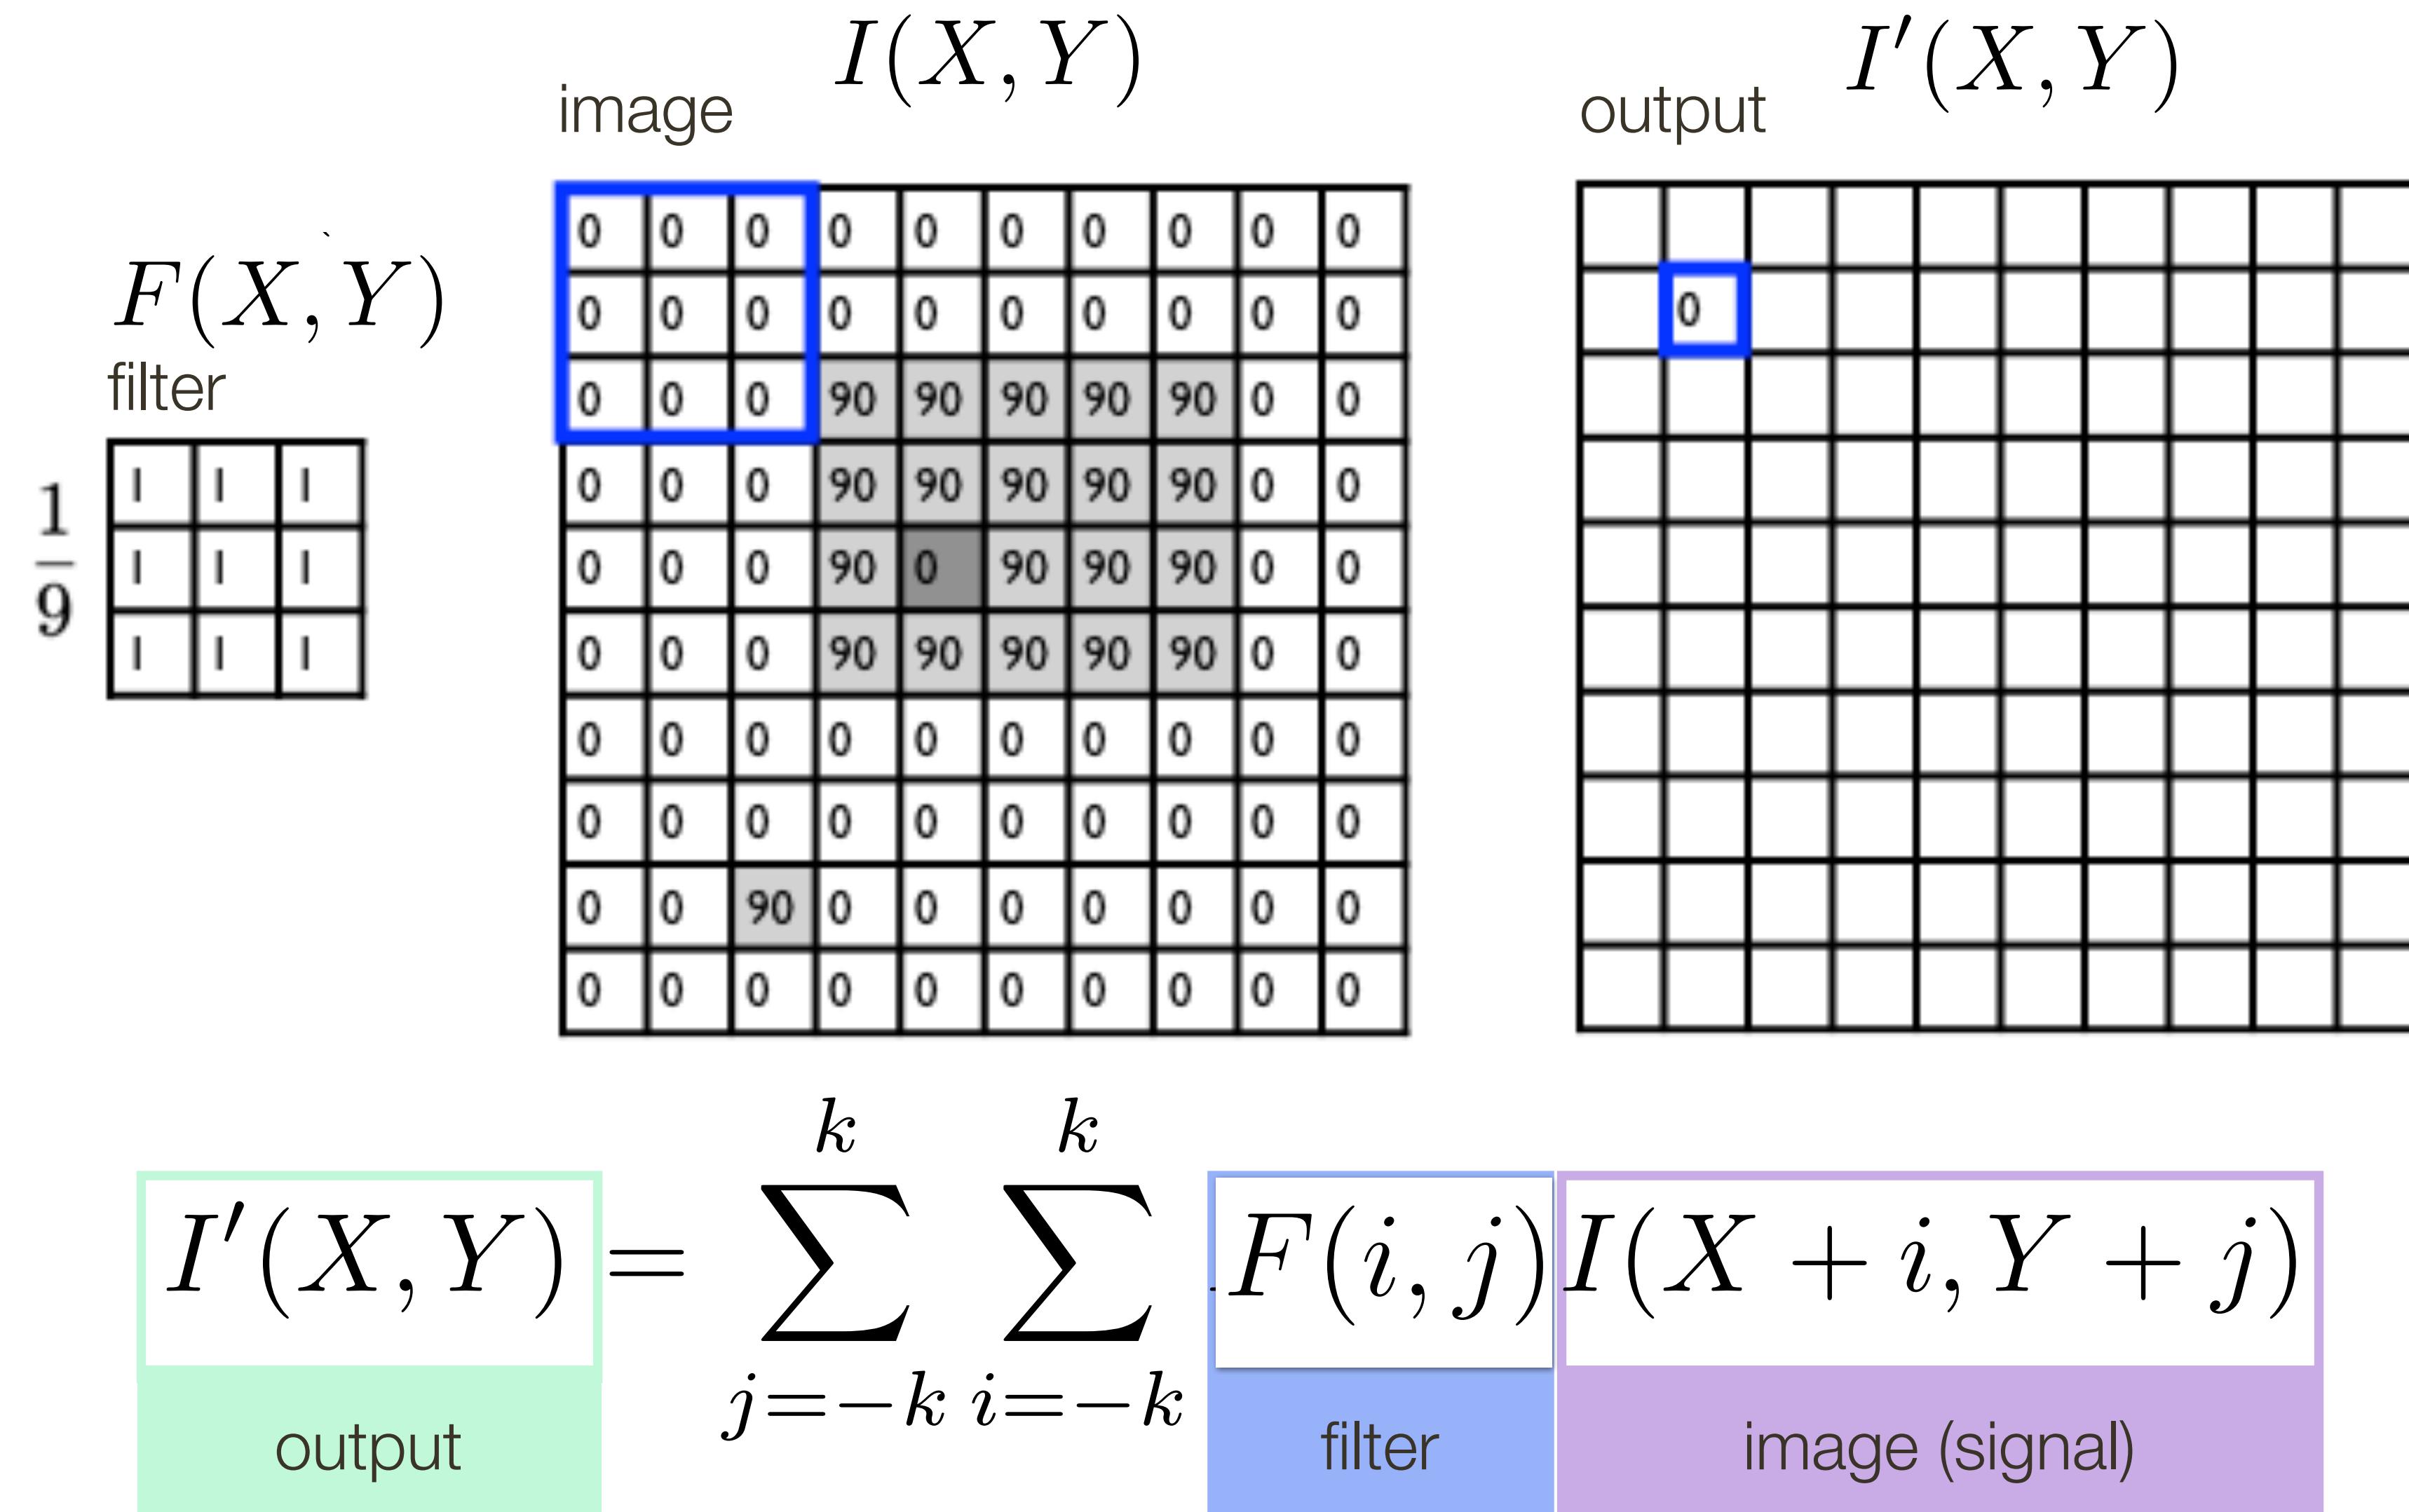
$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

output

filter

image (signal)

# Linear Filter Example

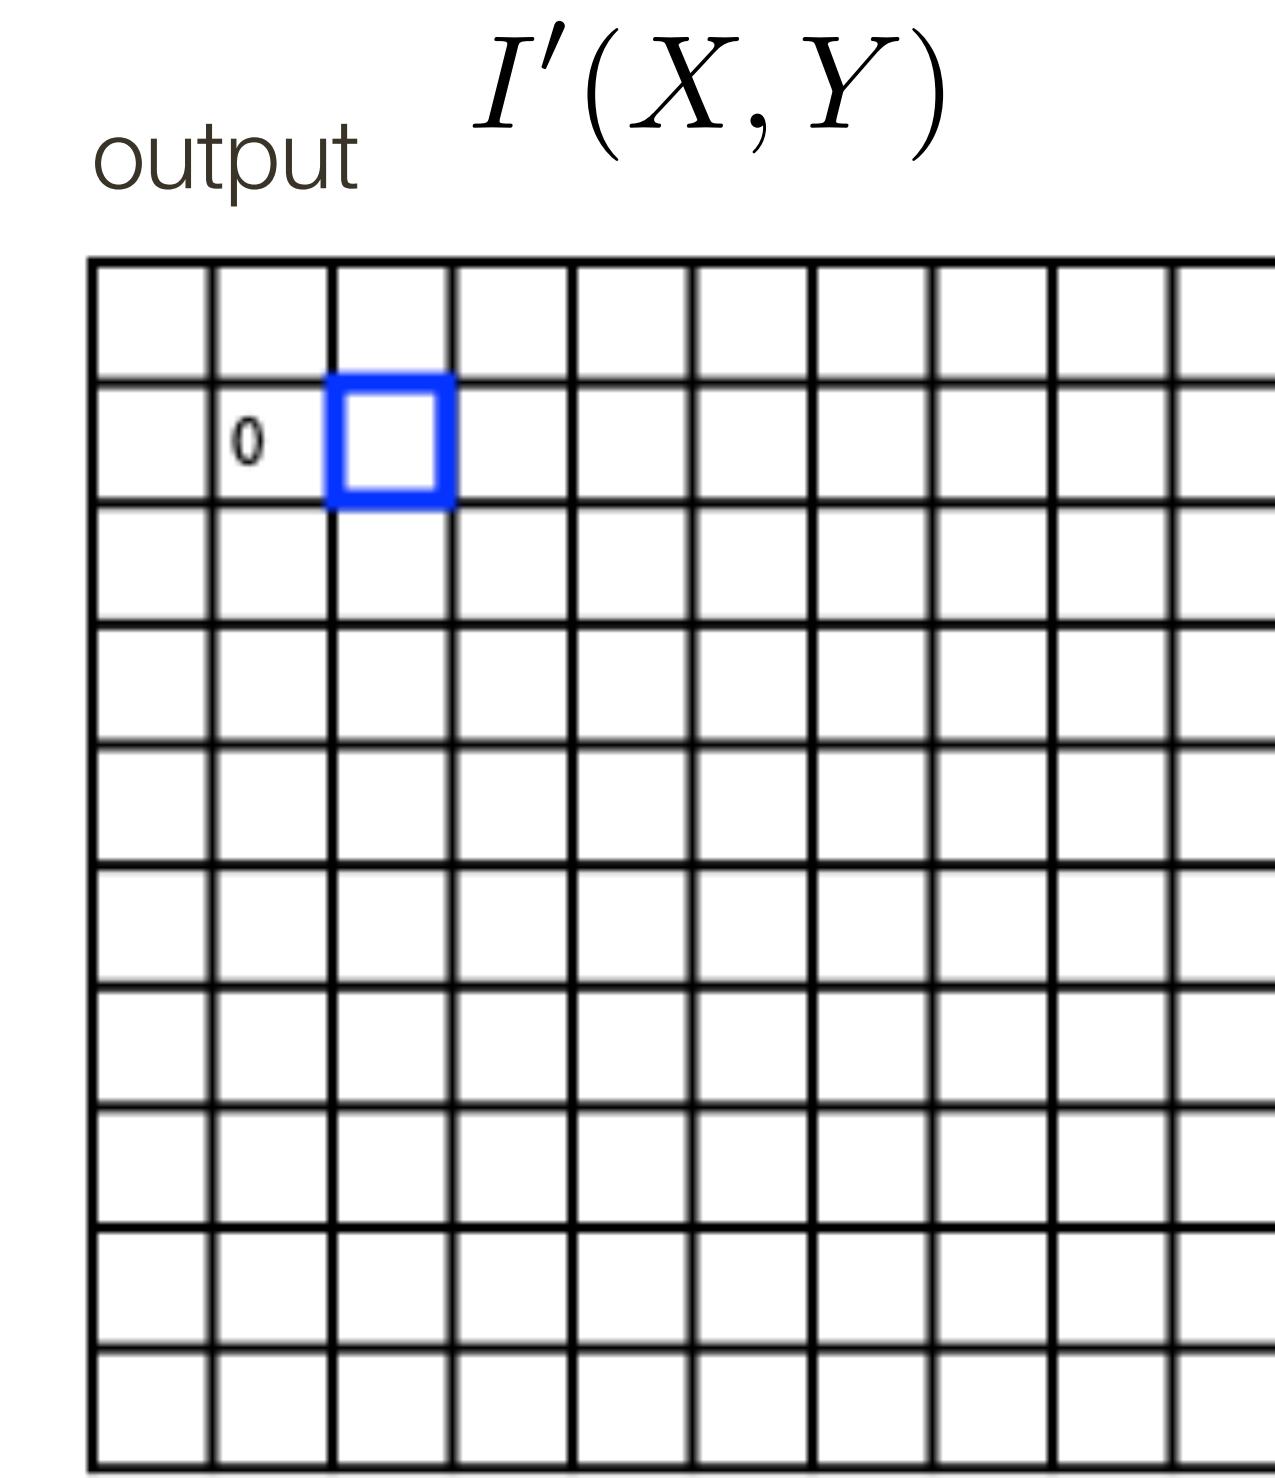
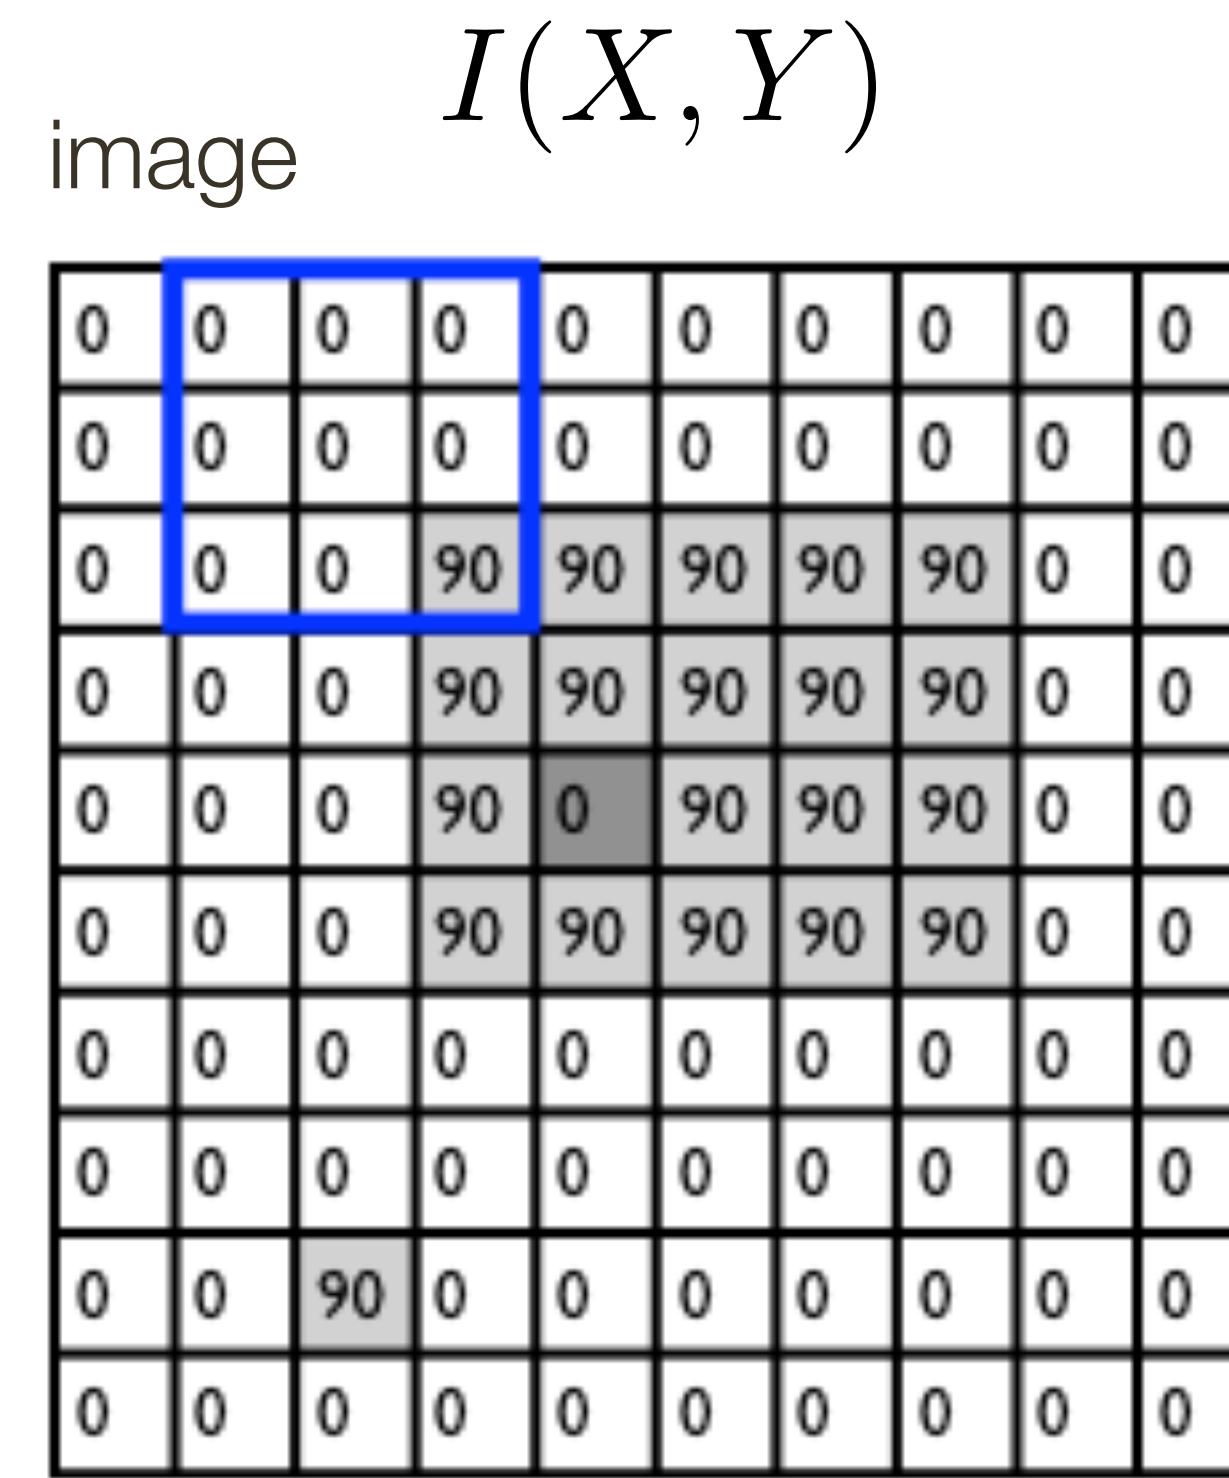


# Linear Filter Example

$$F(X, Y)$$

filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

output

filter

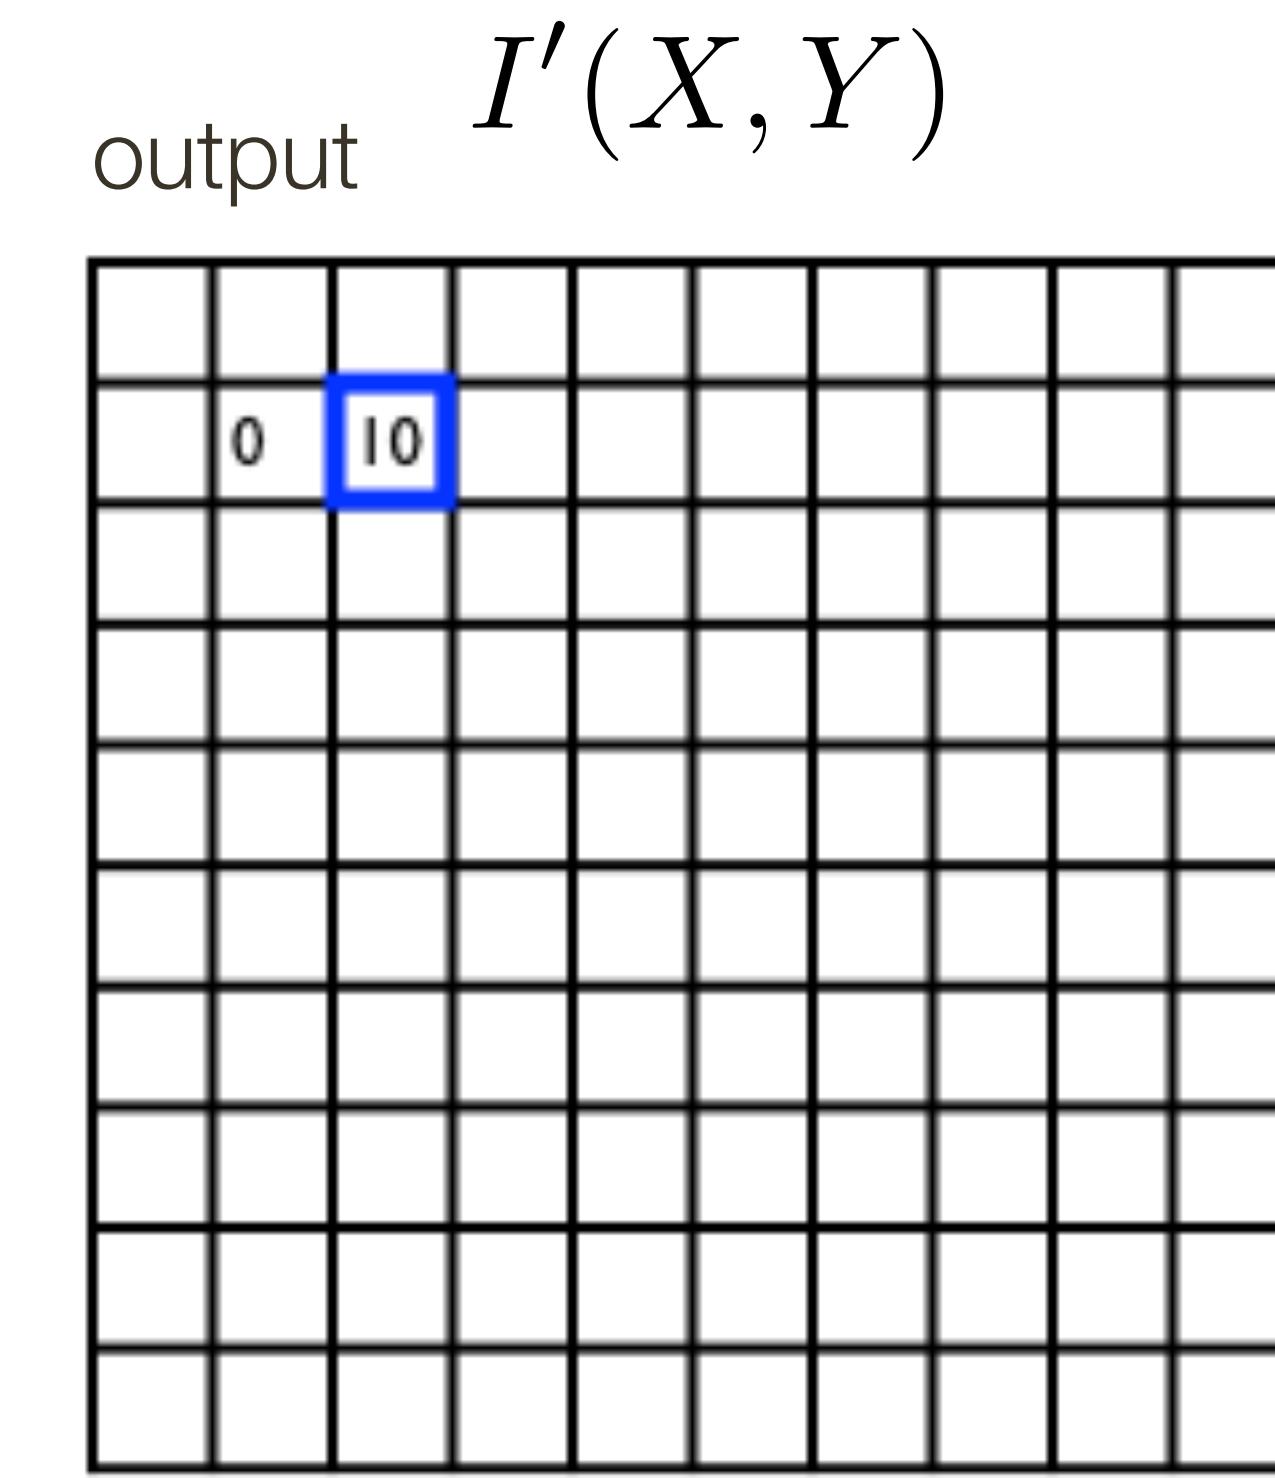
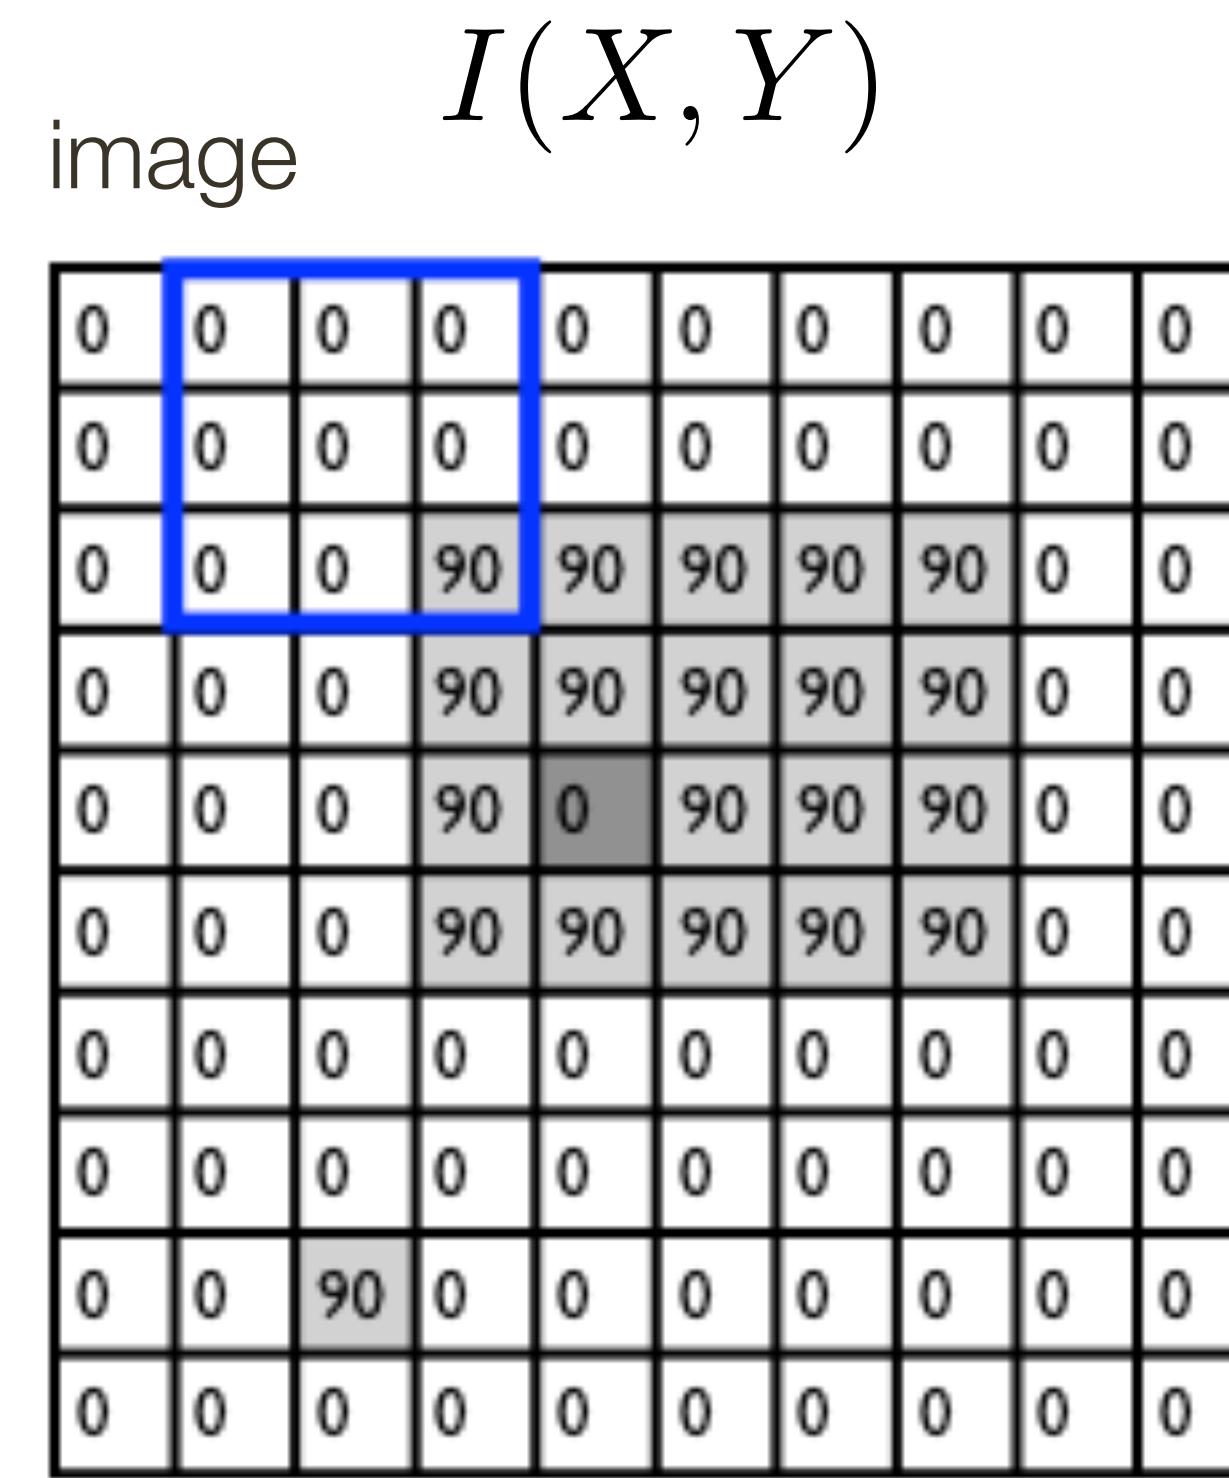
image (signal)

# Linear Filter Example

$$F(X, Y)$$

filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



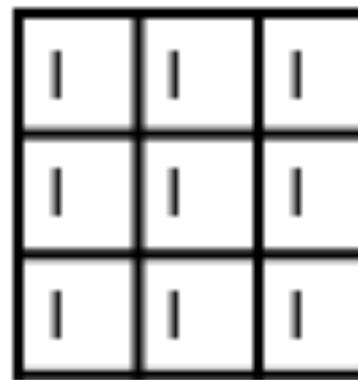
$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

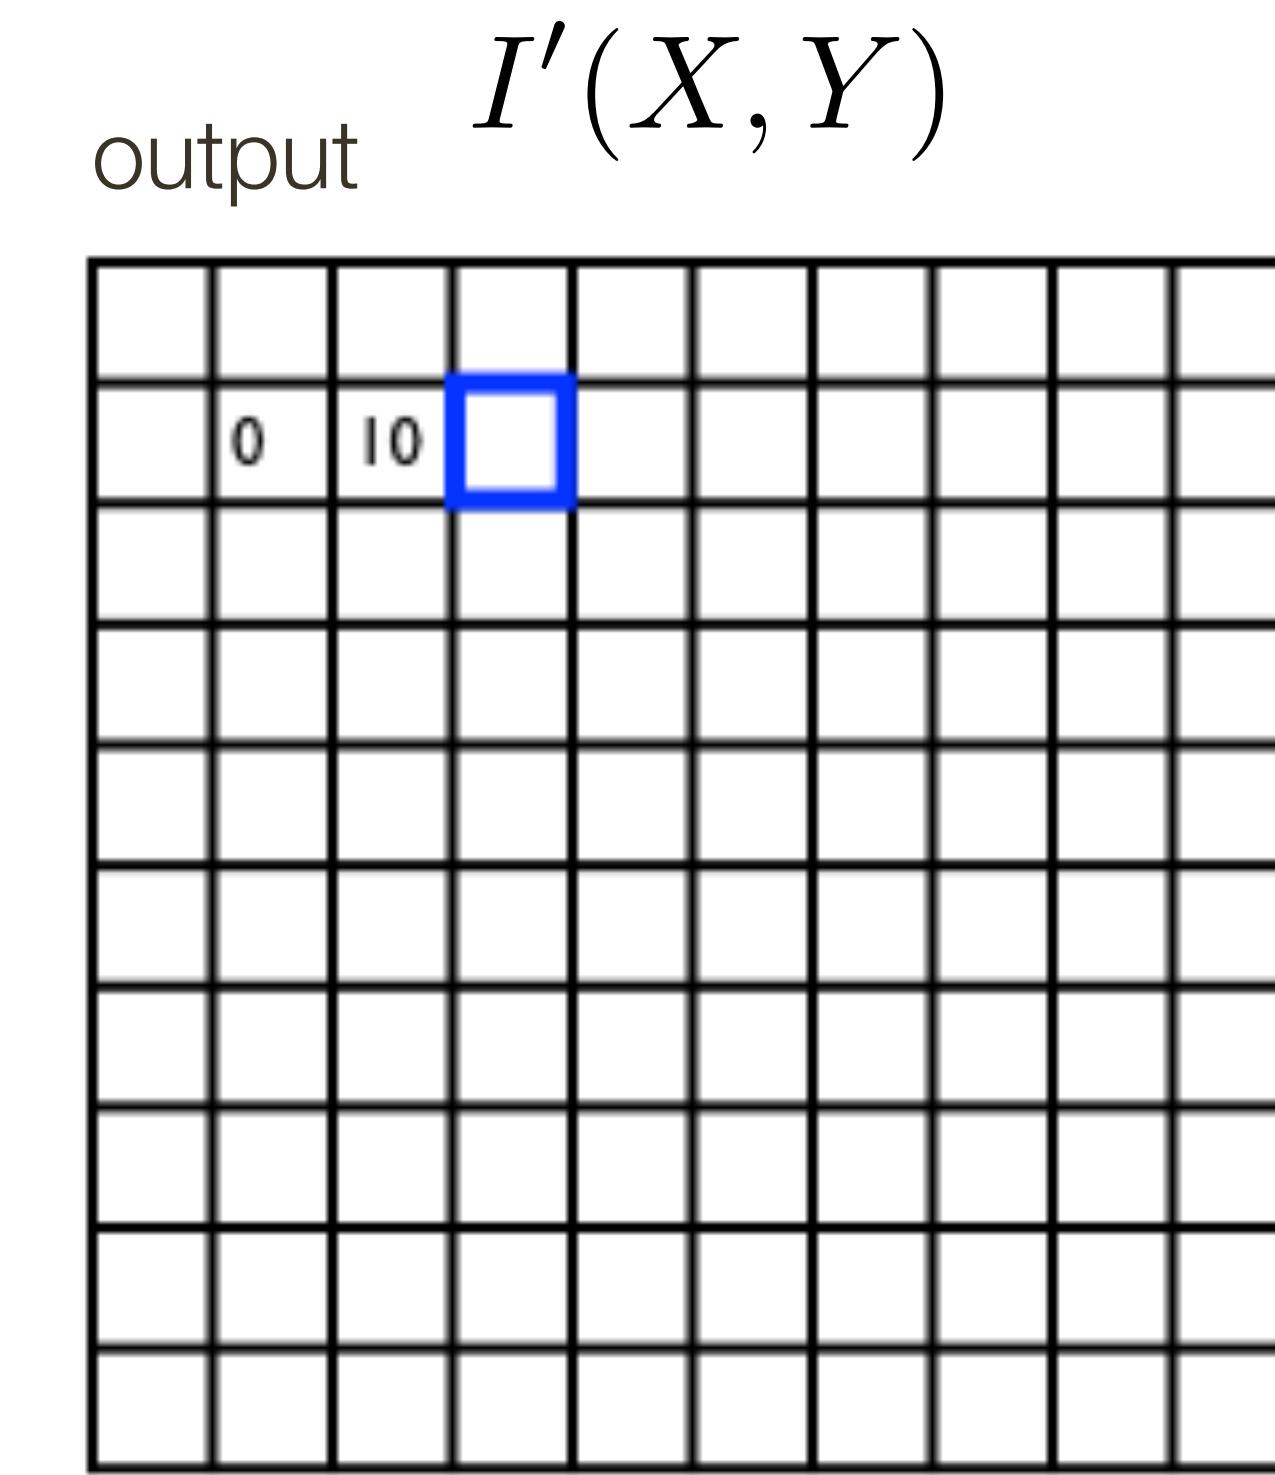
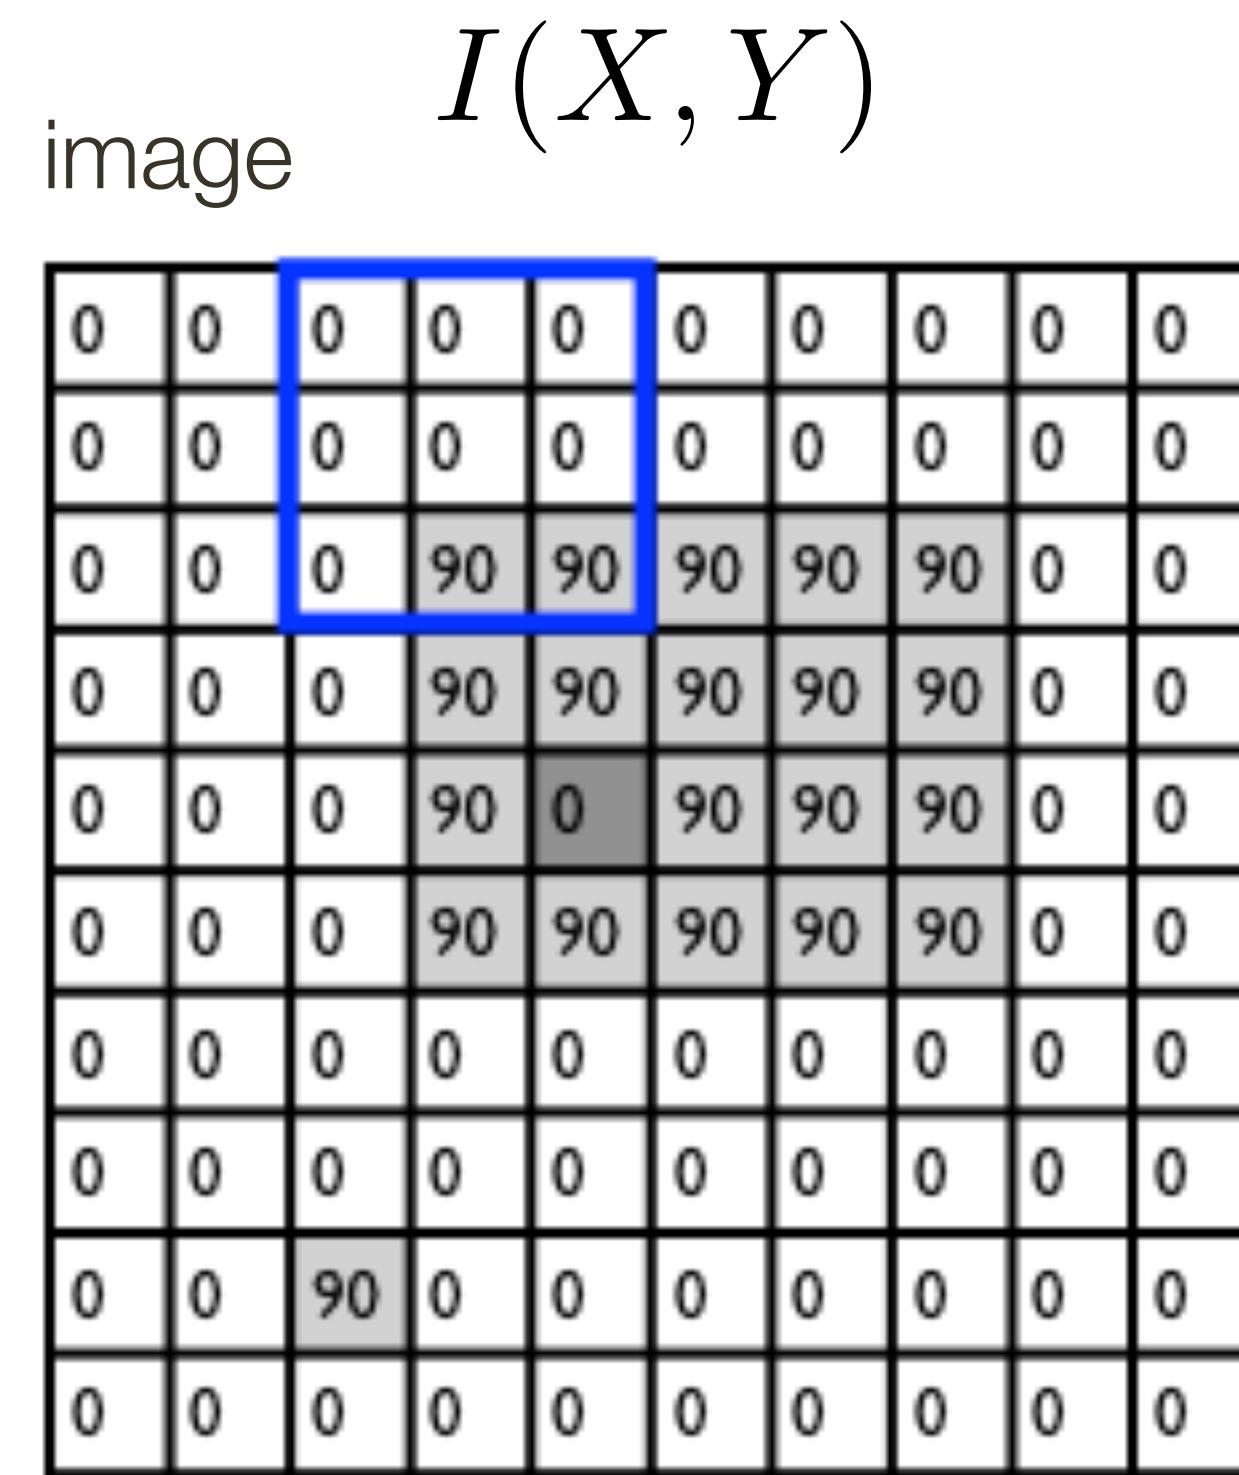
output

filter

image (signal)

# Linear Filter Example

$F(X, Y)$   
filter  
 $\frac{1}{9}$   


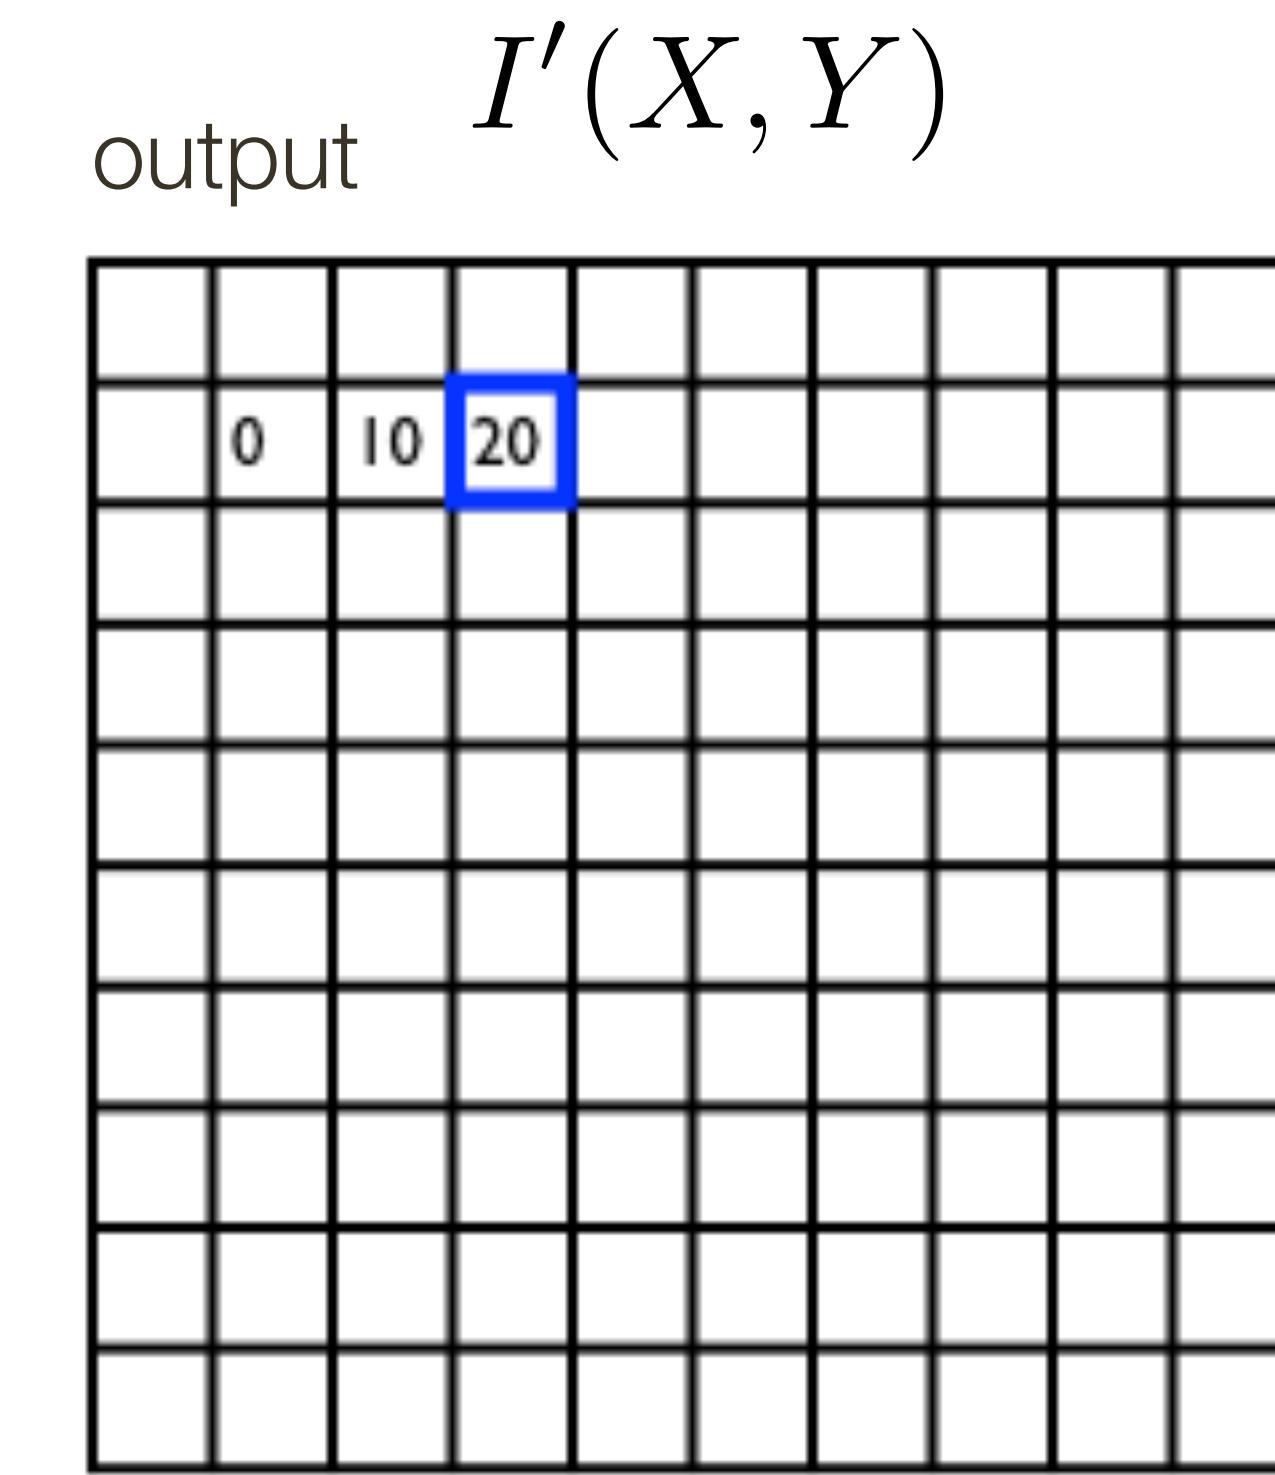
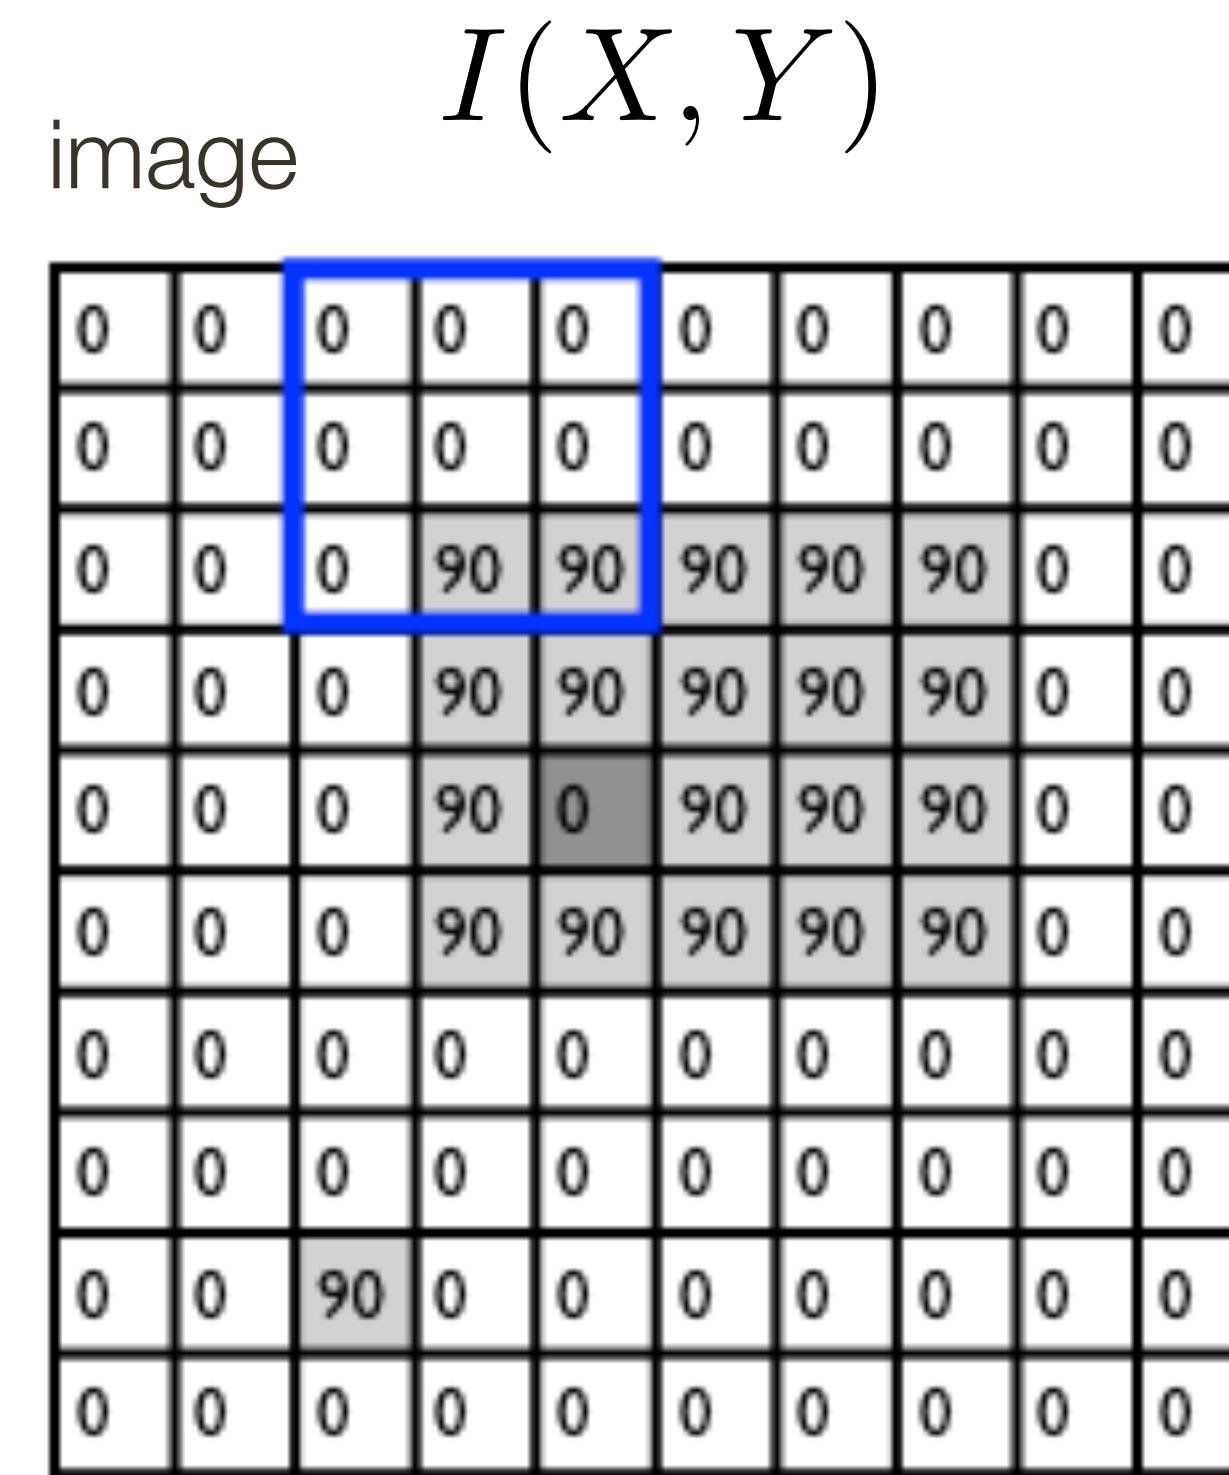


$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

output                      filter                      image (signal)

# Linear Filter Example

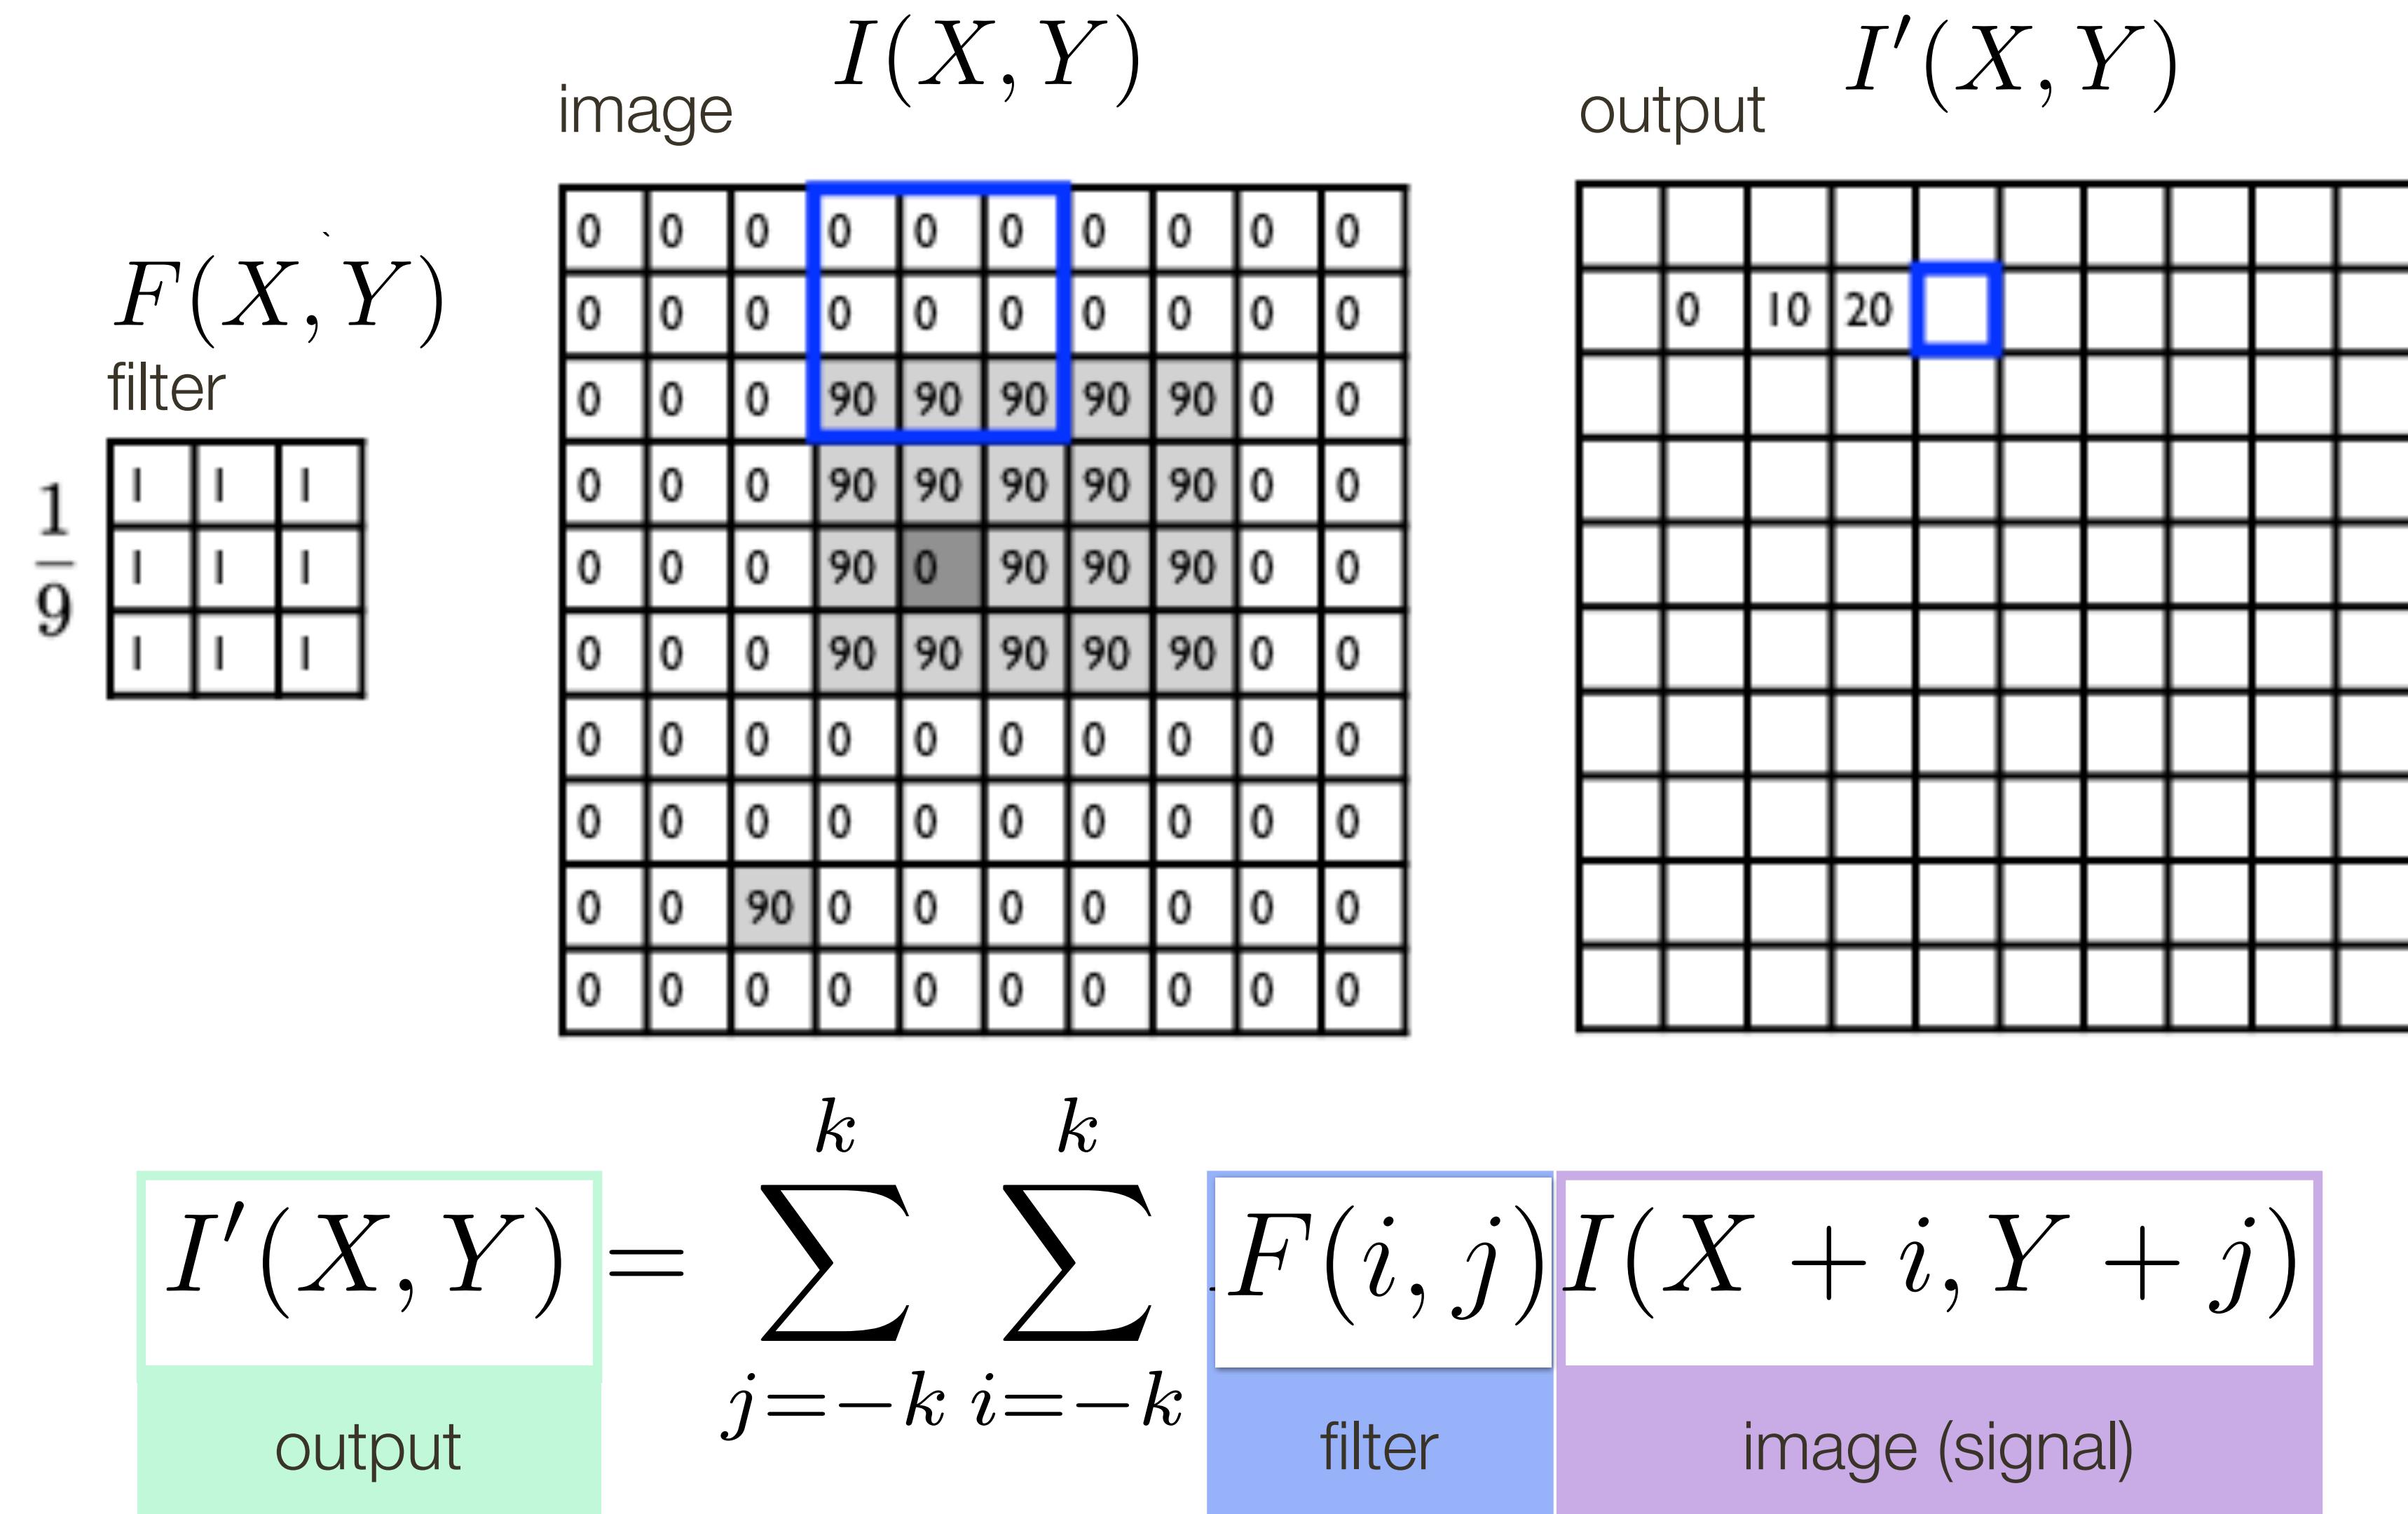
$$F(X, Y) \text{ filter}$$
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

output                          filter                          image (signal)

# Linear Filter Example

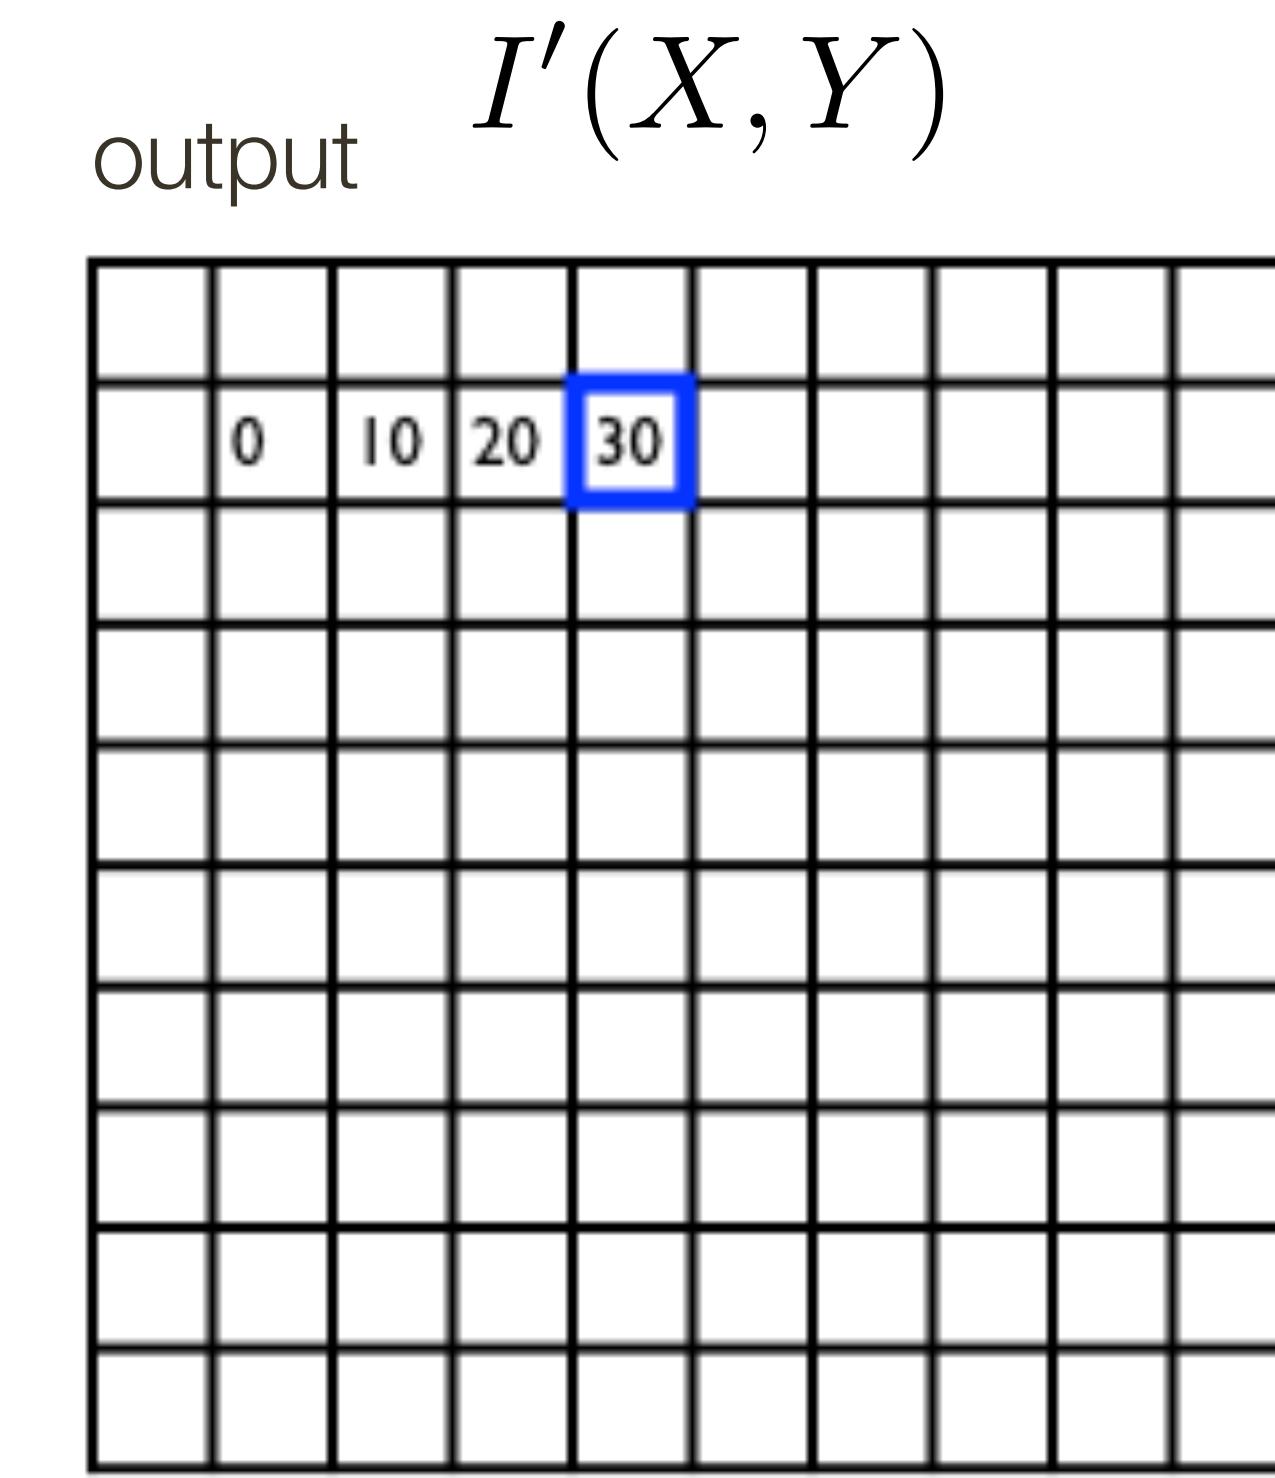
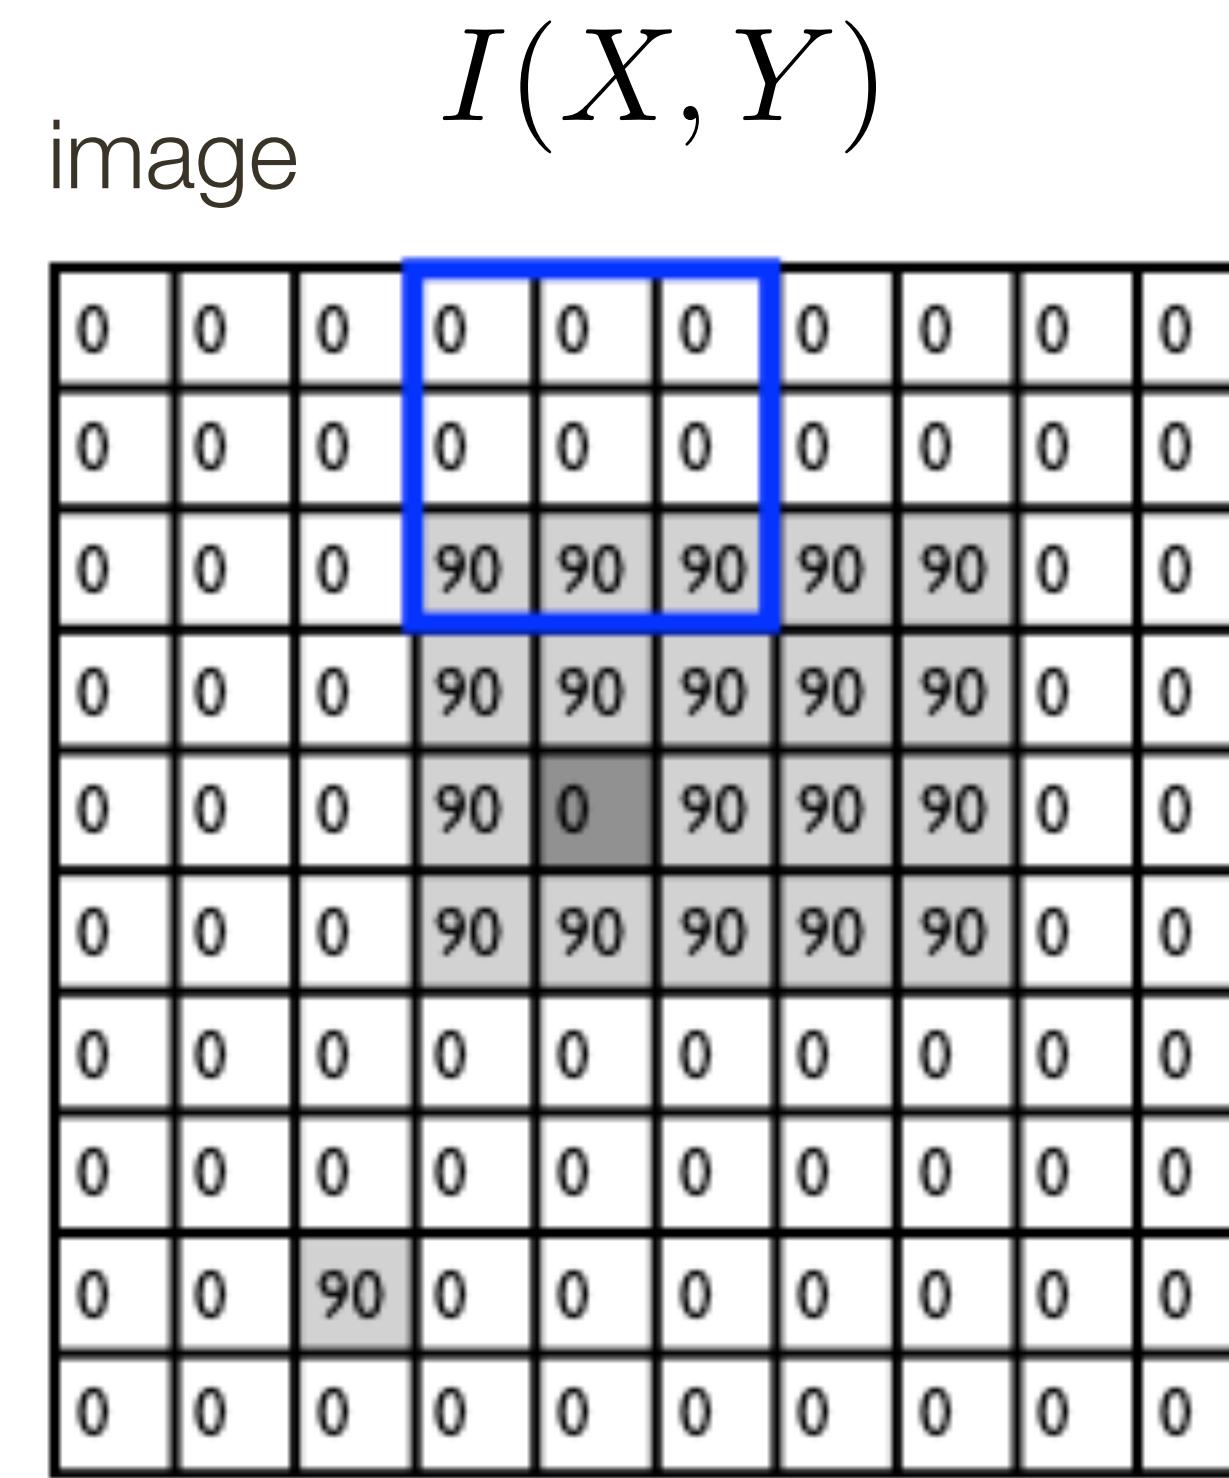


# Linear Filter Example

$$F(X, Y)$$

filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



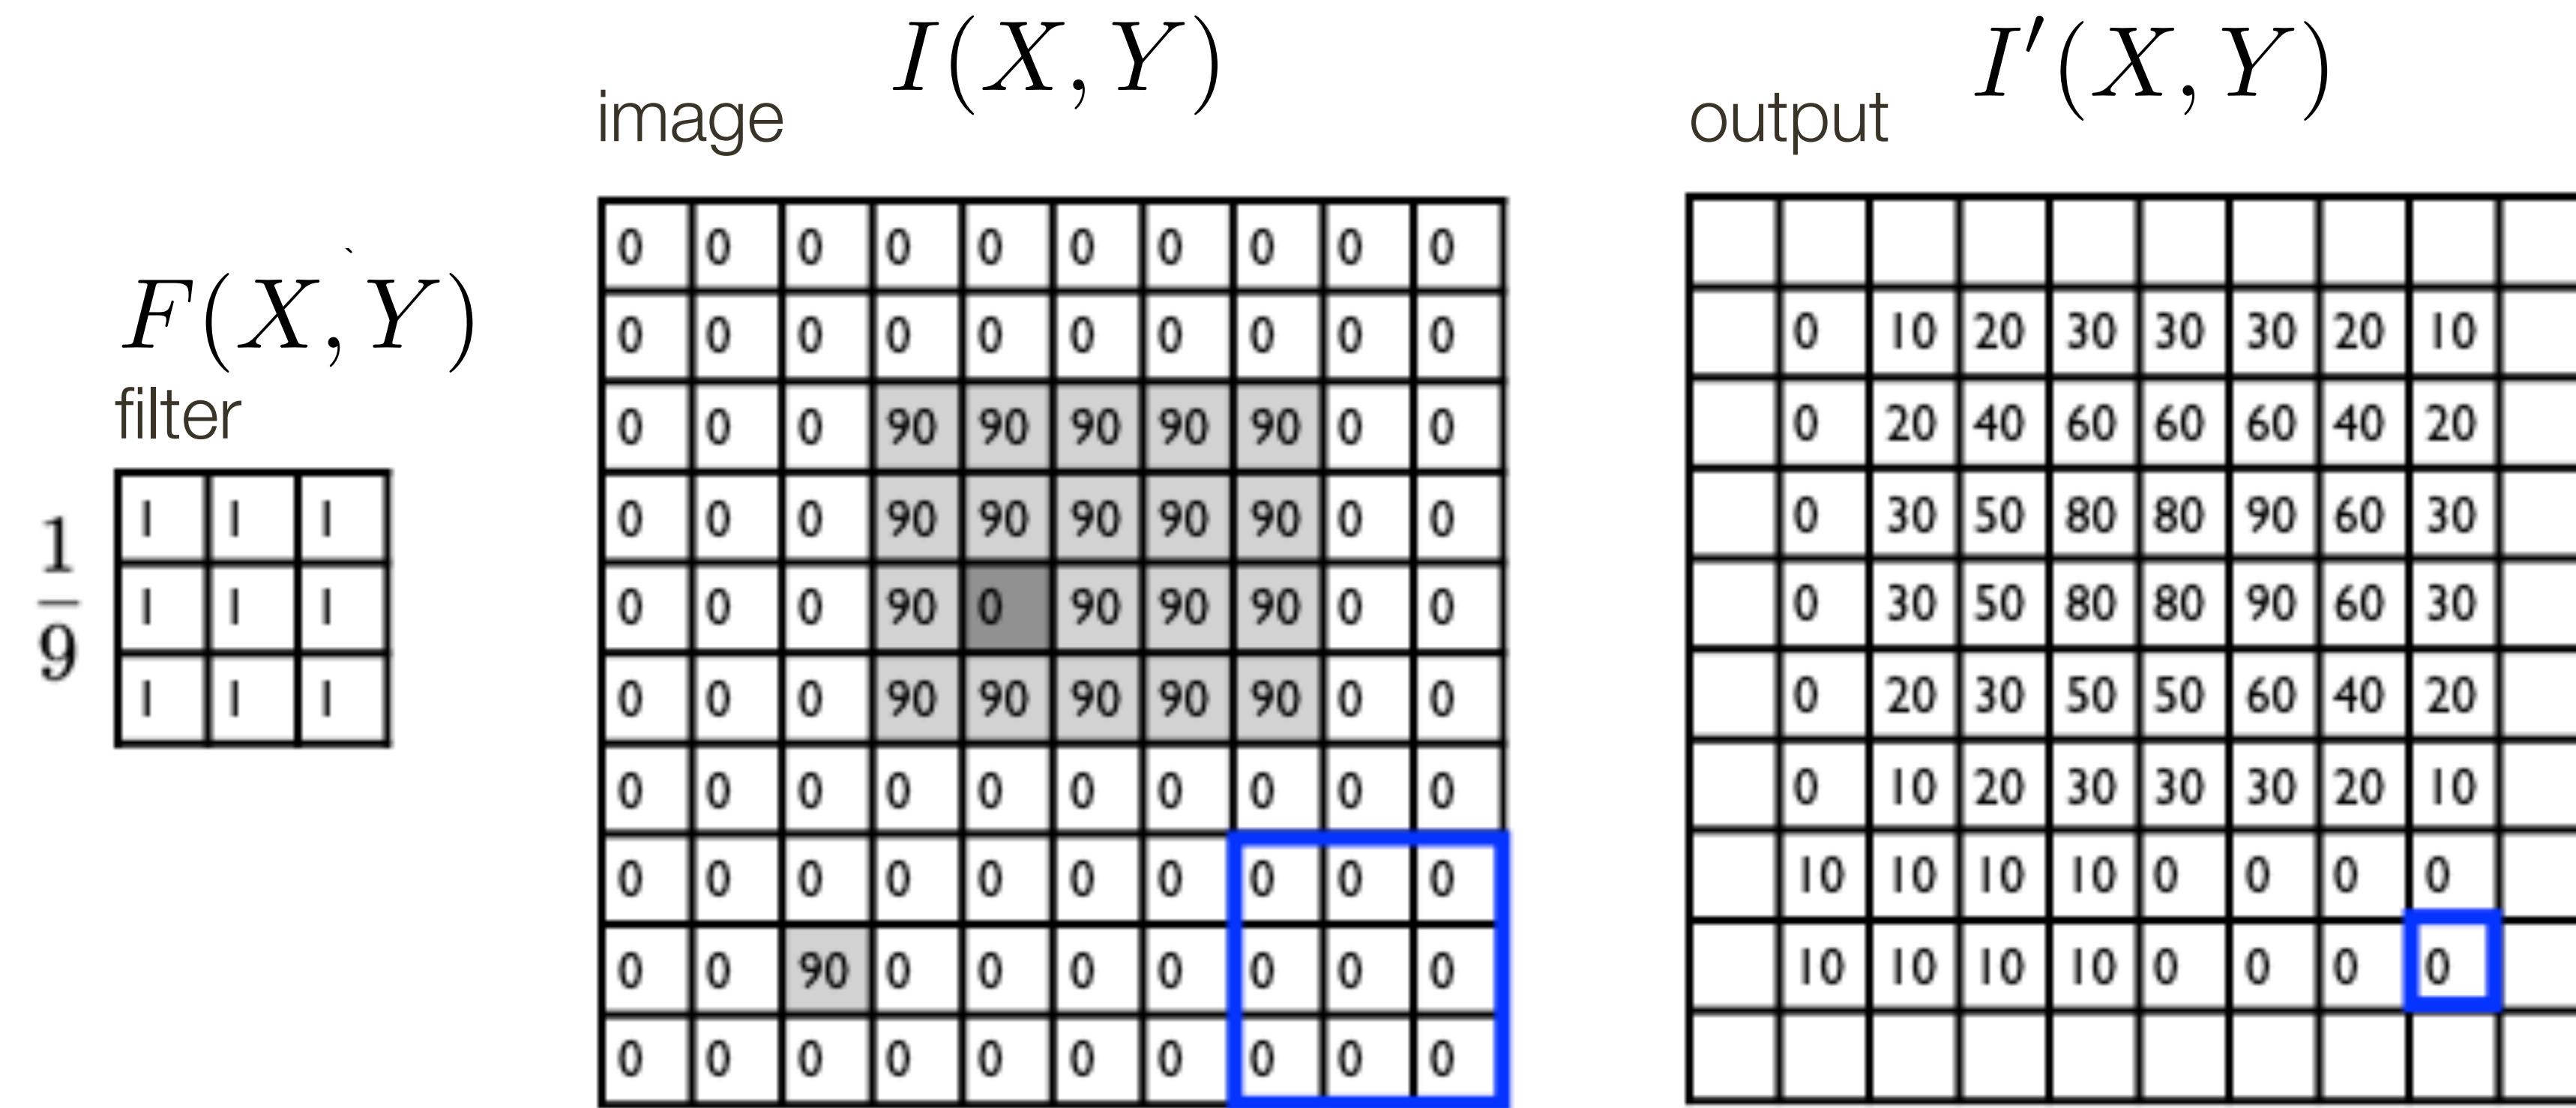
$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

output

filter

image (signal)

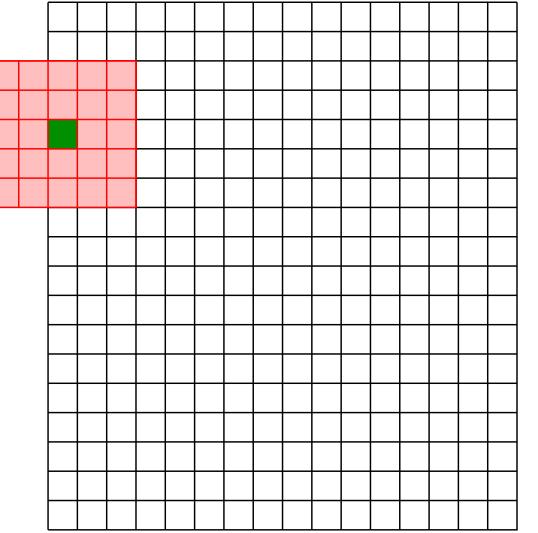
# Linear Filter Example



$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

outputfilterimage (signal)

# Linear Filters: **Boundary** Effects



Three standard ways to deal with boundaries:

1. **Ignore these locations:** Make the computation undefined for the top and bottom  $k$  rows and the leftmost and rightmost  $k$  columns
2. **Pad the image with zeros:** Return zero whenever a value of  $I$  is required at some position outside the defined limits of  $X$  and  $Y$
3. **Assume periodicity:** The top row wraps around to the bottom row; the leftmost column wraps around to the rightmost column

# Point Spread Function

$$\begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & \boxed{1} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & \boxed{1} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{matrix} * \begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{matrix} = \begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 9 & 8 & 7 & 0 & 0 & 0 & 0 & 0 \\ 0 & 6 & 5 & 4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 2 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 9 & 8 & 7 & 0 & 0 \\ 0 & 0 & 0 & 0 & 6 & 5 & 4 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 2 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{matrix}$$

# Point Spread Function

- The point spread function is the correlation kernel rotated by 180° (= the convolution kernel)

# Lecture 4: Re-cap

**Linear** filtering (one interpretation):

- new pixels are a weighted sum of original pixel values
- “filter” defines weights

**Linear** filtering (another interpretation):

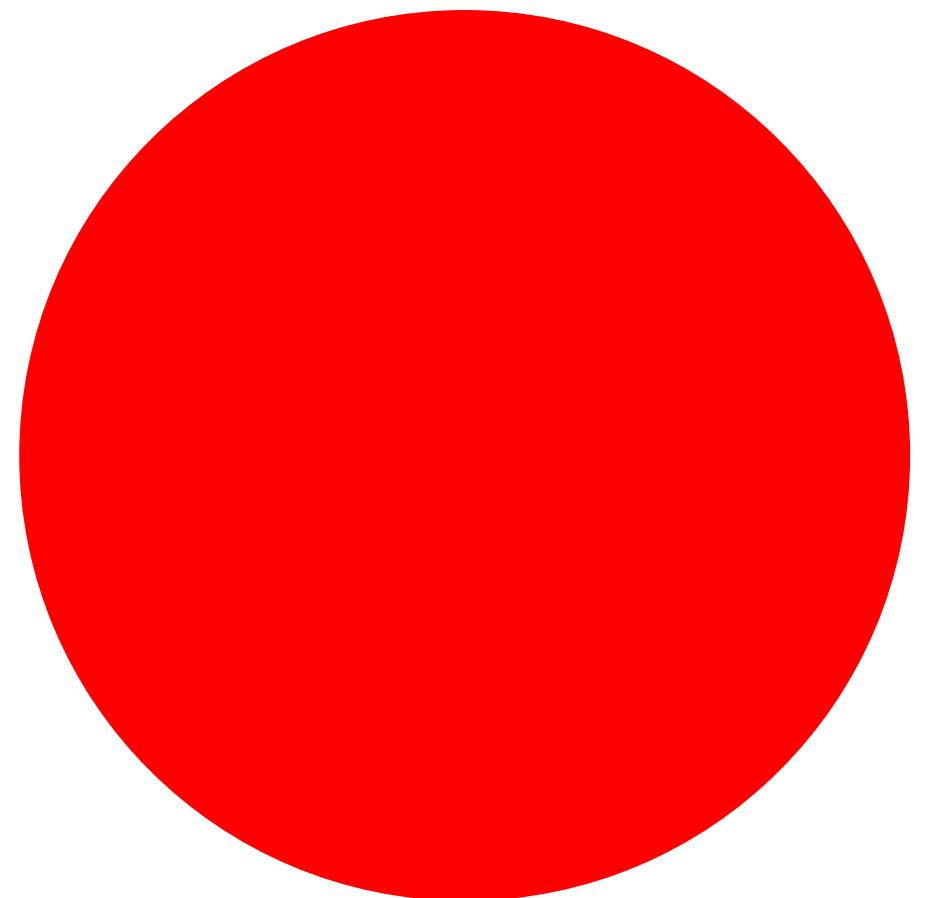
- each pixel creates a scaled copy of point spread function in its location
- “filter” specifies the point spread function

# Low-pass Filtering = “Smoothing”

**Box Filter**

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

**Pillbox Filter**



**Gaussian Filter**

$$\frac{1}{256} \begin{array}{|c|c|c|c|c|} \hline 1 & 4 & 6 & 4 & 1 \\ \hline 4 & 16 & 24 & 16 & 4 \\ \hline 6 & 24 & 36 & 24 & 6 \\ \hline 4 & 16 & 24 & 16 & 4 \\ \hline 1 & 4 & 6 & 4 & 1 \\ \hline \end{array}$$

All of these filters are **Low-pass Filters**

**Low-pass filter:** Low pass filter filters out all of the high frequency content of the image, only low frequencies remain

# Example: Separable Filter

$$\frac{1}{16} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \end{bmatrix} \otimes \frac{1}{16} \begin{bmatrix} 1 \\ 4 \\ 6 \\ 4 \\ 1 \end{bmatrix} = \frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

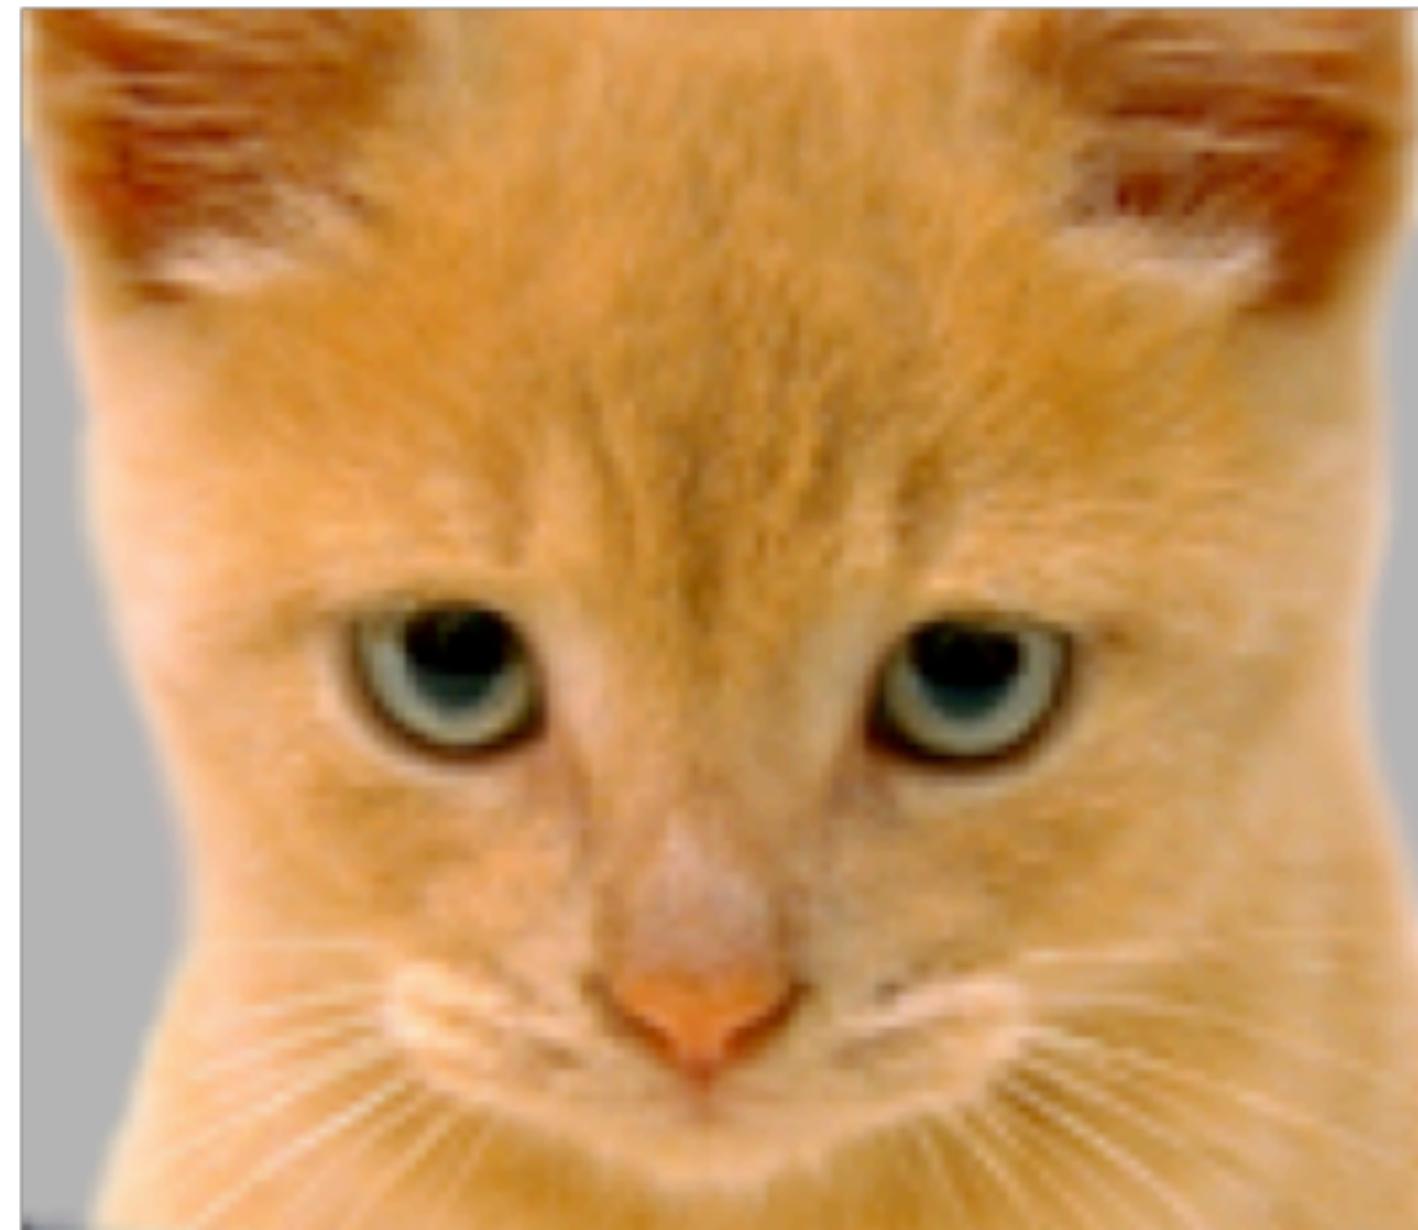
# Gaussian Blur

- 2D Gaussian filter can be thought of as an **outer product** or **convolution** of row and column filters

$$\begin{matrix} \text{Row Filter} \\ \times \\ \text{Column Filter} \end{matrix} = \text{Gaussian Filter}$$

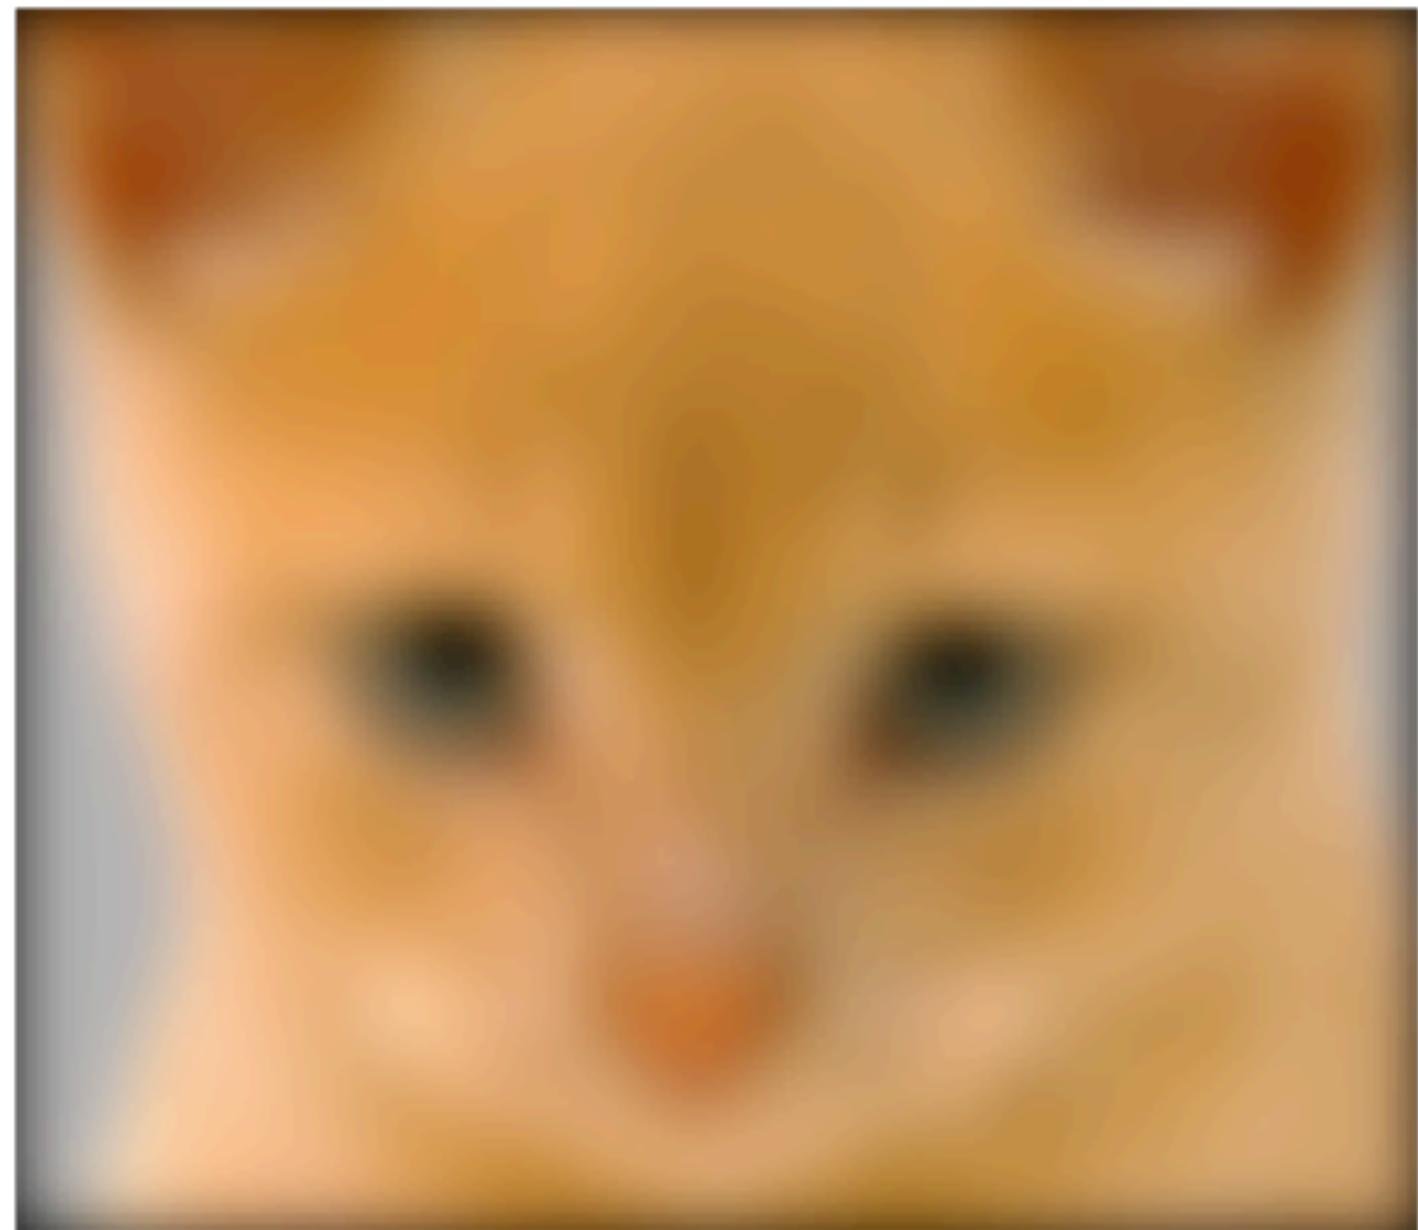
The diagram illustrates the mathematical operation of a 2D Gaussian filter. It shows a vertical column of grayscale blocks representing the Row Filter, followed by a horizontal row of grayscale blocks representing the Column Filter. A multiplication symbol (\*) is placed between the two, indicating the outer product. An equals sign (=) is placed below the row filter, followed by a large square grayscale image representing the resulting Gaussian blur filter, which has a distinct bell-shaped gradient.

# Assignment 1: Low/High Pass Filtering



Original

$$I(x, y)$$



Low-Pass Filter

$$I(x, y) * g(x, y)$$



High-Pass Filter

$$I(x, y) - I(x, y) * g(x, y)$$

**Next:** Please get your **iClickers** –  
**Quiz 1:** 6 questions

# Advanced Convolution Topics

- Multiple filters
- Fourier transforms

# Multiple Filters: Two Box Filters

```
filter = boxfilter(3)
```

```
signal.correlate2d(filter, filter, 'full')
```

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \otimes \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \frac{1}{81} \begin{array}{|c|c|c|c|c|} \hline 1 & 2 & 3 & 2 & 1 \\ \hline 2 & 4 & 6 & 4 & 2 \\ \hline 3 & 6 & 9 & 6 & 3 \\ \hline 2 & 4 & 6 & 4 & 2 \\ \hline 1 & 2 & 3 & 2 & 1 \\ \hline \end{array}$$

3x3 Box                    3x3 Box

# Multiple Filters: Two Box Filters

```
filter = boxfilter(3)
```

```
temp = signal.correlate2d(filter, filter, 'full')
```

```
signal.correlate2d(filter, temp, 'full')
```

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \otimes \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \otimes \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \frac{1}{729}$$

**3x3 Box      3x3 Box      3x3 Box**

1	3	6	7	6	3	1
3	9	18	21	18	9	3
6	18	36	42	36	18	6
7	21	42	49	42	21	7
6	18	36	42	36	18	6
3	9	18	21	18	9	3
1	3	6	7	6	3	1



5.1

# Gaussian: An Additional Property

Let  $\otimes$  denote convolution. Let  $G_{\sigma_1}(x)$  and  $G_{\sigma_2}(x)$  be two 1D Gaussians

$$G_{\sigma_1}(x) \otimes G_{\sigma_2}(x) = G_{\sqrt{\sigma_1^2 + \sigma_2^2}}(x)$$

Convolution of two Gaussians is another Gaussian

**Special case:** Convolving with  $G_\sigma(x)$  twice is equivalent to  $G_{\sqrt{2}\sigma}(x)$

What follows is for fun  
(you will **NOT** be tested on this)

# Convolution using Fourier Transforms

[ Szeliski 3.4 ]

Convolution **Theorem**:

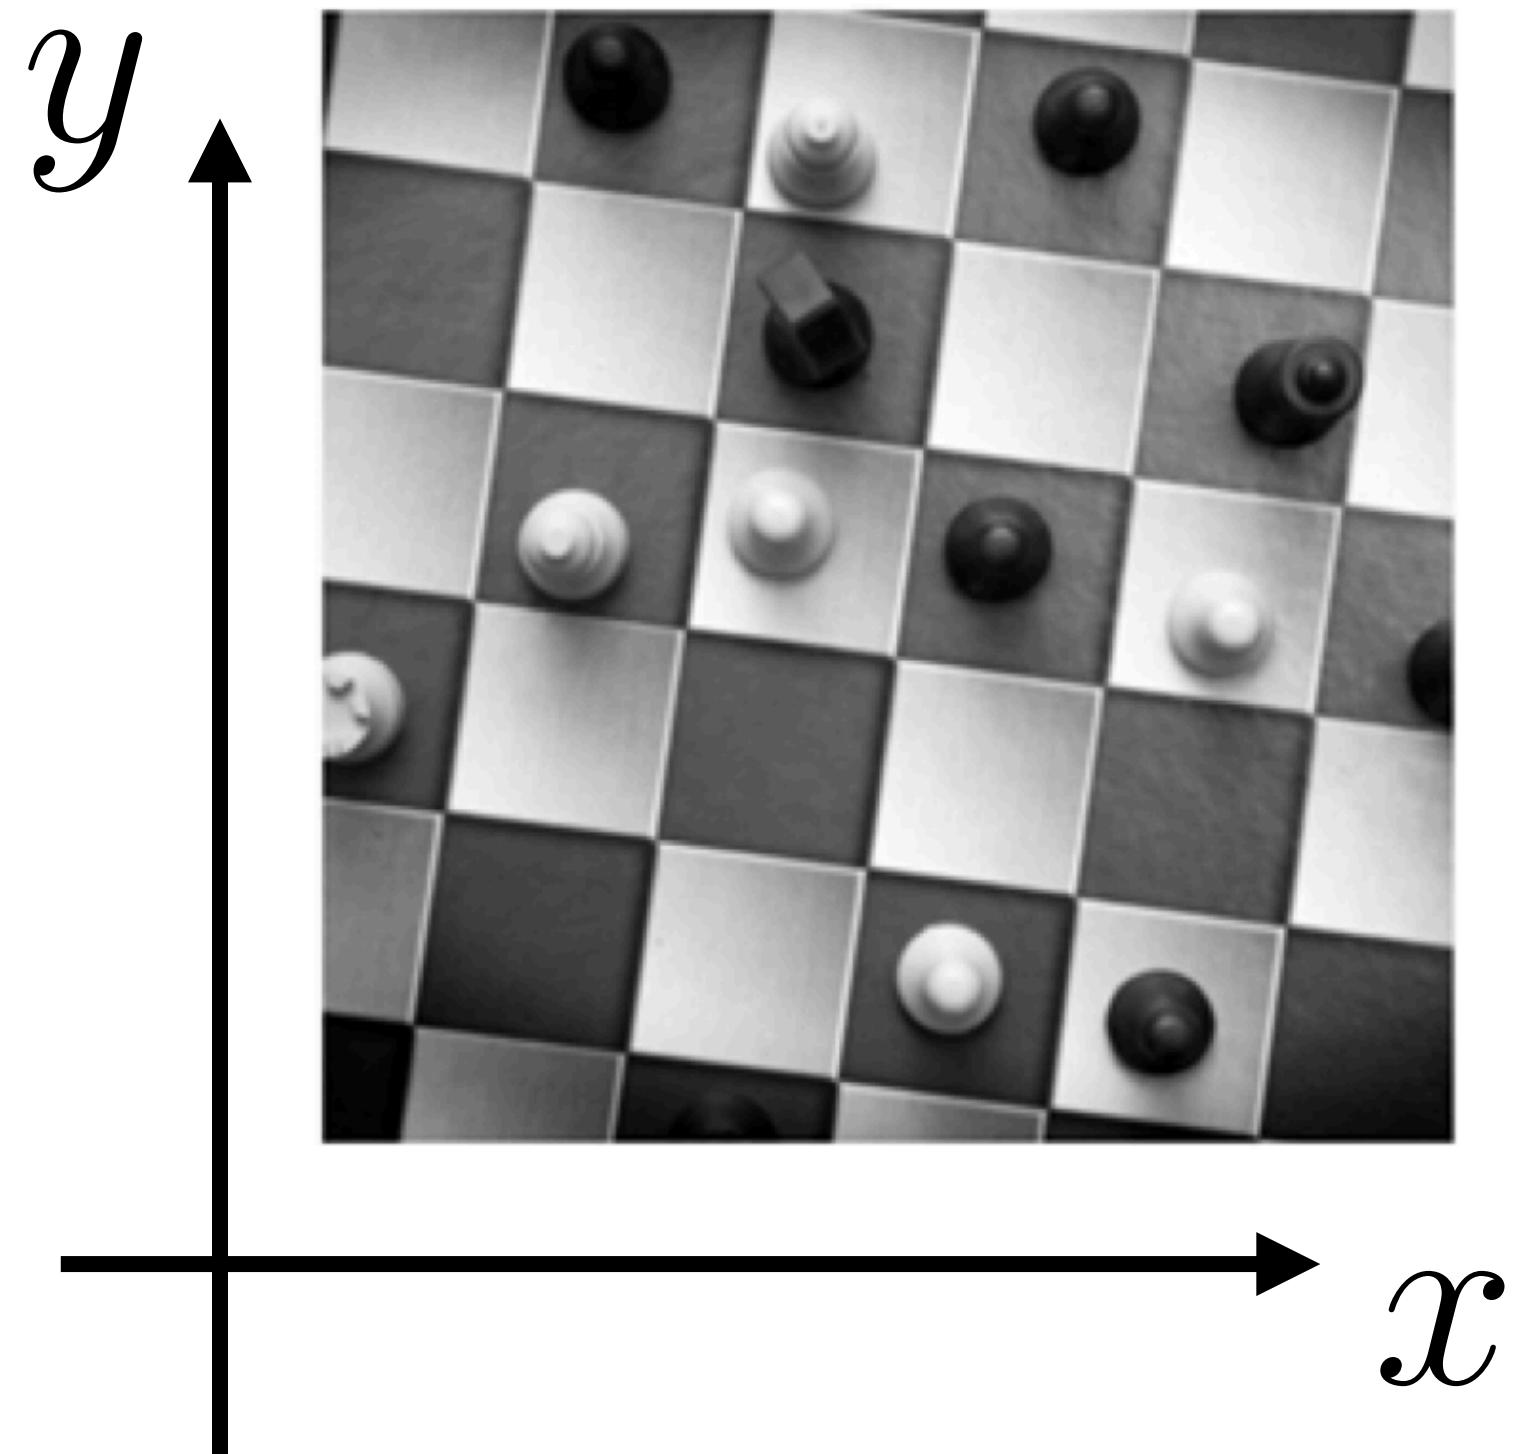
Let  $i'(x, y) = f(x, y) \otimes i(x, y)$

then  $\mathcal{I}'(w_x, w_y) = \mathcal{F}(w_x, w_y) \mathcal{I}(w_x, w_y)$

where  $\mathcal{I}'(w_x, w_y)$ ,  $\mathcal{F}(w_x, w_y)$ , and  $\mathcal{I}(w_x, w_y)$  are Fourier transforms of  $i'(x, y)$ ,  $f(x, y)$  and  $i(x, y)$

At the expense of two **Fourier** transforms and one inverse Fourier transform, convolution can be reduced to (complex) multiplication

# 2D Fourier Transforms: Images

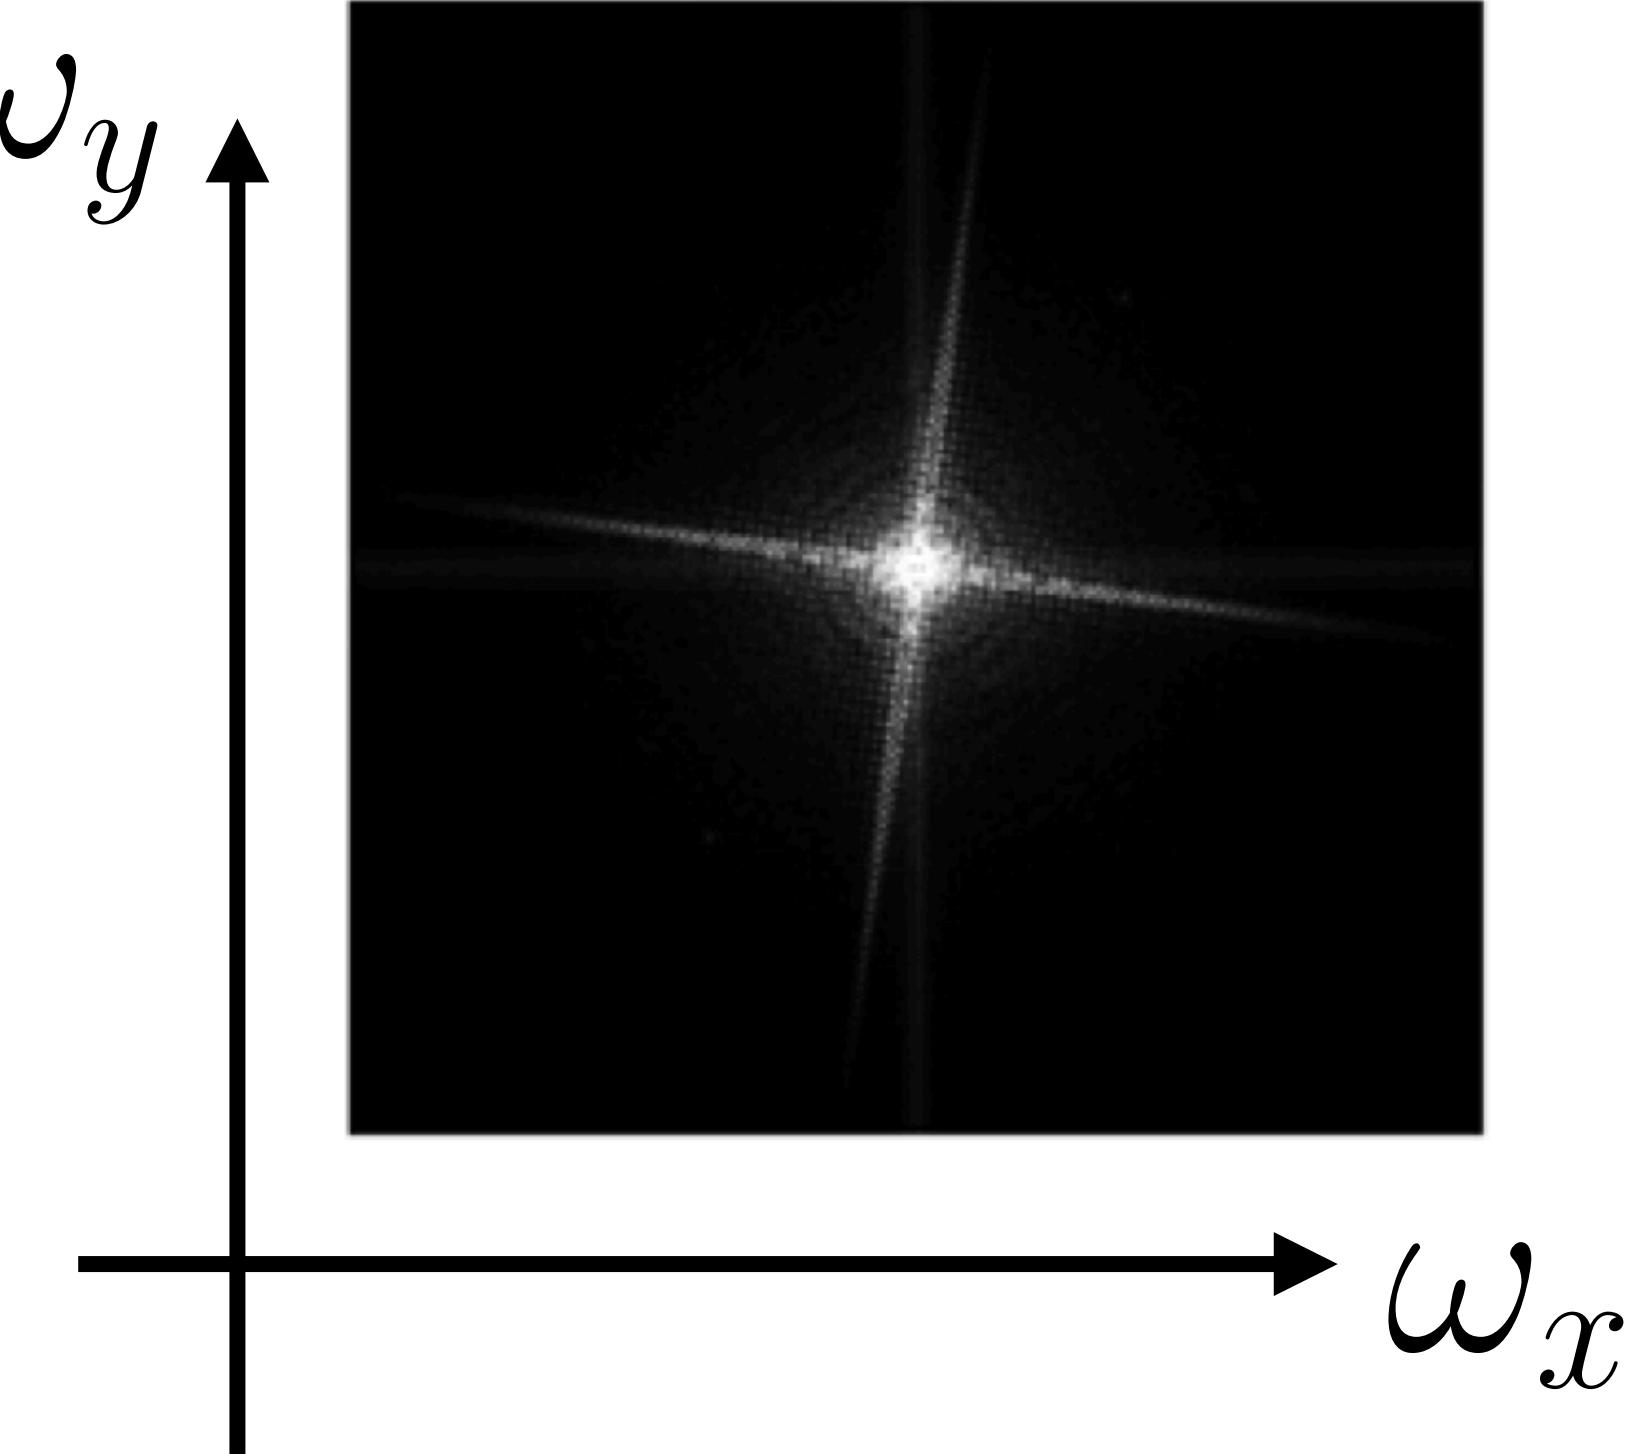


$$f(x, y)$$



5.2

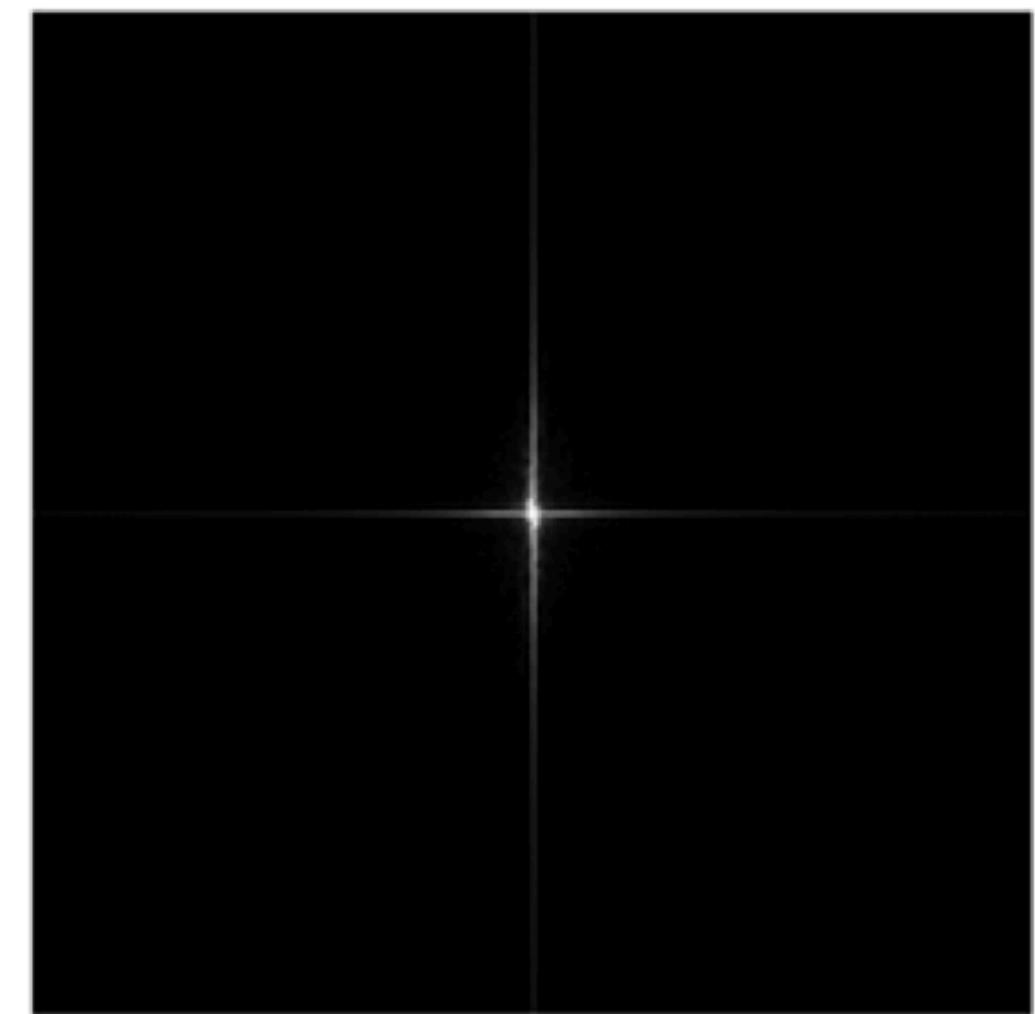
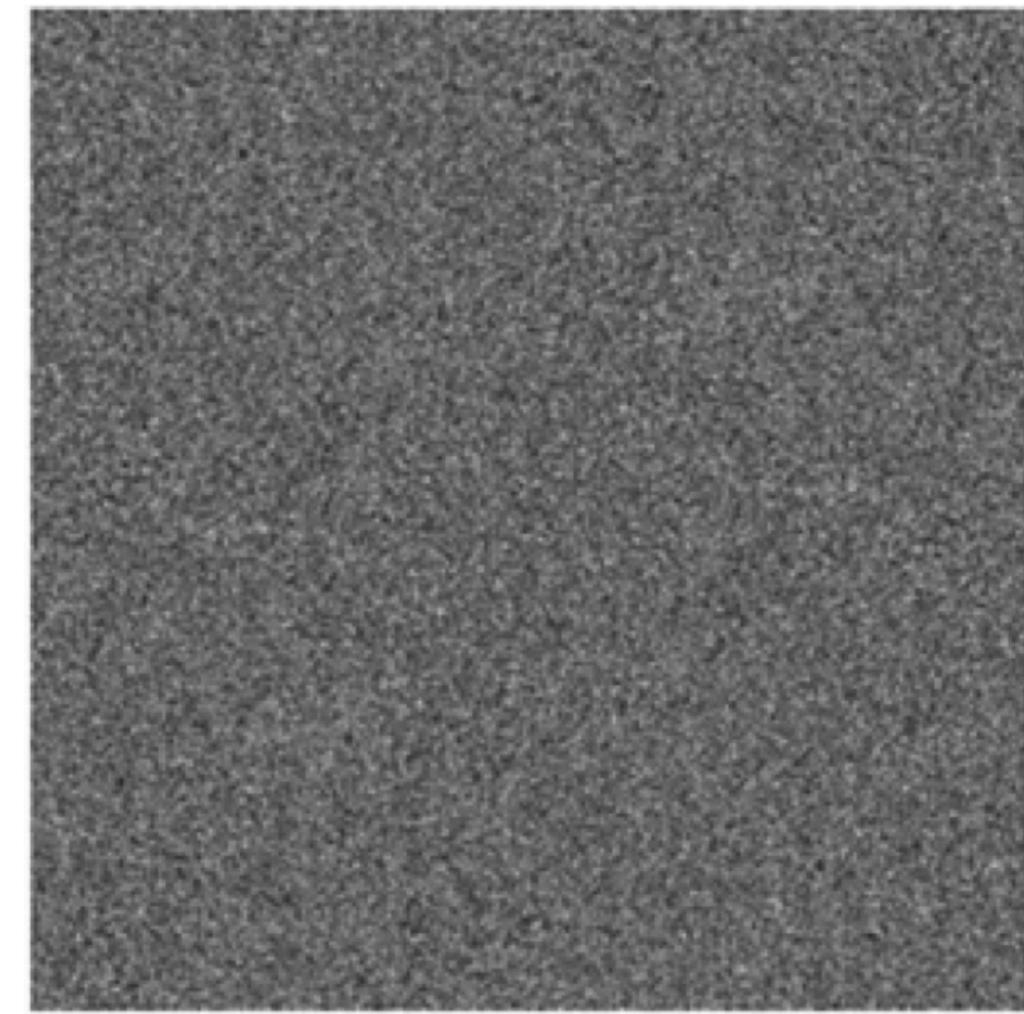
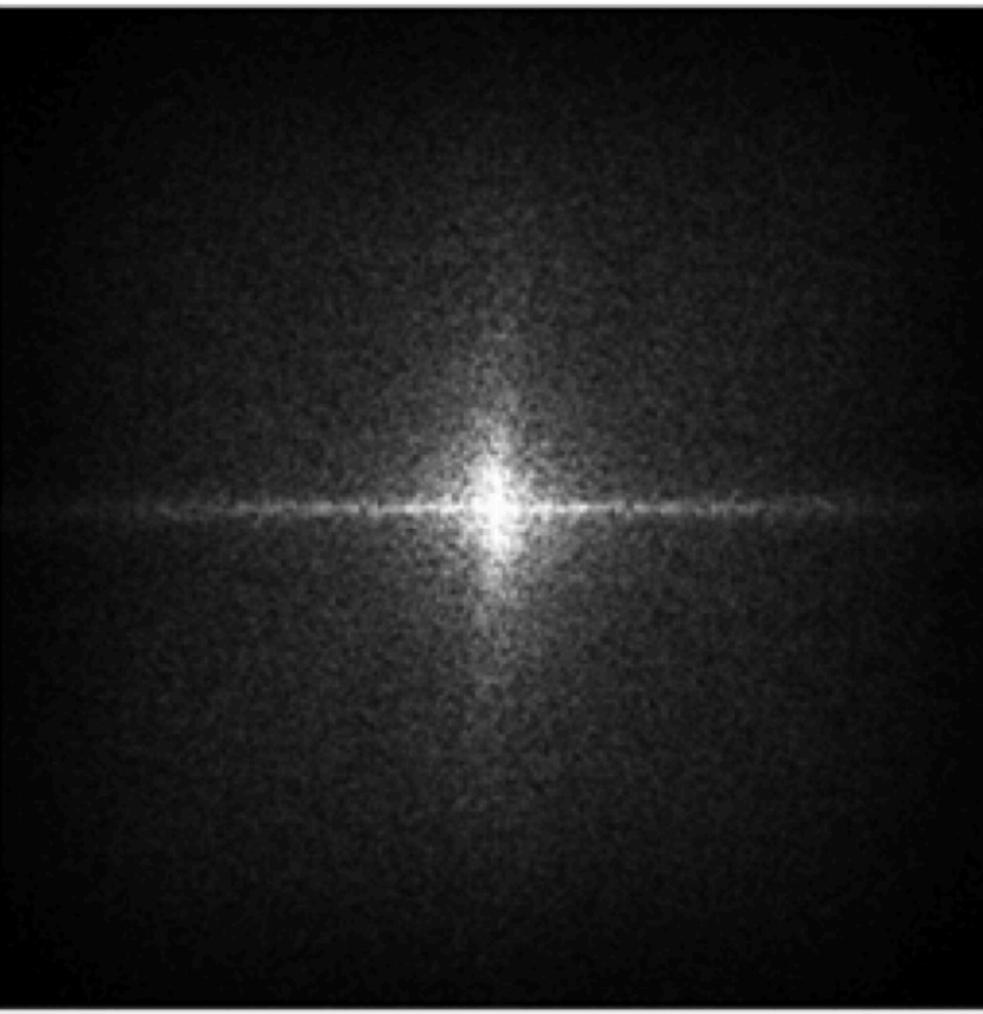
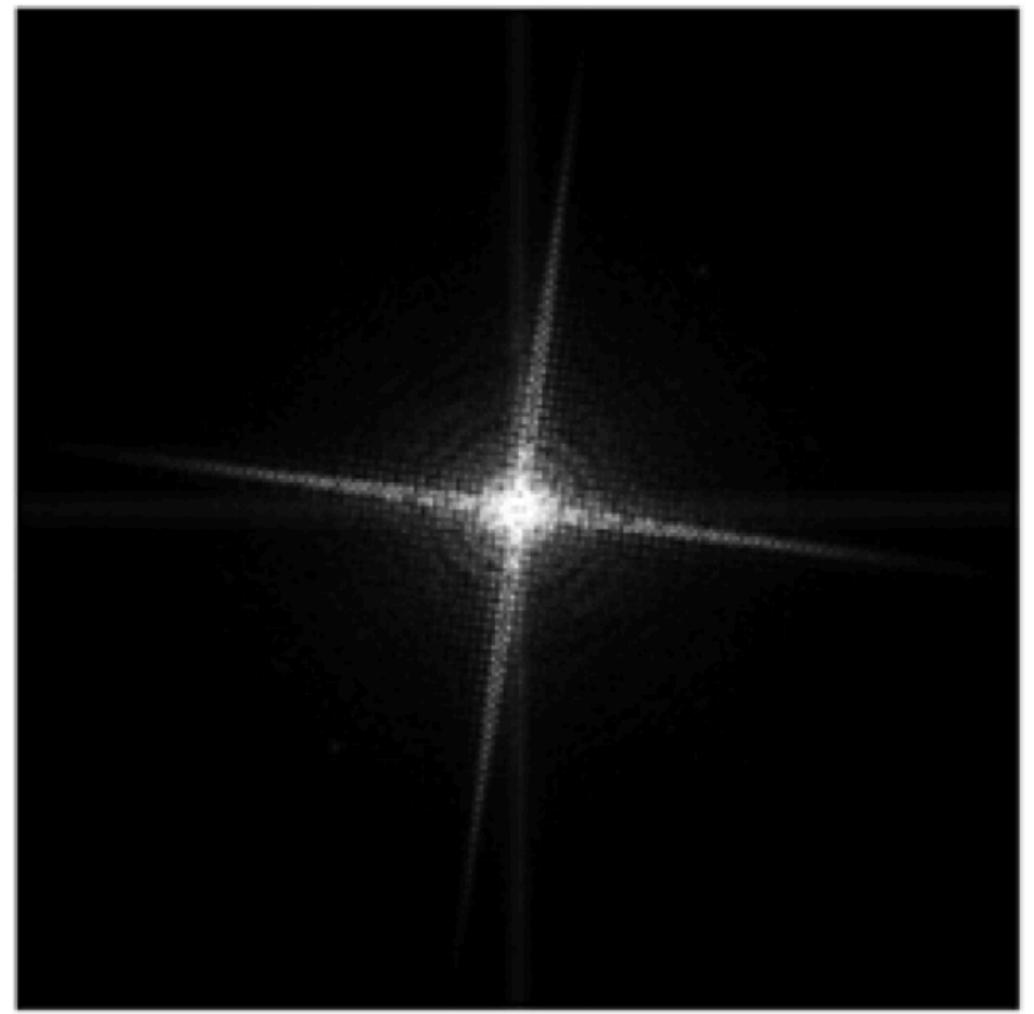
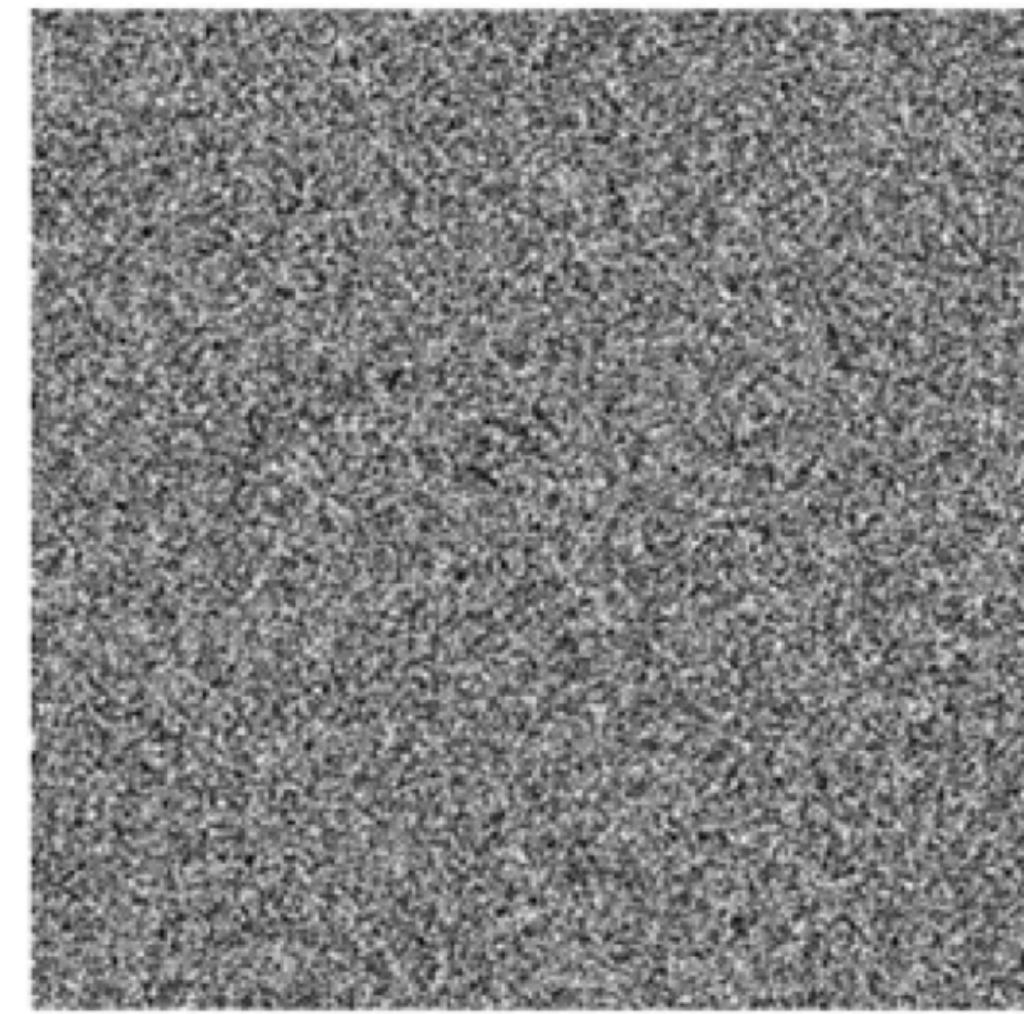
Image



$$F(\omega_x, \omega_y)$$

Fourier Transform

# 2D Fourier Transforms: Images



# Aside: You will not be tested on this ...



**Image**

<https://photo.stackexchange.com/questions/40401/what-does-frequency-mean-in-an-image/40410#40410>

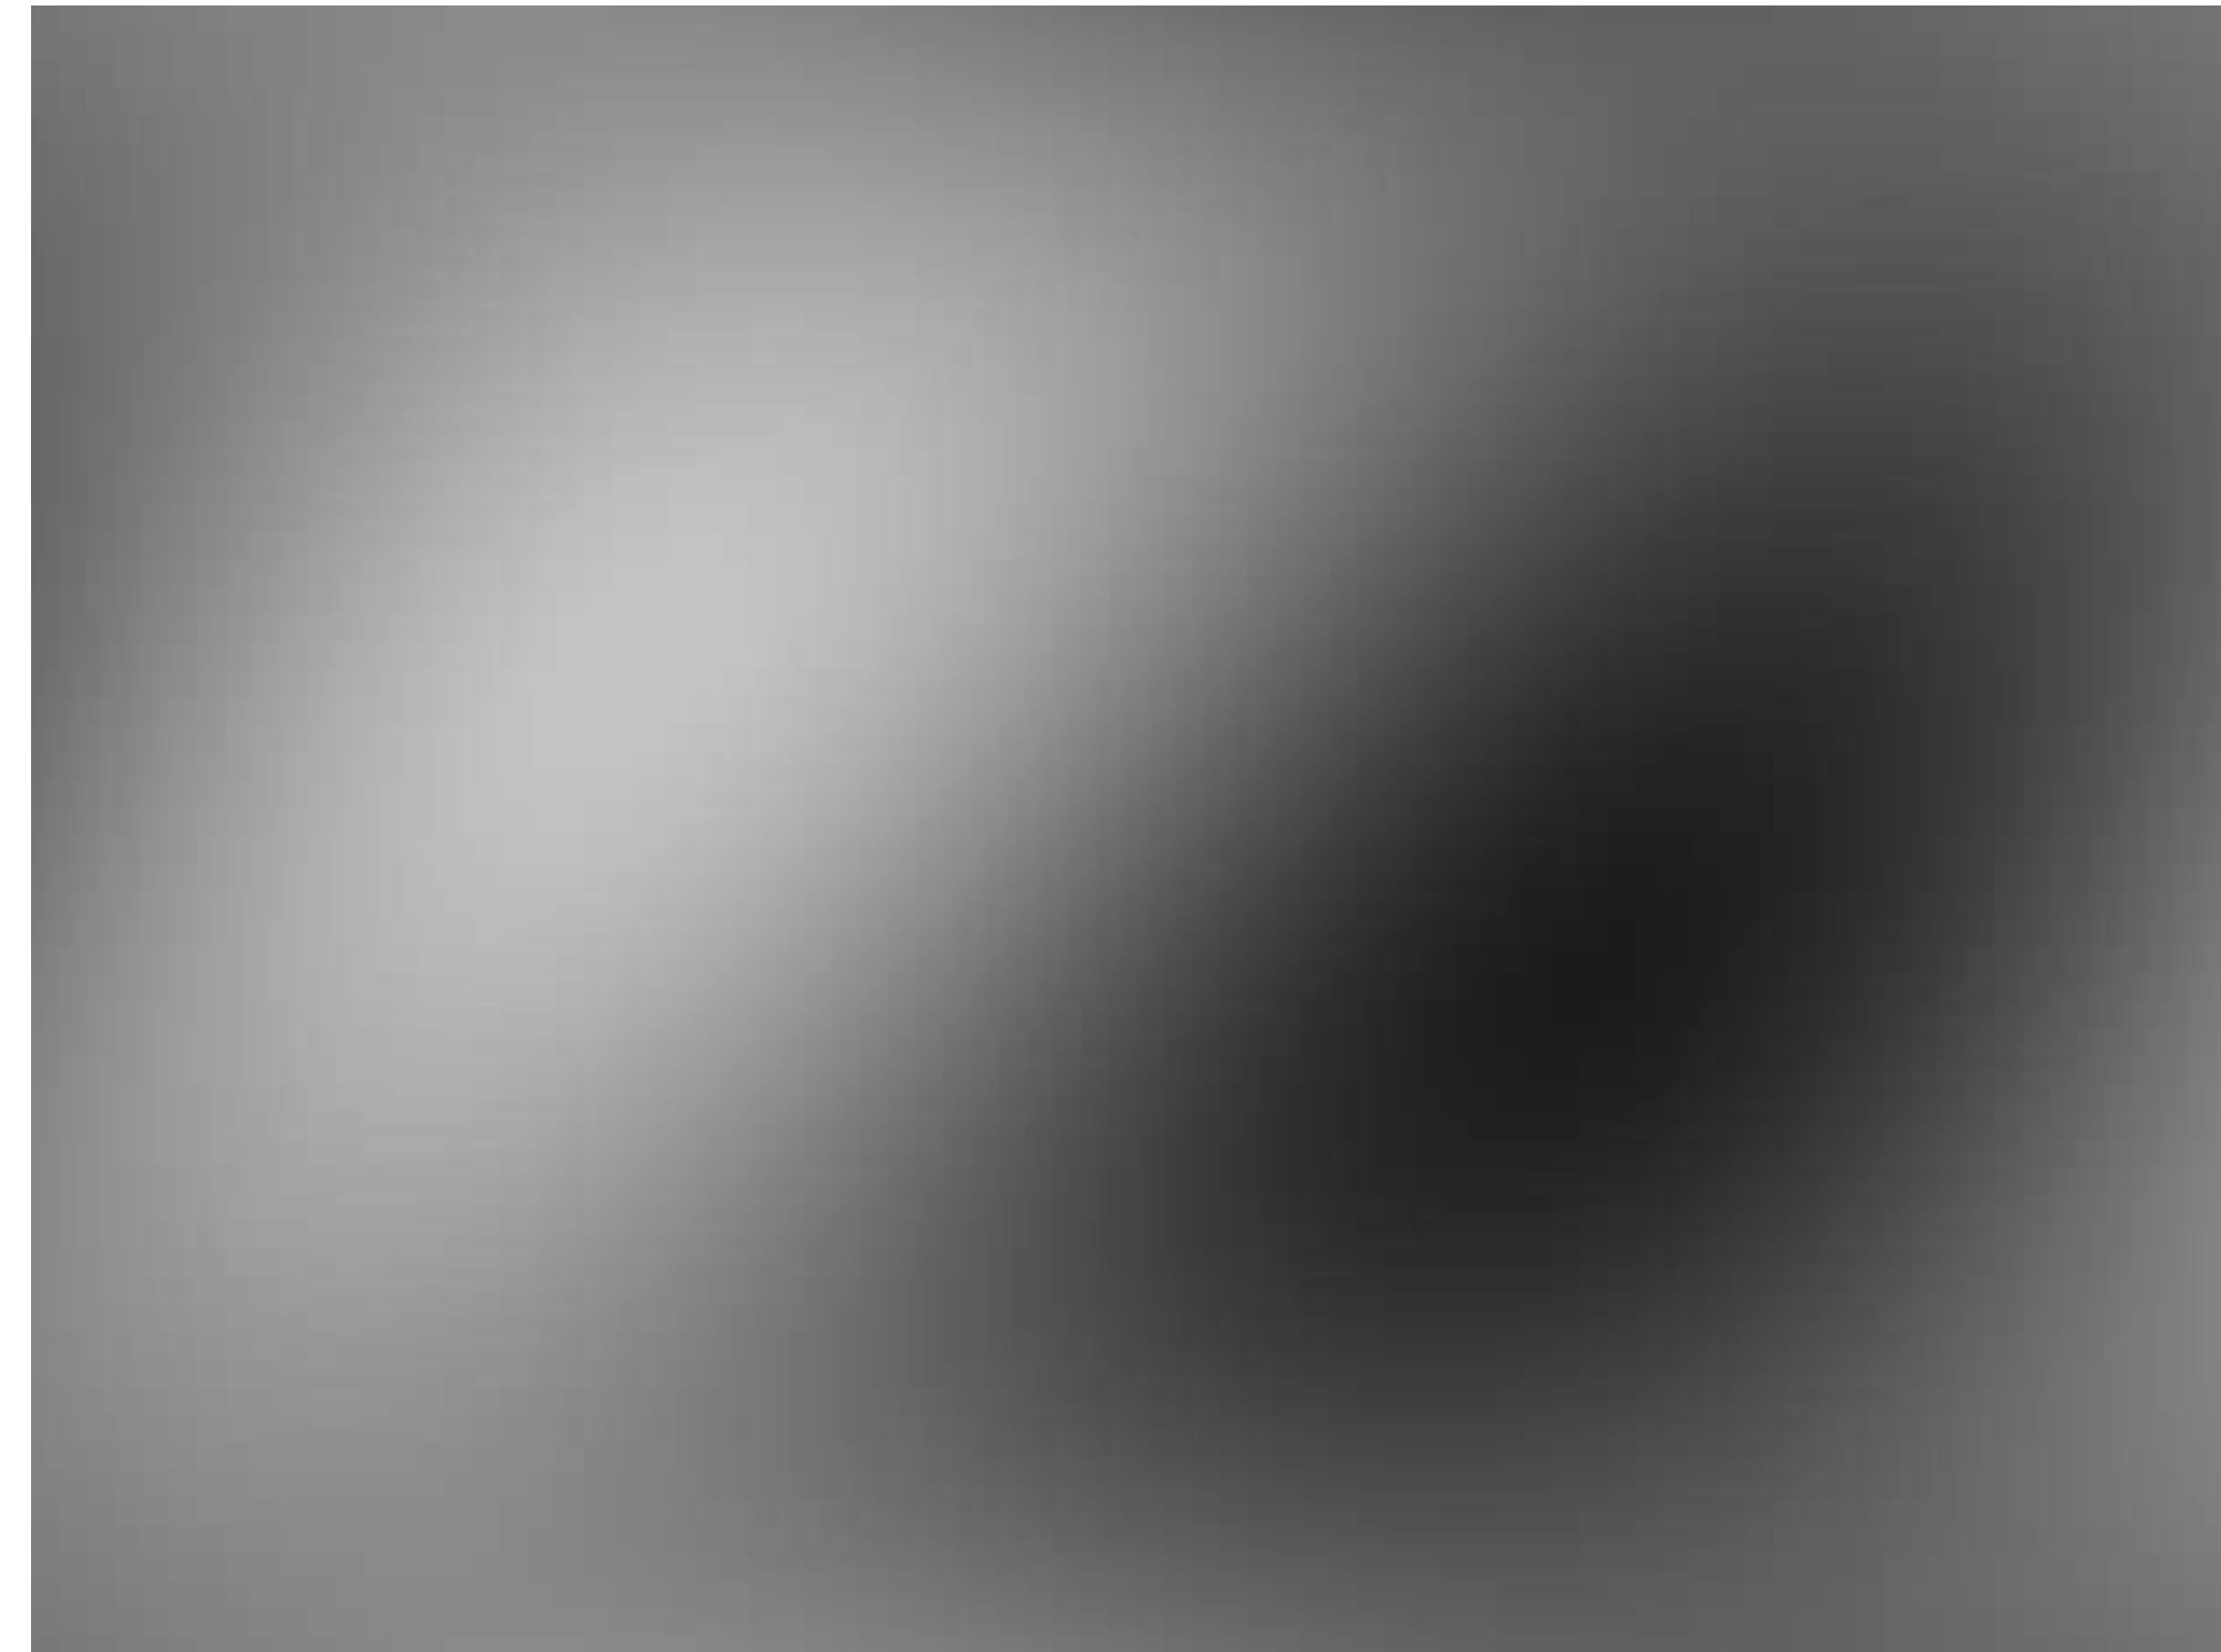
# Aside: You will not be tested on this ...



**First** (lowest) frequency, a.k.a. average

<https://photo.stackexchange.com/questions/40401/what-does-frequency-mean-in-an-image/40410#40410>

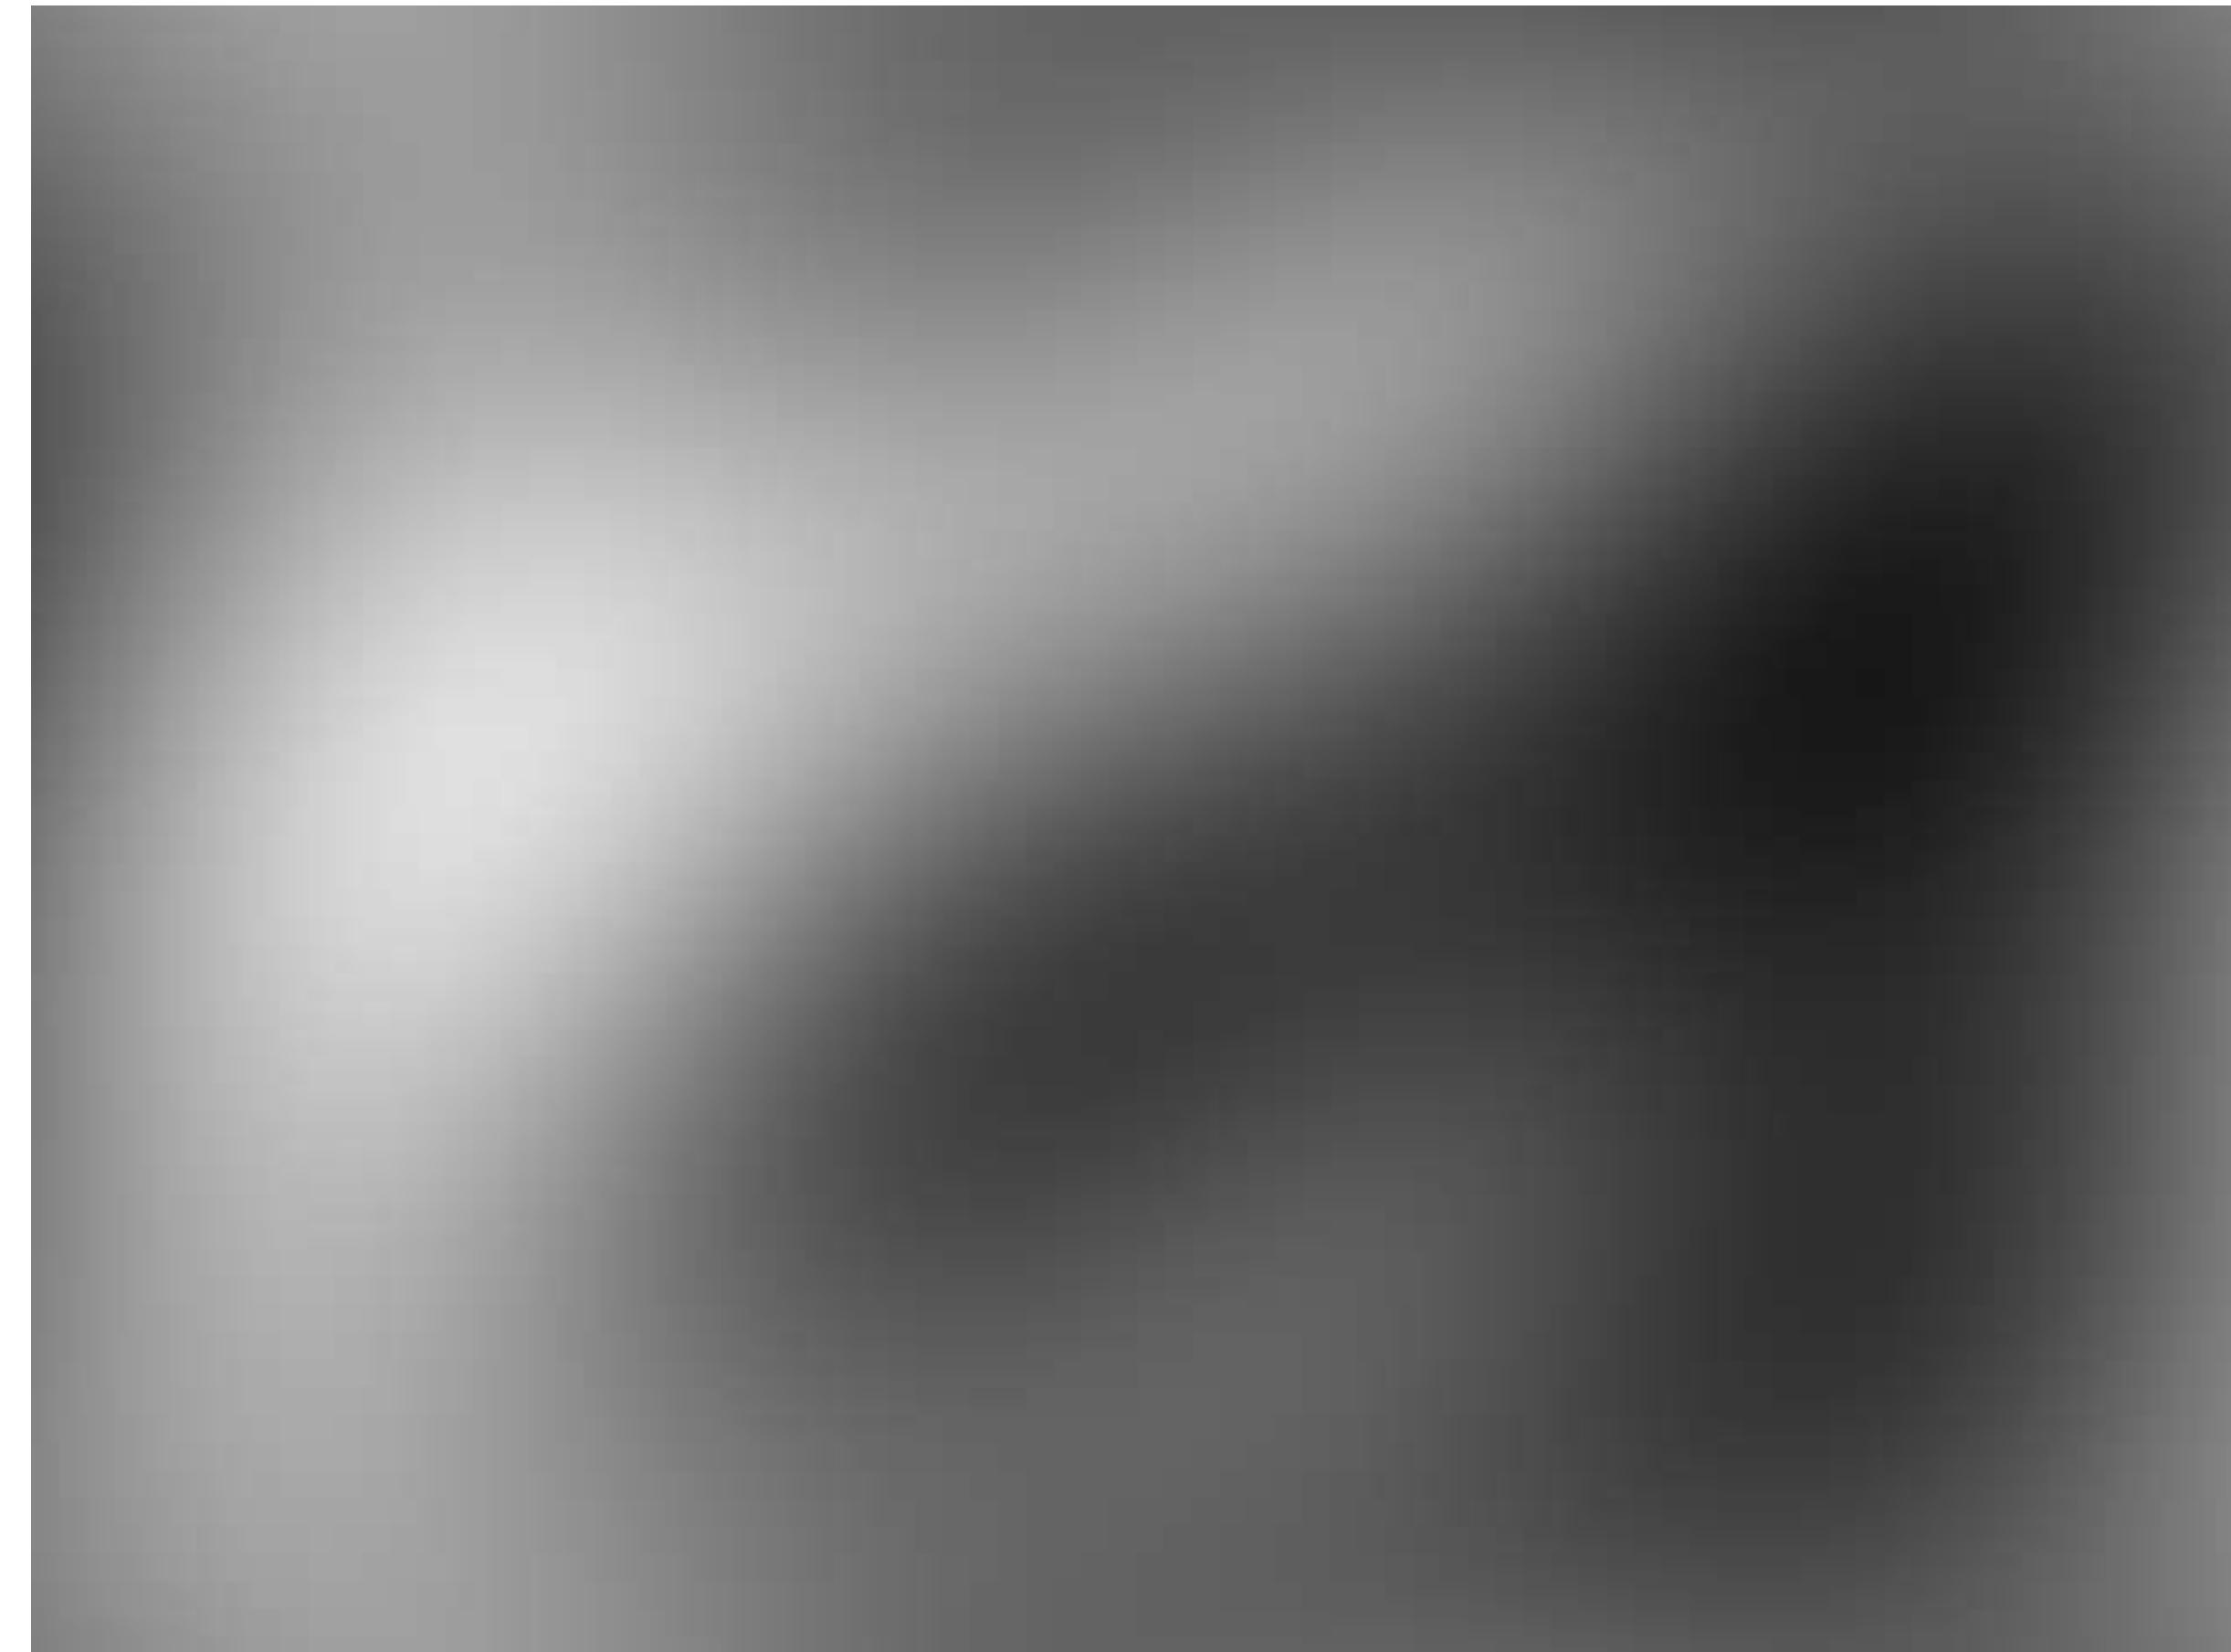
# Aside: You will not be tested on this ...



+ **Second** frequency

<https://photo.stackexchange.com/questions/40401/what-does-frequency-mean-in-an-image/40410#40410>

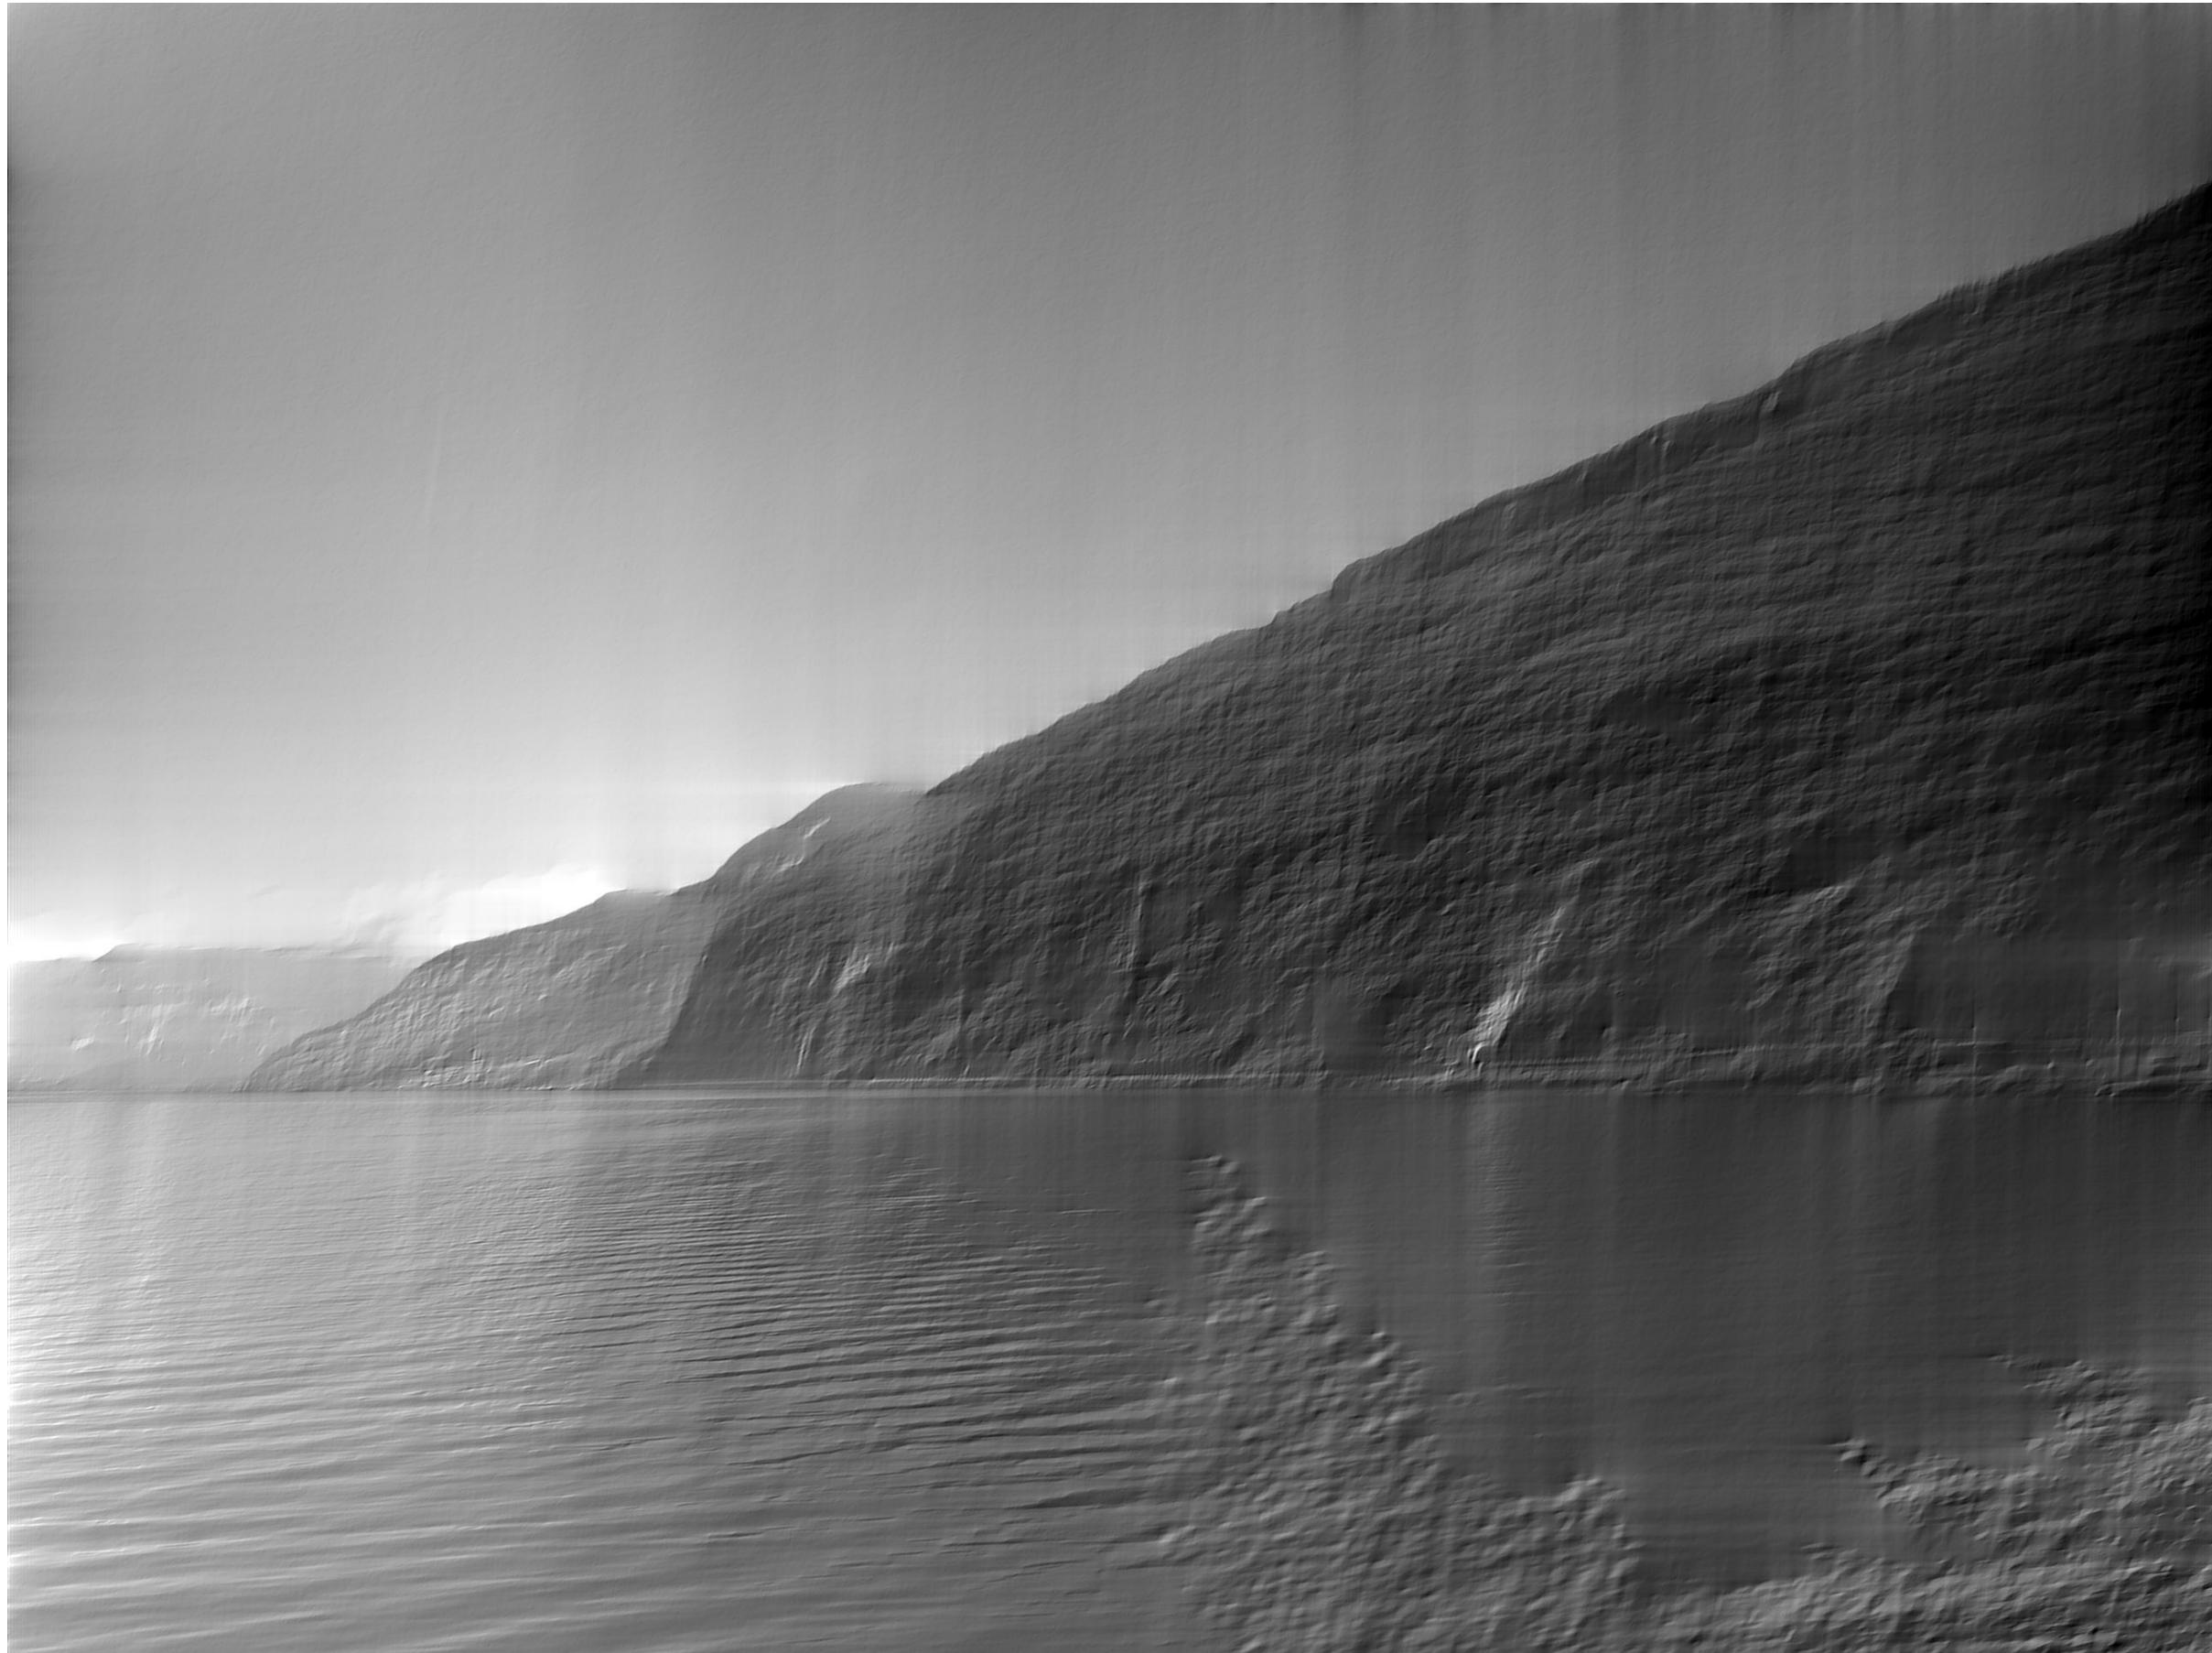
# Aside: You will not be tested on this ...



+ **Third** frequency

<https://photo.stackexchange.com/questions/40401/what-does-frequency-mean-in-an-image/40410#40410>

# Aside: You will not be tested on this ...



+ **50%** of frequencies

<https://photo.stackexchange.com/questions/40401/what-does-frequency-mean-in-an-image/40410#40410>

# Aside: You will not be tested on this ...

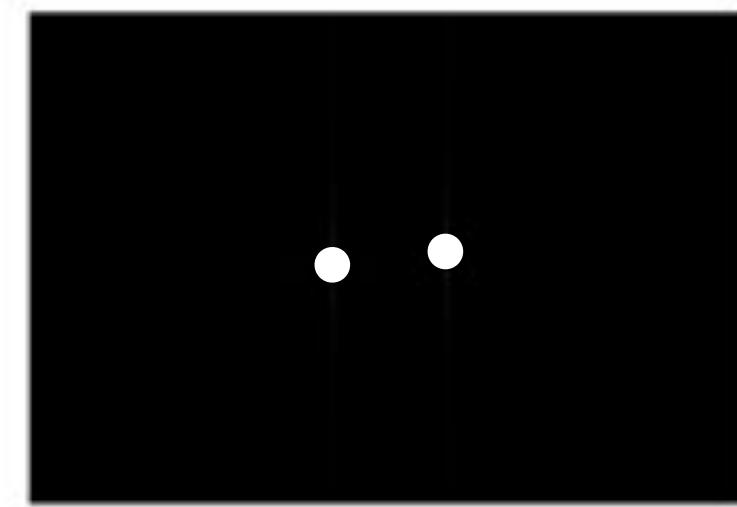
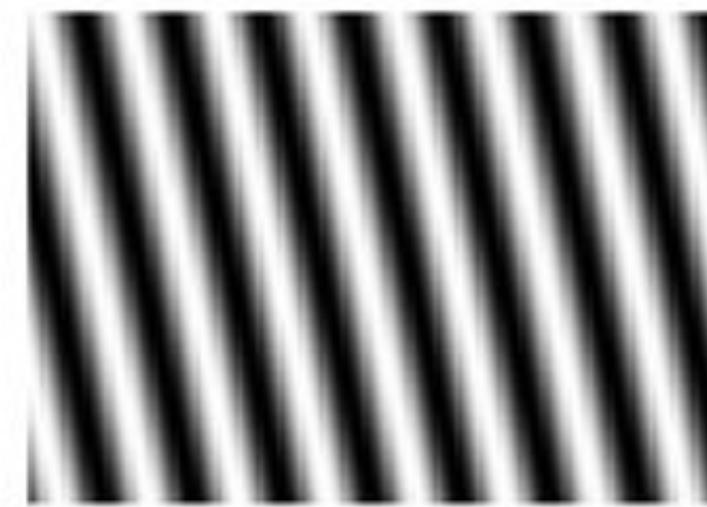


<https://photo.stackexchange.com/questions/40401/what-does-frequency-mean-in-an-image/40410#40410>

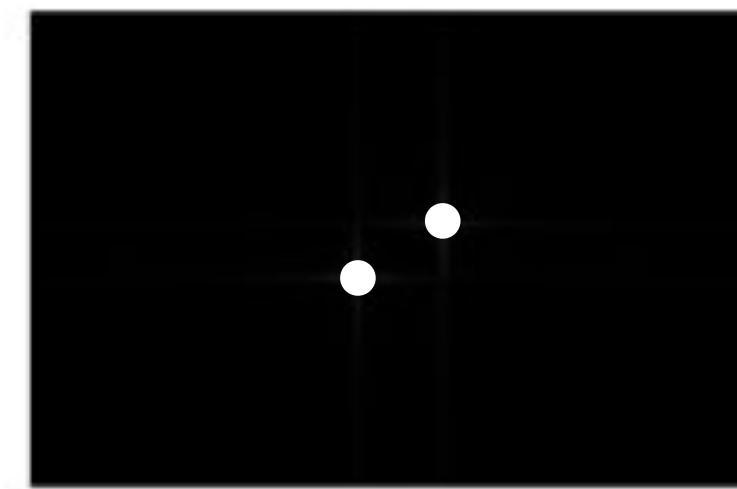
# Aside: You will not be tested on this ...

What are “frequencies” in an image?

**Spatial** frequency



$\Theta=30^\circ$



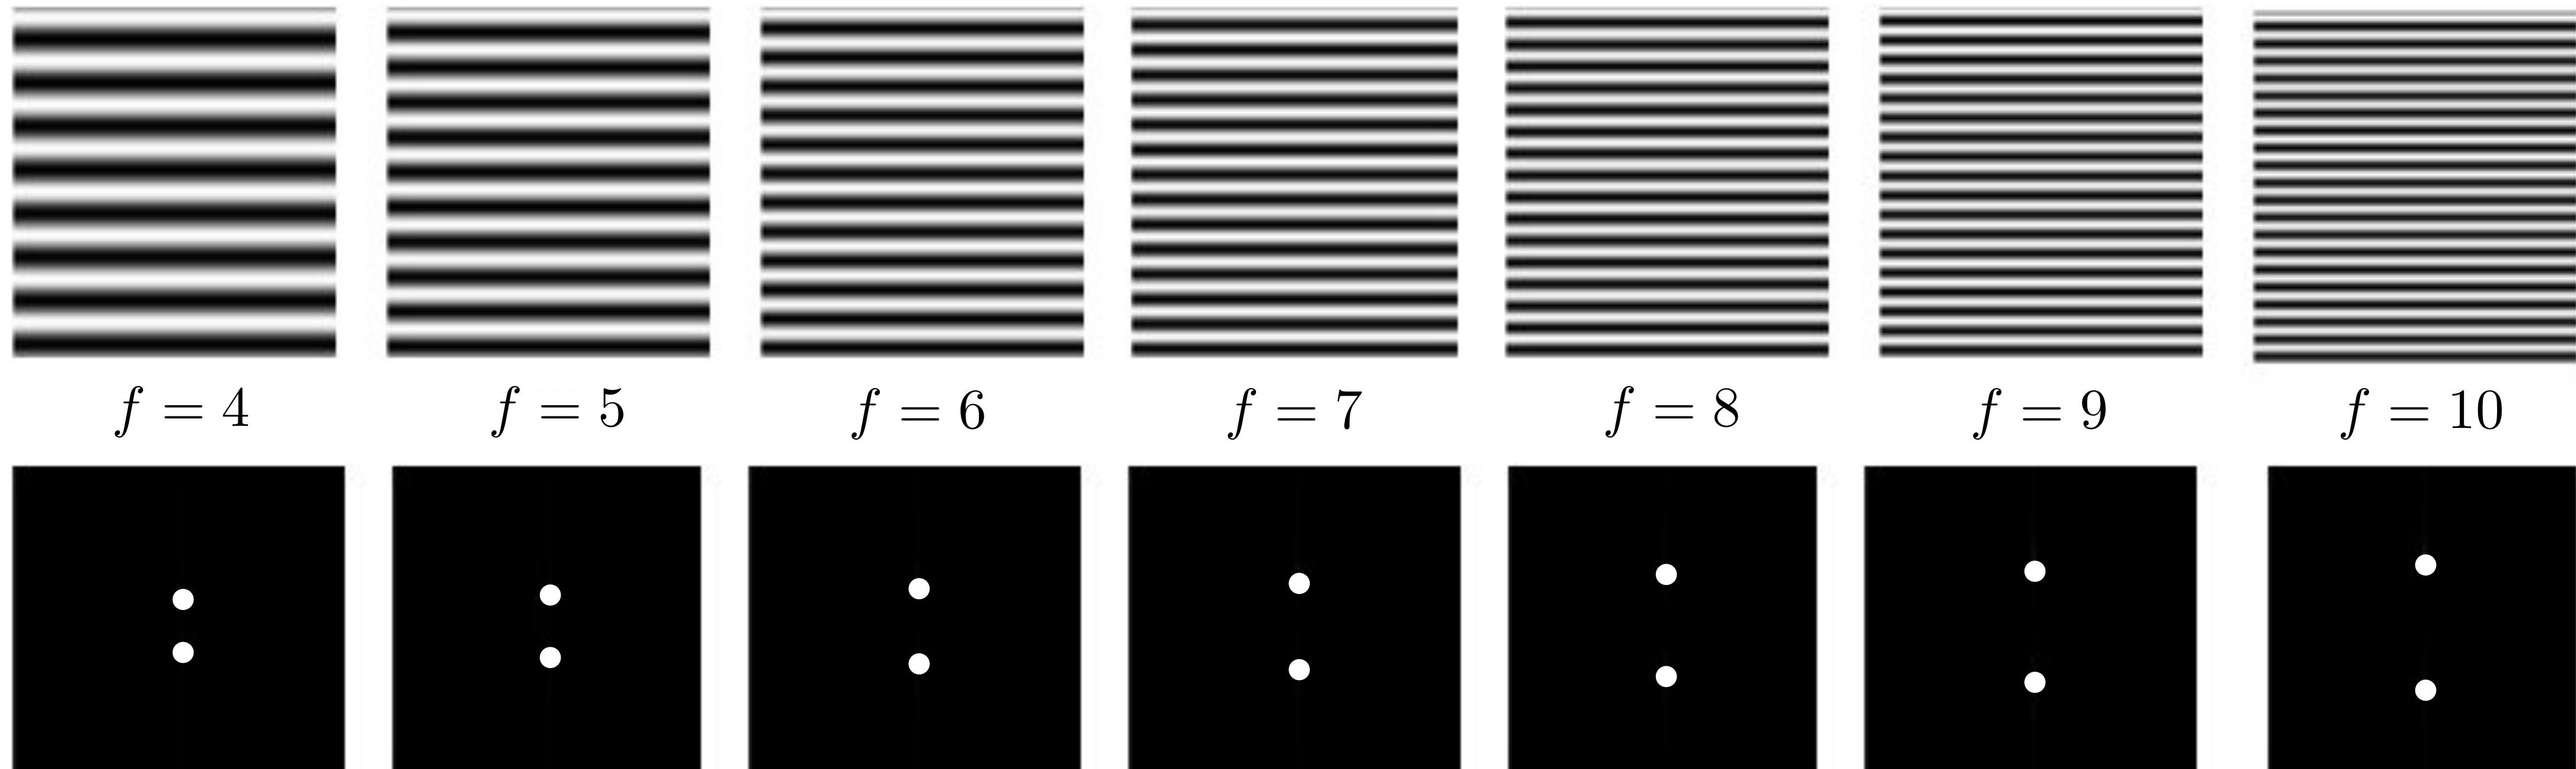
$\Theta=150^\circ$

Points (actually point pairs) in Fourier Space correspond to sinusoids with a given frequency and direction

# Aside: You will not be tested on this ...

What are “frequencies” in an image?

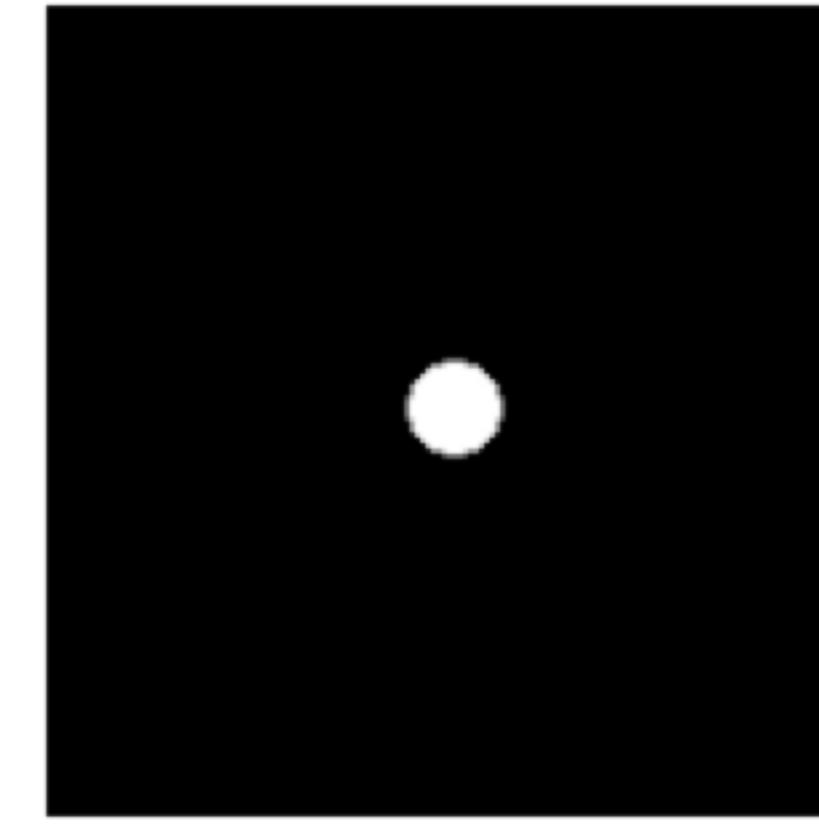
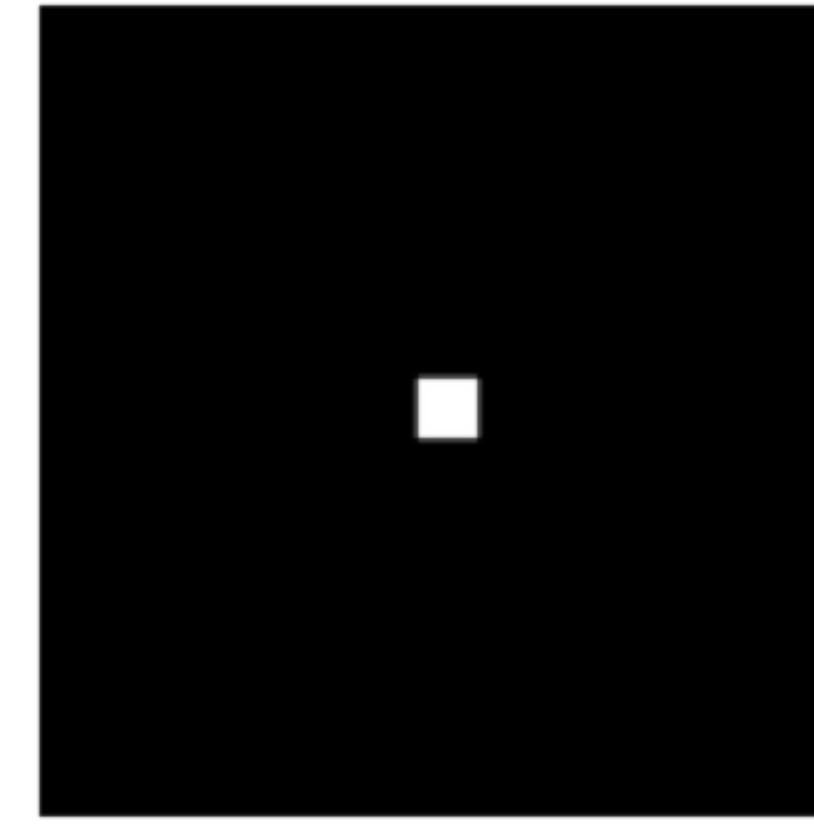
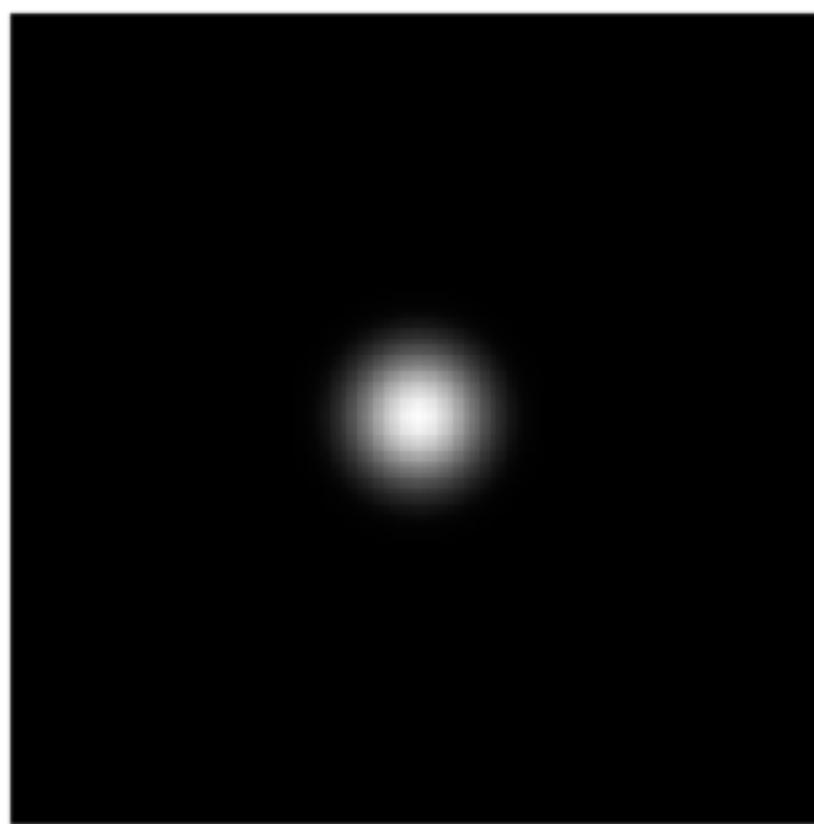
**Spatial** frequency



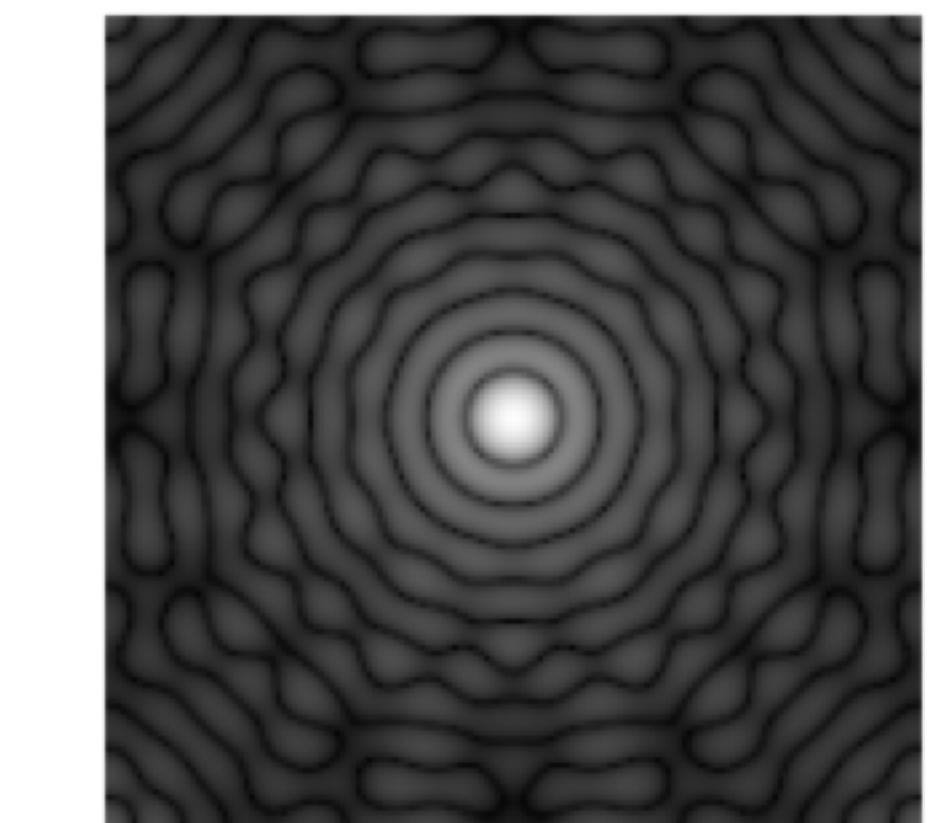
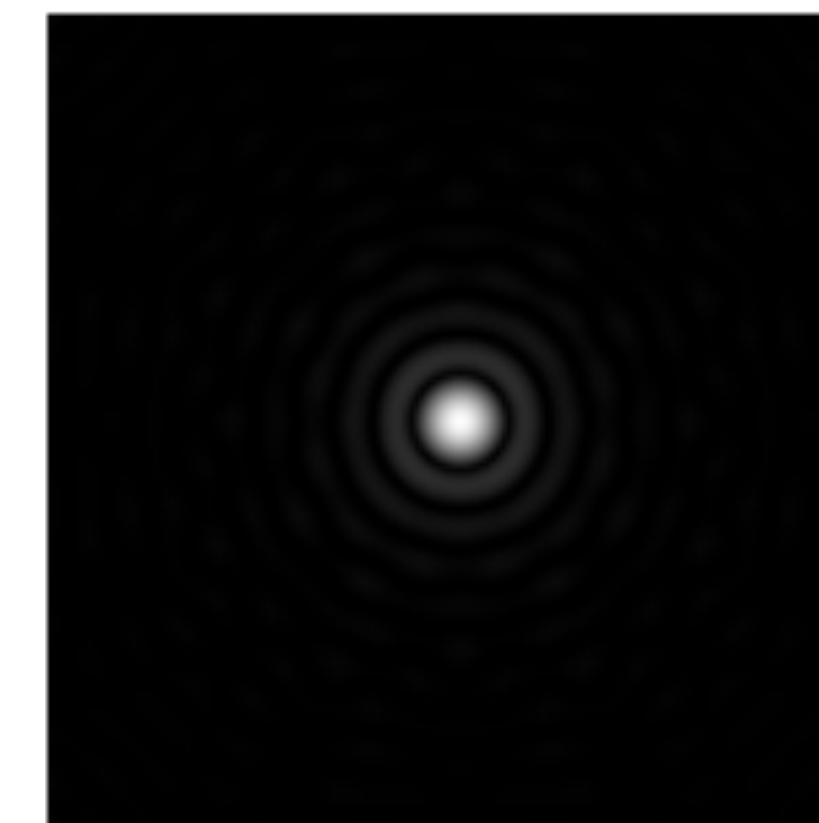
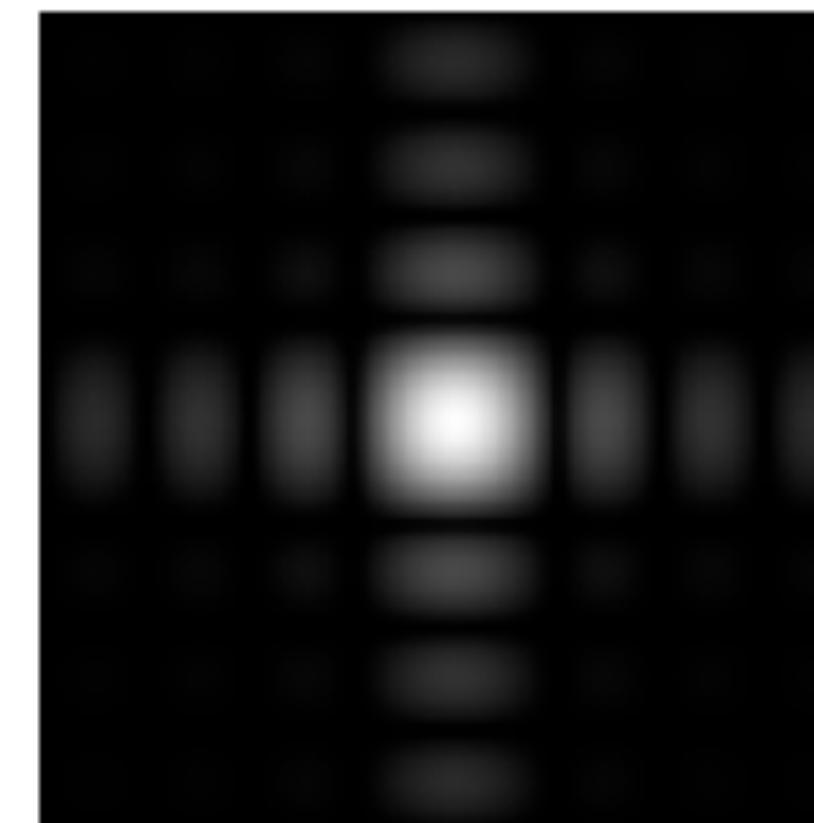
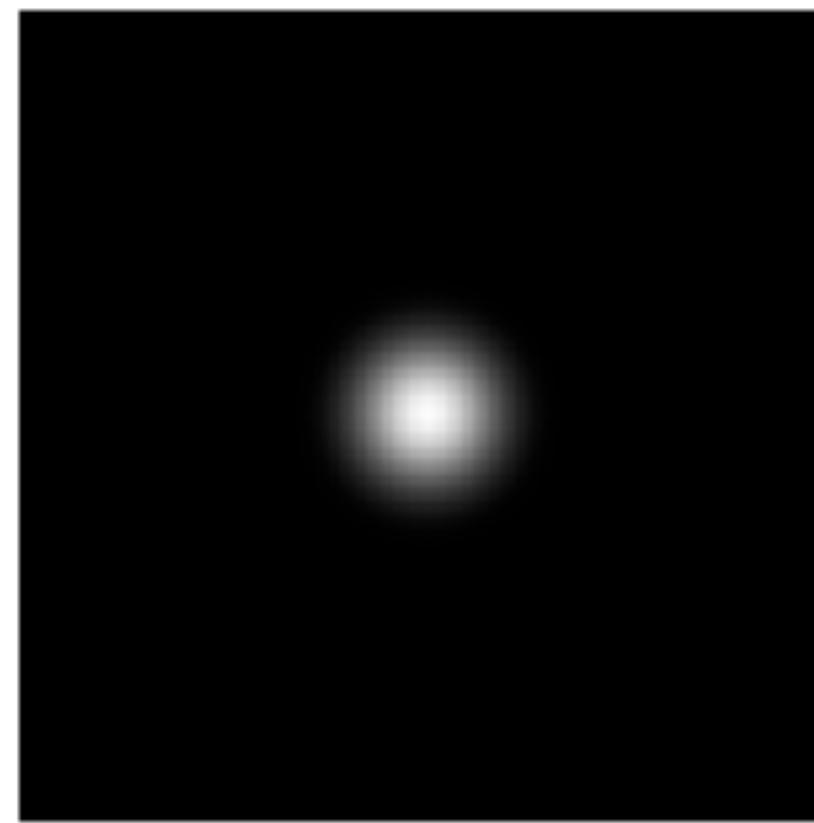
**Observation:** low frequencies close to the center

# 2D Fourier Transforms: Kernels

$f(x, y)$



$F(\omega_x, \omega_y)$

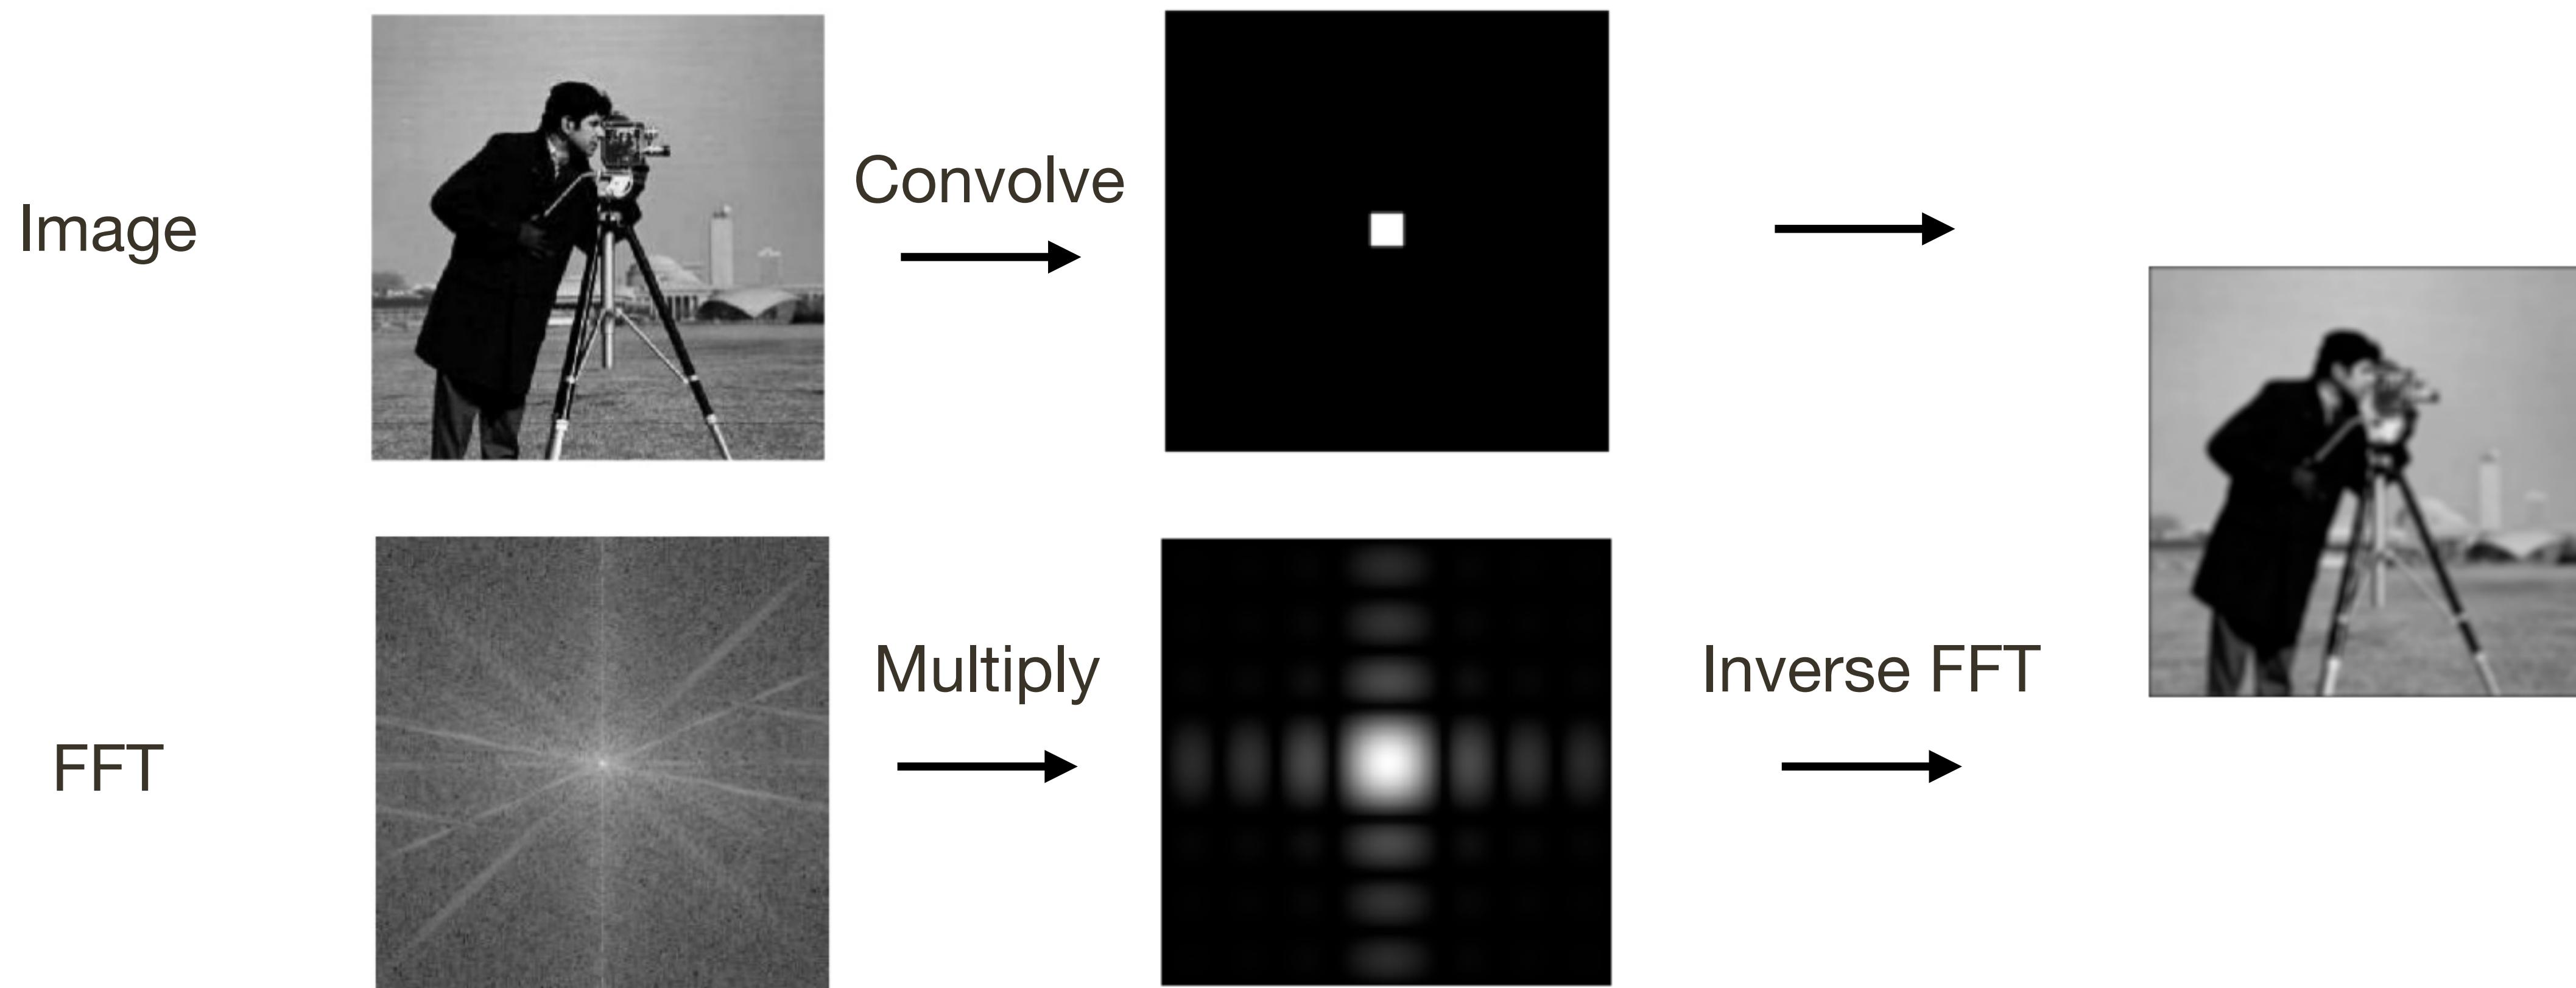


Compress power 0.5  
to exaggerate lobes  
(just for visualization)

# Convolution using Fourier Transforms

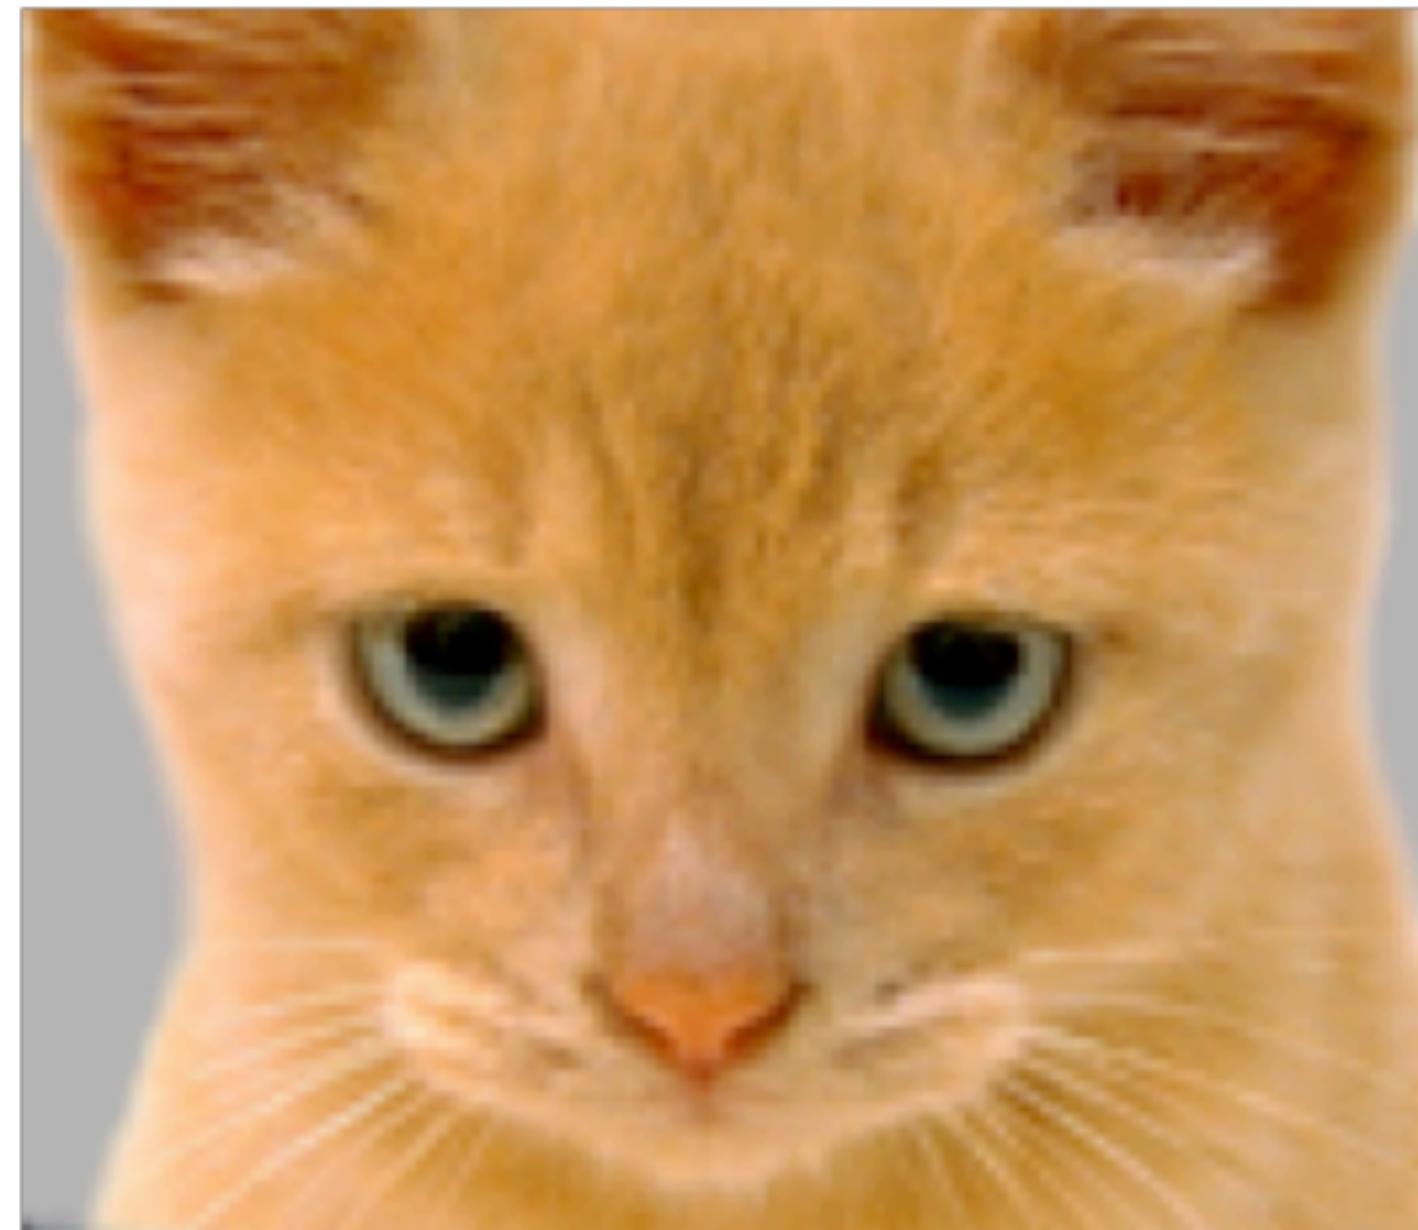
Convolution **Theorem**:  $i'(x, y) = f(x, y) \otimes i(x, y)$

$$\mathcal{I}'(w_x, w_y) = \mathcal{F}(w_x, w_y) \mathcal{I}(w_x, w_y)$$



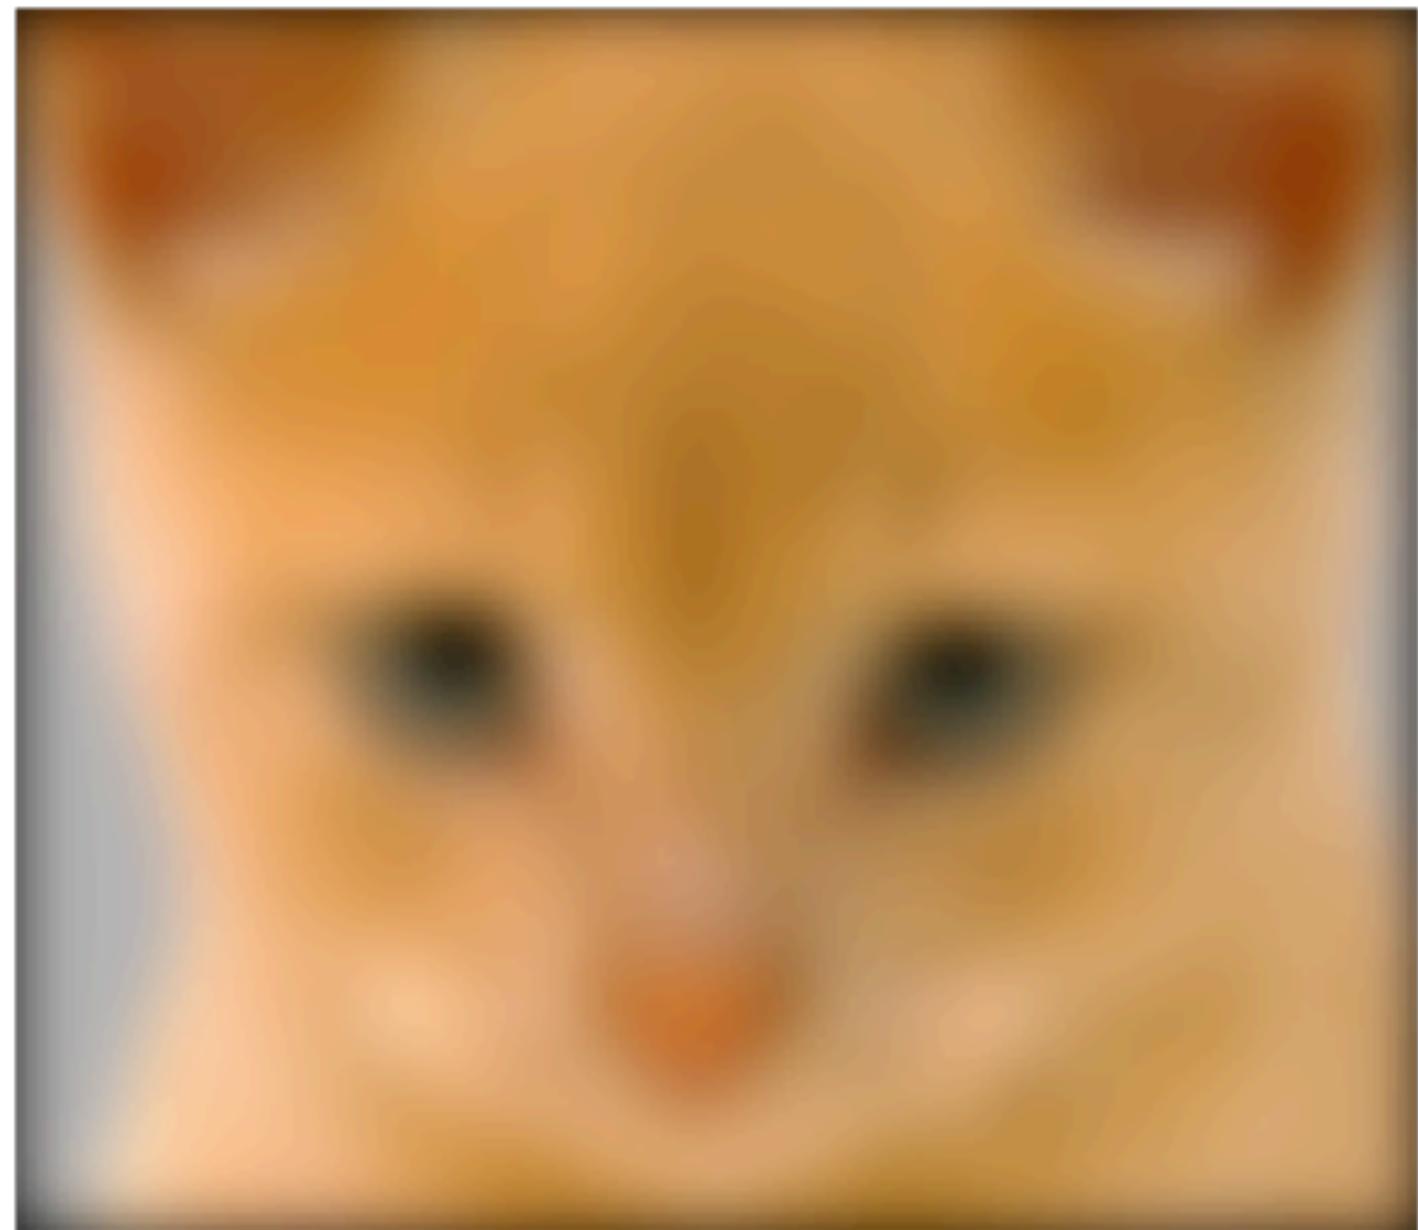
What preceded was for fun  
(you will **NOT** be tested on it)

# Assignment 1: Low/High Pass Filtering



Original

$$I(x, y)$$



Low-Pass Filter

$$I(x, y) * g(x, y)$$



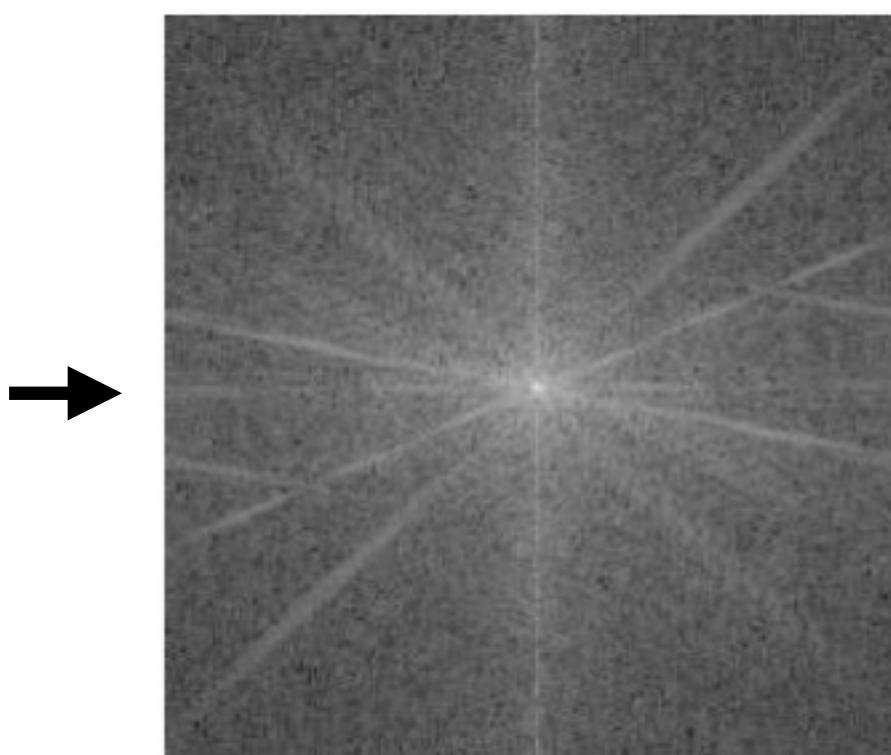
High-Pass Filter

$$I(x, y) - I(x, y) * g(x, y)$$

# Aside: You will not be tested on this ...

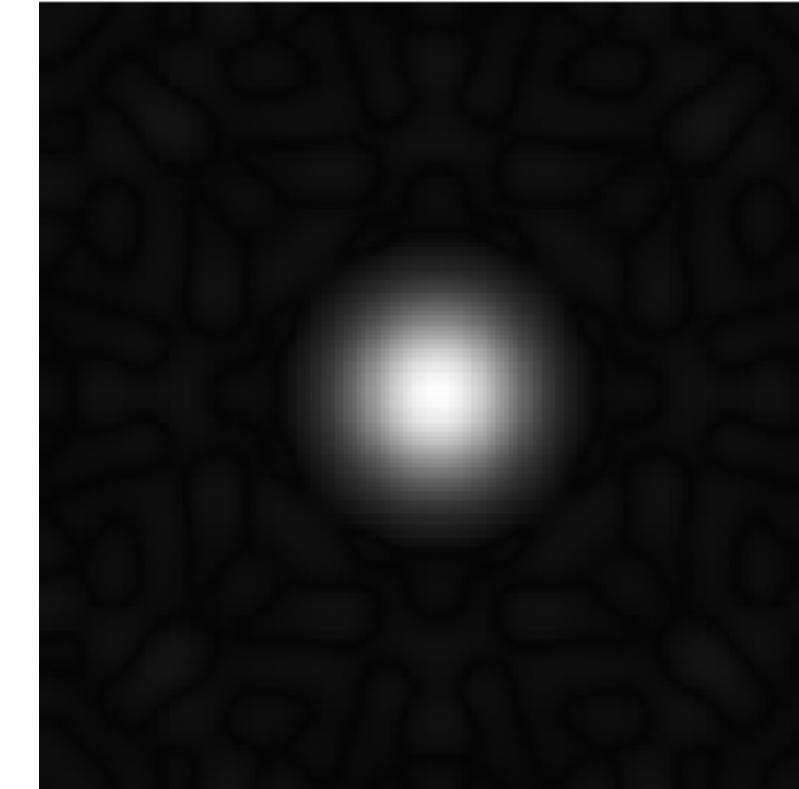


image



FFT (Mag)

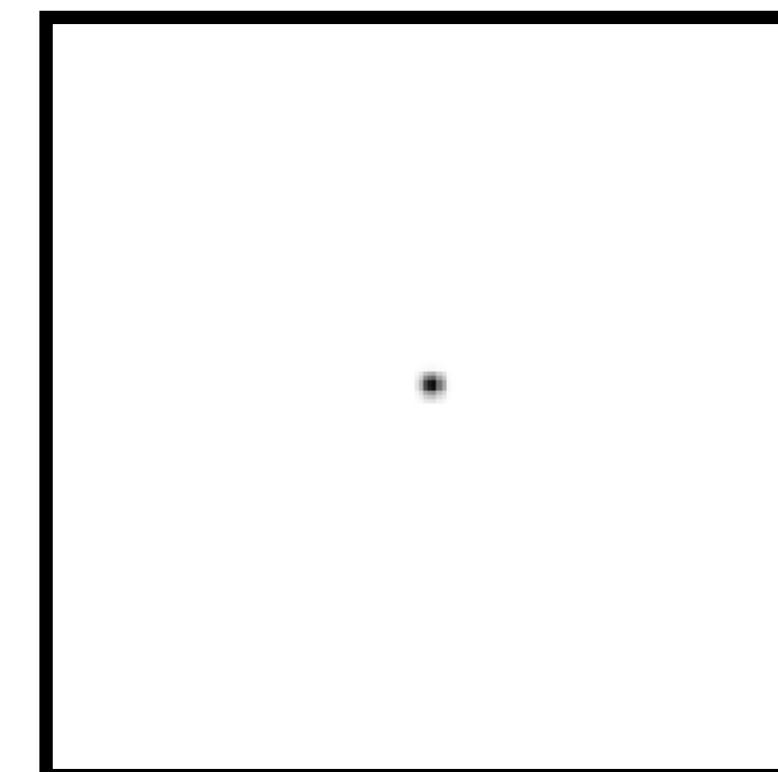
complex  
element-wise  
multiplication



Low pass



filtered image



High pass



filtered image

# Convolution using Fourier Transforms

**General** implementation of **convolution**:

At each pixel,  $(X, Y)$ , there are  $m \times m$  multiplications

There are  $n \times n$  pixels in  $(X, Y)$

---

**Total:**  $m^2 \times n^2$  multiplications

**Convolution** in FFT space:

Cost of FFT/IFFT for image:  $\mathcal{O}(n^2 \log n)$

Cost of FFT/IFFT for filter:  $\mathcal{O}(m^2 \log m)$

Worthwhile if image and kernel are **both** large

# Non-linear Filters

We've seen that **linear filters** can perform a variety of image transformations

- shifting
- smoothing
- sharpening

In some applications, better performance can be obtained by using **non-linear filters**.

For example, the median filter selects the **median** value from each pixel's neighborhood.

# Non-linear Filtering

- Example: Median filter



“shot” noise



gaussian blurred



median filtered

# Median Filter

Take the **median value** of the pixels under the filter:

5	13	5	221
4	16	7	34
24	54	34	23
23	75	89	123
54	25	67	12

Image

4	5	5	7	13	16	24	34	54
---	---	---	---	----	----	----	----	----



	13		

Output

# Median Filter

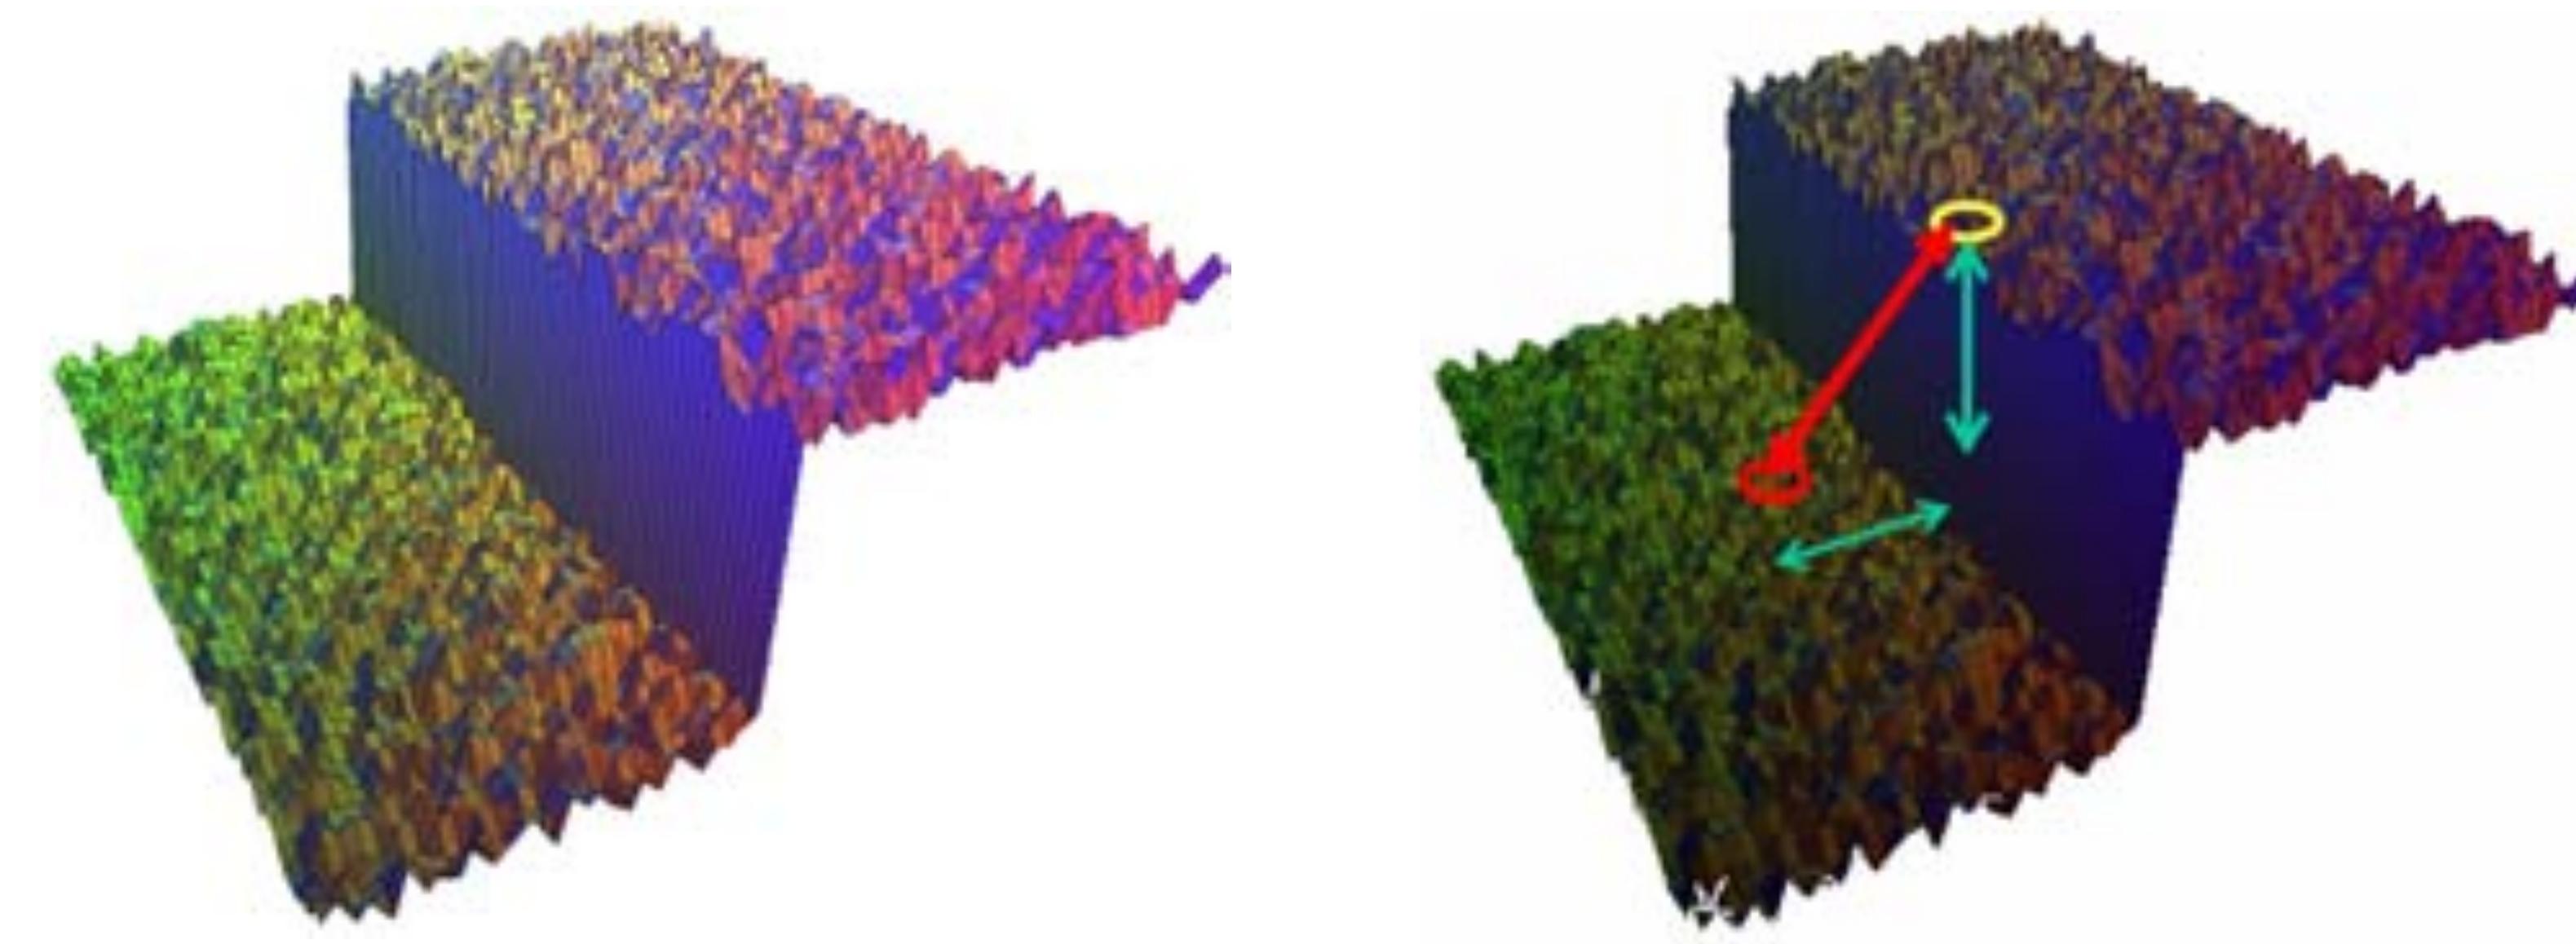
Effective at reducing certain kinds of noise, such as impulse noise (a.k.a ‘salt and pepper’ noise or ‘shot’ noise)

The median filter forces points with distinct values to be more like their neighbors



**Image credit:** [https://en.wikipedia.org/wiki/Median\\_filter#/media/File:Medianfilterp.png](https://en.wikipedia.org/wiki/Median_filter#/media/File:Medianfilterp.png)

# Bilateral Filter



Suppose we want to smooth a noisy step function

A Gaussian kernel performs a weighted average of points over a spatial neighbourhood..

But this averages points both at the top and bottom of the step – blurring

**Bilateral Filter** idea: look at distances in **range** (value) as well as **space** x,y

# Bilateral Filter

An edge-preserving non-linear filter

**Like** a Gaussian filter:

- The filter weights depend on spatial distance from the center pixel
- Pixels nearby (in space) should have greater influence than pixels far away

**Unlike** a Gaussian filter:

- The filter weights also depend on range distance from the center pixel
- Pixels with similar brightness value should have greater influence than pixels with dissimilar brightness value

# Bilateral Filter

**Gaussian** filter: weights of neighbor at a spatial offset  $(x, y)$  away from the center pixel  $I(X, Y)$  given by:

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2+y^2}{2\sigma^2}}$$

(with appropriate normalization)

# Bilateral Filter

**Gaussian** filter: weights of neighbor at a spatial offset  $(x, y)$  away from the center pixel  $I(X, Y)$  given by:

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2+y^2}{2\sigma^2}}$$

(with appropriate normalization)

**Bilateral** filter: weights of neighbor at a spatial offset  $(x, y)$  away from the center pixel  $I(X, Y)$  given by a product:

$$\exp^{-\frac{x^2+y^2}{2\sigma_d^2}} \exp^{-\frac{(I(X+x, Y+y) - I(X, Y))^2}{2\sigma_r^2}}$$

(with appropriate normalization)

# Bilateral Filter

**Gaussian** filter: weights of neighbor at a spatial offset  $(x, y)$  away from the center pixel  $I(X, Y)$  given by:

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2+y^2}{2\sigma^2}}$$

(with appropriate normalization)

**Bilateral** filter: weights of neighbor at a spatial offset  $(x, y)$  away from the center pixel  $I(X, Y)$  given by a product:

The diagram illustrates the bilateral filter kernel as a product of two components: a domain kernel and a range kernel. The domain kernel is represented by a green box containing the expression  $\exp^{-\frac{x^2+y^2}{2\sigma_d^2}}$ . The range kernel is represented by a blue box containing the expression  $\exp^{-\frac{(I(X+x, Y+y) - I(X, Y))^2}{2\sigma_r^2}}$ . The two boxes are joined together by a horizontal line.

$$\text{domain kernel } \exp^{-\frac{x^2+y^2}{2\sigma_d^2}} \times \text{range kernel } \exp^{-\frac{(I(X+x, Y+y) - I(X, Y))^2}{2\sigma_r^2}}$$

(with appropriate normalization)

# Bilateral Filter

image  $I(X, Y)$

25	0	25	255	255	255
0	0	0	230	255	255
0	25	25	255	230	255
0	0	25	255	255	255



Normalised

image  $I(X, Y)$

0.1	0	0.1	1	1	1
0	0	0	0.9	1	1
0	0.1	0.1	1	0.9	1
0	0	0.1	1	1	1

**Domain Kernel**  
 $\sigma_d = 0.45$

0.08	0.12	0.08
0.12	0.20	0.12
0.08	0.12	0.08

# Bilateral Filter

image  $I(X, Y)$

25	0	25	255	255	255
0	0	0	230	255	255
0	25	25	255	230	255
0	0	25	255	255	255

Normalised

image  $I(X, Y)$

0.1	0	0.1	1	1	1
0	0	0	0.9	1	1
0	0.1	0.1	1	0.9	1
0	0	0.1	1	1	1

**Domain Kernel**  
 $\sigma_d = 0.45$

0.08	0.12	0.08
0.12	0.20	0.12
0.08	0.12	0.08

**Range Kernel**

$$\sigma_r = 0.45$$

0.98	0.98	0.2
1	1	0.1
0.98	1	0.1

(differences based on  
**centre pixel**)

# Bilateral Filter

image  $I(X, Y)$

25	0	25	255	255	255
0	0	0	230	255	255
0	25	25	255	230	255
0	0	25	255	255	255

Normalised

image  $I(X, Y)$

0.1	0	0.1	1	1	1
0	0	0	0.9	1	1
0	0.1	0.1	1	0.9	1
0	0	0.1	1	1	1

**Domain Kernel**  
 $\sigma_d = 0.45$

0.08	0.12	0.08
0.12	0.20	0.12
0.08	0.12	0.08

**Range Kernel**

$$\sigma_r = 0.45$$

0.98	0.98	0.2
1	1	0.1
0.98	1	0.1

**Range \* Domain Kernel**

multiply

0.08	0.12	0.02
0.12	0.20	0.01
0.08	0.12	0.01

(differences based on  
**centre pixel**)

# Bilateral Filter

image  $I(X, Y)$

25	0	25	255	255	255
0	0	0	230	255	255
0	25	25	255	230	255
0	0	25	255	255	255

Normalised

image  $I(X, Y)$

0.1	0	0.1	1	1	1
0	0	0	0.9	1	1
0	0.1	0.1	1	0.9	1
0	0	0.1	1	1	1

**Domain Kernel**  
 $\sigma_d = 0.45$

0.08	0.12	0.08
0.12	0.20	0.12
0.08	0.12	0.08

**Range Kernel**

$$\sigma_r = 0.45$$

0.98	0.98	0.2
1	1	0.1
0.98	1	0.1

**Range \* Domain Kernel**

multiply

0.08	0.12	0.02
0.12	0.20	0.01
0.08	0.12	0.01

sum to 1

0.11	0.16	0.03
0.16	0.26	0.01
0.11	0.16	0.01

(differences based on  
**centre pixel**)

# Bilateral Filter

image  $I(X, Y)$

25	0	25	255	255	255
0	0	0	230	255	255
0	25	25	255	230	255
0	0	25	255	255	255
0	0	25	255	255	255

Normalised

image  $I(X, Y)$

0.1	0	0.1	1	1	1
0	0	0	0.9	1	1
0	0.1	0.1	1	0.9	1
0	0	0.1	1	1	1
0	0	0.1	1	1	1

**Domain Kernel**  
 $\sigma_d = 0.45$

0.08	0.12	0.08
0.12	0.20	0.12
0.08	0.12	0.08

**Range Kernel**

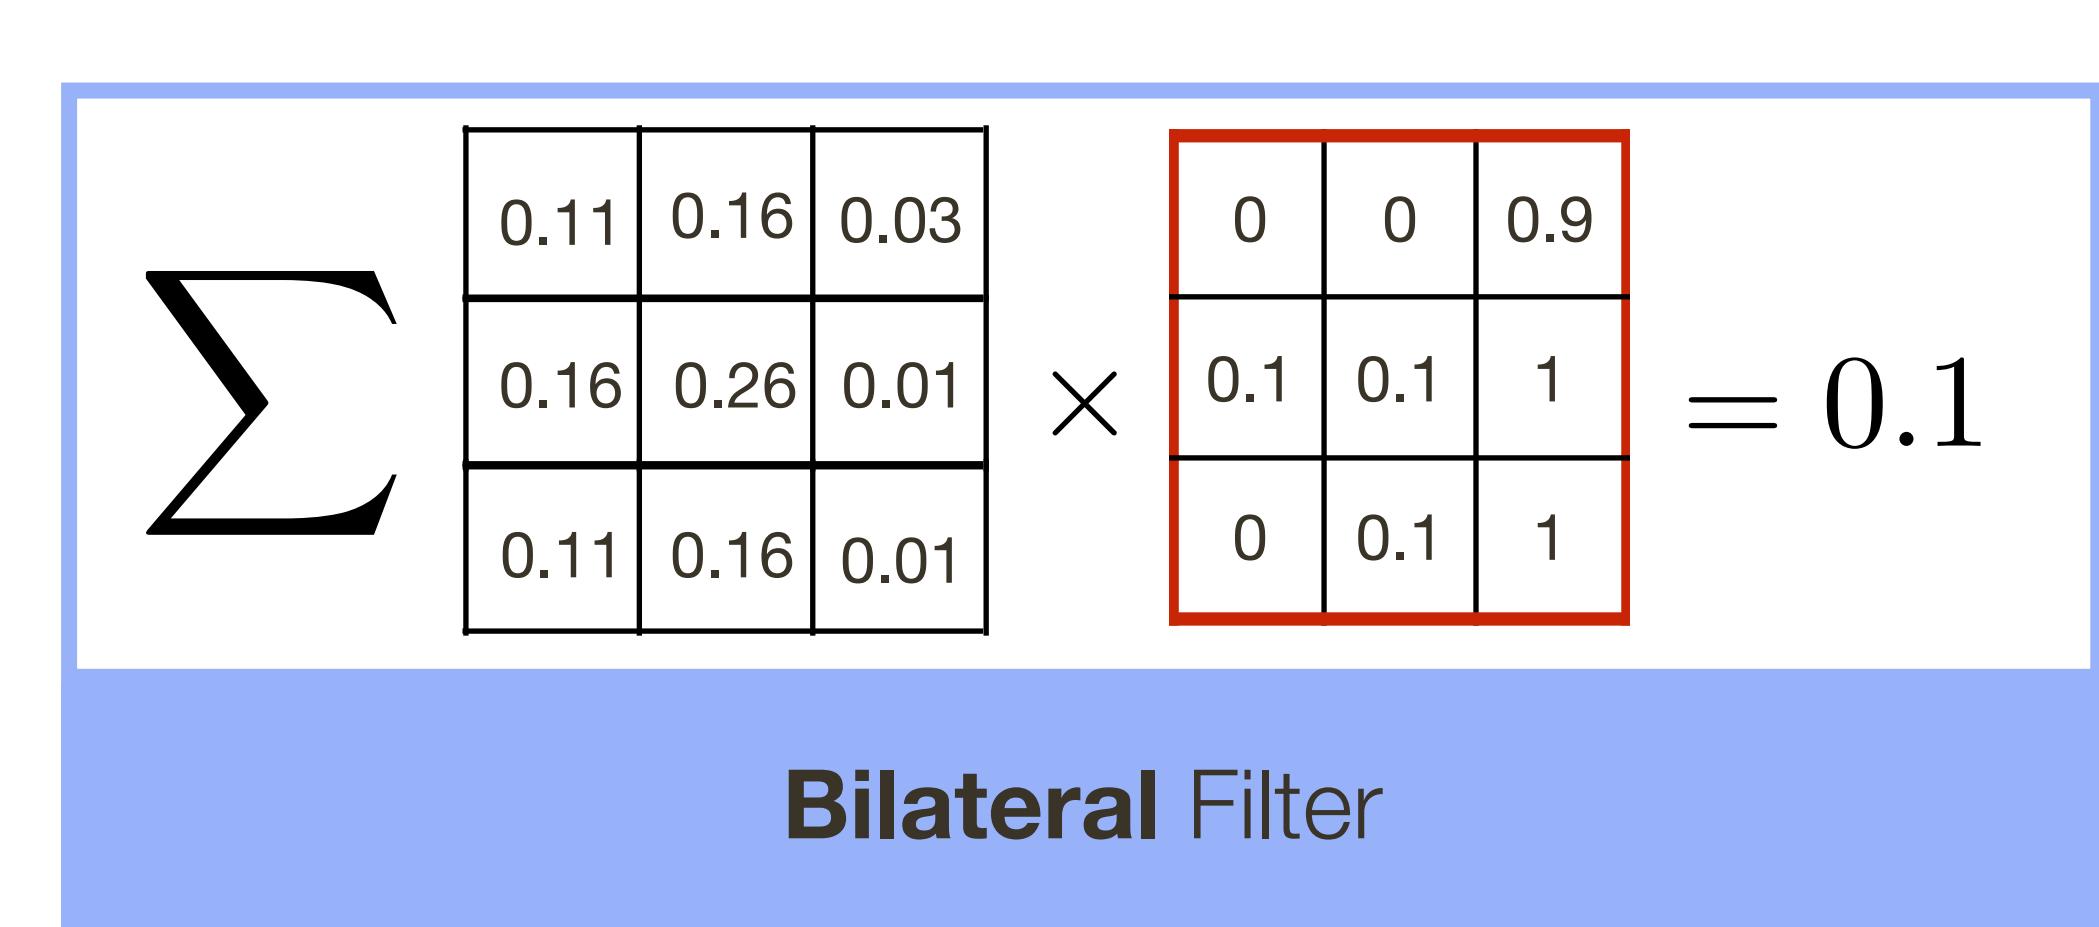
$$\sigma_r = 0.45$$

0.98	0.98	0.2
1	1	0.1
0.98	1	0.1

multiply  
→

0.08	0.12	0.02
0.12	0.20	0.01
0.08	0.12	0.01

(differences based on  
**centre pixel**)



# Bilateral Filter

image  $I(X, Y)$

25	0	25	255	255	255
0	0	0	230	255	255
0	25	25	255	230	255
0	0	25	255	255	255

Normalised

image  $I(X, Y)$

0.1	0	0.1	1	1	1
0	0	0	0.9	1	1
0	0.1	0.1	1	0.9	1
0	0	0.1	1	1	1



**Domain Kernel**  
 $\sigma_d = 0.45$

$$\sum \begin{array}{|c|c|c|} \hline 0.08 & 0.12 & 0.08 \\ \hline 0.12 & 0.20 & 0.12 \\ \hline 0.08 & 0.12 & 0.08 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 0 & 0 & 0.9 \\ \hline 0.1 & 0.1 & 1 \\ \hline 0 & 0.1 & 1 \\ \hline \end{array} = 0.3$$

**Gaussian Filter (only)**

**Range Kernel**

$$\sigma_r = 0.45$$

0.98	0.98	0.2
1	1	0.1
0.98	1	0.1

multiply  
→

0.08	0.12	0.02
0.12	0.20	0.01
0.08	0.12	0.01

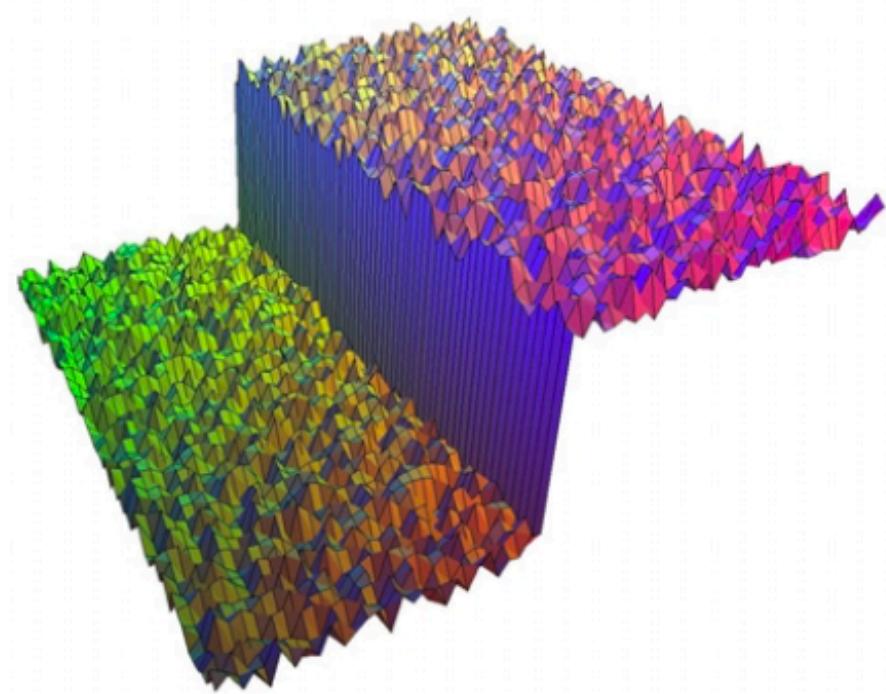
(differences based on  
**centre pixel**)

**Range \* Domain Kernel**

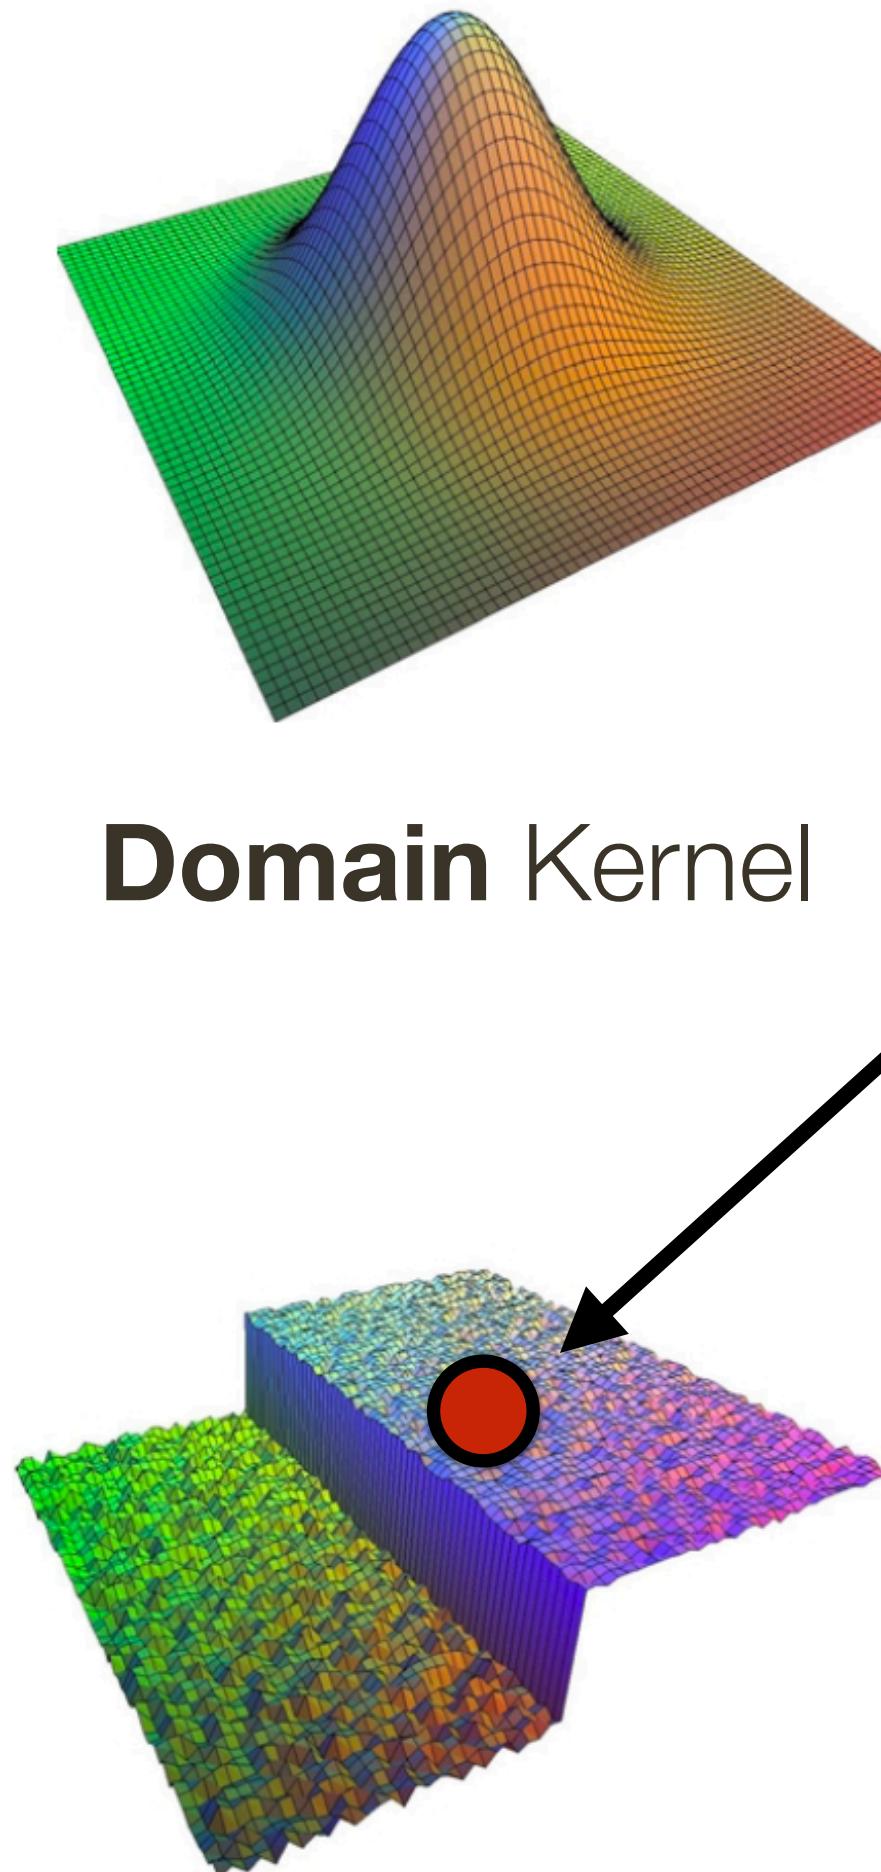
$$\sum \begin{array}{|c|c|c|} \hline 0.11 & 0.16 & 0.03 \\ \hline 0.16 & 0.26 & 0.01 \\ \hline 0.11 & 0.16 & 0.01 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 0 & 0 & 0.9 \\ \hline 0.1 & 0.1 & 1 \\ \hline 0 & 0.1 & 1 \\ \hline \end{array} = 0.1$$

**Bilateral Filter**

# Bilateral Filter



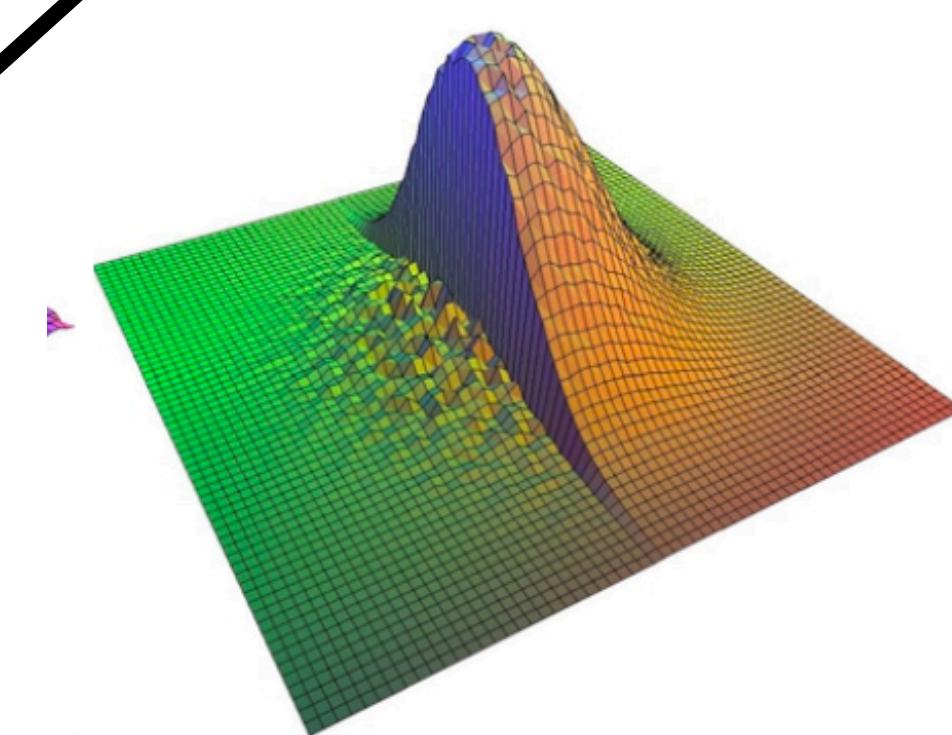
**Input**



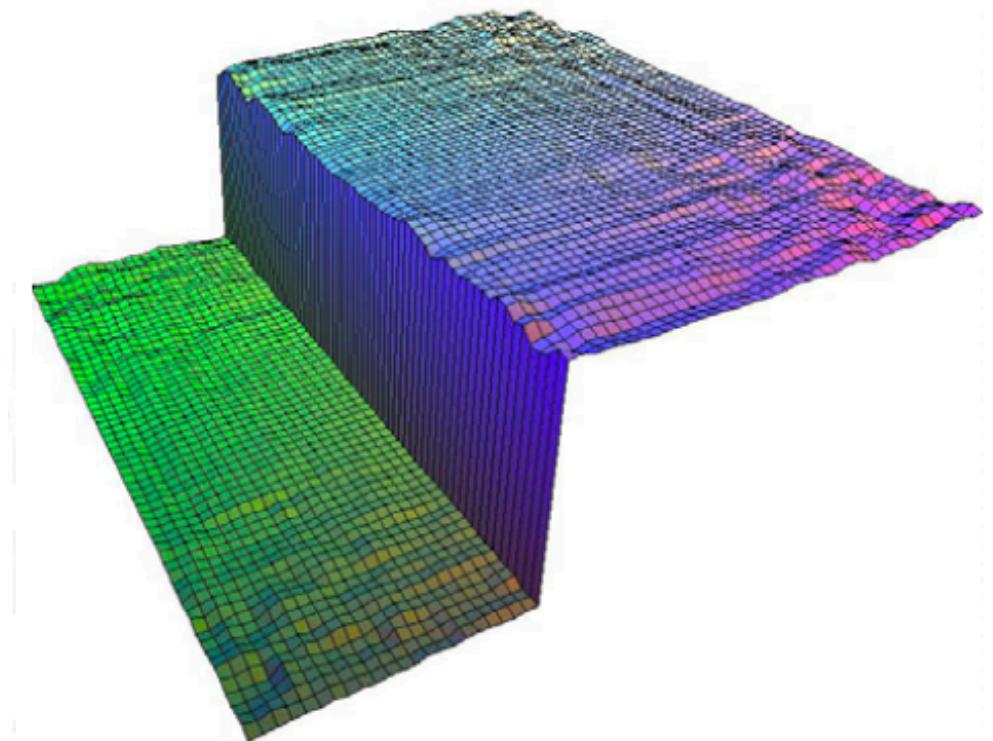
**Range Kernel Influence**

This example:  
weights for point  
on top of edge

**Domain Kernel**



**Bilateral Filter**  
(domain \* range)



**Output**

**Images from:** Durand and Dorsey, 2002

# Bilateral Filter Application: Denoising



**Noisy Image**



**Gaussian Filter**



**Bilateral Filter**

# Bilateral Filter Application: Cartooning



Original Image



After 5 iterations of **Bilateral Filter**

# Bilateral Filter Application: Flash Photography

Non-flash images taken under low light conditions often suffer from excessive **noise** and **blur**

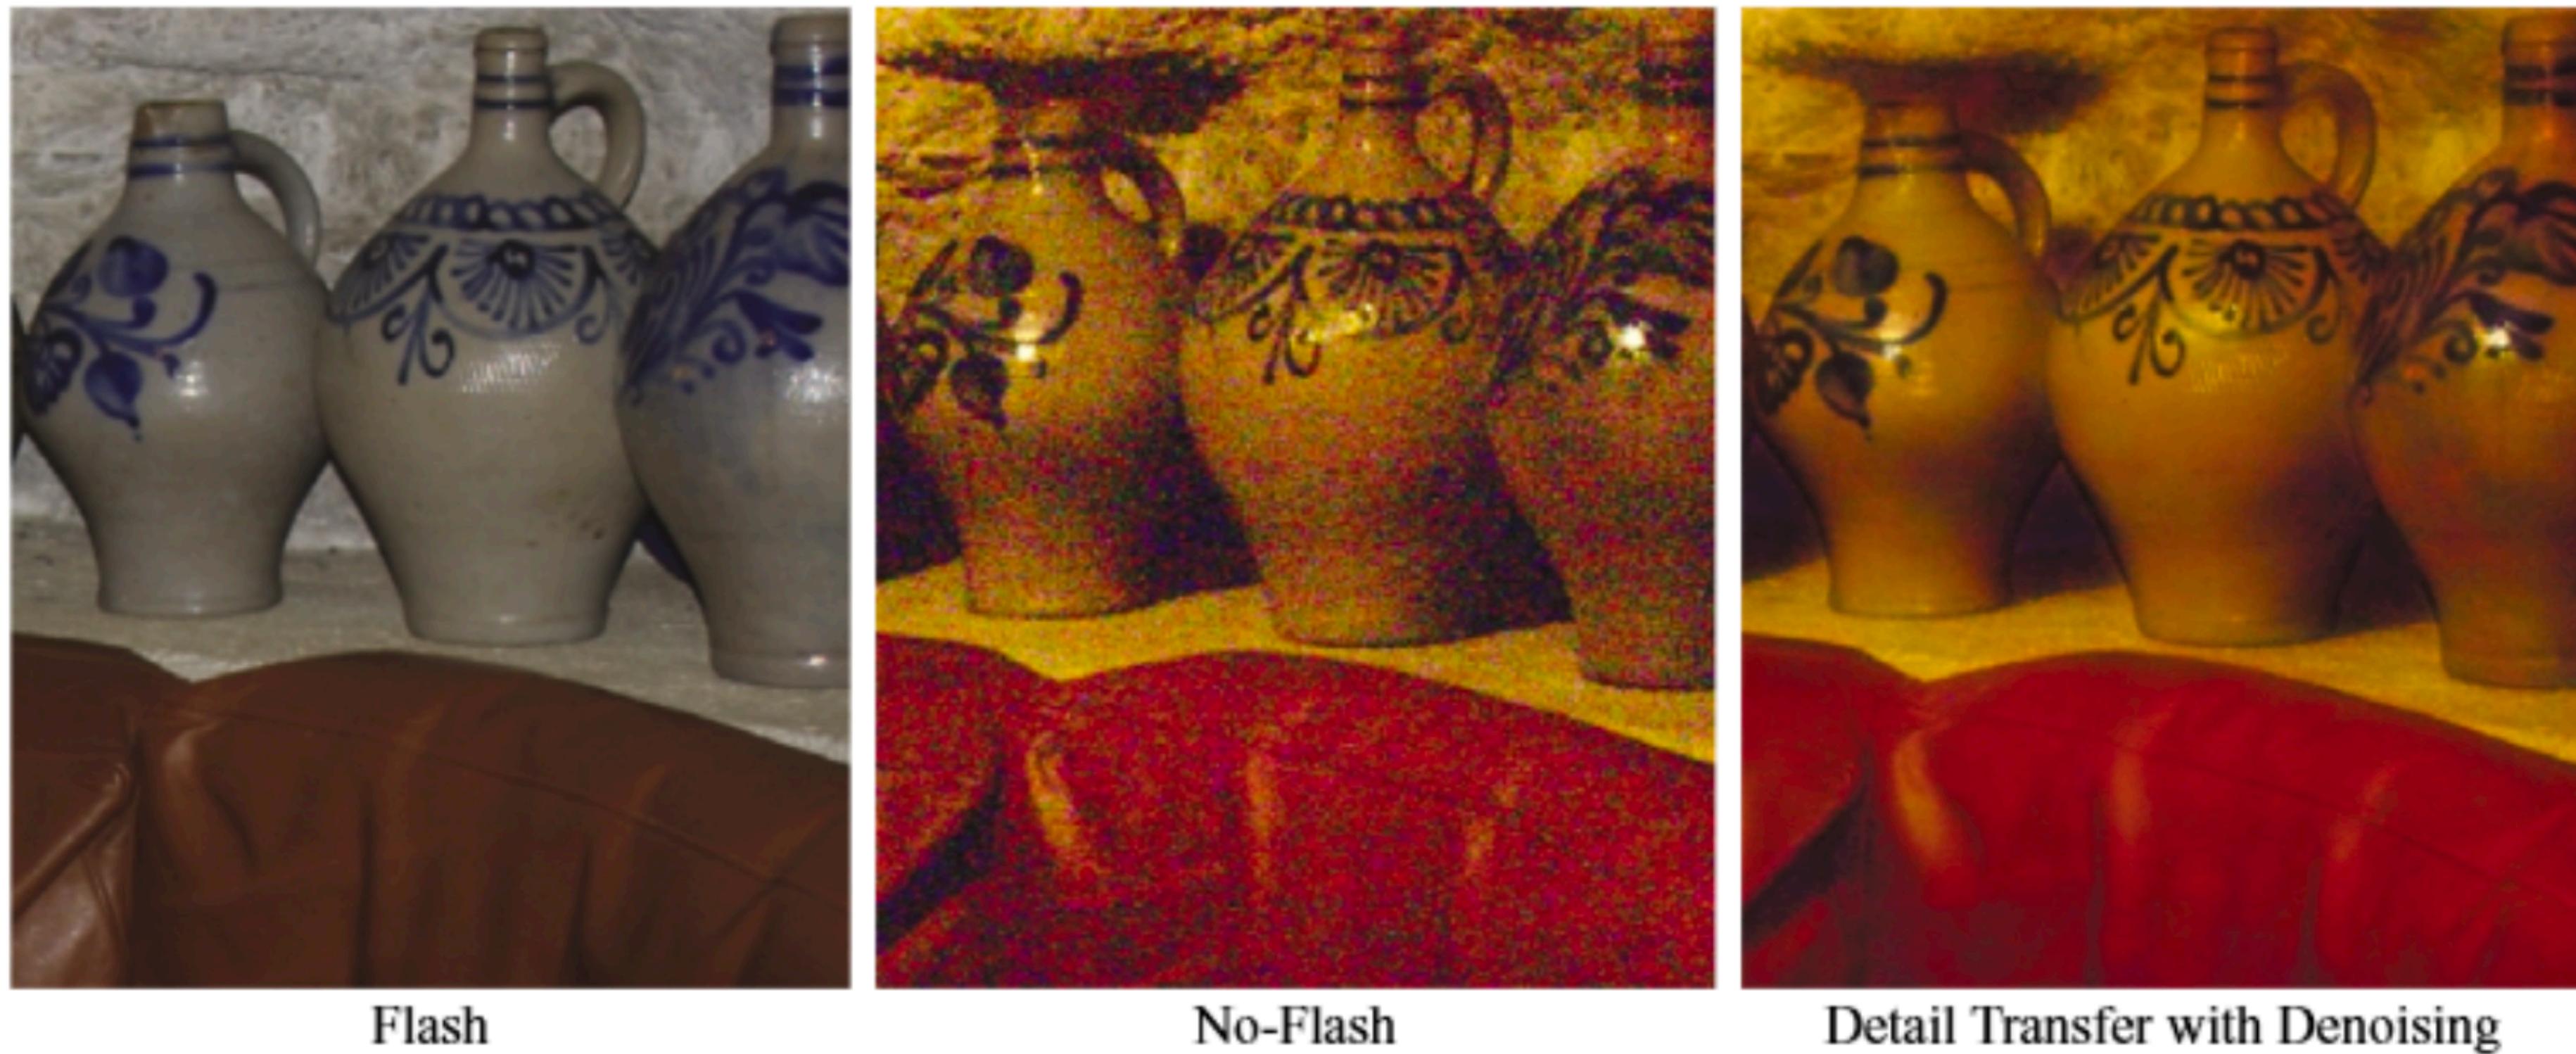
But there are problems with **flash images**:

- colour is often unnatural
- there may be strong shadows or specularities

**Idea:** Combine flash and non-flash images to achieve better exposure and colour balance, and to reduce noise

# Bilateral Filter Application: Flash Photography

System using ‘joint’ or ‘cross’ bilateral filtering:

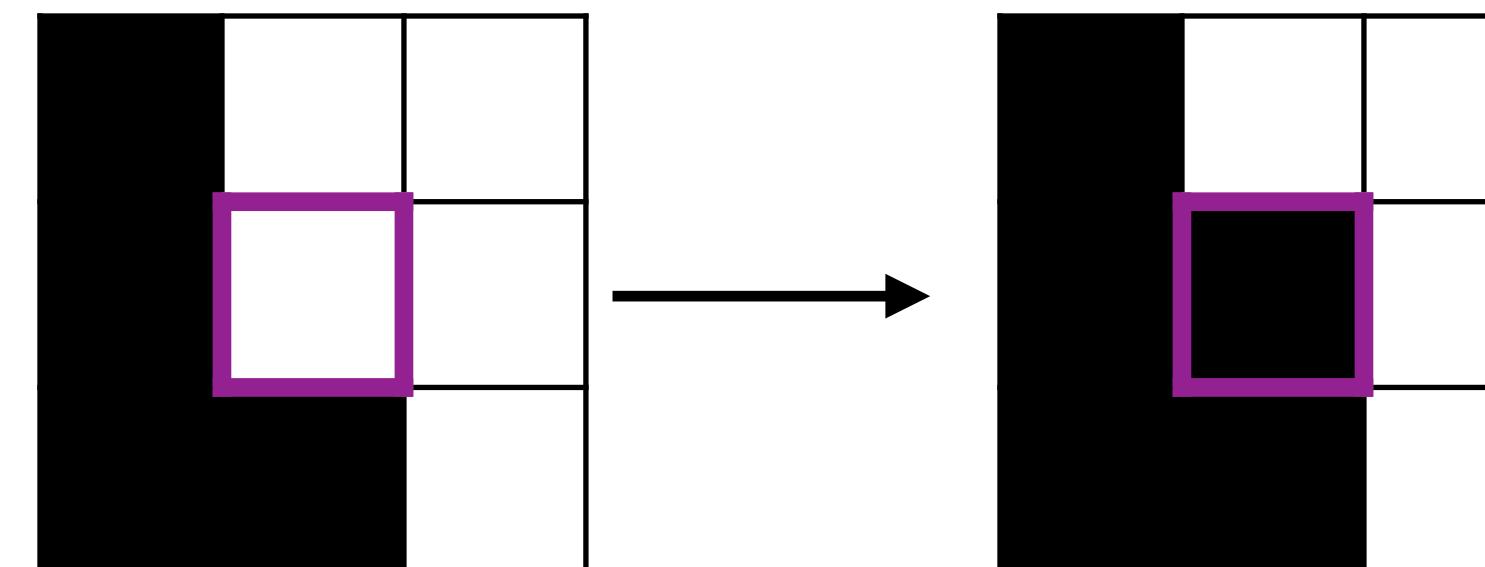
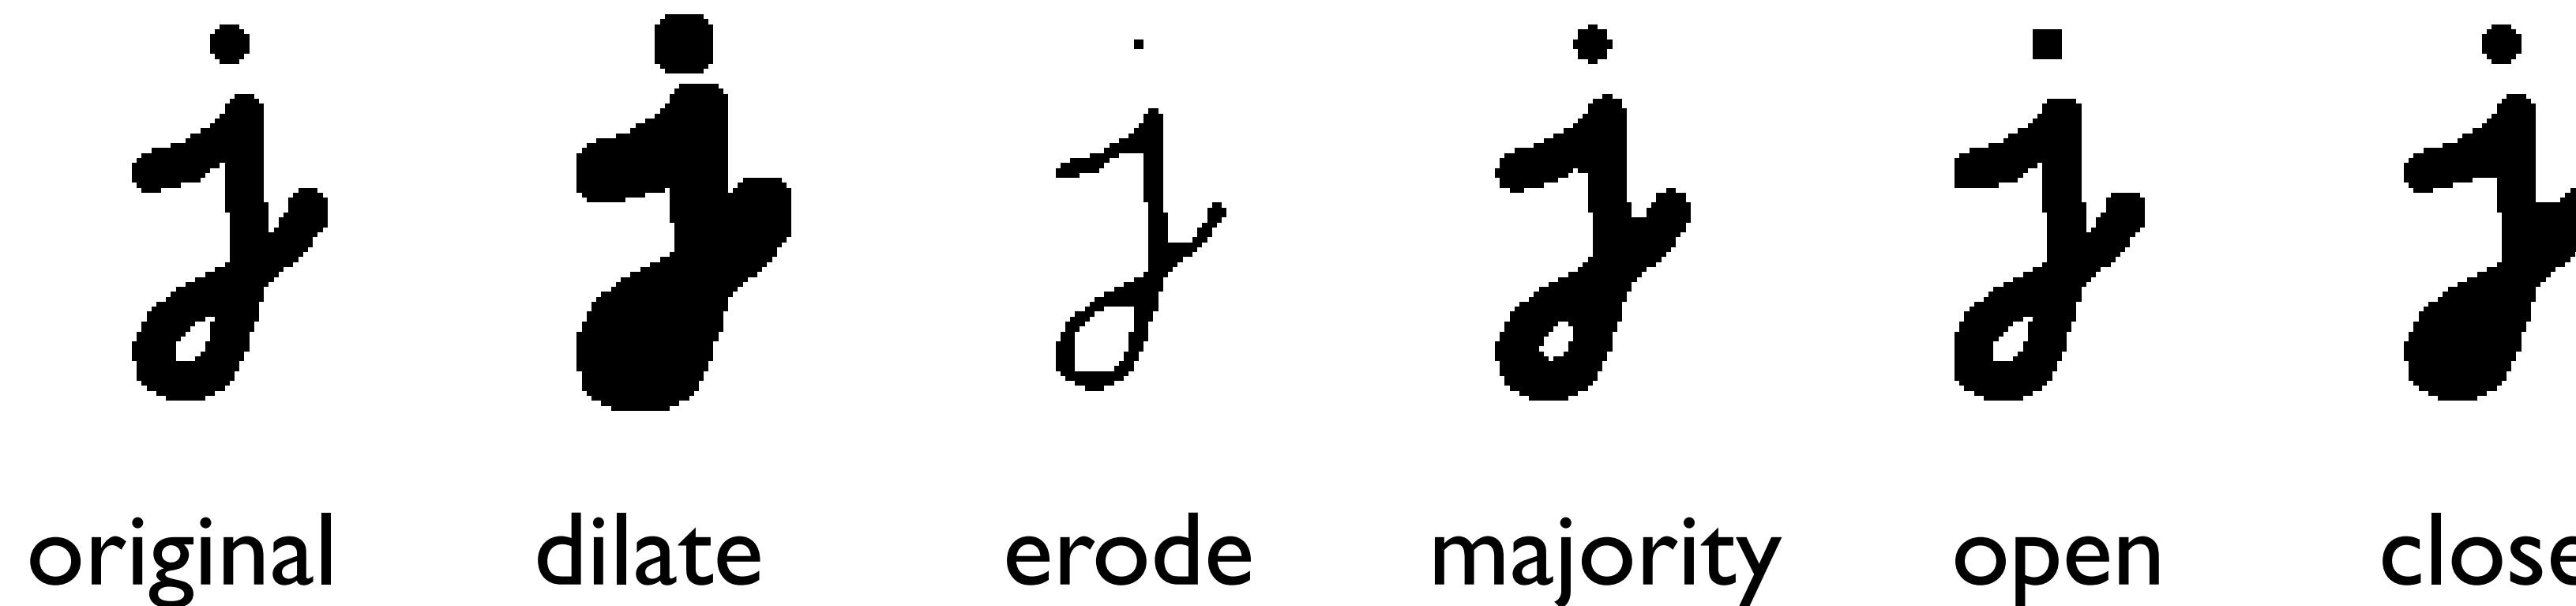


**‘Joint’ or ‘Cross’ bilateral:** Range kernel is computed using a separate guidance image instead of the input image

**Figure Credit:** Petschnigg et al., 2004

# Morphology

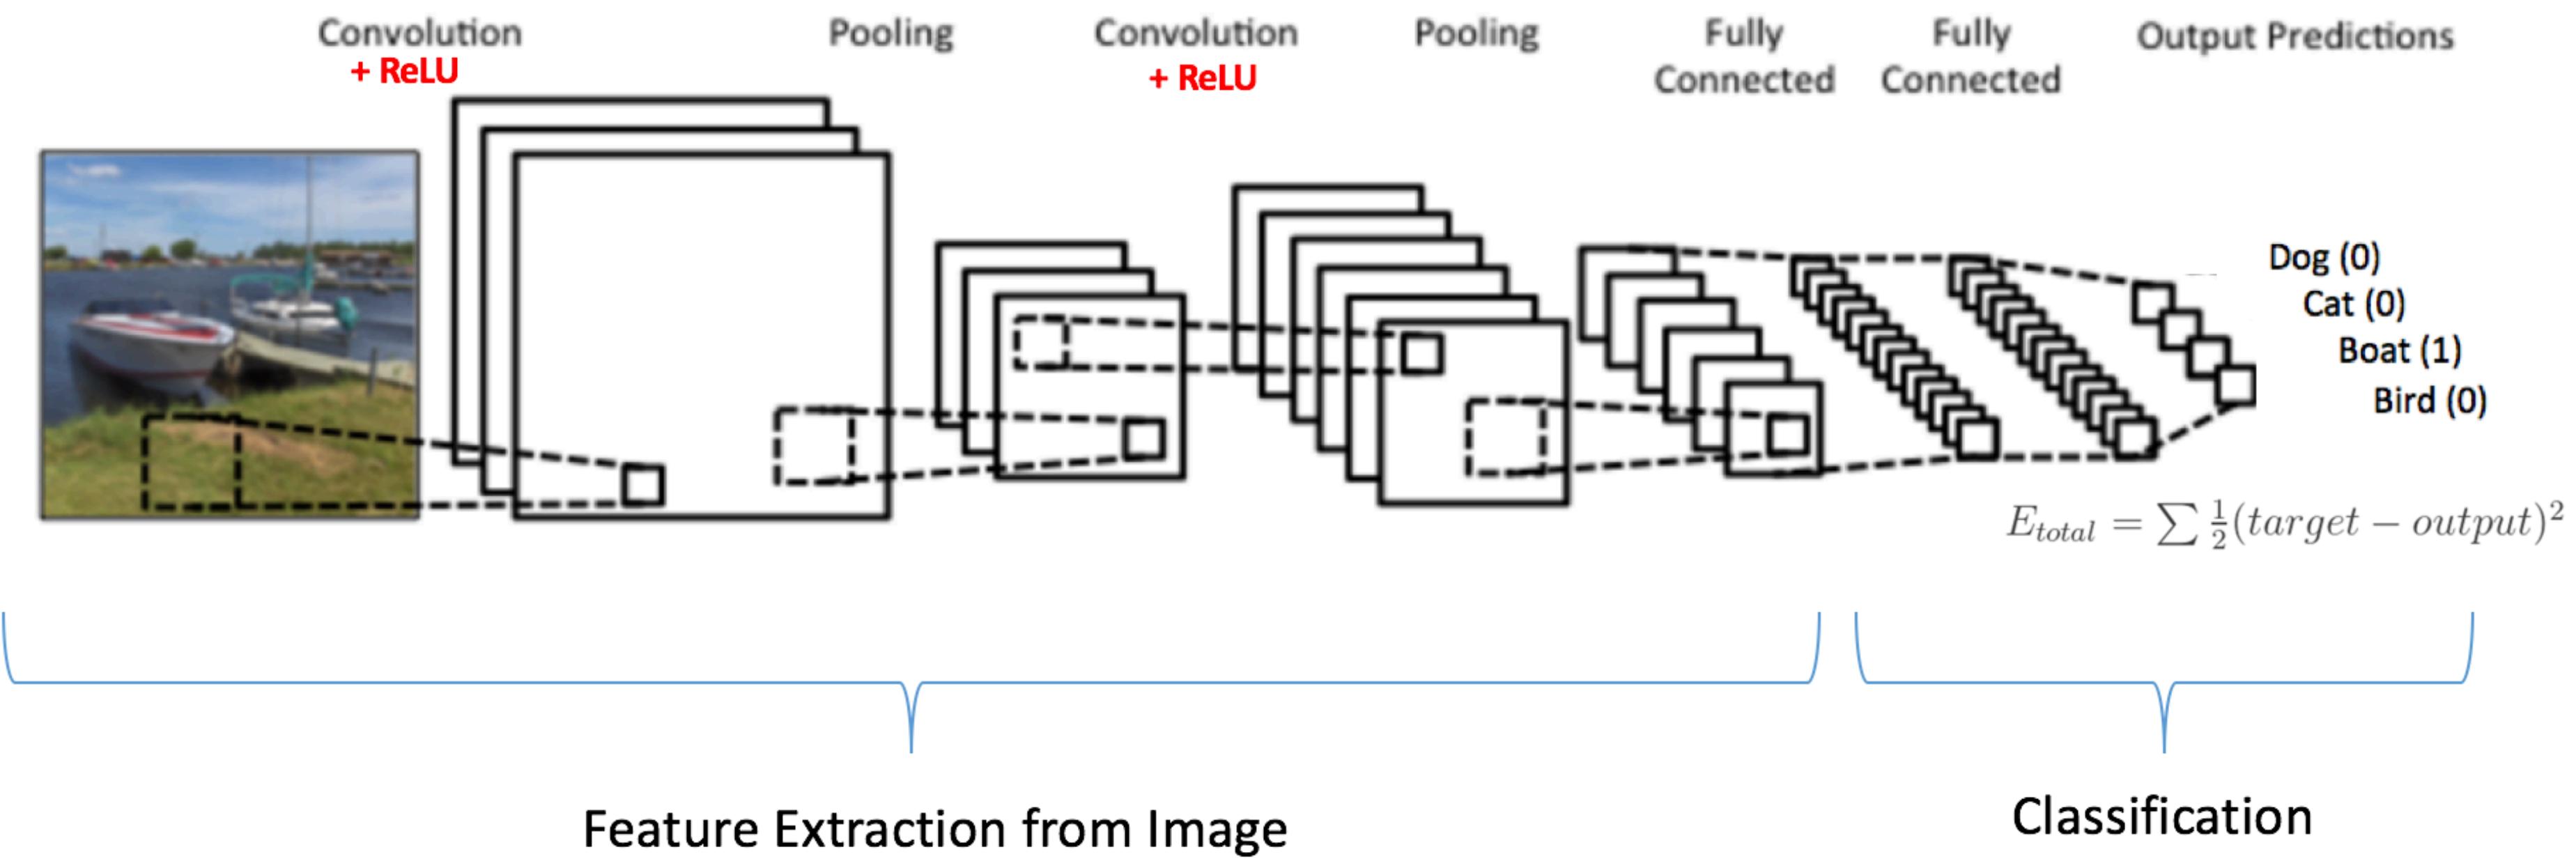
- Non-linear binary image operations



Threshold function  
in local structuring  
element

$\text{close}(\cdot) = \text{erode}(\text{dilate}(\cdot))$  etc., see Szeliski 3.3.2

# Aside: Linear Filter with ReLU



9	3	5	-8
-6	2	-3	1
1	3	4	1
3	-4	5	1



9	3	5	0
0	2	0	1
1	3	4	1
3	0	5	1

Result of: Linear Image Filtering

After Non-linear ReLU

# Summary

We covered two three **non-linear filters**: Median, Bilateral, ReLU

The **median filter** is a non-linear filter that selects the median in the neighbourhood

The **bilateral filter** is a non-linear filter that considers both spatial distance and range (intensity) distance, and has edge-preserving properties

**Speeding-up Convolution** can be achieved using separable filters or Fourier Transforms if the filter and image are both large

**Fourier Transforms** give us a way to think about image processing operations in Frequency Space, e.g., low pass filter = removing high frequency components

# Menu for Today

## Topics:

- **Linear Filtering** recap
- Efficient convolution, Fourier aside
- **Quiz 1**
- **Non-linear Filters:**  
Median, ReLU, Bilateral Filter

## Readings:

- **Today's Lecture:** Szeliski 3.3-3.4, Forsyth & Ponce (2nd ed.) 4.4

## Reminders:

- **Assignment 1:** Image Filtering and Hybrid Images due **September 28th**