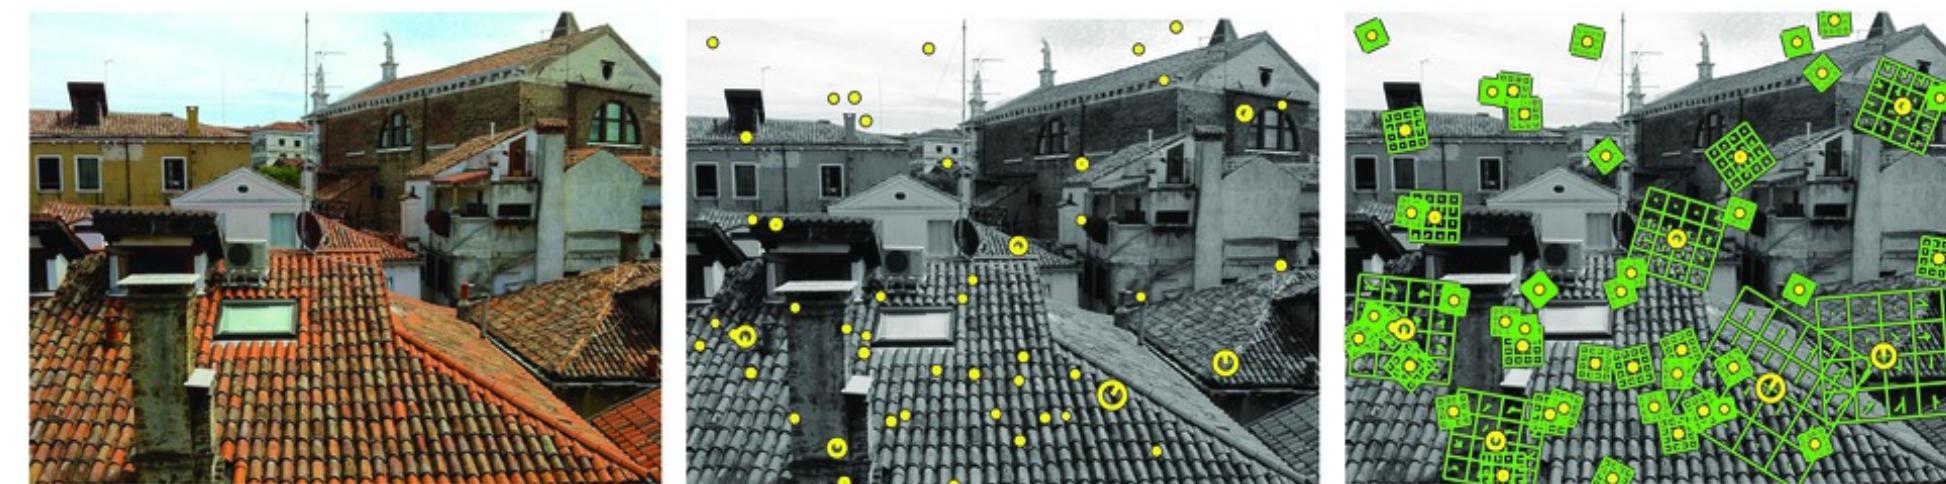




CPSC 425: Computer Vision



Lecture 13: Planar Geometry and RANSAC

Menu for Today

Topics:

- **Planar Geometry**
- **Image Alignment**, Object Recognition
- **RANSAC**

Readings:

- **Today's Lecture:** Szeliski 2.1, 8.1, Forsyth & Ponce 10.4.2

Reminders:

- **Assignment 4:** RANSAC and Panorama Stitching – **now available**

Image Alignment

- Aim: warp our images together using a 2D transformation



Image Alignment

- Aim: warp our images together using a 2D transformation



Image Alignment

- Find corresponding (matching) points between the images

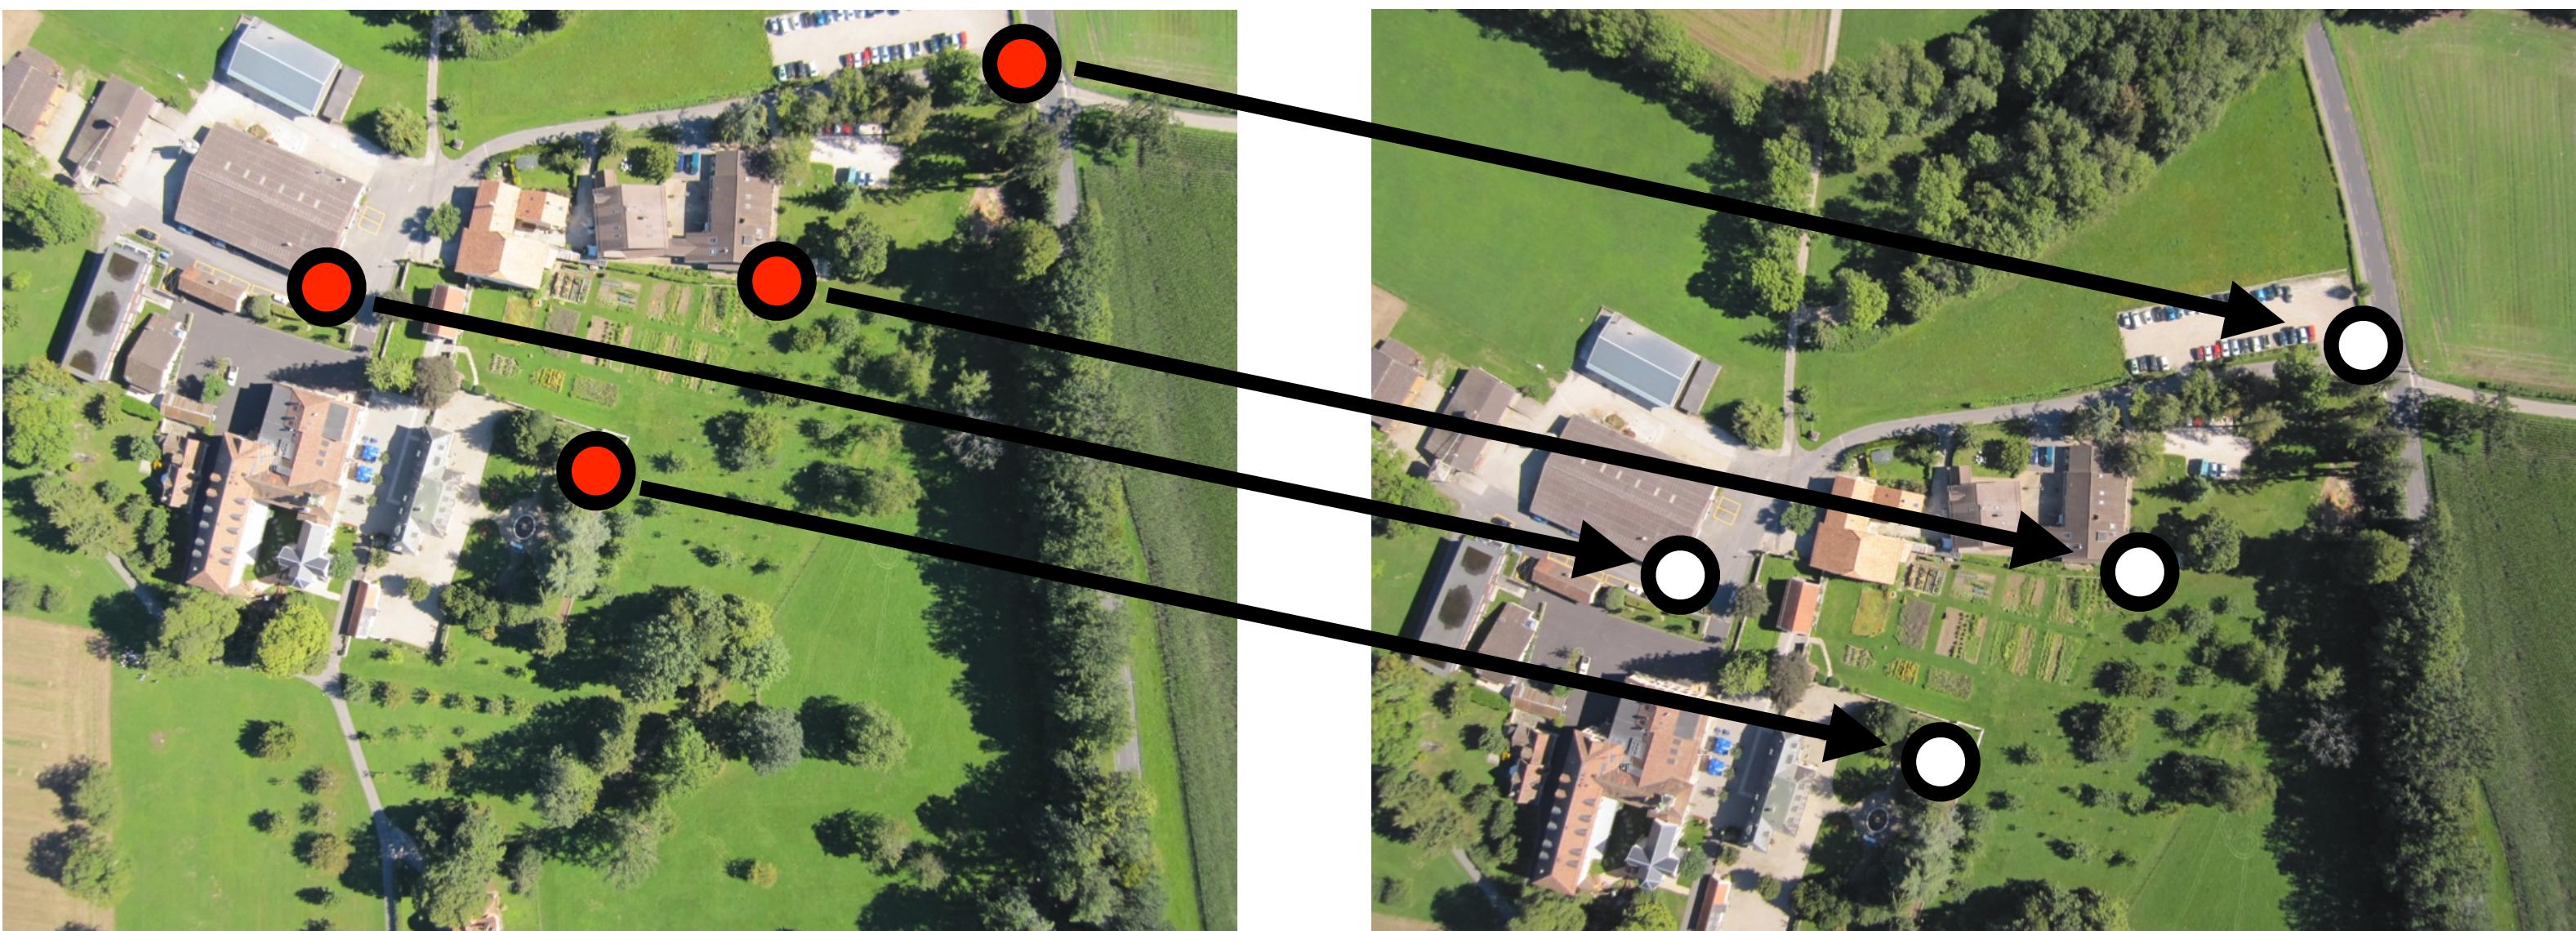


Image Alignment

- Compute the transformation to align the points



Image Alignment

- We can also use this transformation to reject outliers

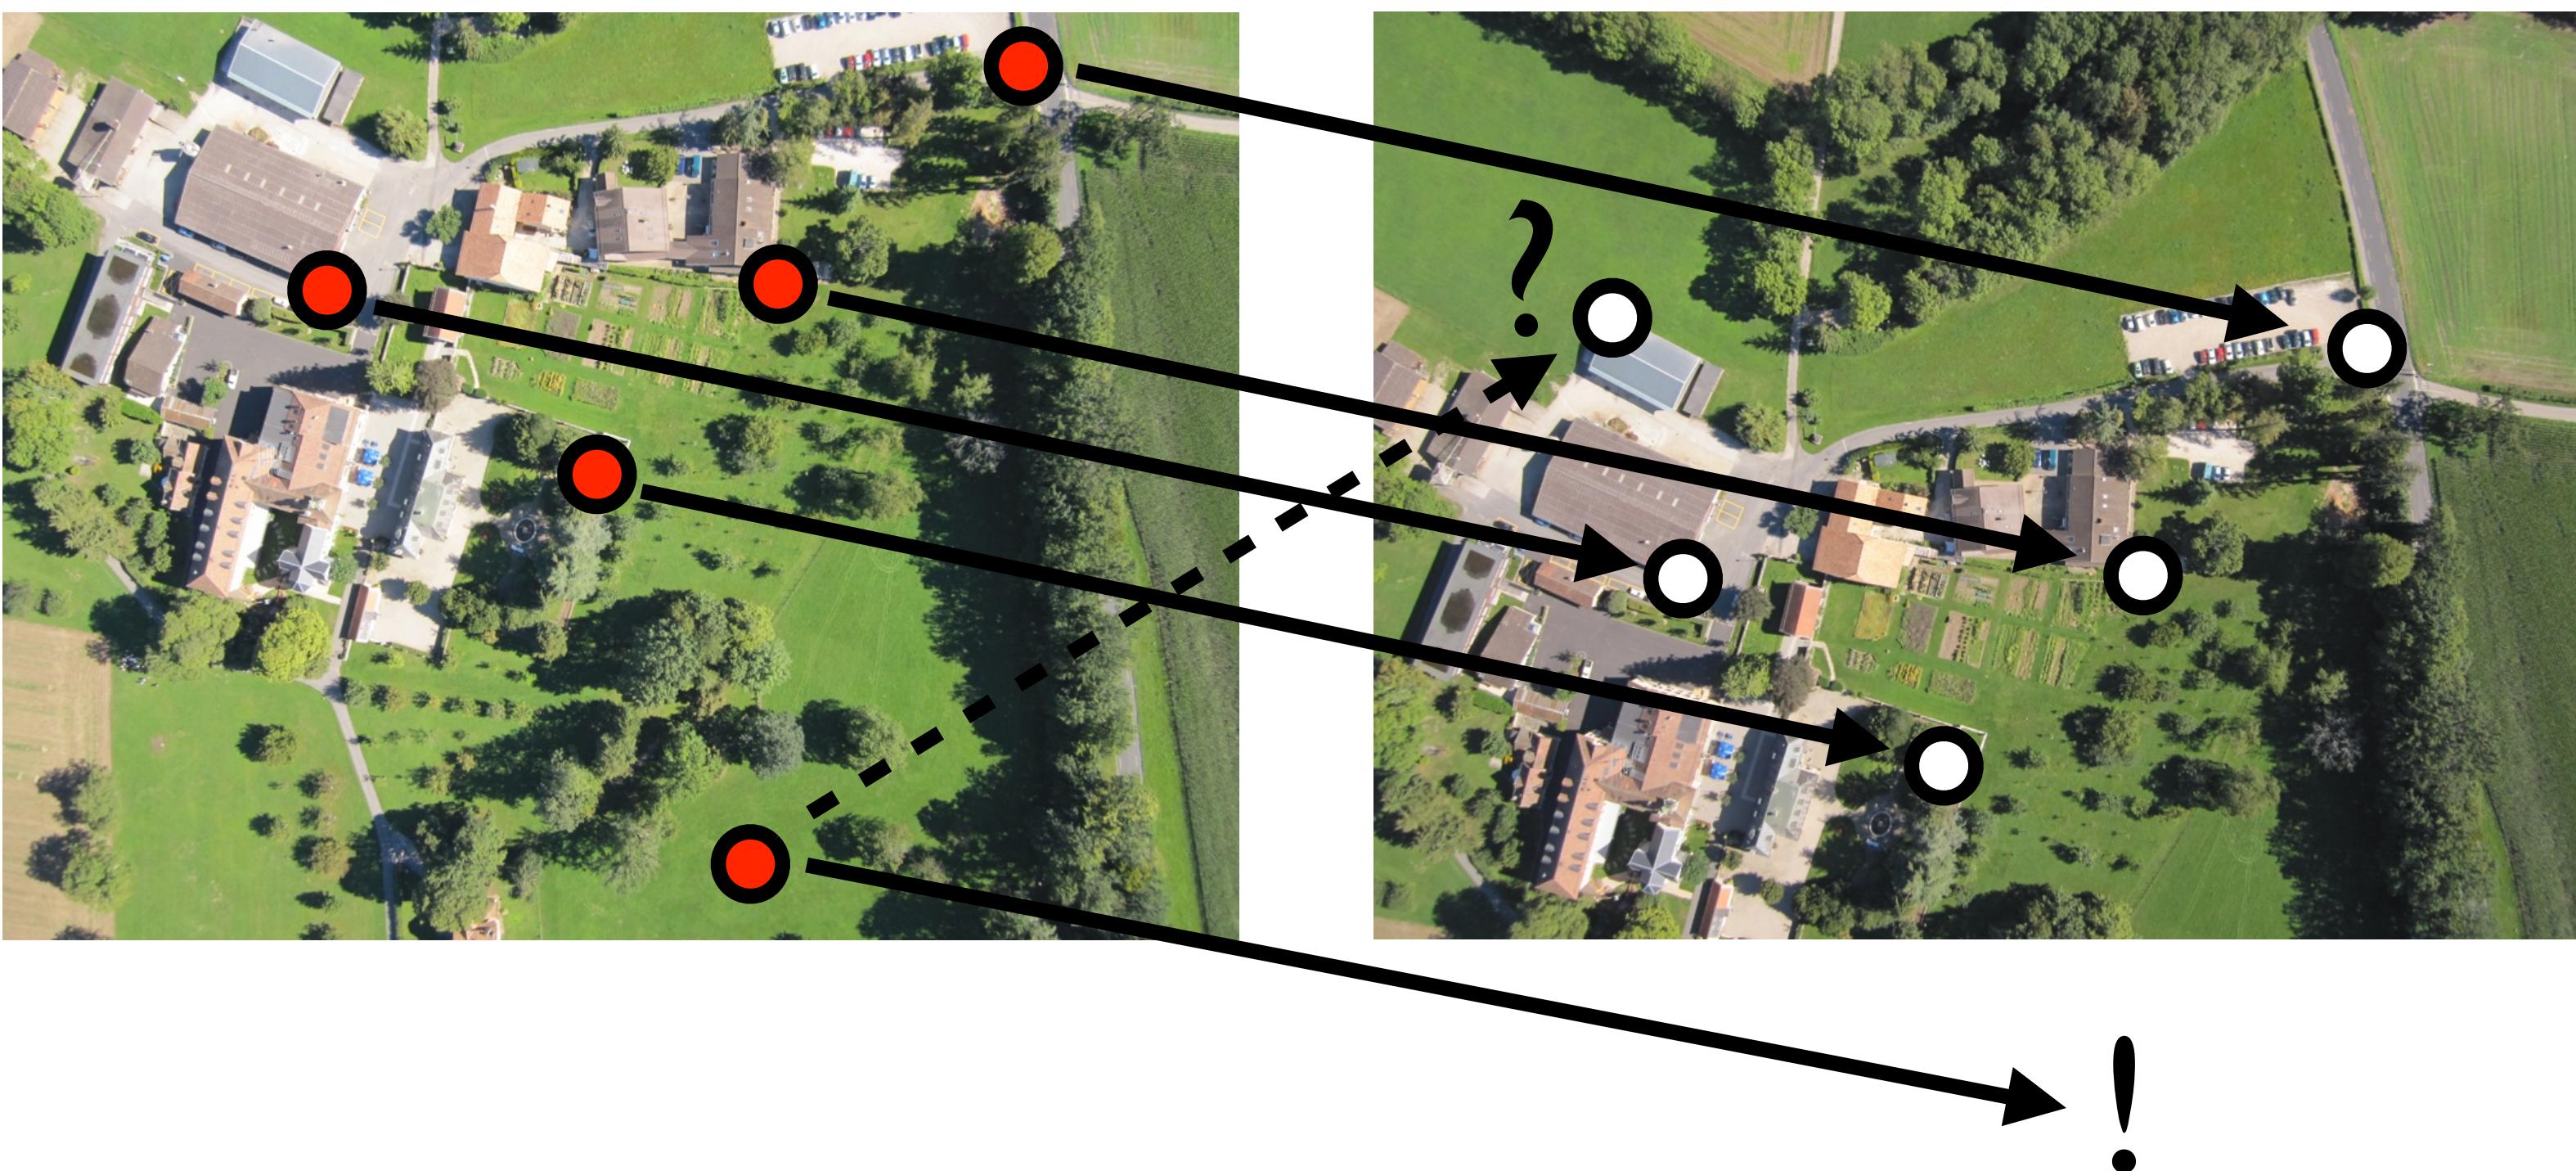
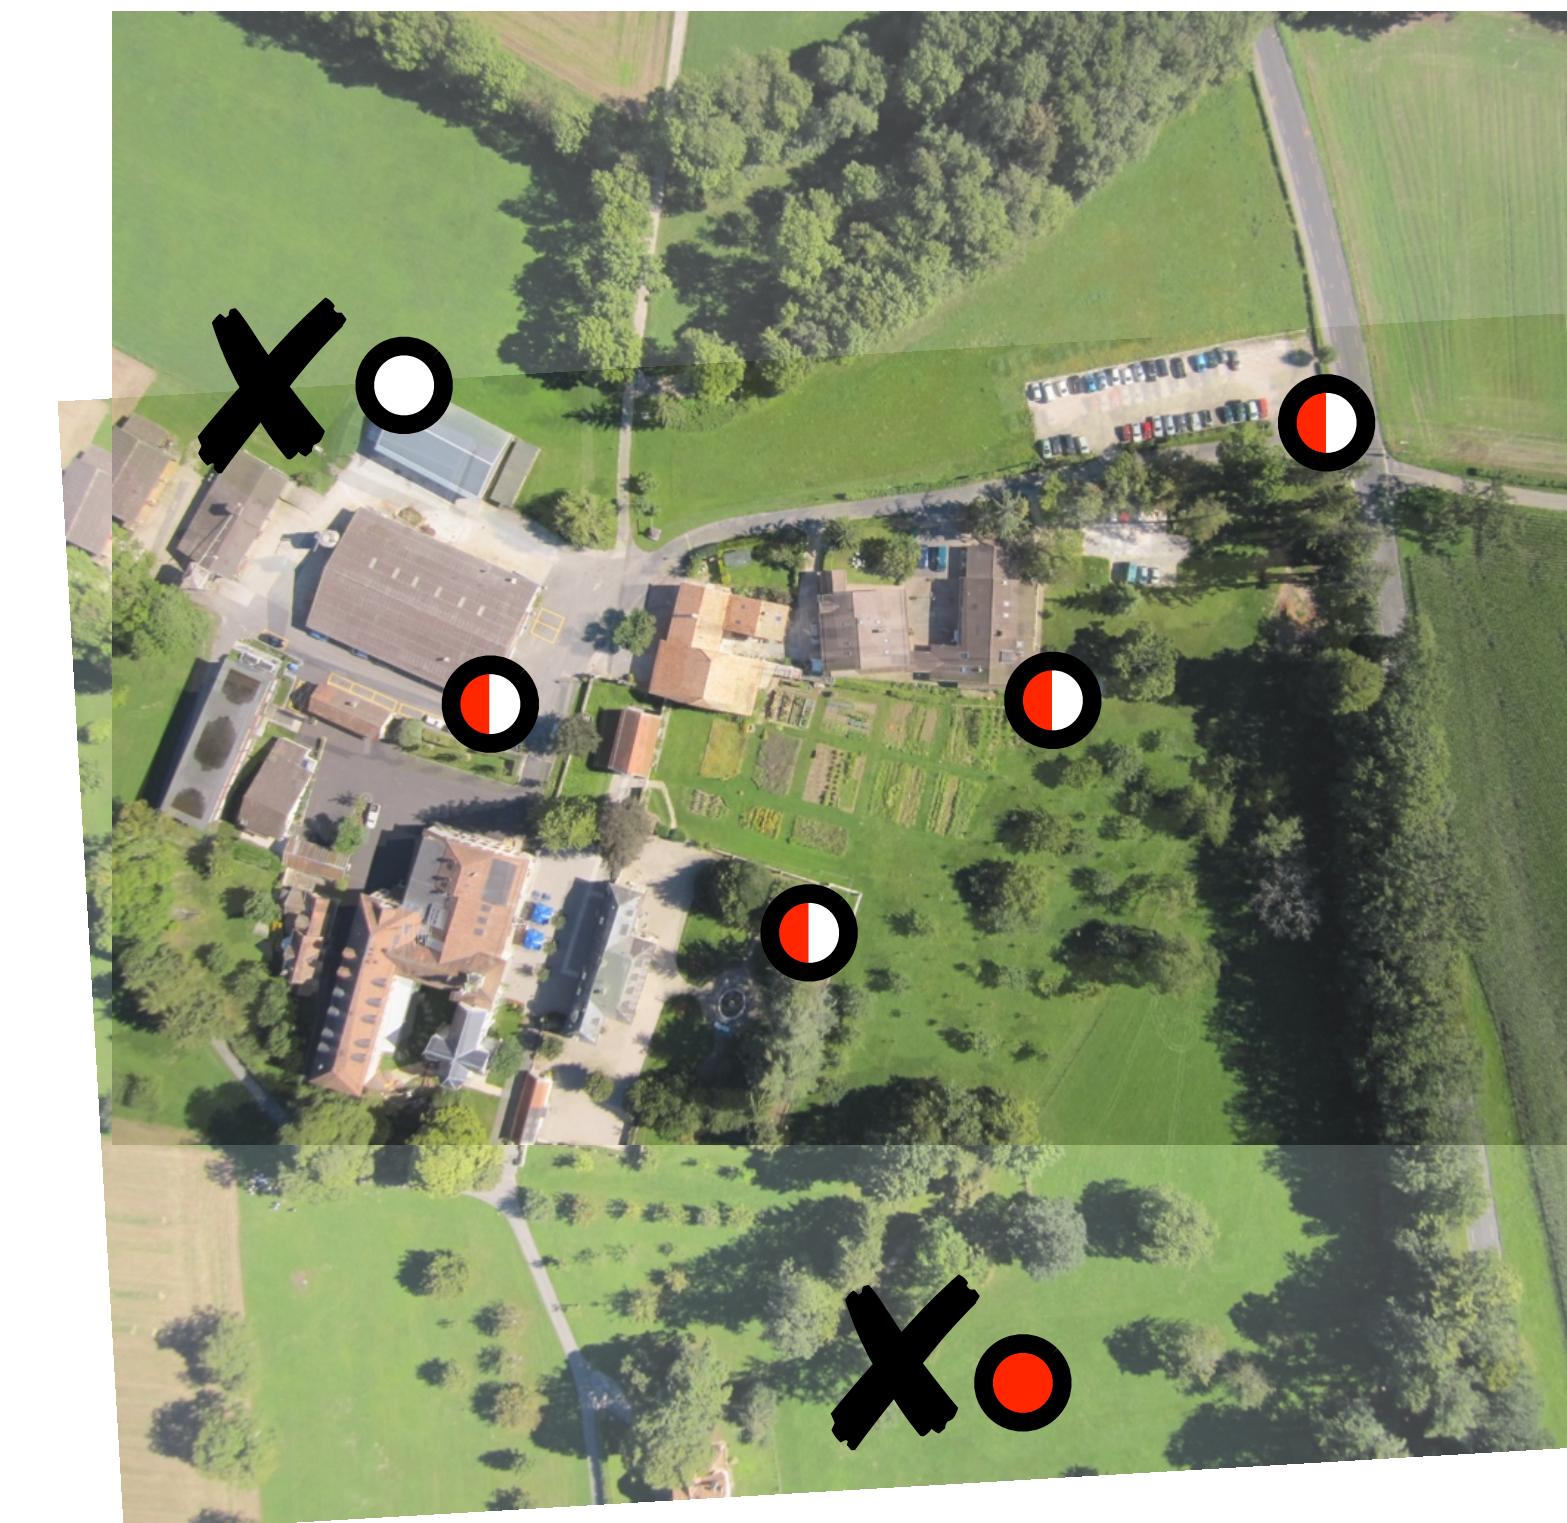


Image Alignment

- We can also use this transformation to reject outliers

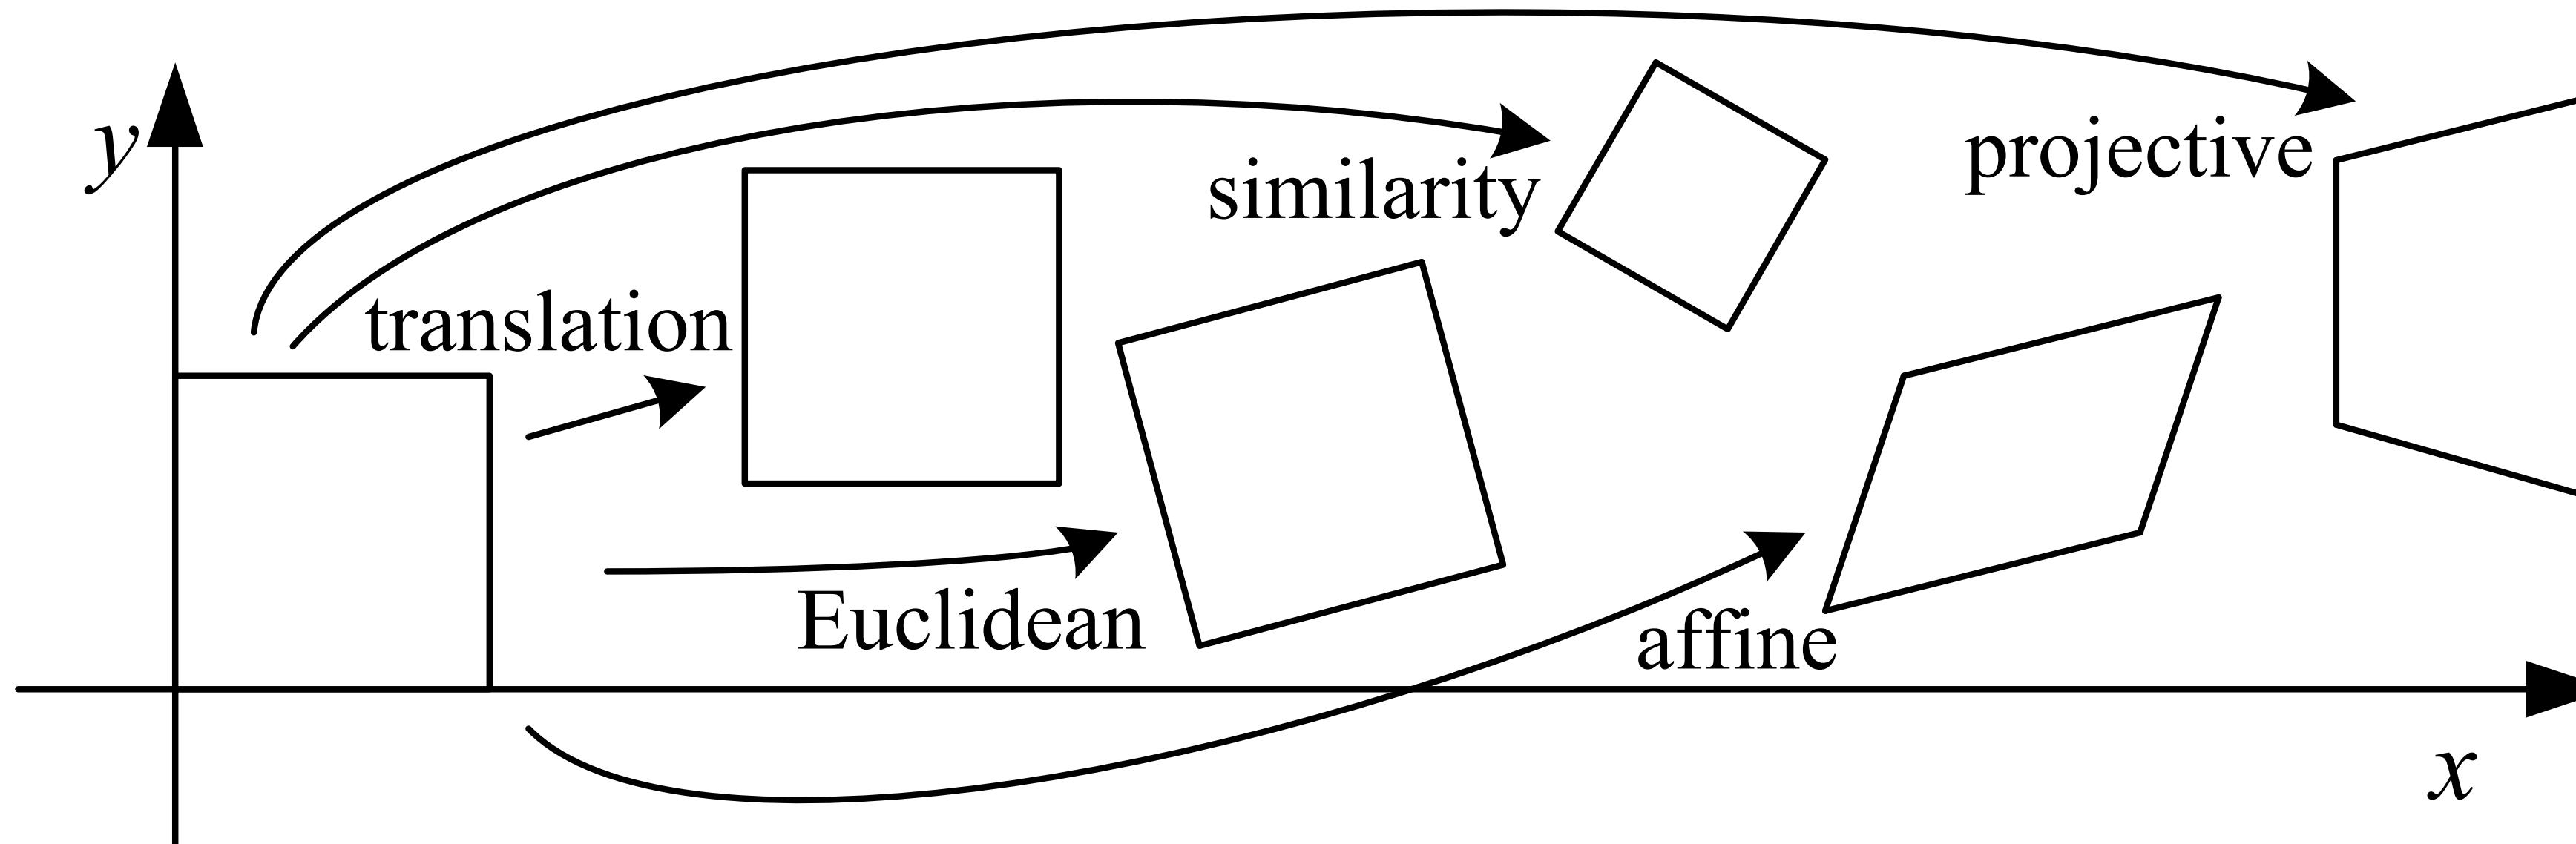


Planar Geometry

- 2D Linear + Projective transformations
 - Euclidean, Similarity, Affine, Homography
- Robust Estimation and RANSAC
 - Estimating 2D transforms with noisy correspondences

2D Transformations

- We will look at a family that can be represented by 3×3 matrices



This group represents perspective projections
of **planar surfaces** in the world

Affine Transformations

- Transformed points are a linear function of the input points

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} a_{13} \\ a_{23} \end{bmatrix}$$

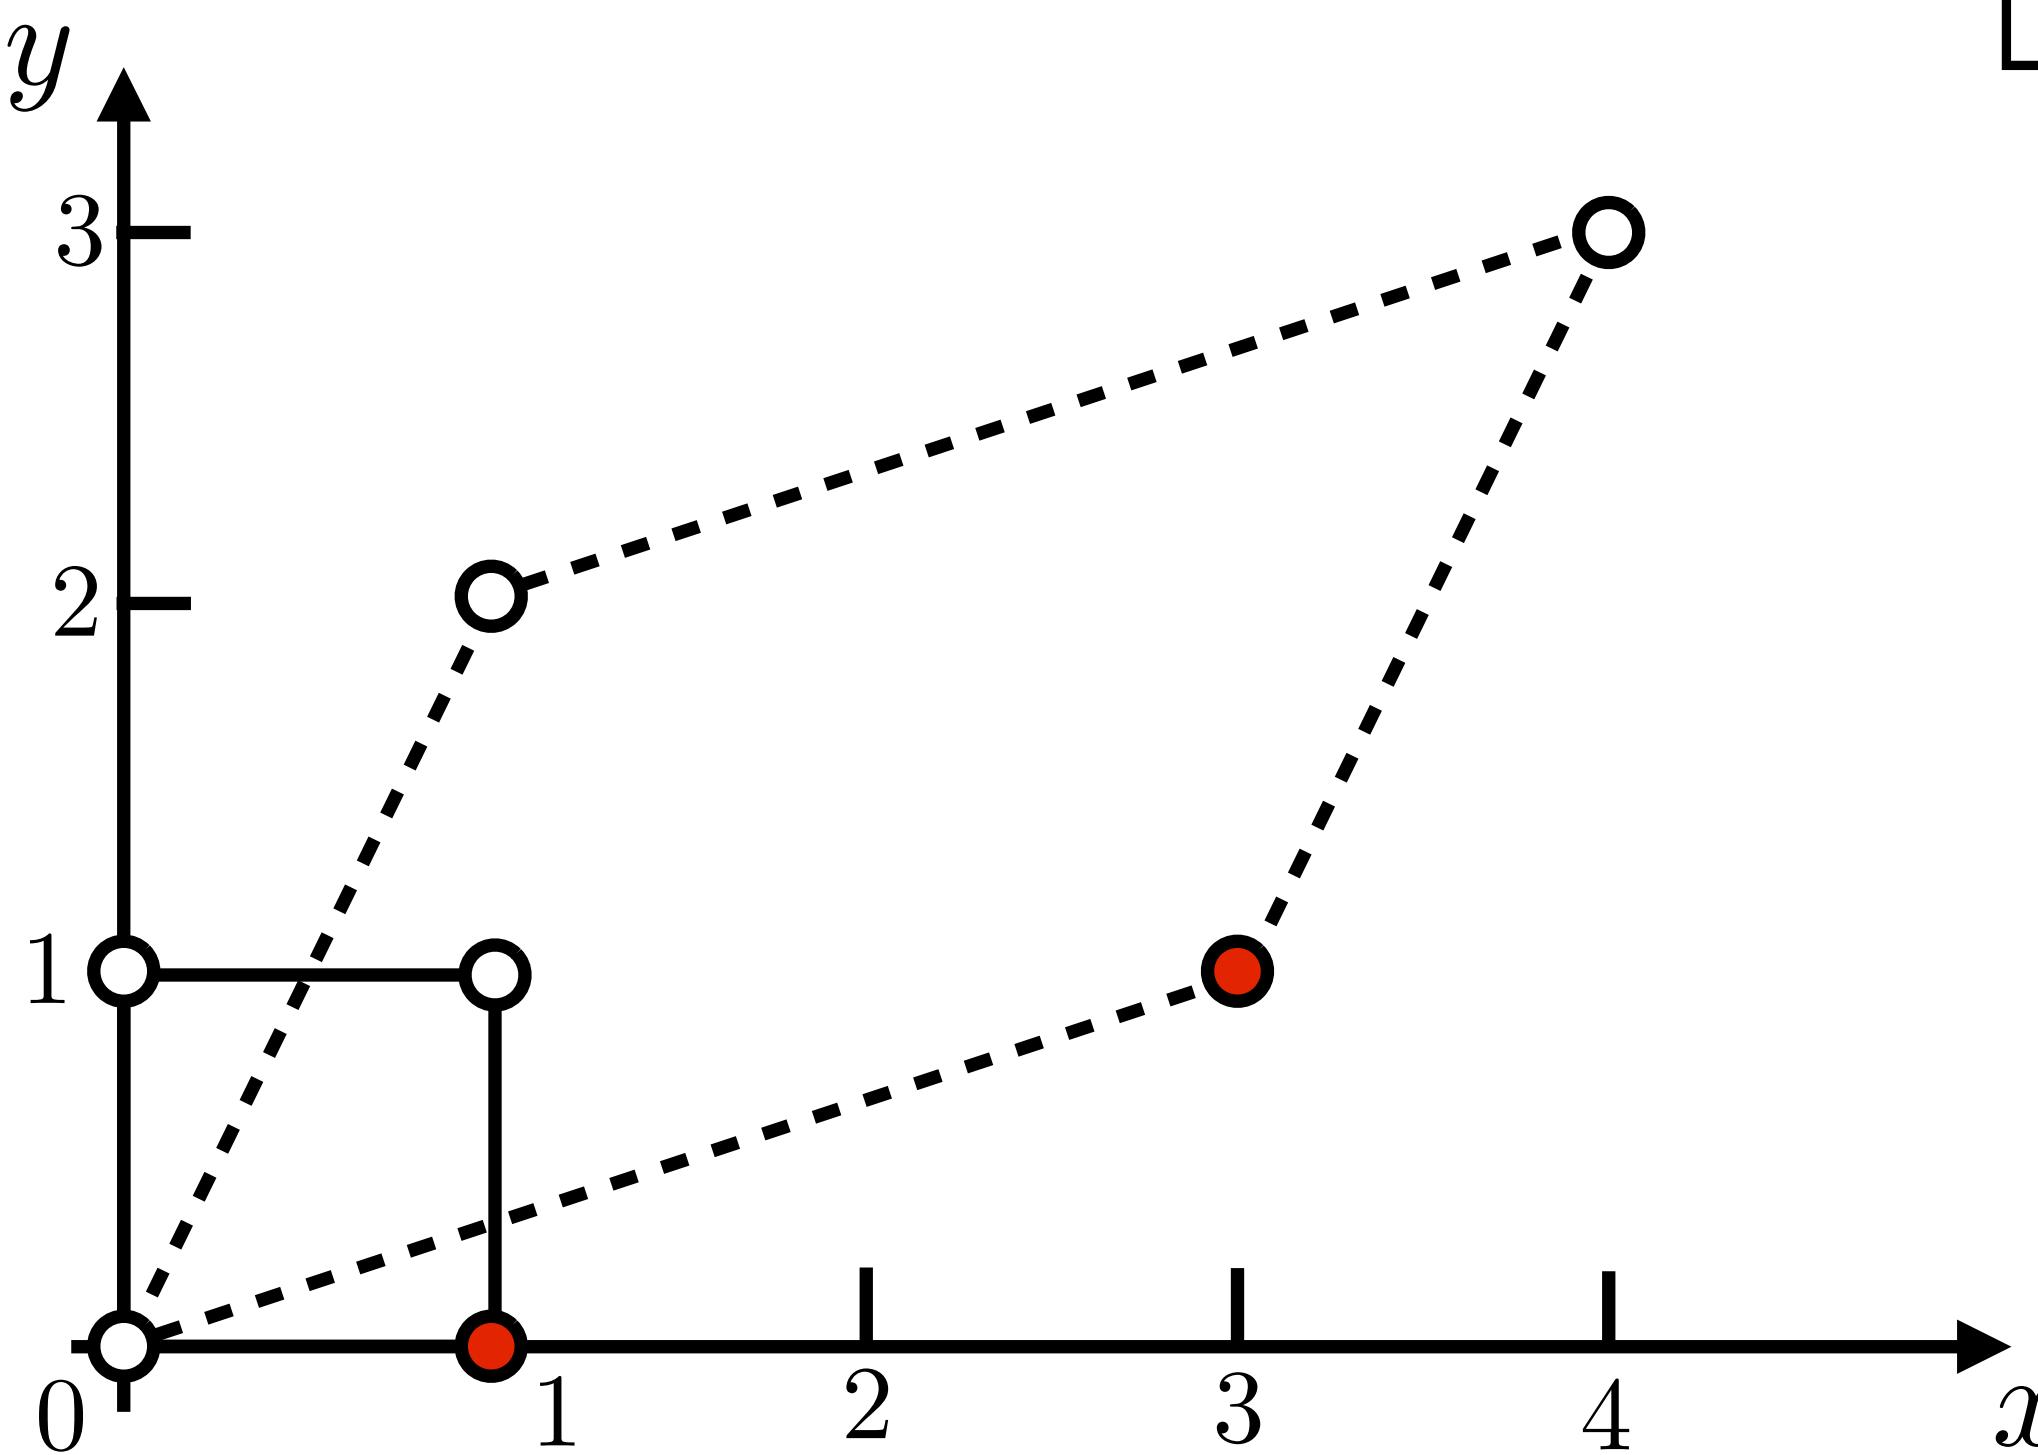
- This can be written as a single matrix multiplication



Linear Transformations

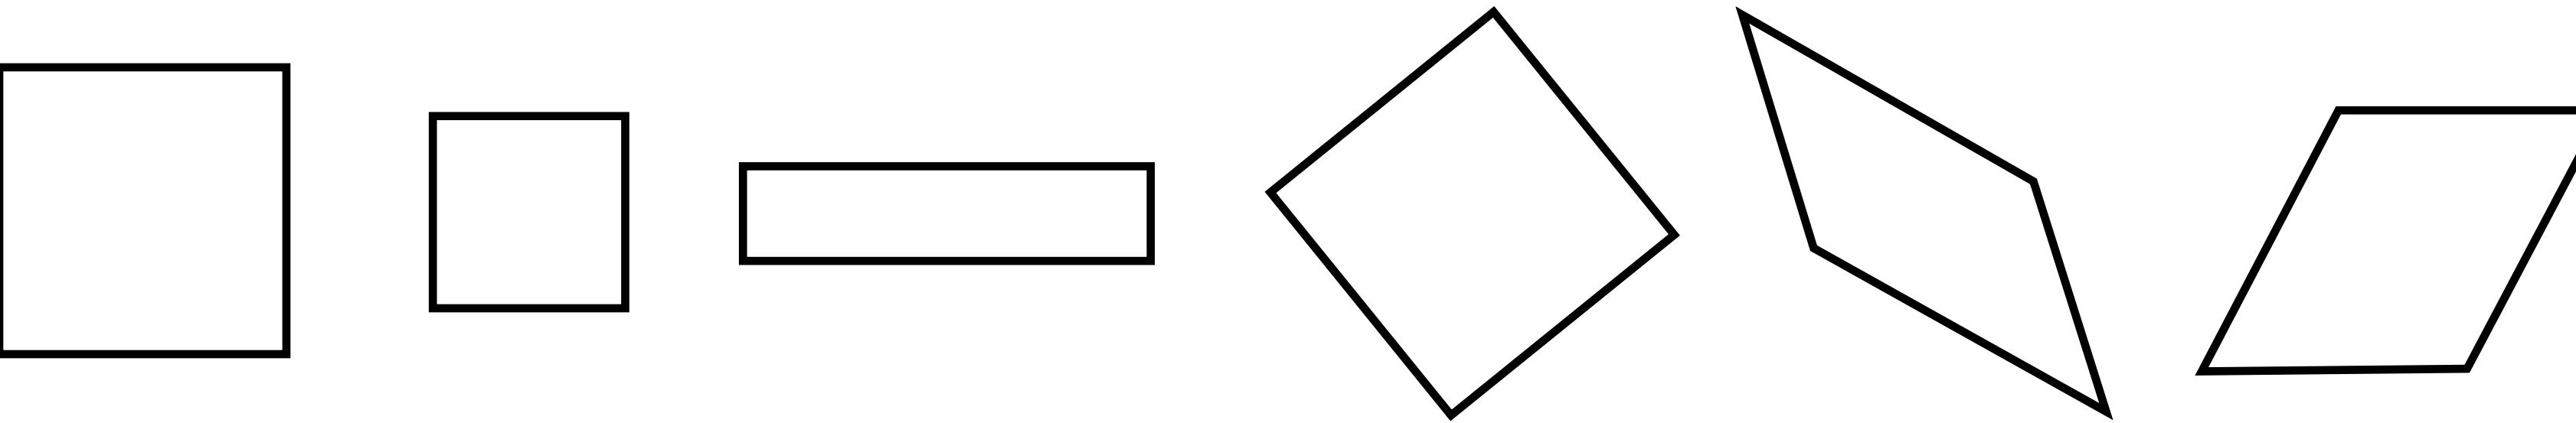
- Consider the action of the unit square under

$$\begin{bmatrix} 3 & 1 & 0 \\ 1 & 2 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

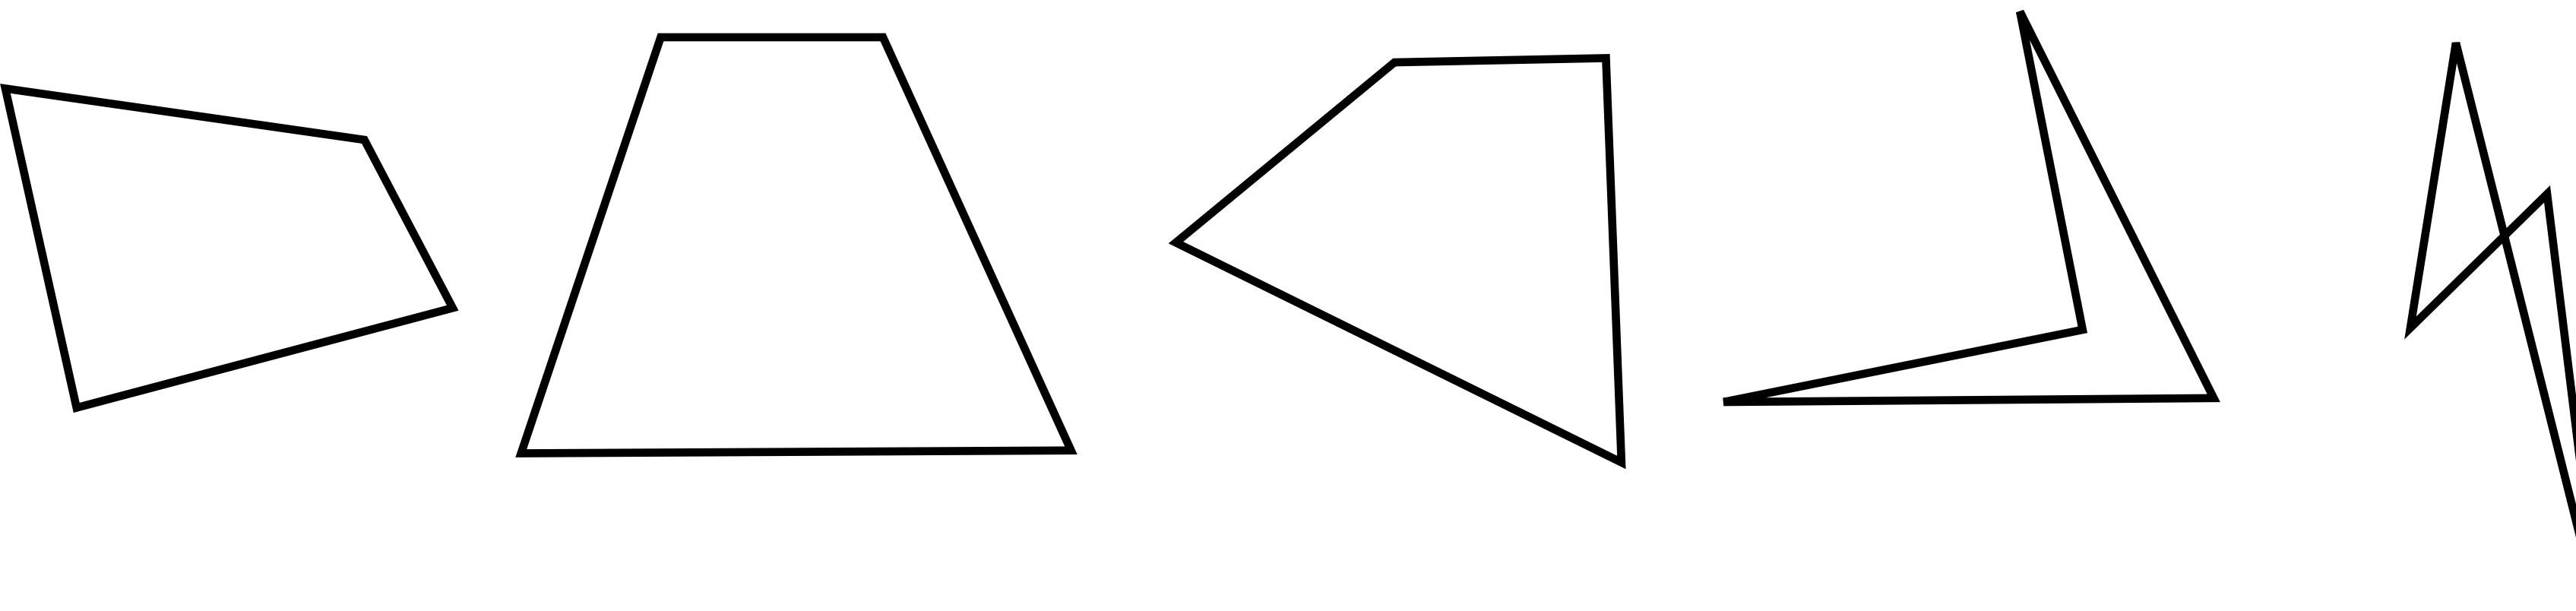


13.2

Linear Transform Examples



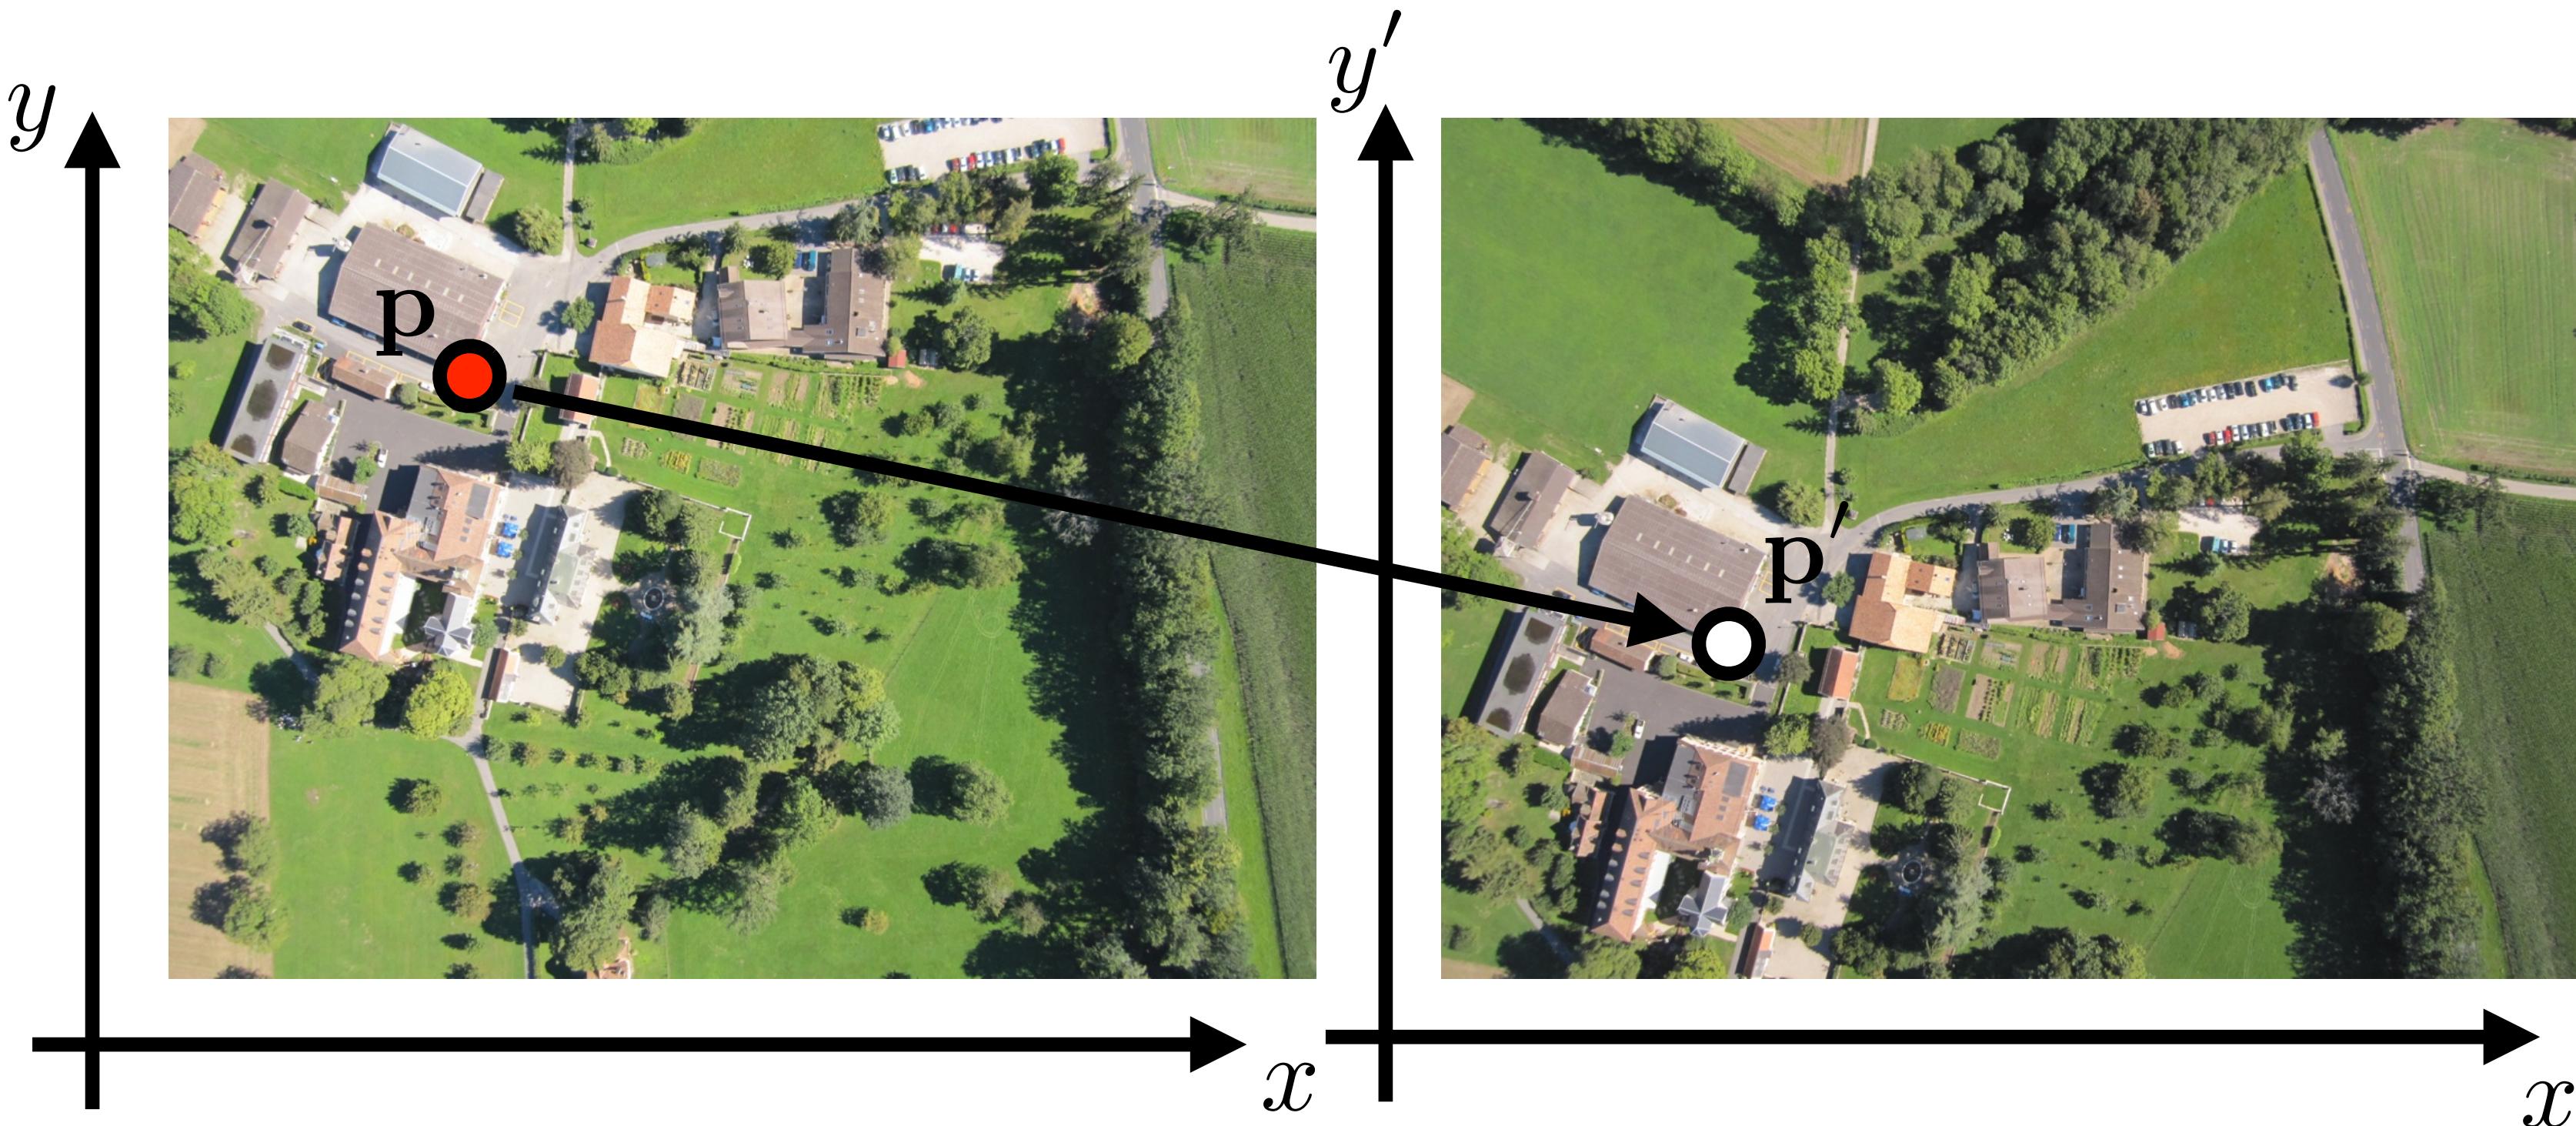
Translation, rotation, scale, shear (parallel lines preserved)



These transforms are not affine (parallel lines not preserved)

Linear Transformations

- Consider a single point correspondence



$$\begin{bmatrix} x'_1 \\ y'_1 \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$



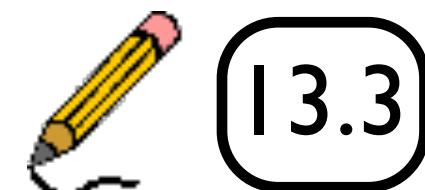
How many points are needed to solve for **a**?

Computing Affine Transforms

- Lets compute an affine transform from correspondences:

$$\begin{bmatrix} x'_1 \\ y'_1 \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

- Re-arrange unknowns into a vector



Computing Affine Transforms

- Linear system in the unknown parameters \mathbf{a}

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_1 & y_1 & 1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_2 & y_2 & 1 \\ x_3 & y_3 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_3 & y_3 & 1 \end{bmatrix} \begin{bmatrix} a_{11} \\ a_{12} \\ a_{13} \\ a_{21} \\ a_{22} \\ a_{23} \end{bmatrix} = \begin{bmatrix} x'_1 \\ y'_1 \\ x'_2 \\ y'_2 \\ x'_3 \\ y'_3 \end{bmatrix}$$

- Of the form

$$\mathbf{M}\mathbf{a} = \mathbf{y}$$

Solve for \mathbf{a} using Gaussian Elimination

Computing Affine Transforms

- We can now map any other points between the two images



$$\begin{bmatrix} x'_1 \\ y'_1 \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Computing Affine Transforms

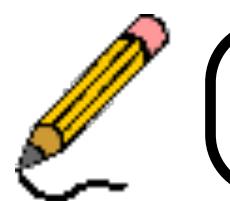
- Or resample one image in the coordinate system of the other

This allows us to “stitch”
the two images



Linear Transformations

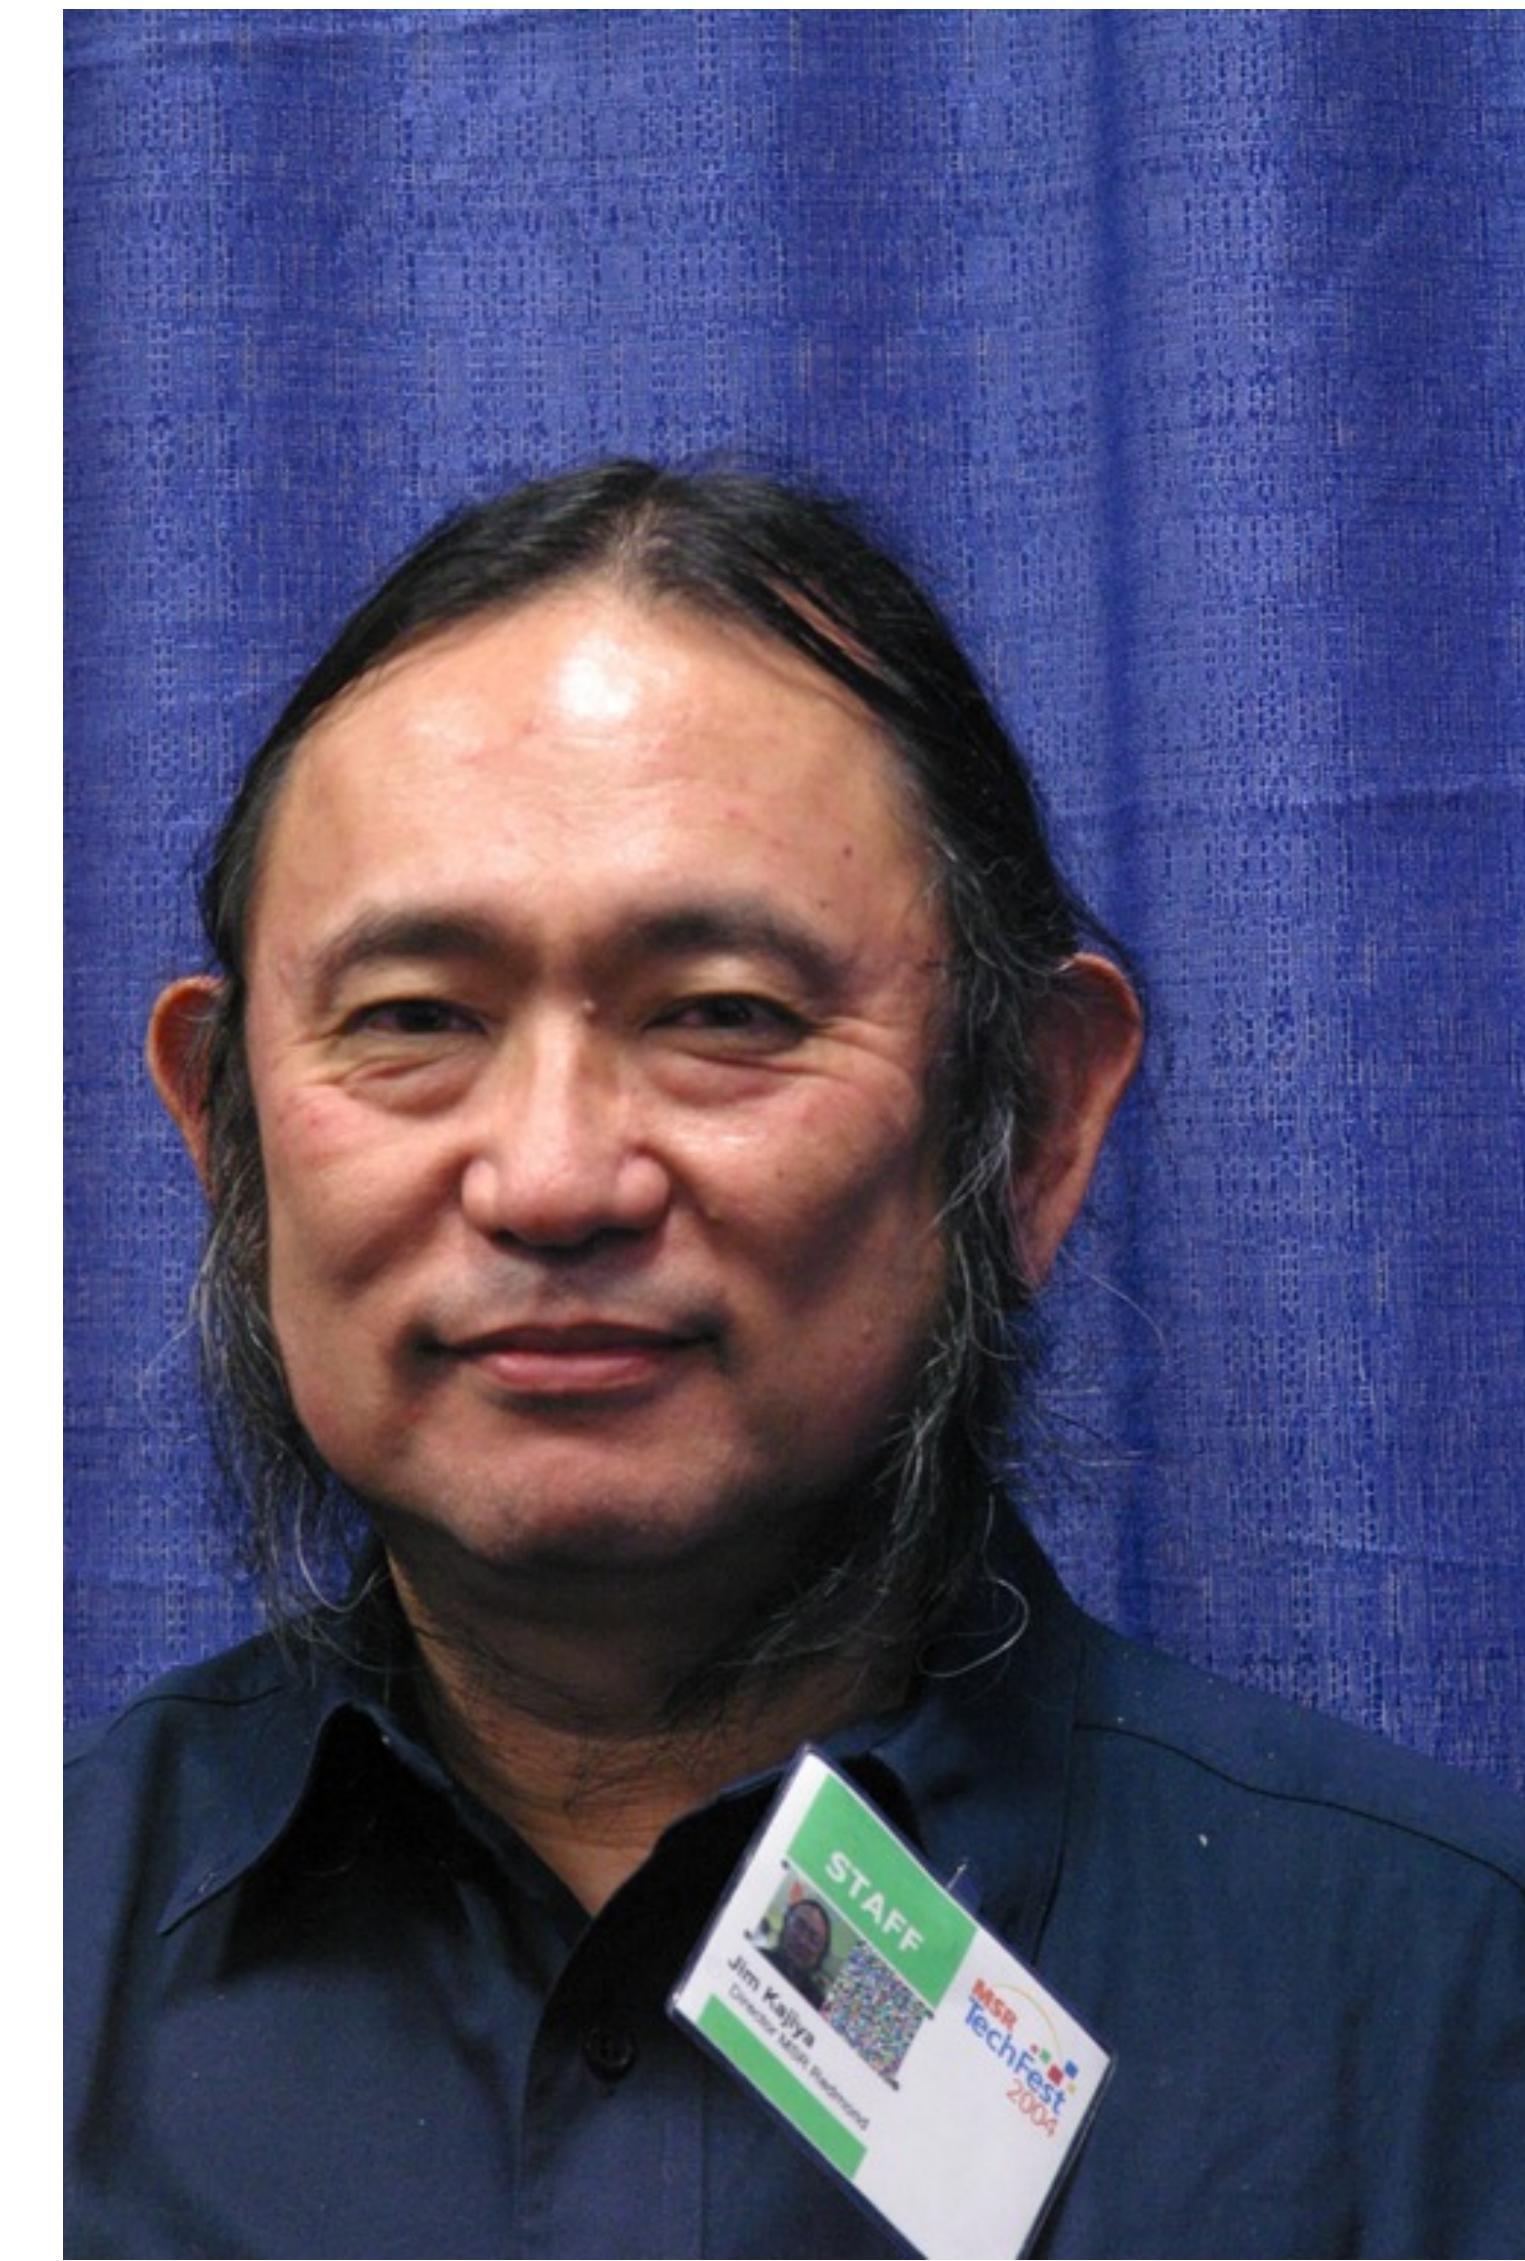
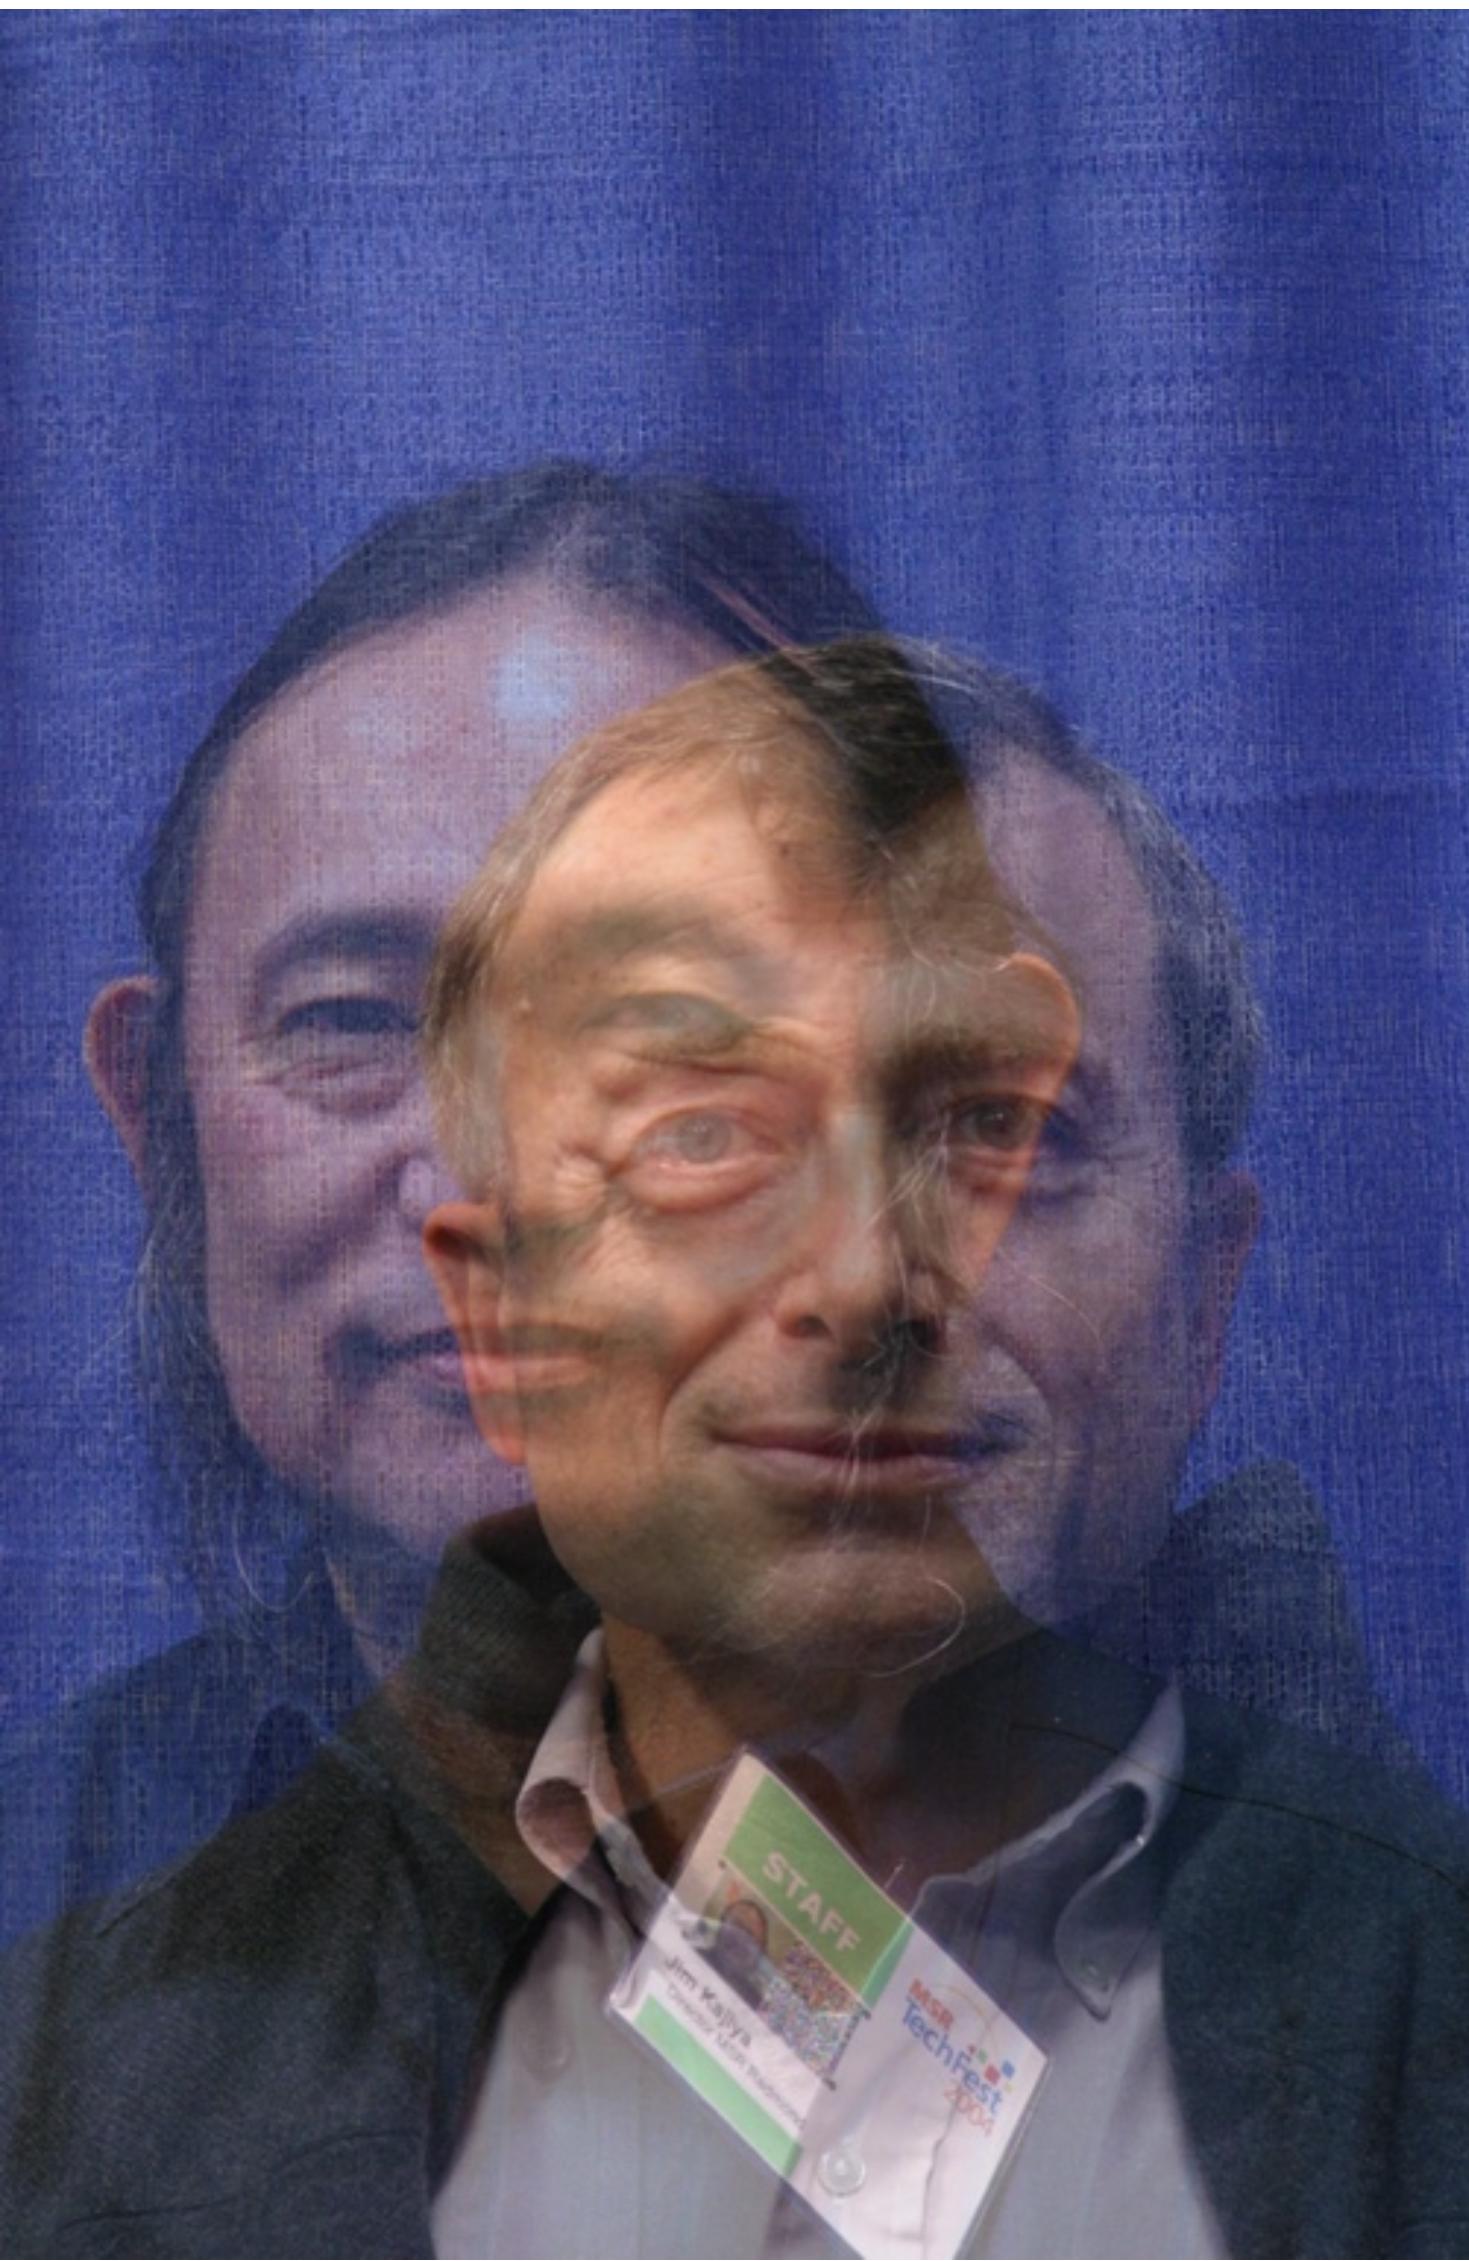
- Other linear transforms are special cases of affine



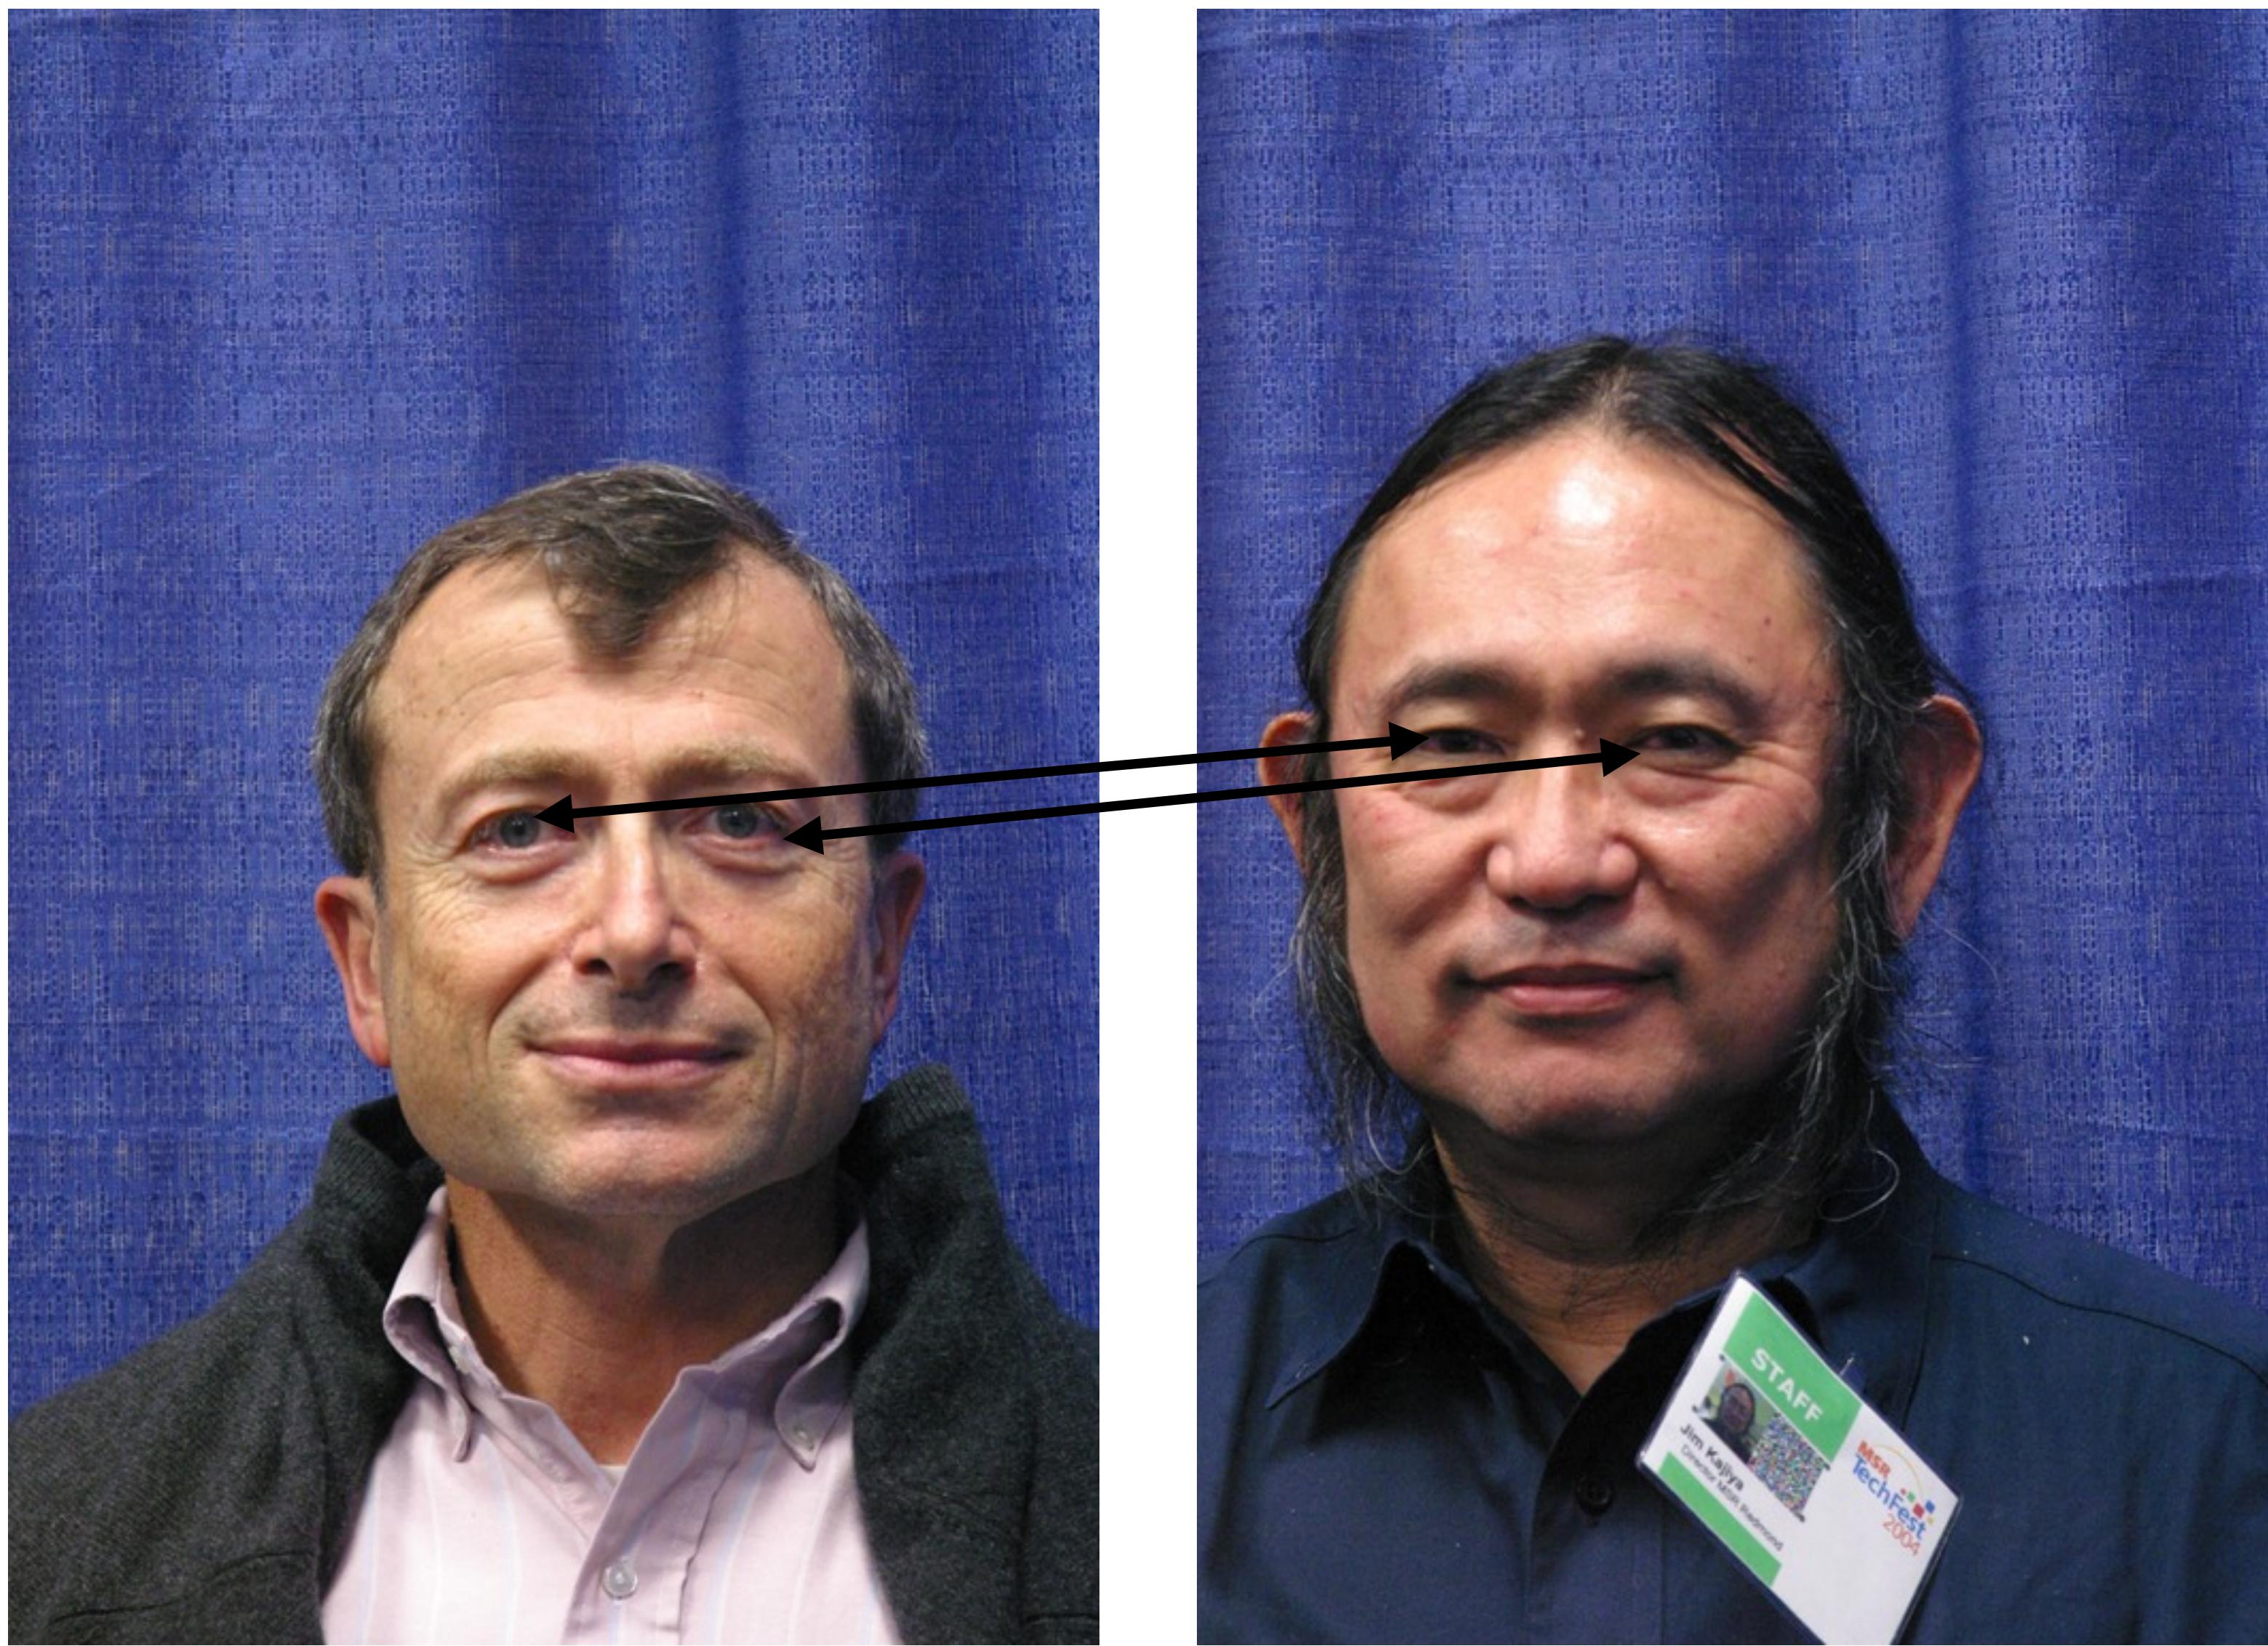
13.4

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ 0 & 0 & 1 \end{bmatrix}$$

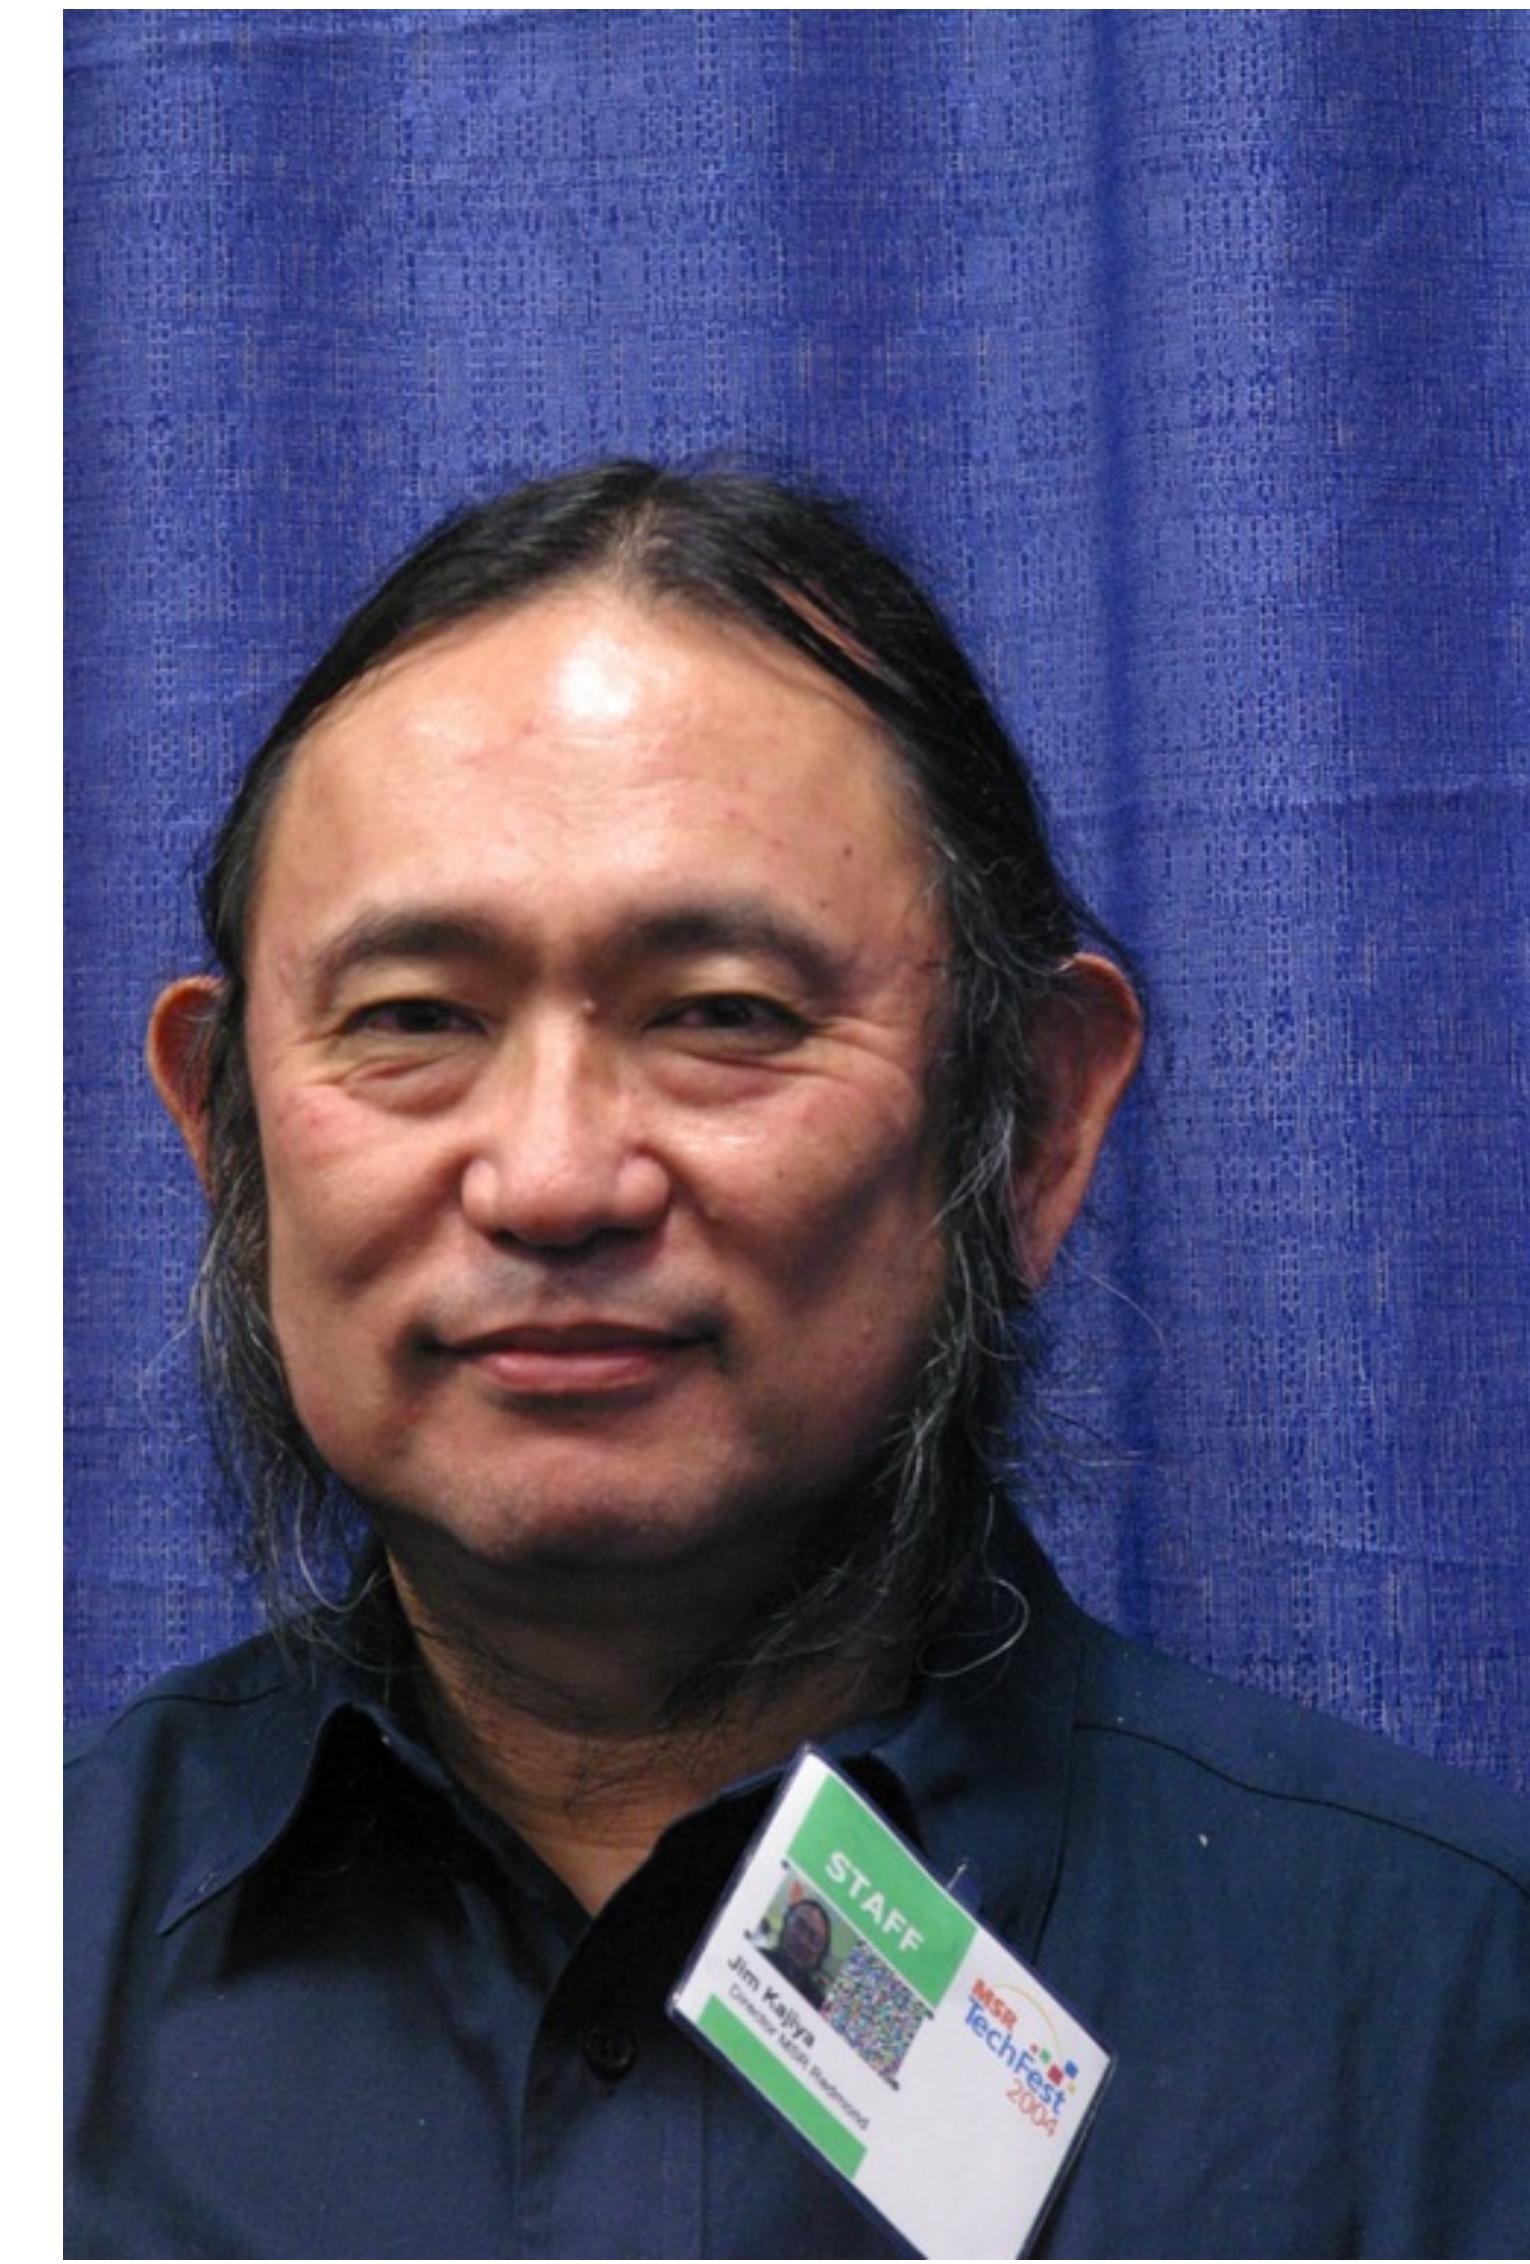
Face Alignment



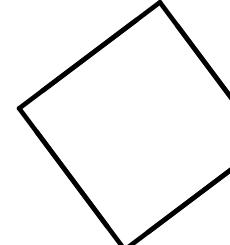
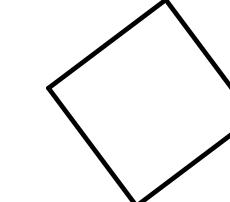
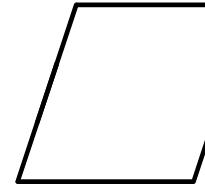
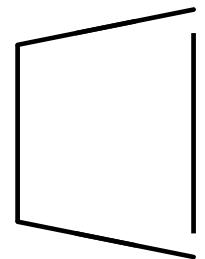
Face Alignment



Face Alignment

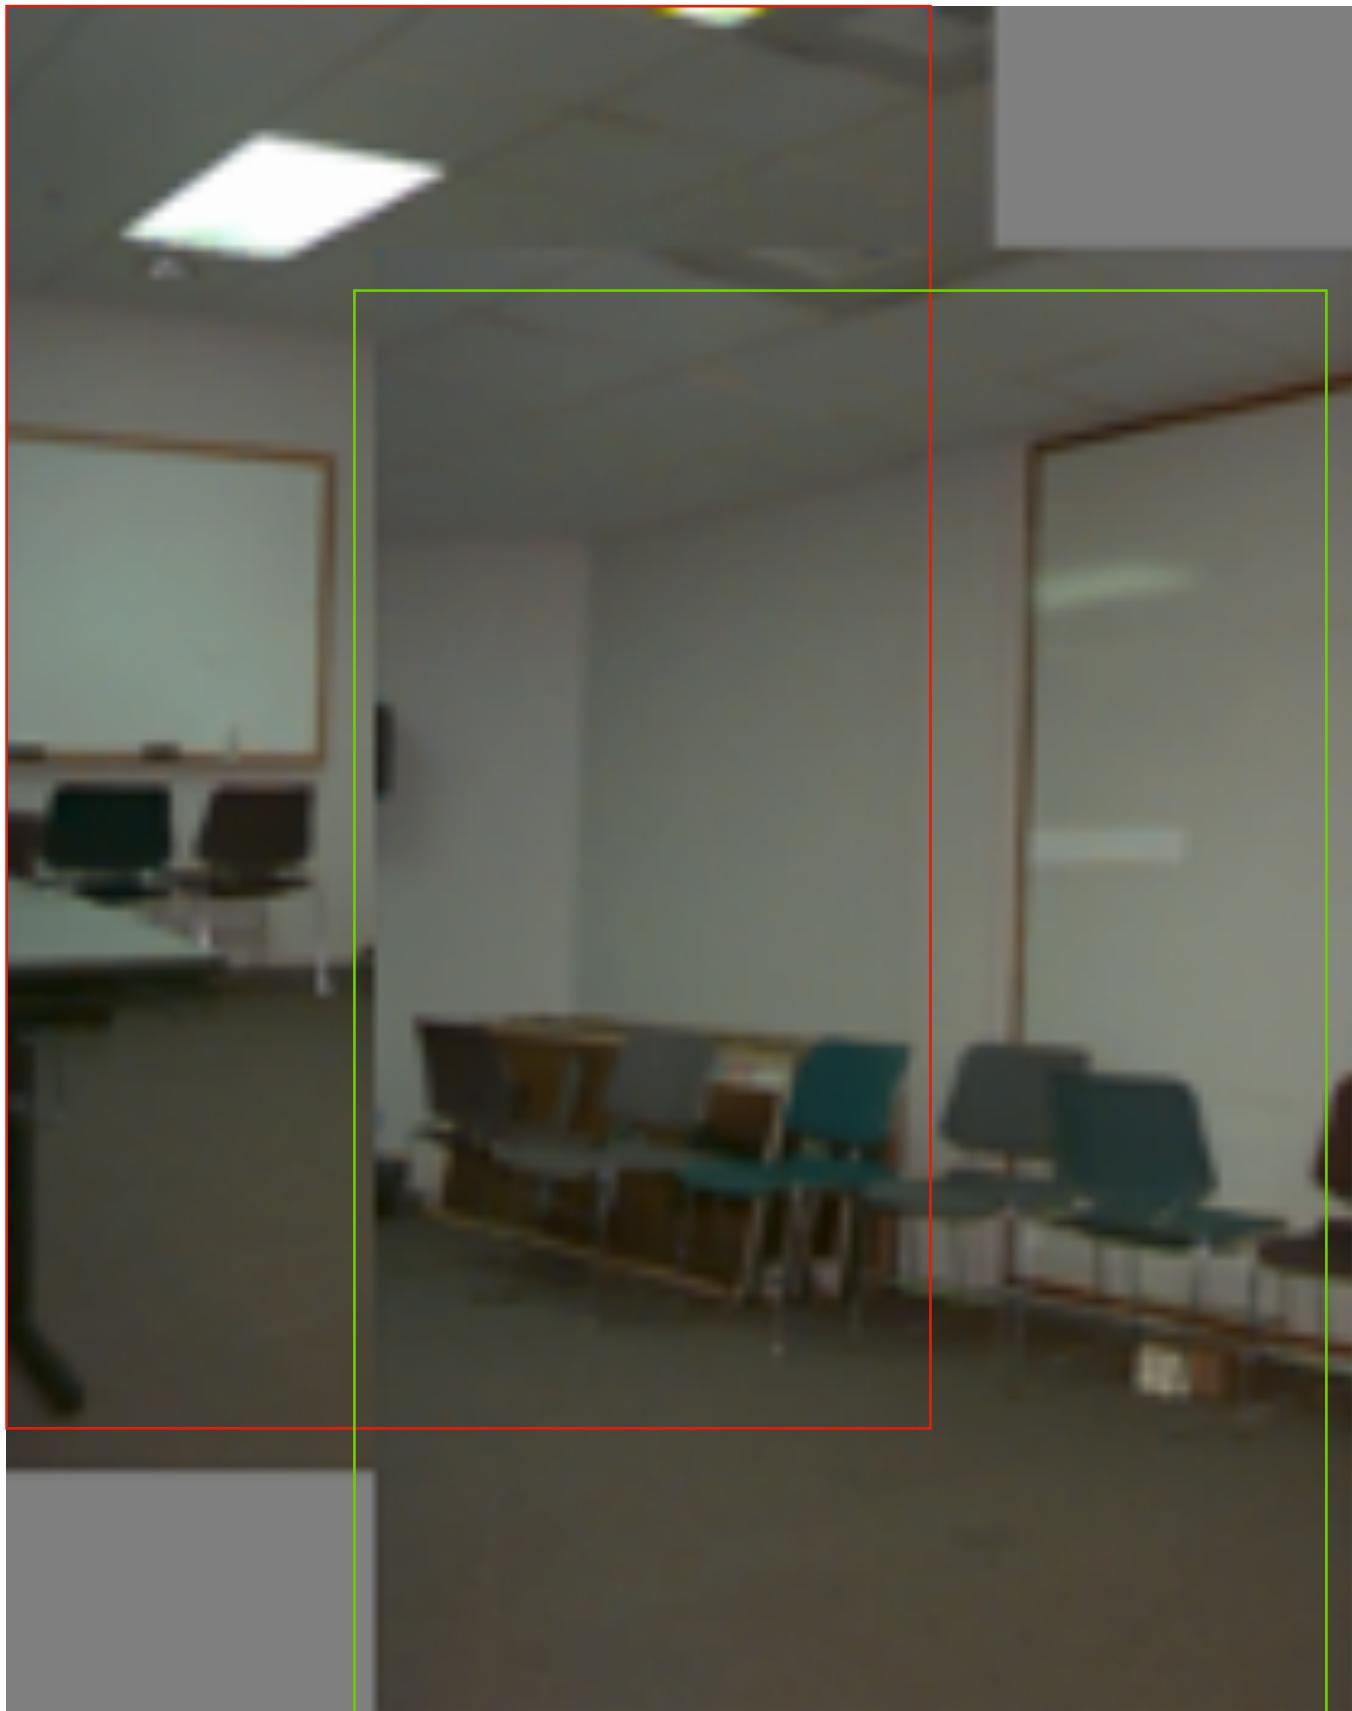


2D Transformations

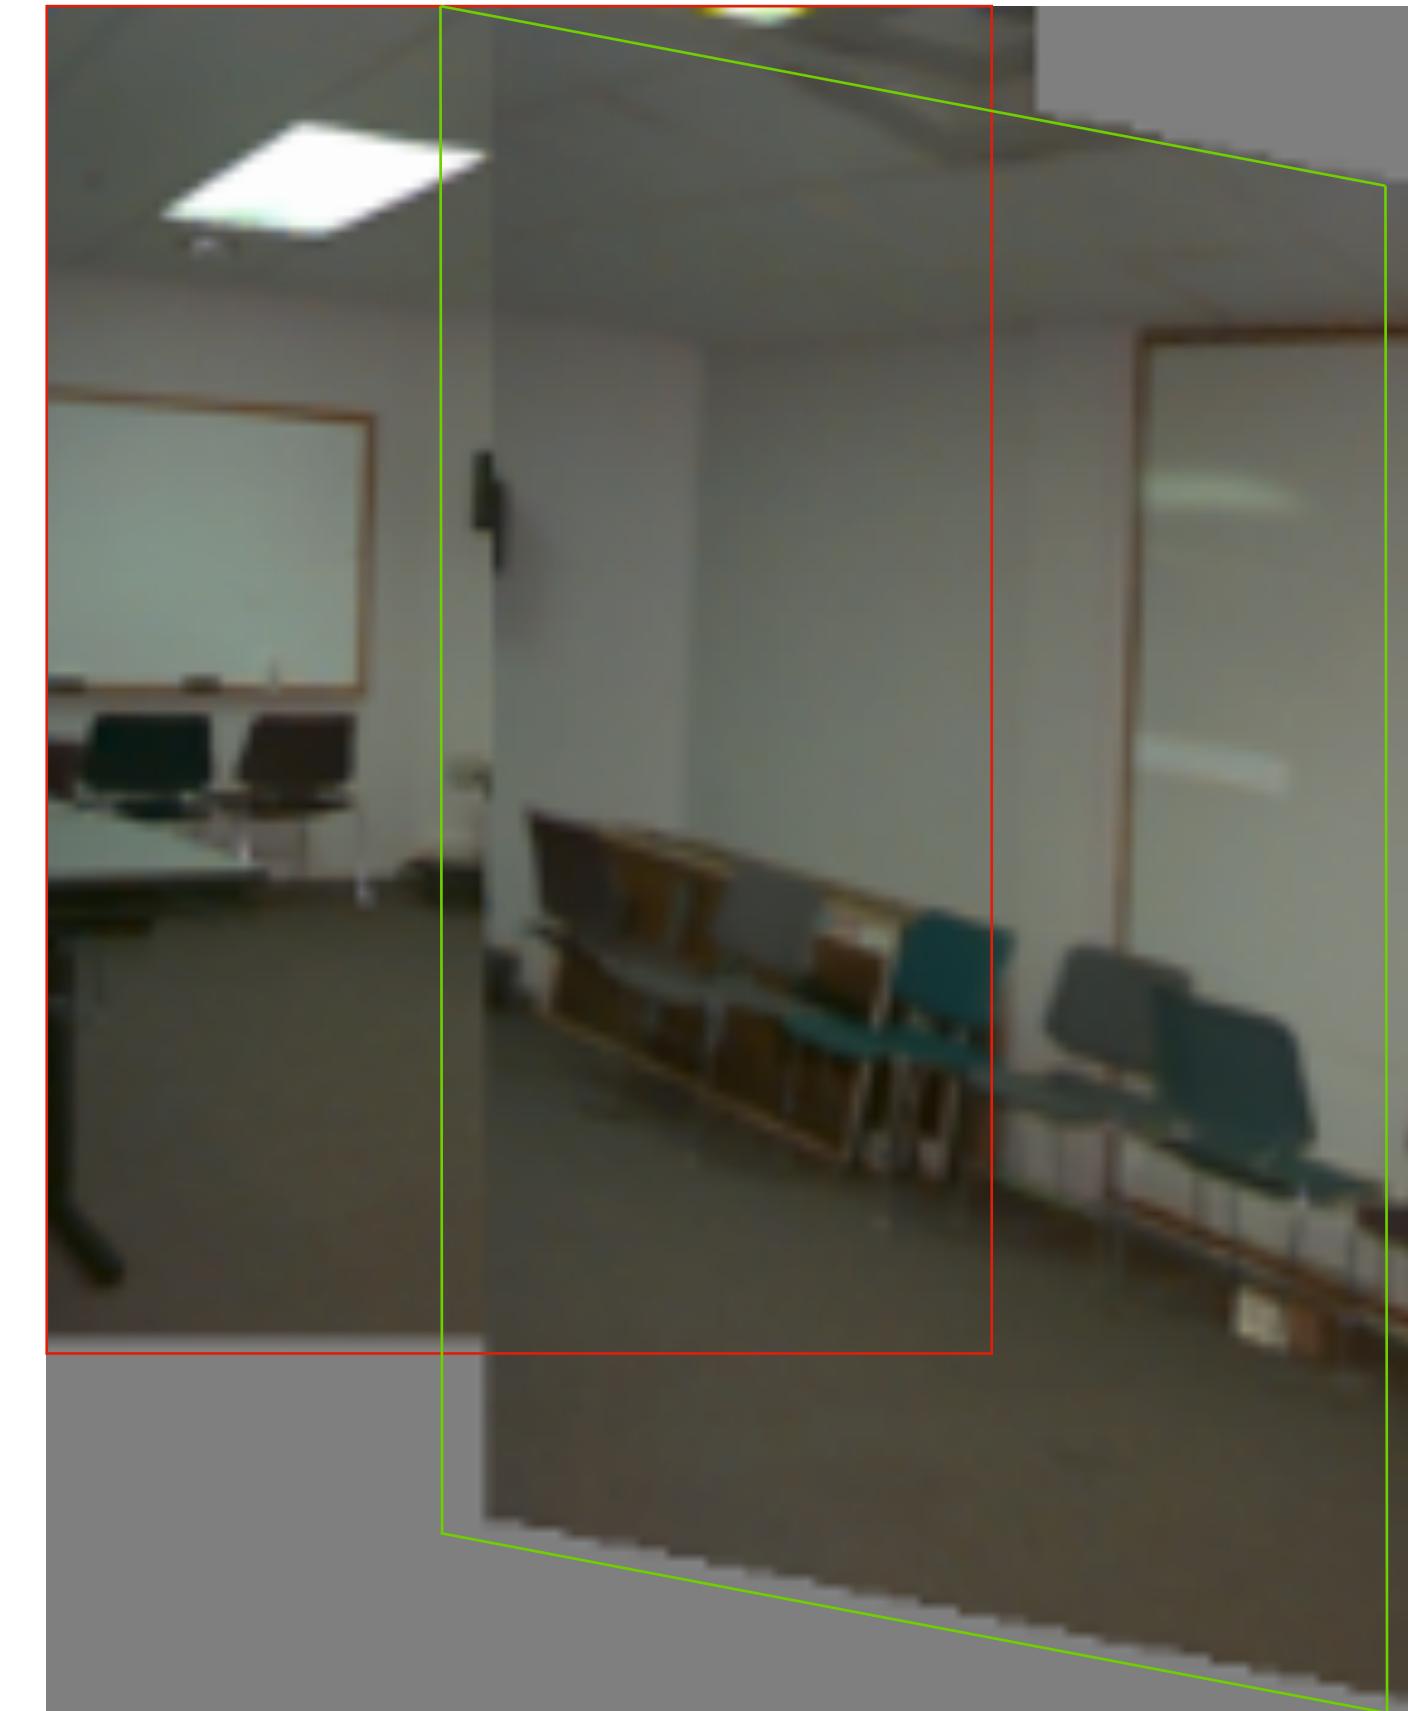
Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

Example: Warping with Different Transformations

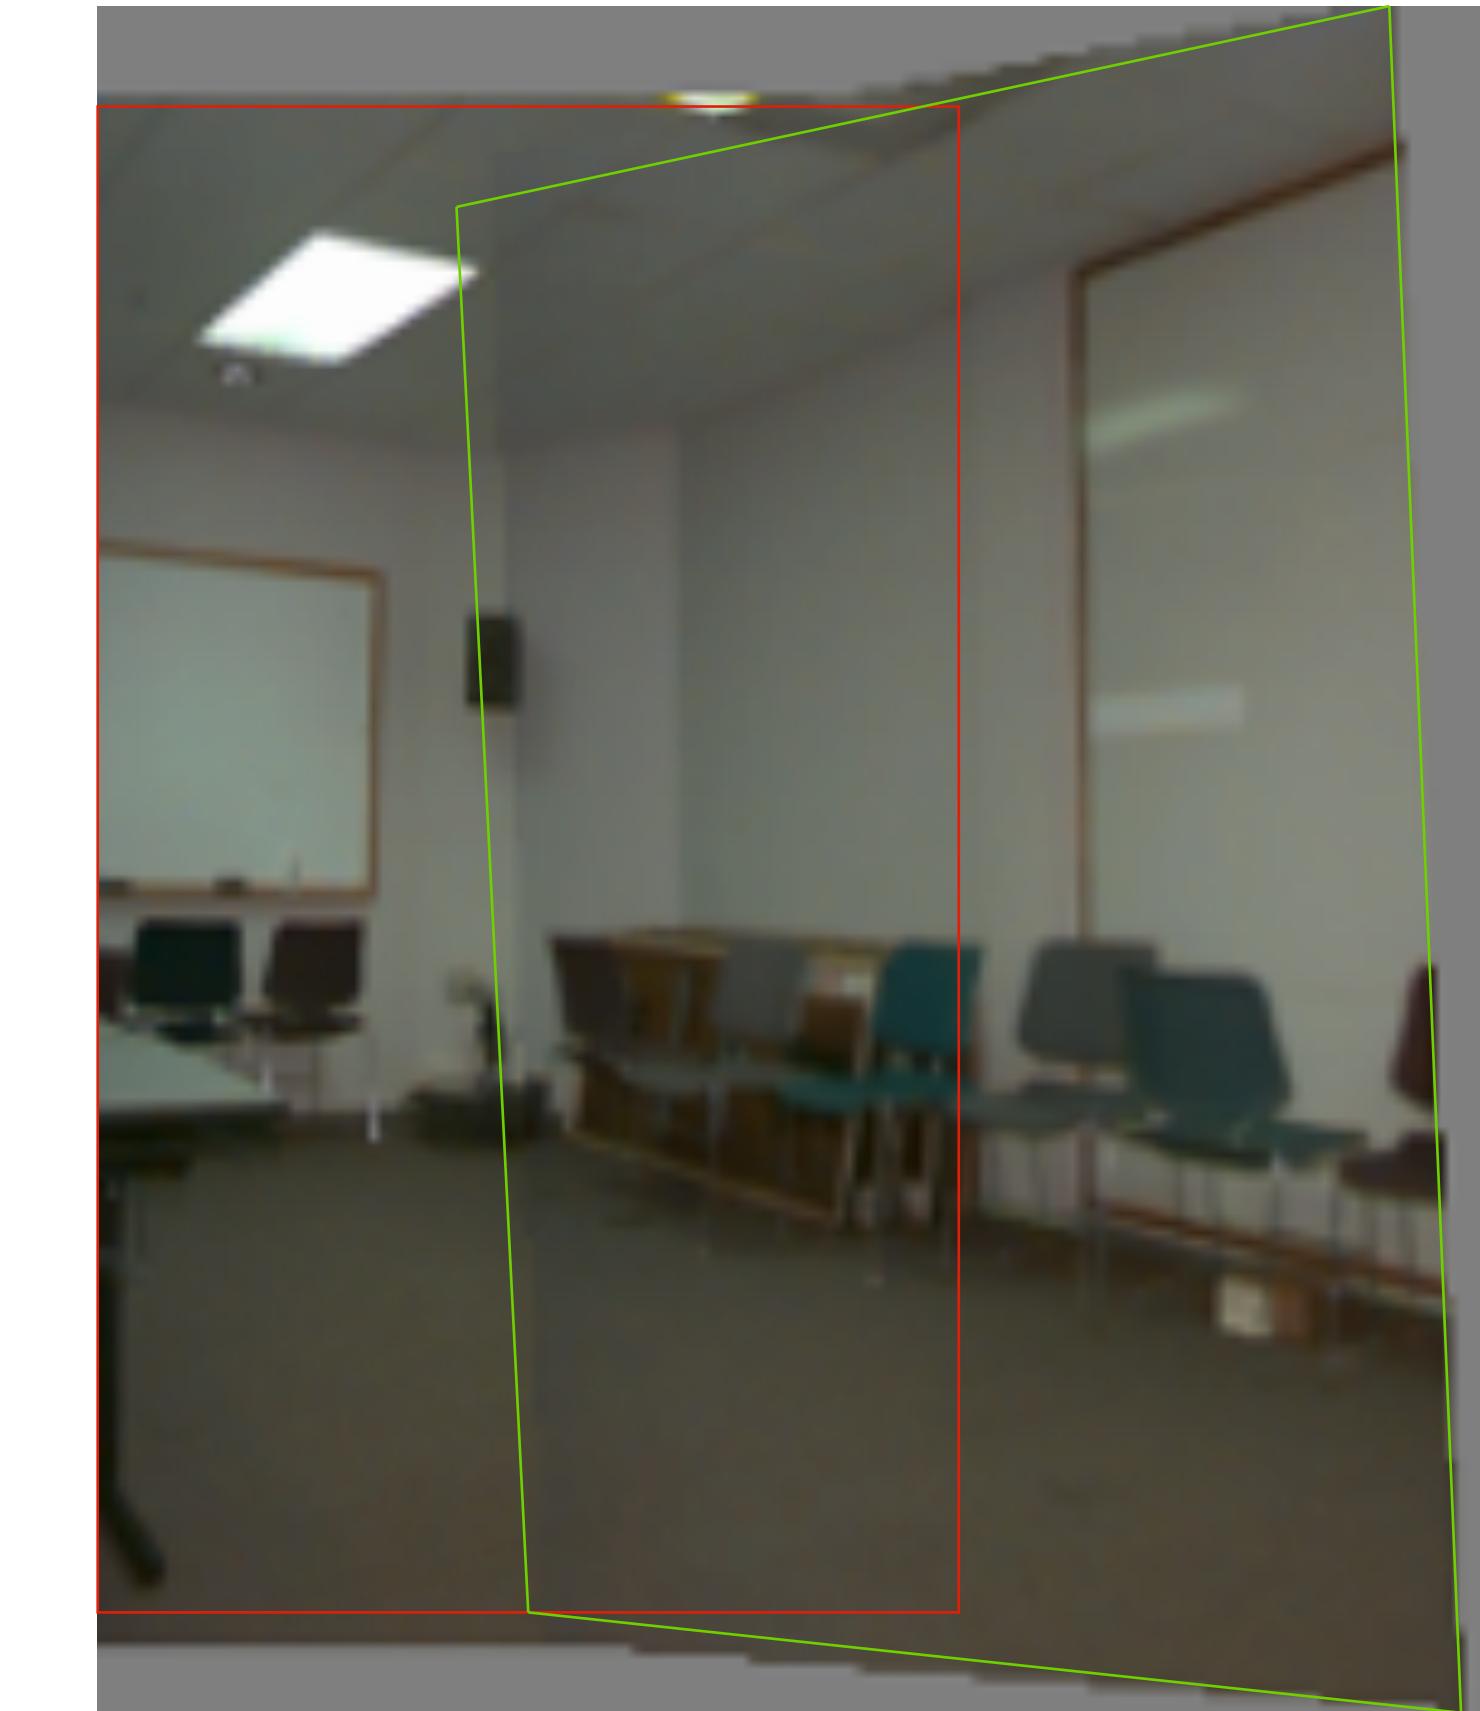
Translation



Affine



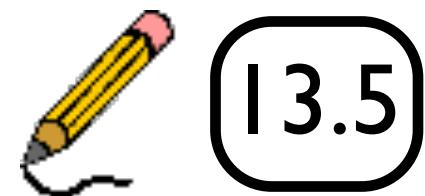
Projective
(homography)



Projective Transformation

- General 3x3 matrix transformation (note need scale factor)

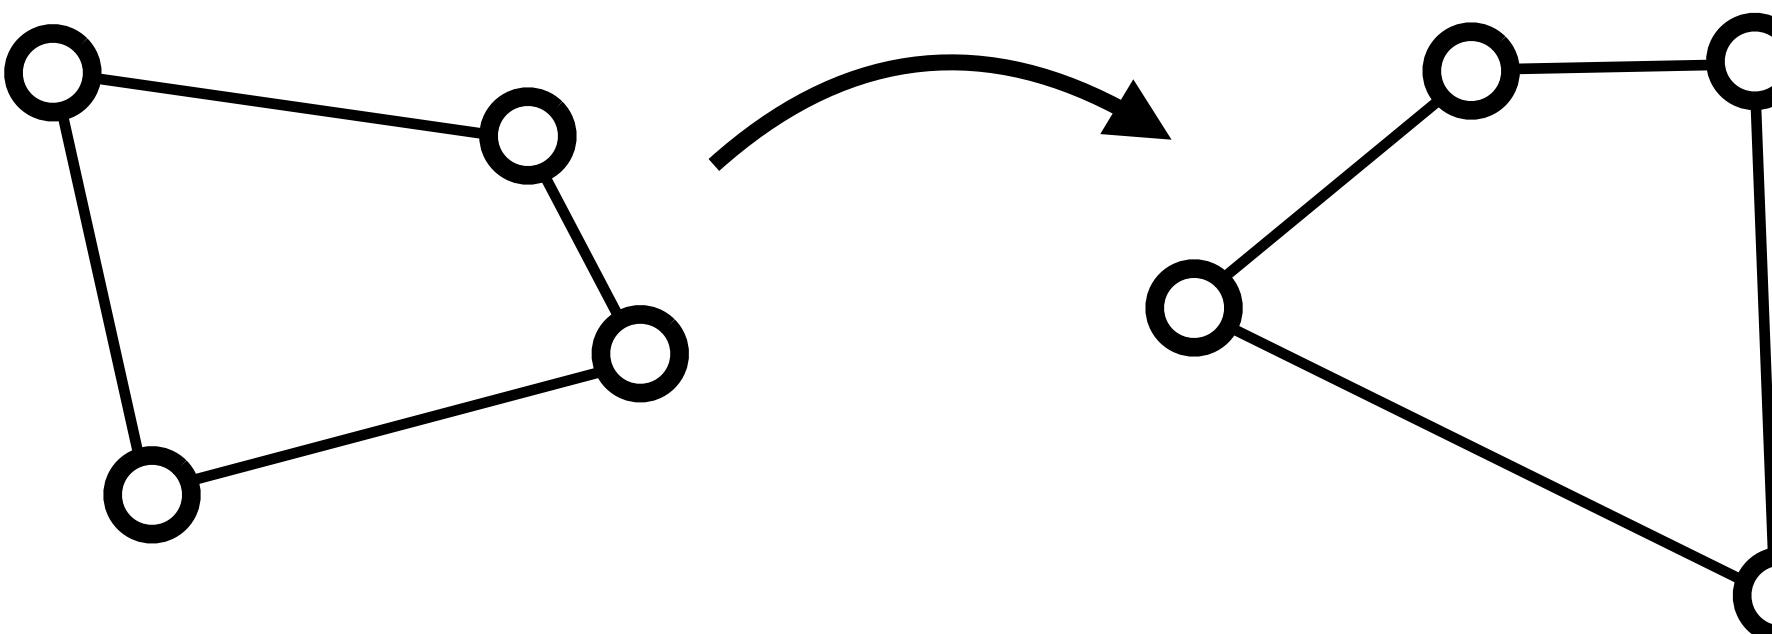
$$s \begin{bmatrix} x'_1 \\ y'_1 \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$



Compute H from Correspondences

- Each match gives 2 equations to solve for 8 parameters (arbitrary scale)

$$s \begin{bmatrix} x'_1 \\ y'_1 \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$



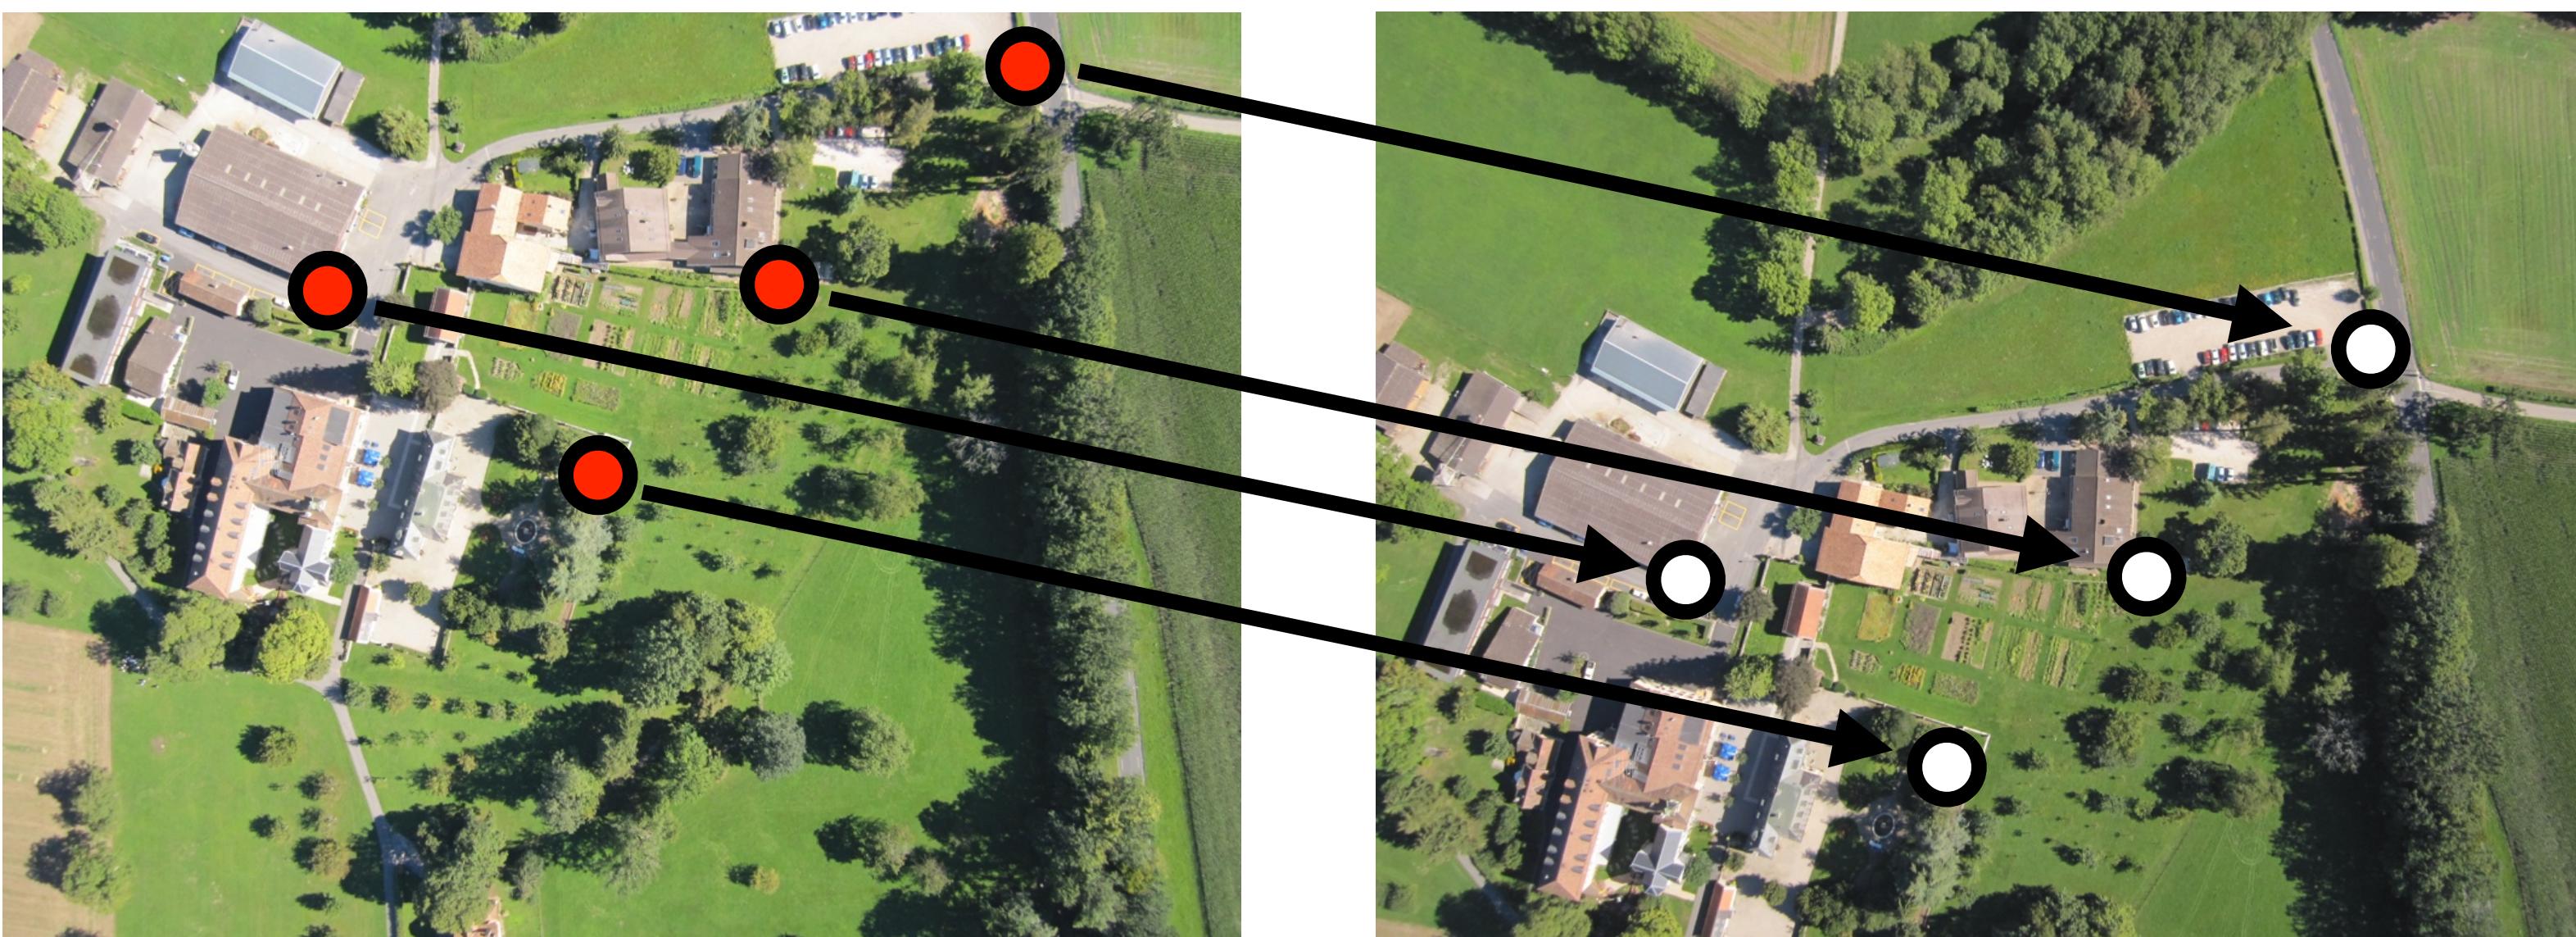
- → 4 correspondences to solve for H matrix
- Solution uses Singular Value Decomposition (SVD)
- In Assignment 4 you can compute this using cv2.findHomography

2-view Alignment + RANSAC

- 2-view alignment: linear equations
- Least squares and outliers
- Robust estimation via sampling

Image Alignment

- Find corresponding (matching) points between the images

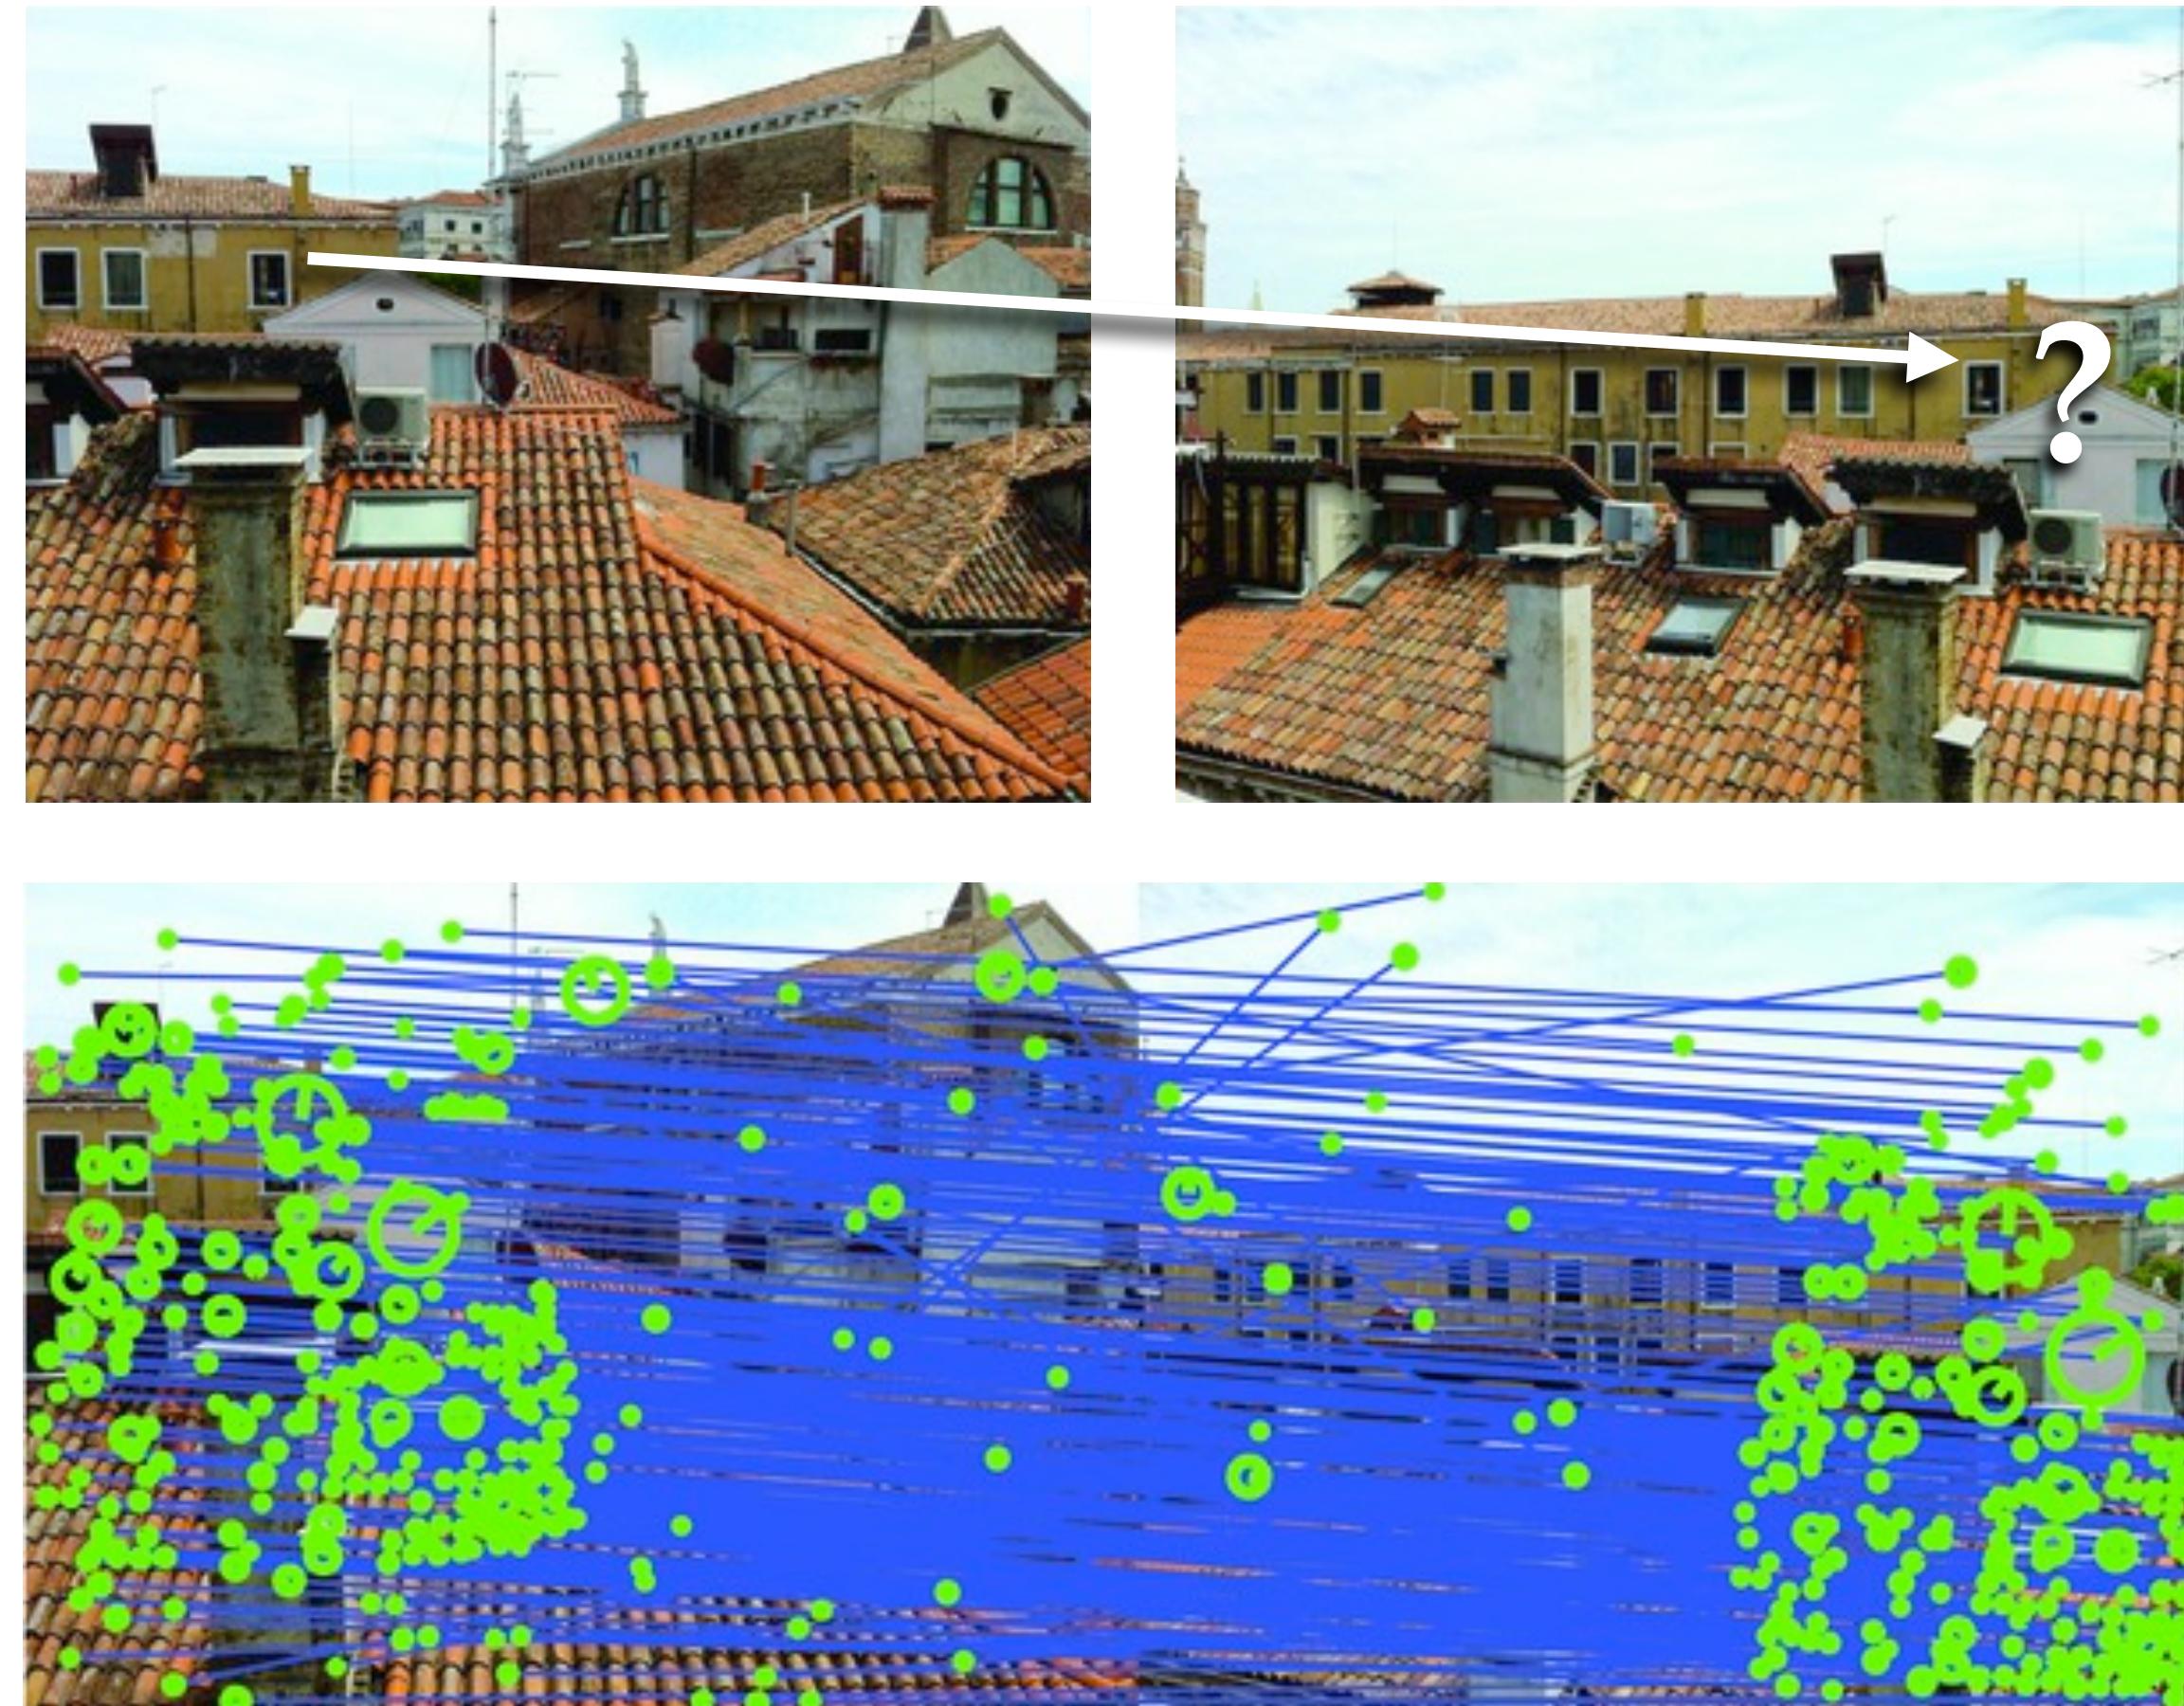


$$\mathbf{u} = \mathbf{Hx}$$

2 points for Similarity
3 for Affine
4 for Homography

Image Alignment

- In practice we have many noisy correspondences + **outliers**



Linear Equations

- e.g., for an affine transform we have a linear system in the unknown parameters \mathbf{a} :

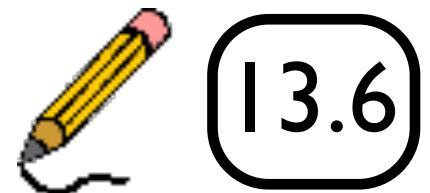
$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_1 & y_1 & 1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_2 & y_2 & 1 \\ x_3 & y_3 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_3 & y_3 & 1 \end{bmatrix} \begin{bmatrix} a_{11} \\ a_{12} \\ a_{13} \\ a_{21} \\ a_{22} \\ a_{23} \end{bmatrix} = \begin{bmatrix} x'_1 \\ y'_1 \\ x'_2 \\ y'_2 \\ x'_3 \\ y'_3 \\ \vdots \\ \vdots \end{bmatrix}$$

- It is **overconstrained** (more equations than unknowns)
- and subject to **outliers** (some rows are completely wrong)

Let's deal with these problems in a simpler context..

Robust Line Fitting

- Consider fitting a line to noisy points



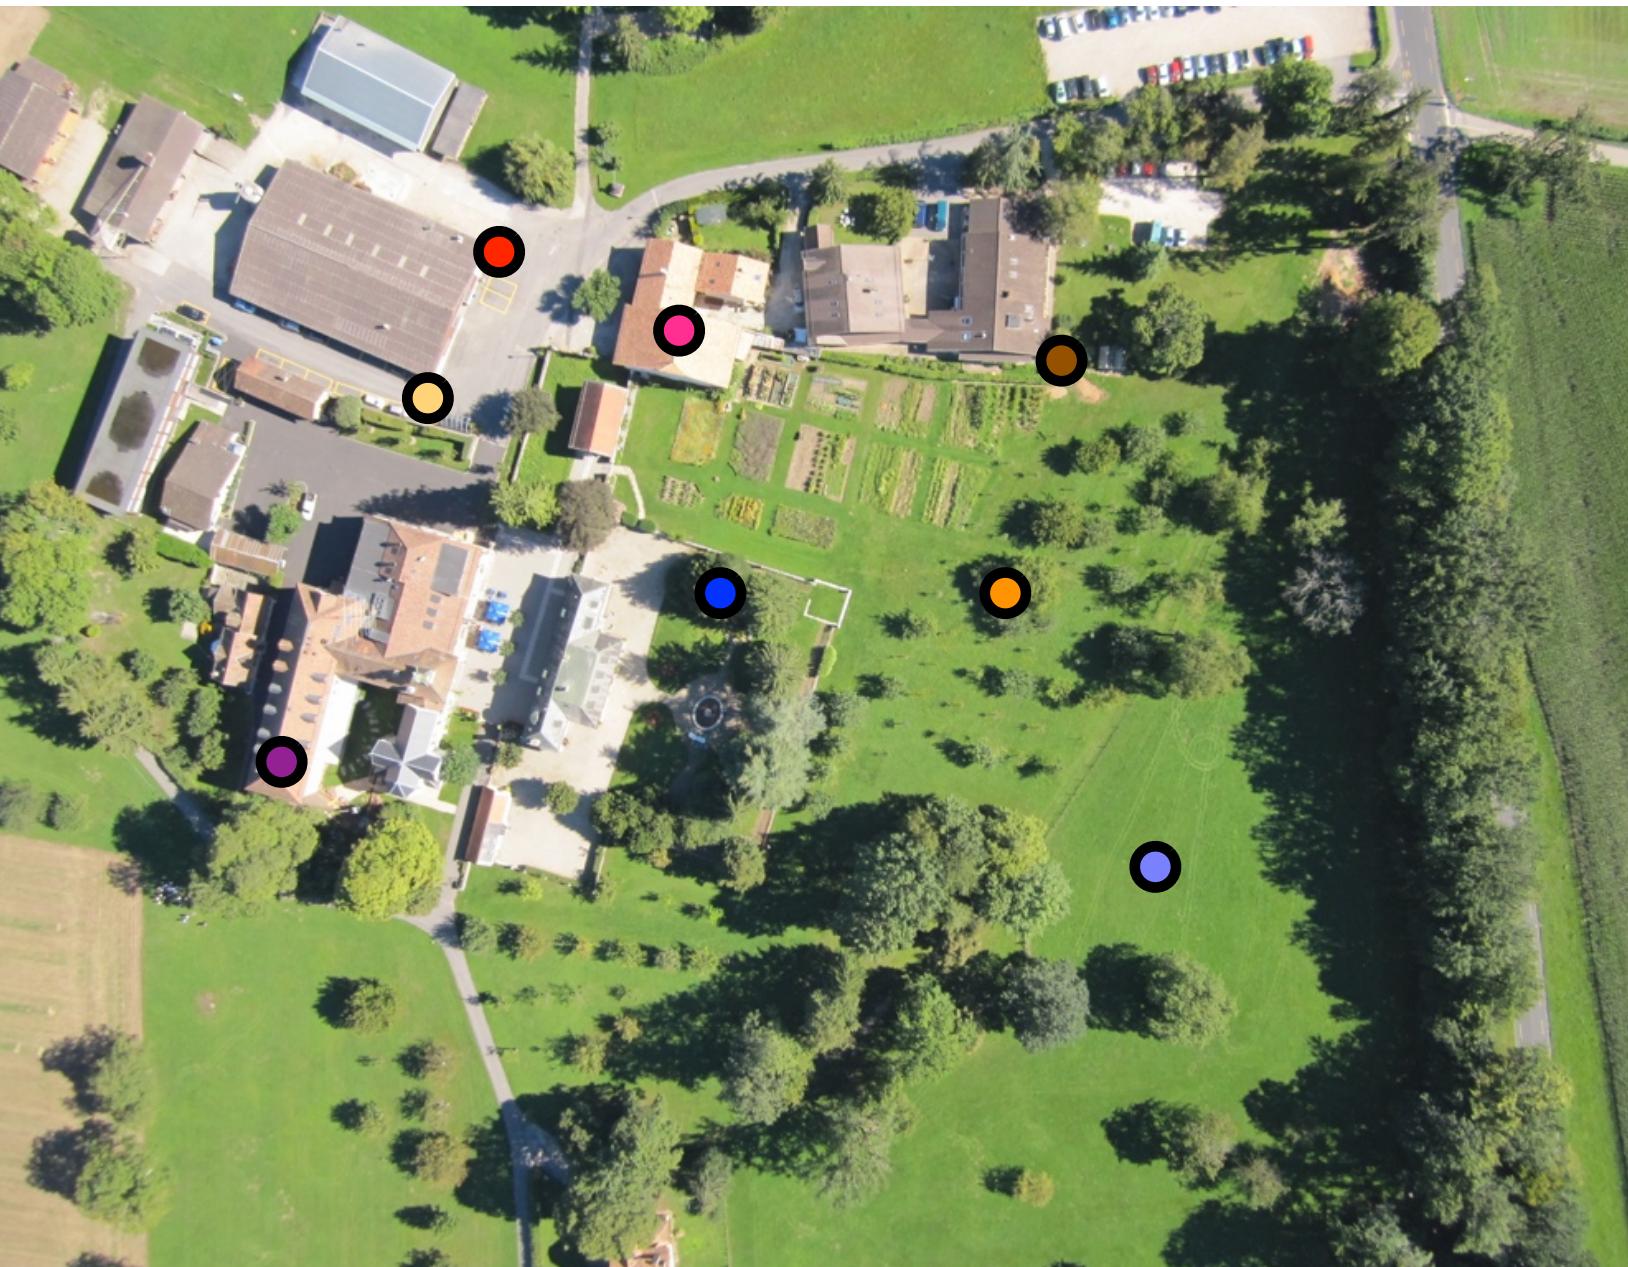
RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



4 inliers (red, yellow, orange, brown),

RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



4 outliers (blue, light blue, purple, pink)

RANSAC Example

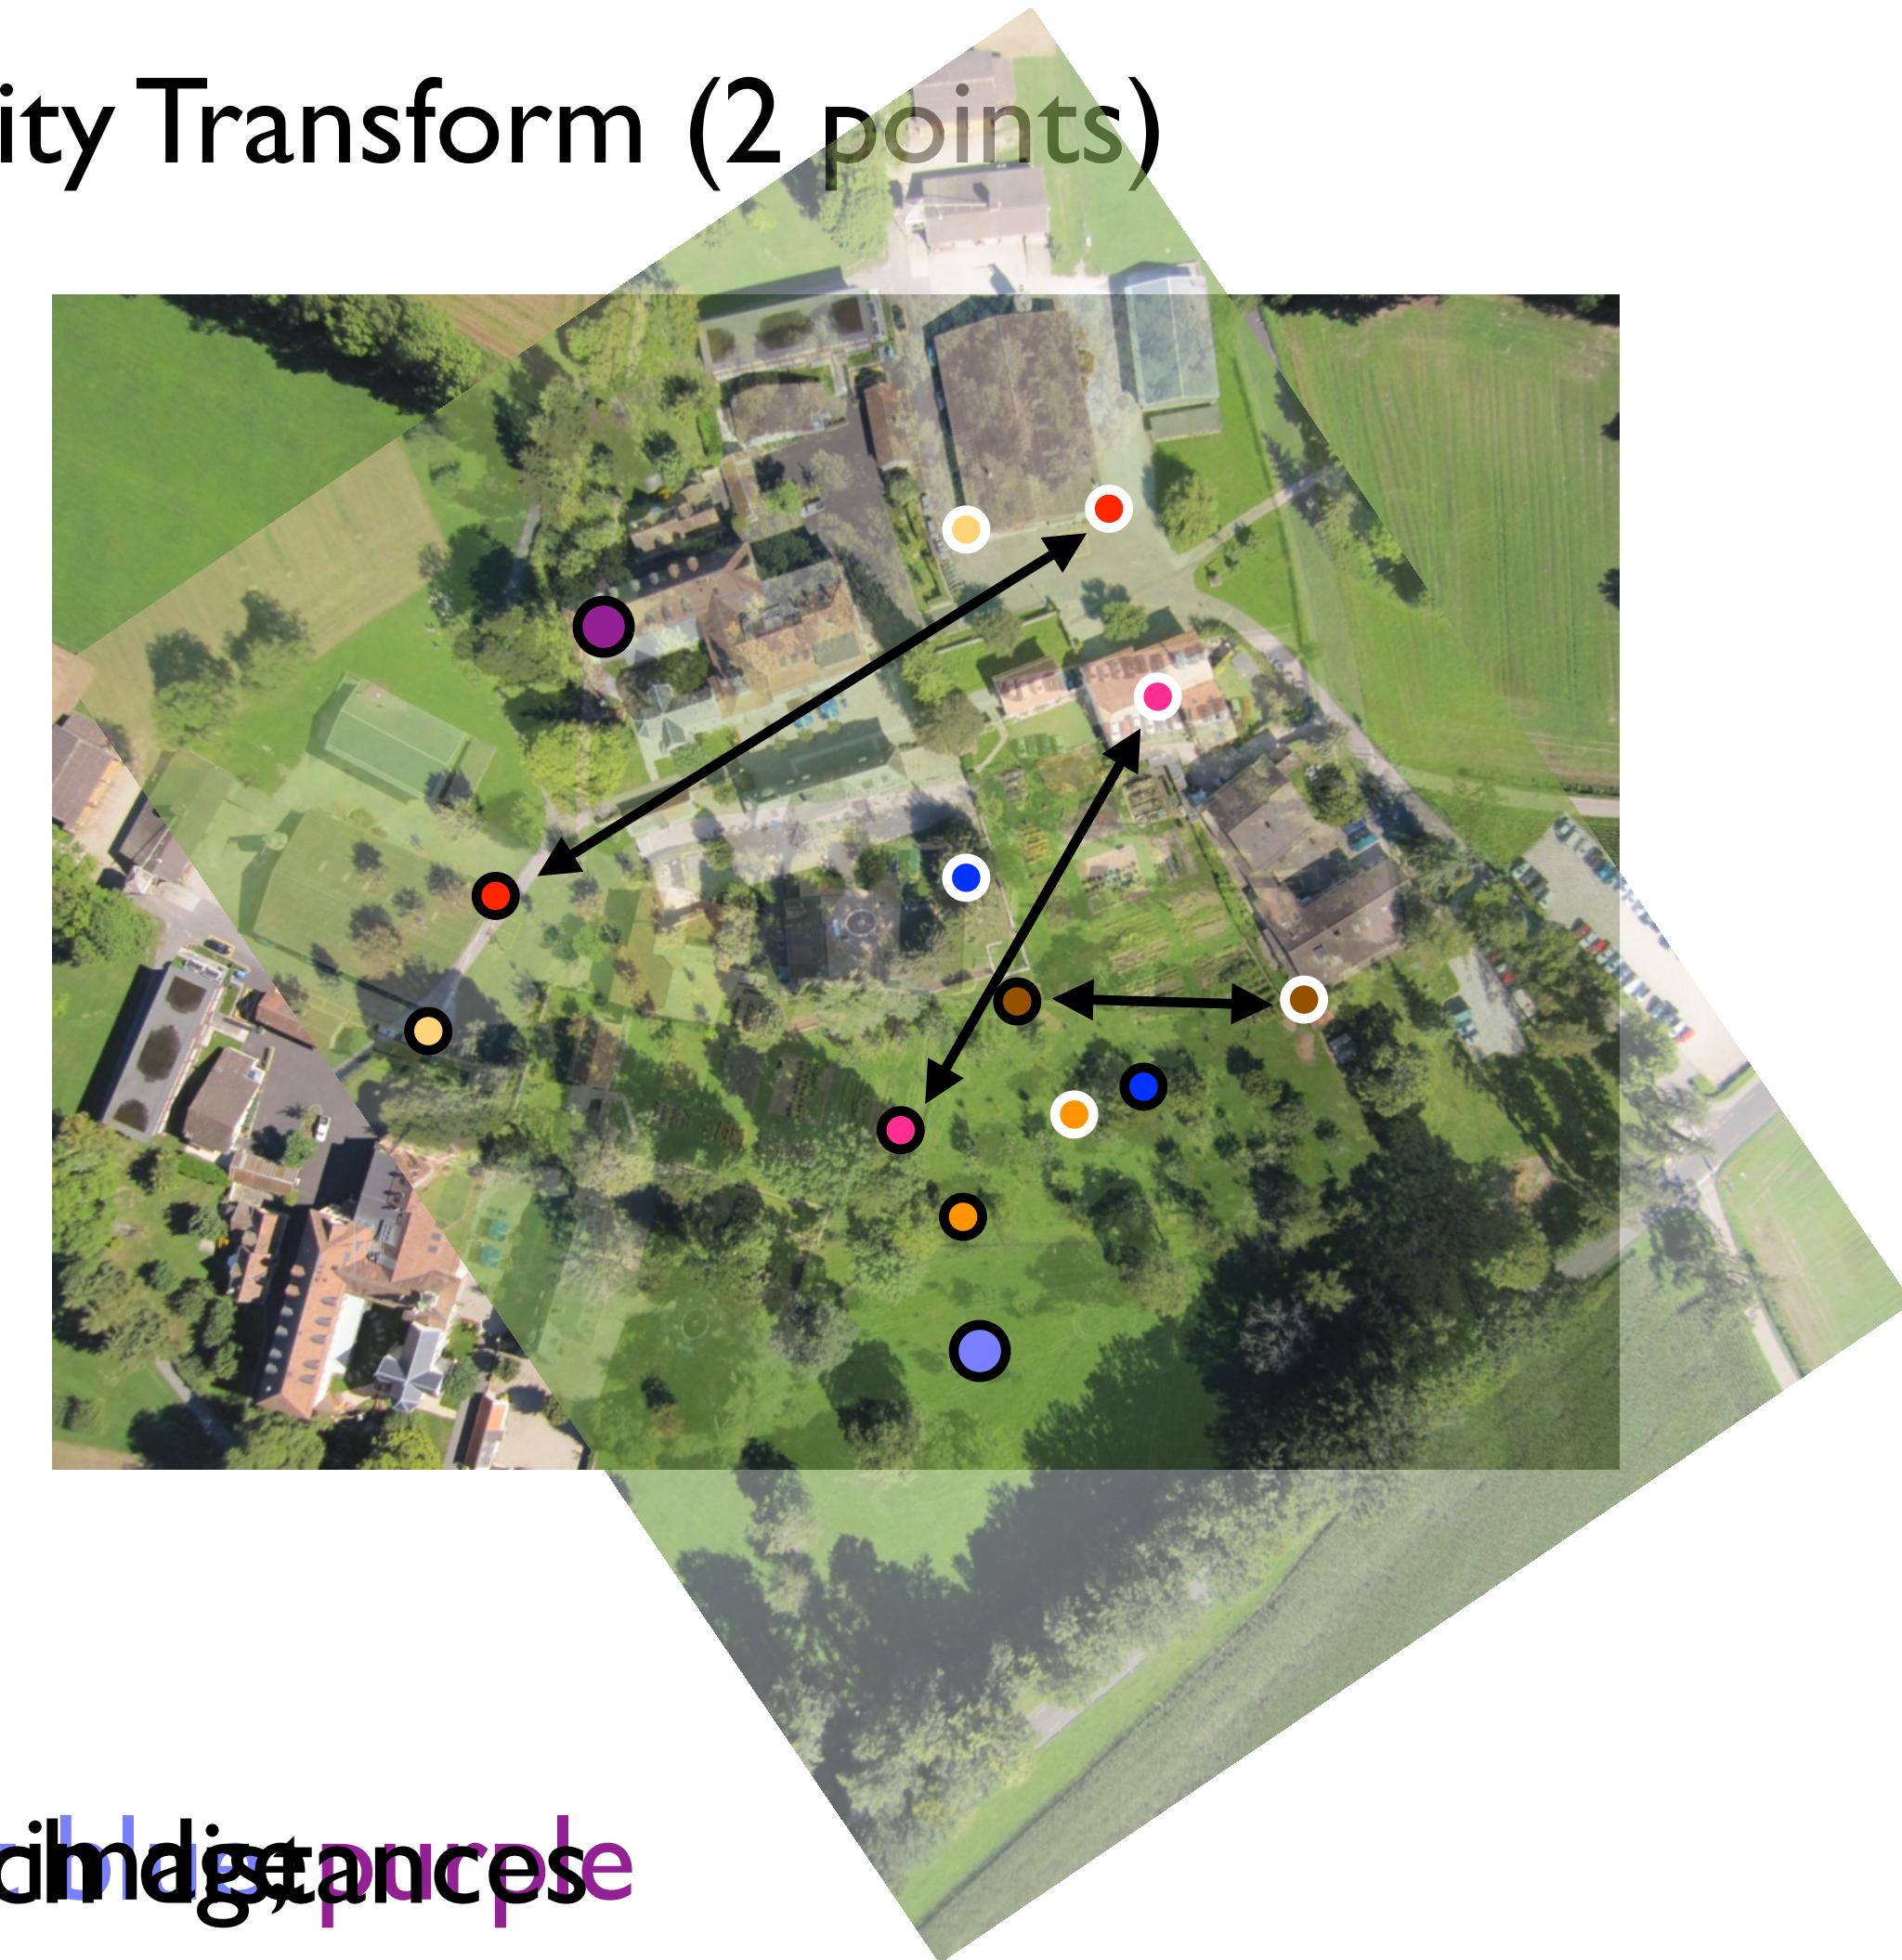
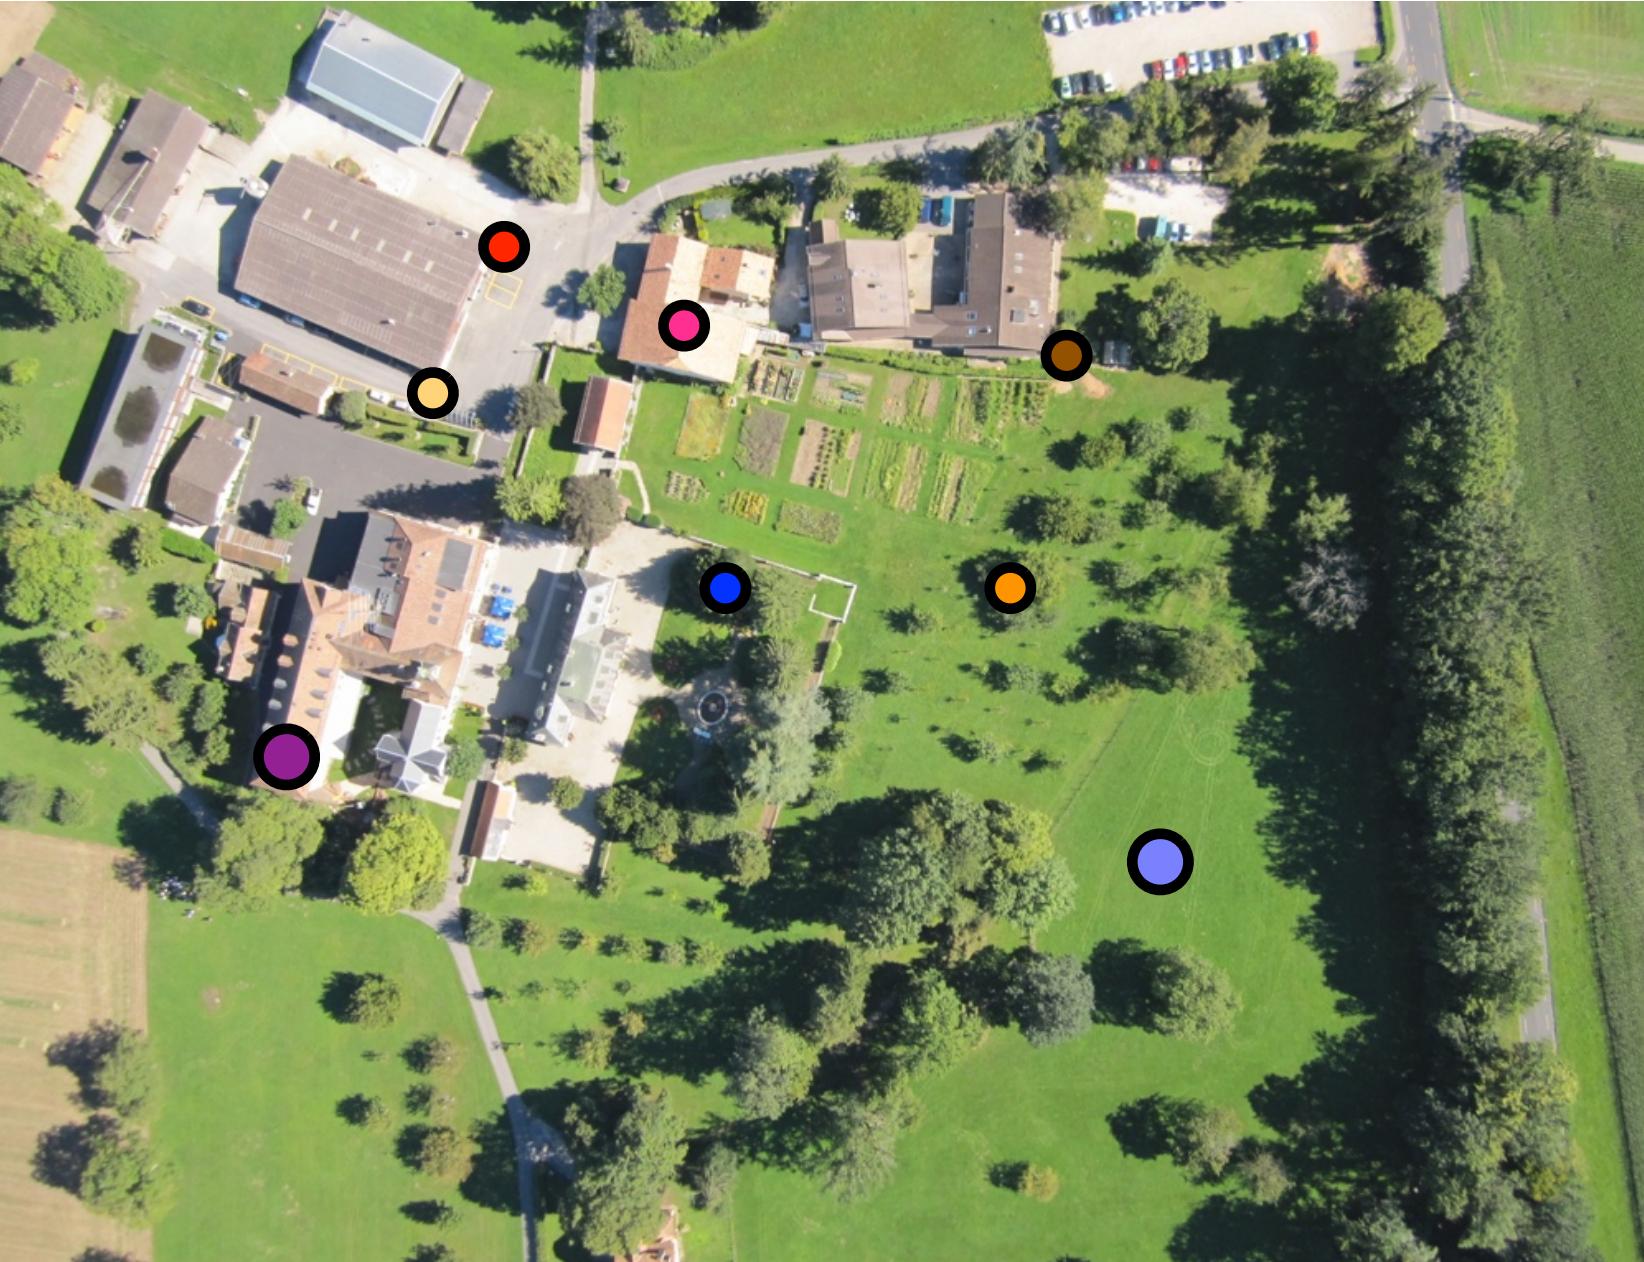
- RANSAC solution for Similarity Transform (2 points)



4 inliers (red, yellow, orange, brown),
4 outliers (blue, light blue, purple, pink)

RANSAC Example

- RANSAC solution for Similarity Transform (2 points)

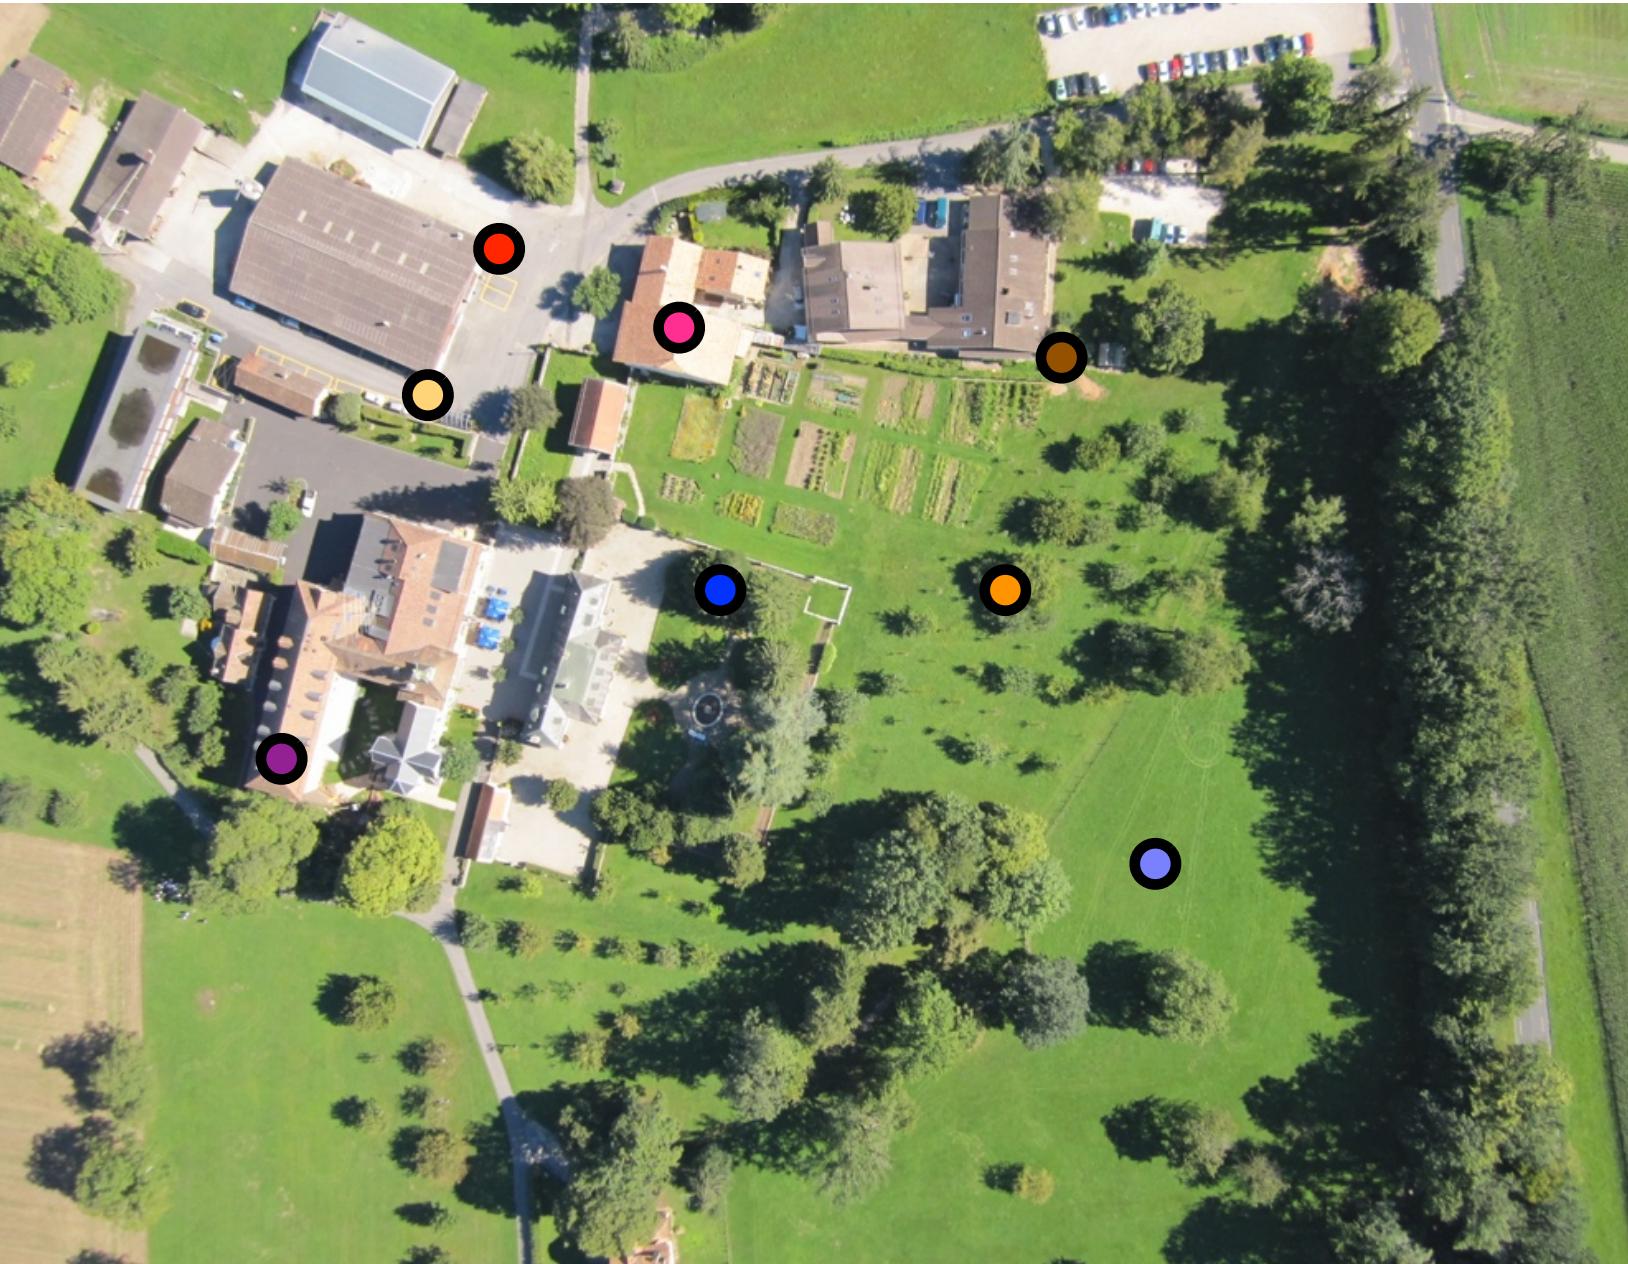


chbe slow light blue principle

#inliers = 2

RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



chebkoosappinkblue

#inliers = 2

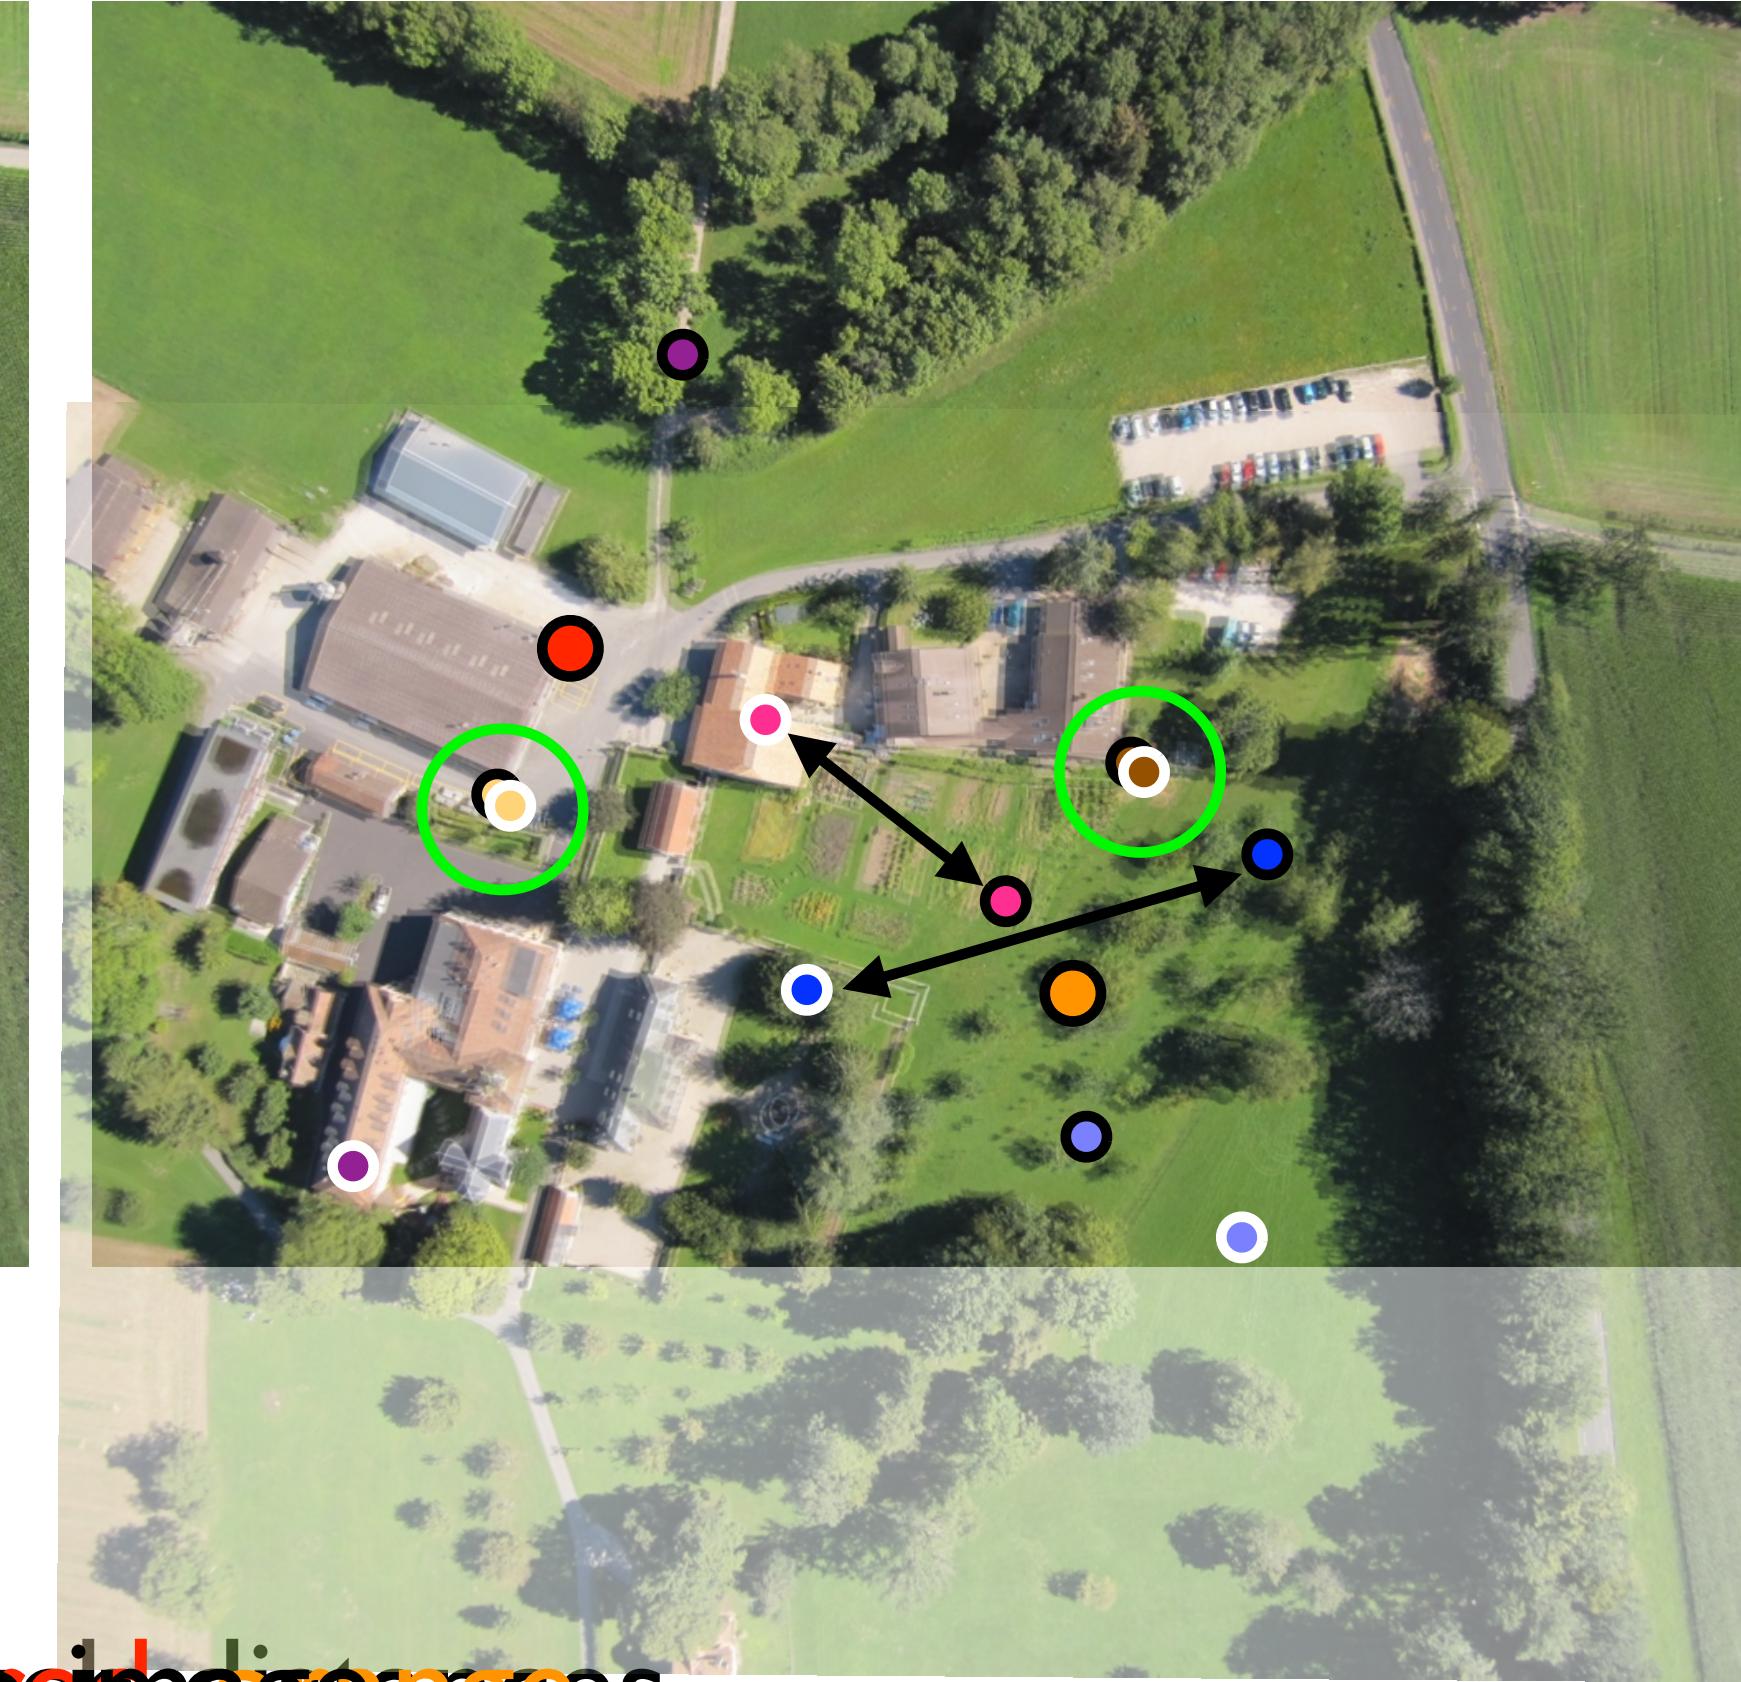
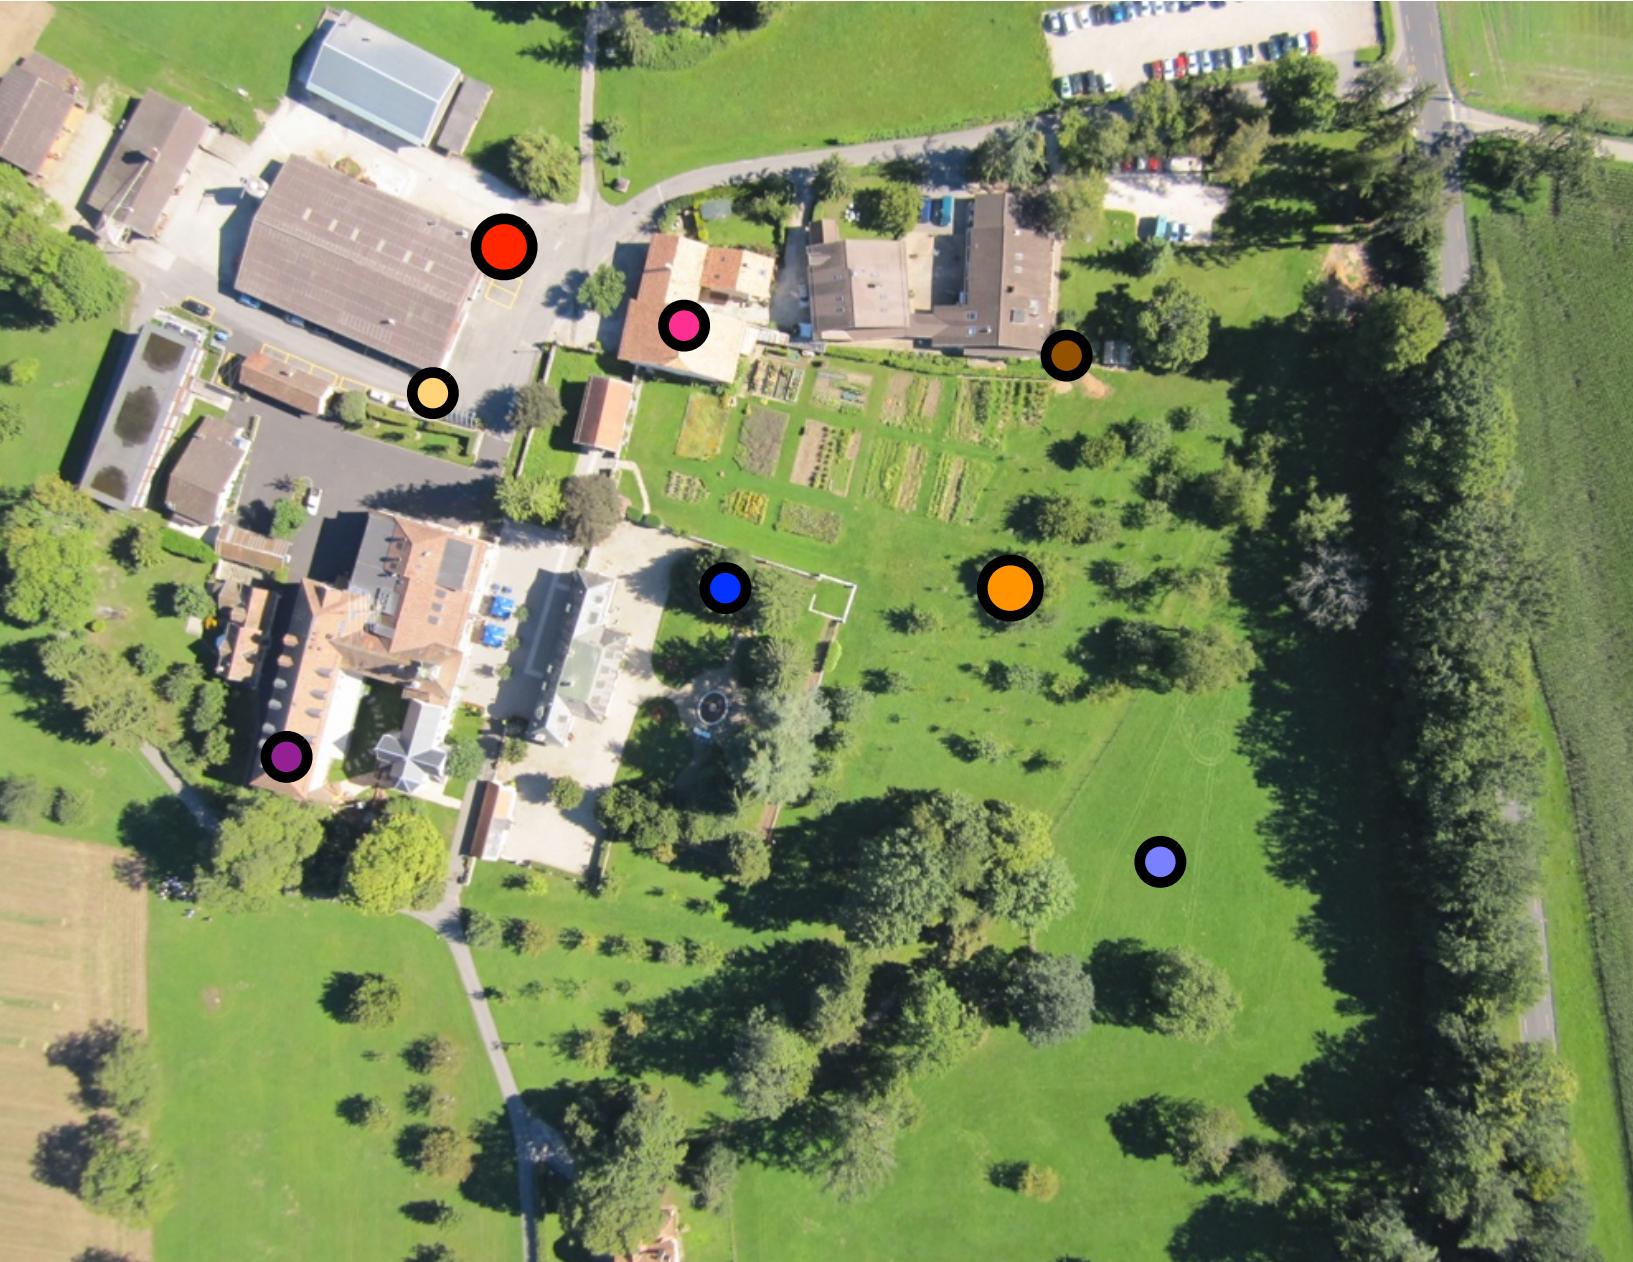
RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



RANSAC Example

- RANSAC solution for Similarity Transform (2 points)



check sample in image 2

#inliers = 4

RANSAC Example

- RANSAC solution for Similarity Transform (2 points)

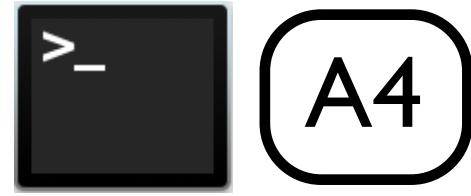


RANSAC algorithm

1. Match feature points between 2 views
2. Select minimal subset of matches*
3. Compute transformation T using minimal subset
4. Check consistency of all points with T — compute projected position and count #inliers with distance $<$ threshold
5. Repeat steps 2-4 to maximise #inliers

* Similarity transform = 2 points, Affine = 3, Homography = 4

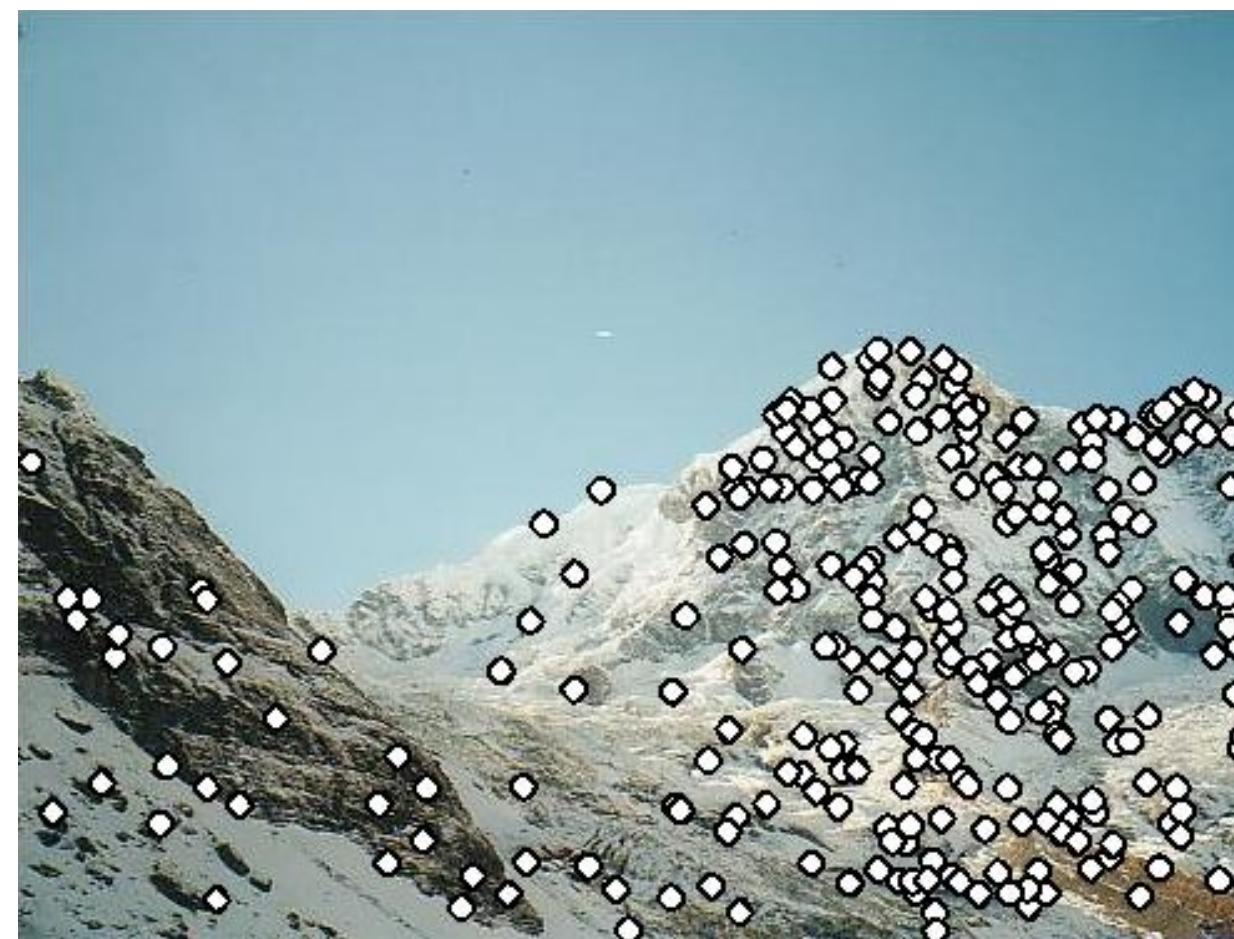
Assignment 4



- Try out the **RANSAC Implementation** section in Assignment 4

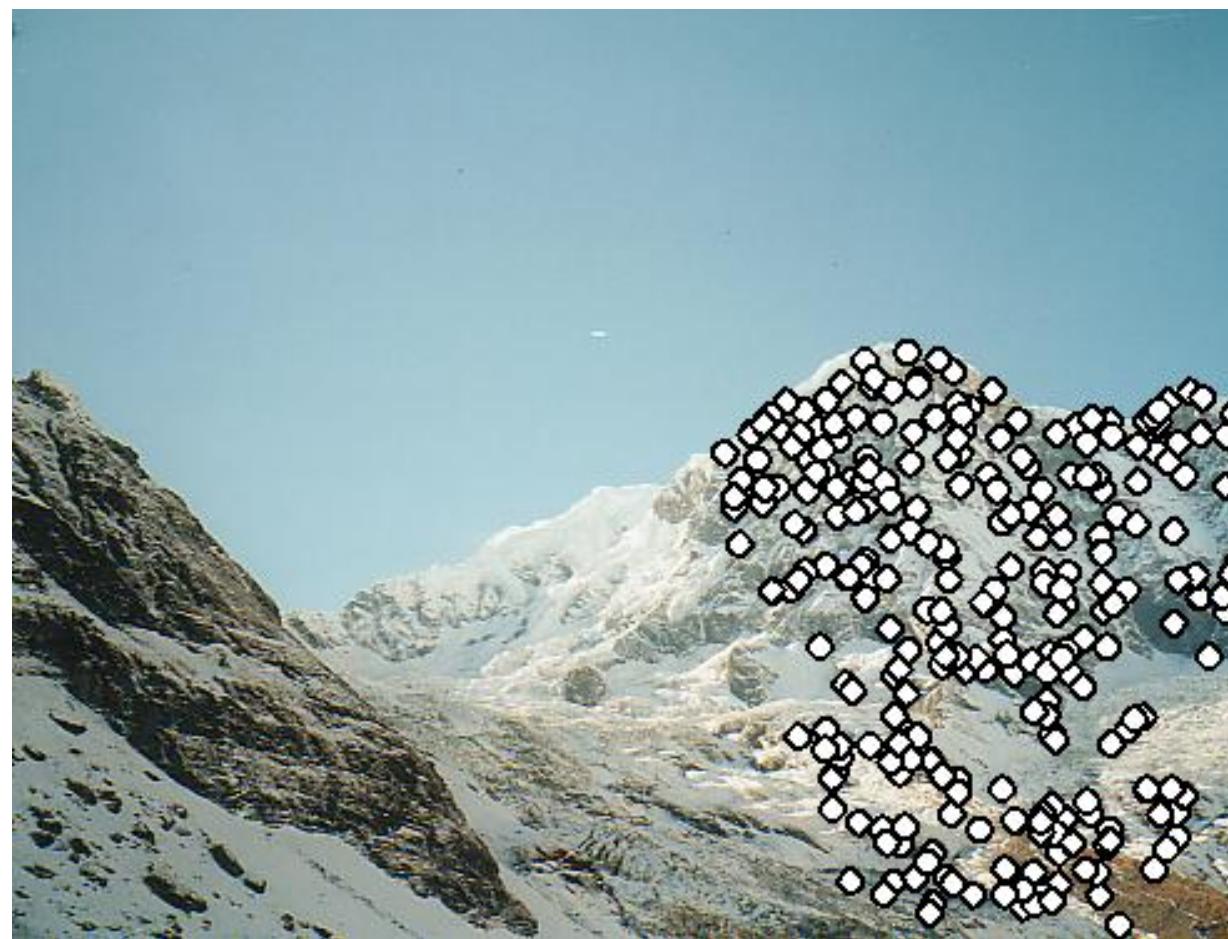
2-view Rotation Estimation

- Find features + raw matches, use RANSAC to find Similarity



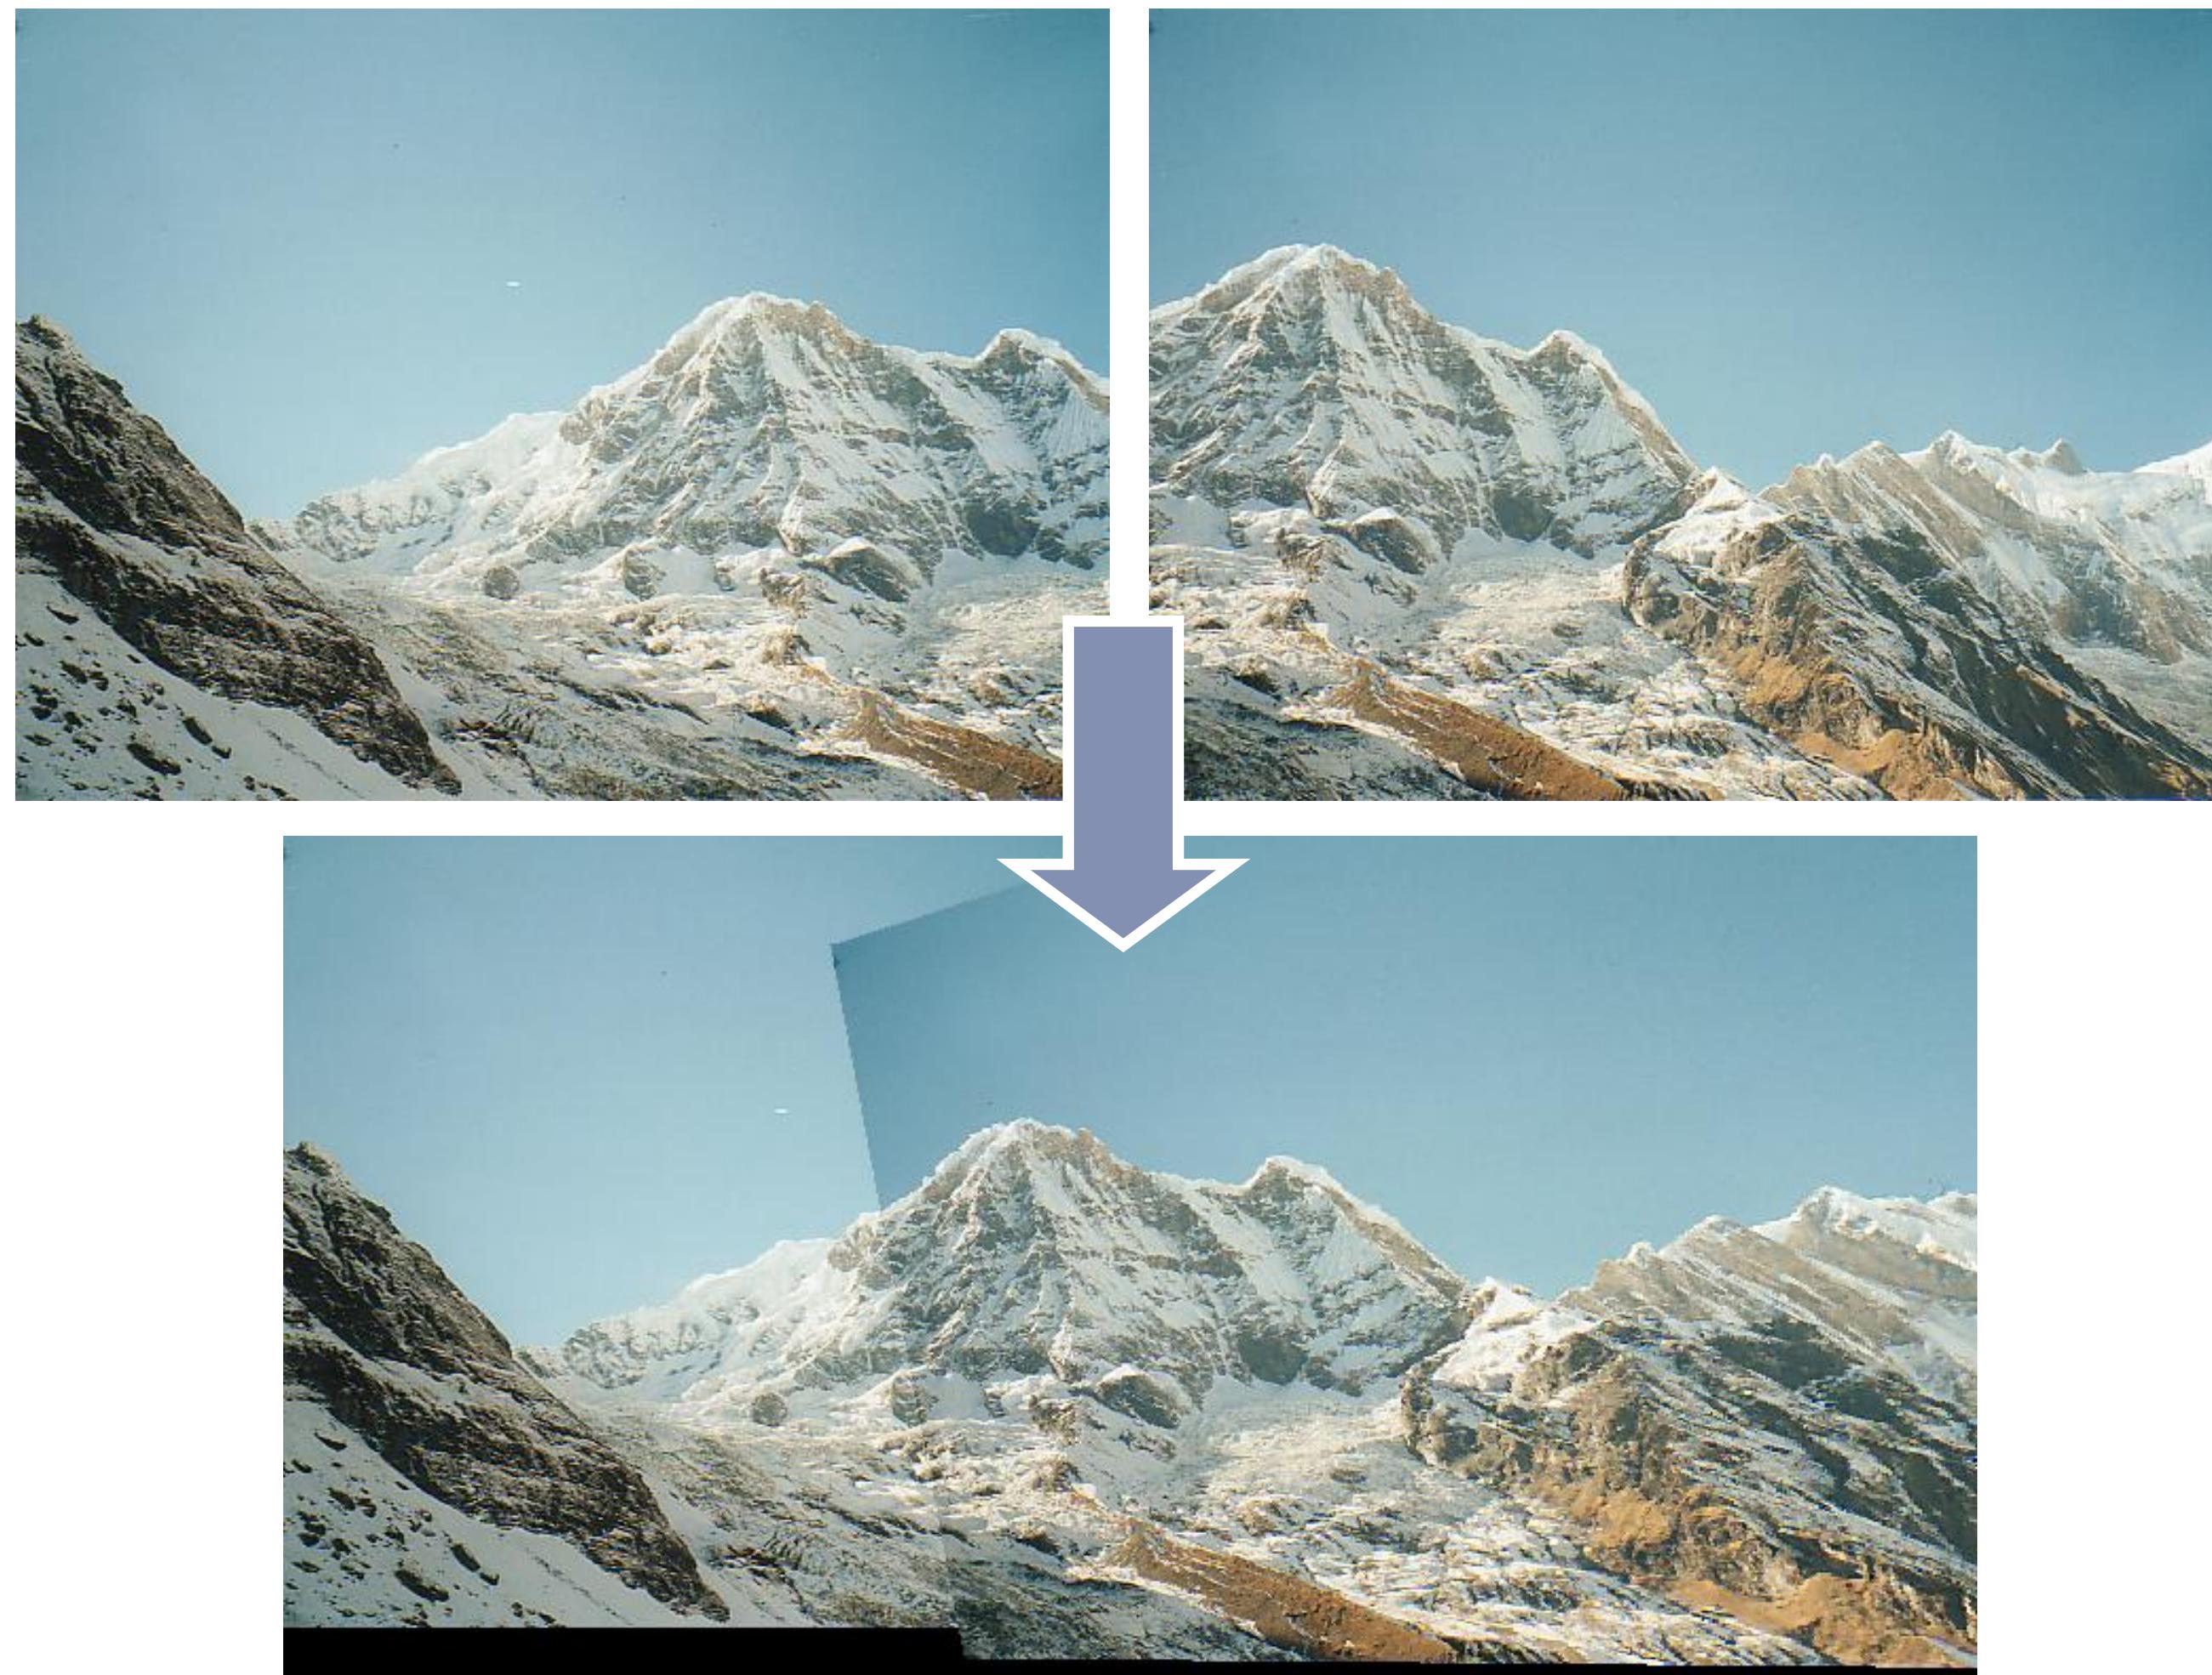
2-view Rotation Estimation

- Remove outliers, can now solve for R using least squares



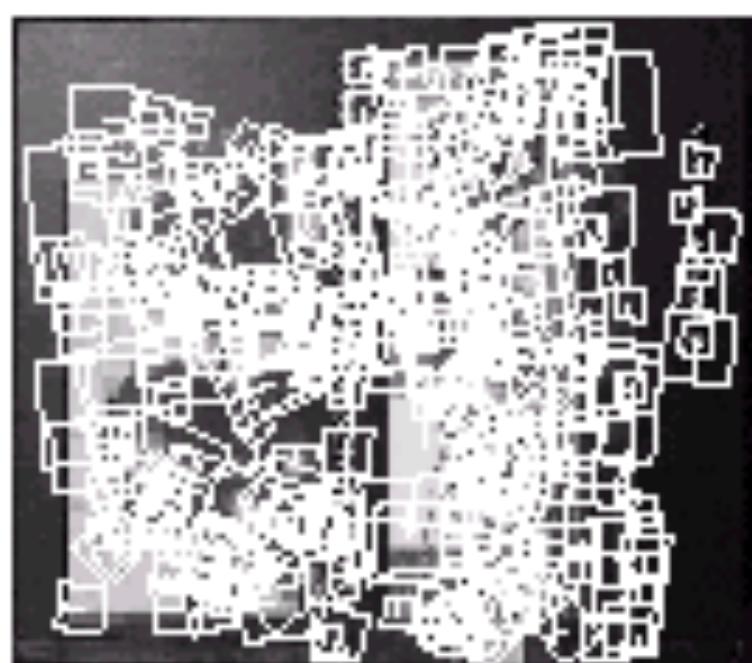
2-view Rotation Estimation

- Final rotation estimation

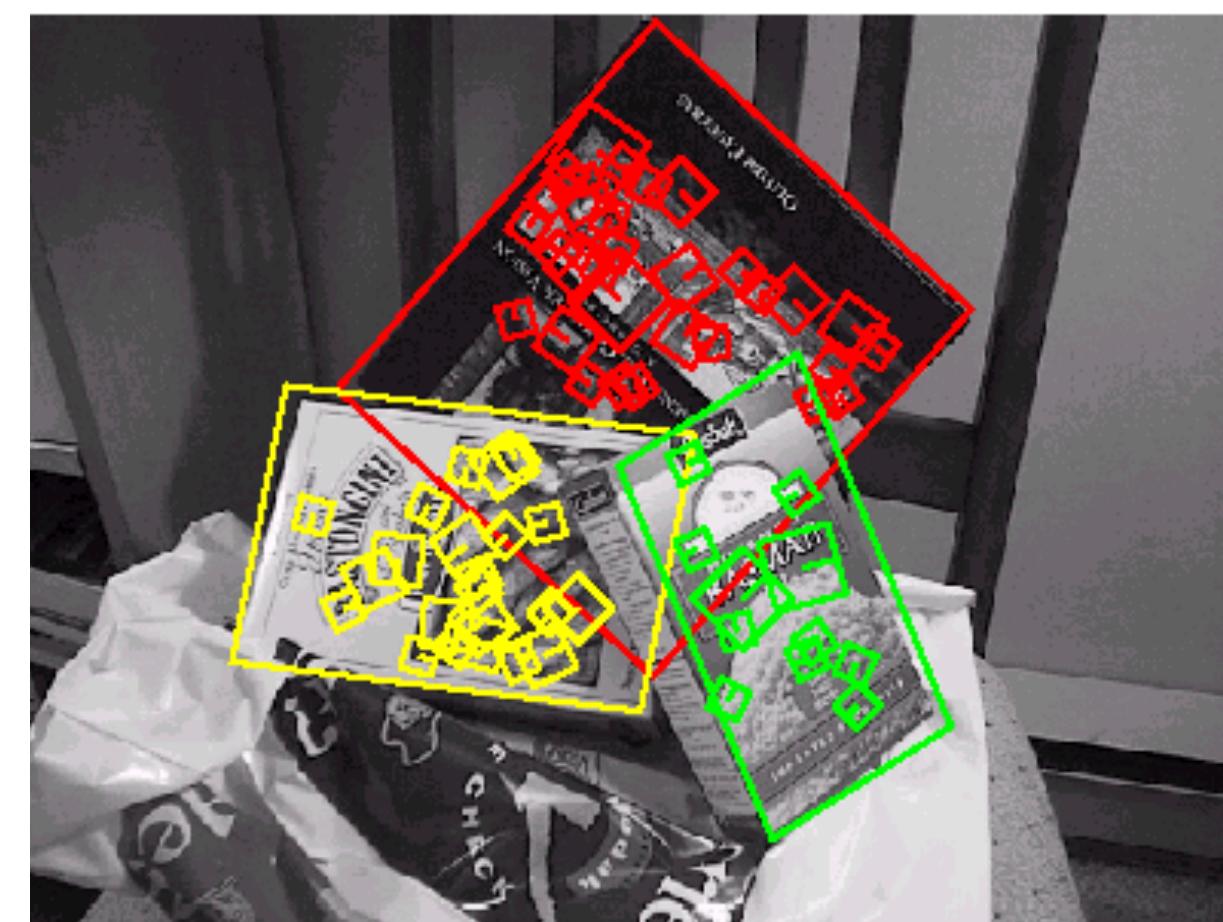


Object Instance Recognition

Database of planar objects



Instance recognition



Object Instance Recognition with SIFT

Match SIFT descriptors between **query image** and a database of known keypoints extracted from **training examples**

- use fast (approximate) nearest neighbour matching
- threshold based on ratio of distances between 1NN and 2NN

Use **RANSAC** to find a **subset of matches** that all agree on an object and geometric transform (e.g., **affine transform**)

Optionally **refine pose estimate** by recomputing the transformation using all the RANSAC inliers

Fitting a Model to Noisy Data

Suppose we are **fitting a line** to a dataset that consists of 50% outliers

We can fit a line using two points

If we draw pairs of points uniformly at random, what fraction of pairs will consist entirely of ‘good’ data points (inliers)?



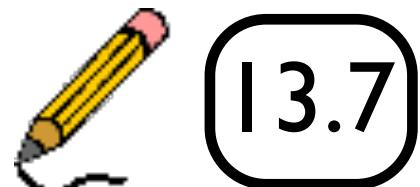
RANSAC: How many samples?

Let p_0 be the fraction of outliers (i.e., points on line)

Let n be the number of points needed to define hypothesis
($n = 2$ for a line in the plane)

Suppose k samples are chosen

How many samples do we need to find a good solution?



RANSAC: k Samples Chosen ($p = 0.99$)

Sample size	Proportion of outliers							
	5%	10%	20%	25%	30%	40%	50%	
n								
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

Figure Credit: Hartley & Zisserman

Lecture 14: Re-cap RANSAC

RANSAC is a technique to fit data to a model

- divide data into inliers and outliers
- estimate model from minimal set of inliers
- improve model estimate using all inliers
- alternate fitting with re-classification as inlier/outlier

RANSAC is a general method suited for a wide range of model fitting problems

- easy to implement
- easy to estimate/control failure rate

RANSAC only handles a moderate percentage of outliers without cost blowing up

Menu for Today

Topics:

- **Planar Geometry**
- **Image Alignment**, Object Recognition
- **RANSAC**

Readings:

- **Today's Lecture:** Szeliski 2.1, 8.1, Forsyth & Ponce 10.4.2

Reminders:

- **Assignment 4:** RANSAC and Panorama Stitching – **now available**